

MACHINE LEARNING AND AI



Pablo Doval

DATA PONTIFEX @Plain Concepts

I work with code and data, but don't tell my mom; she thinks I'm a piano player in a whorehouse.



A BIT OF STORY



A.I.

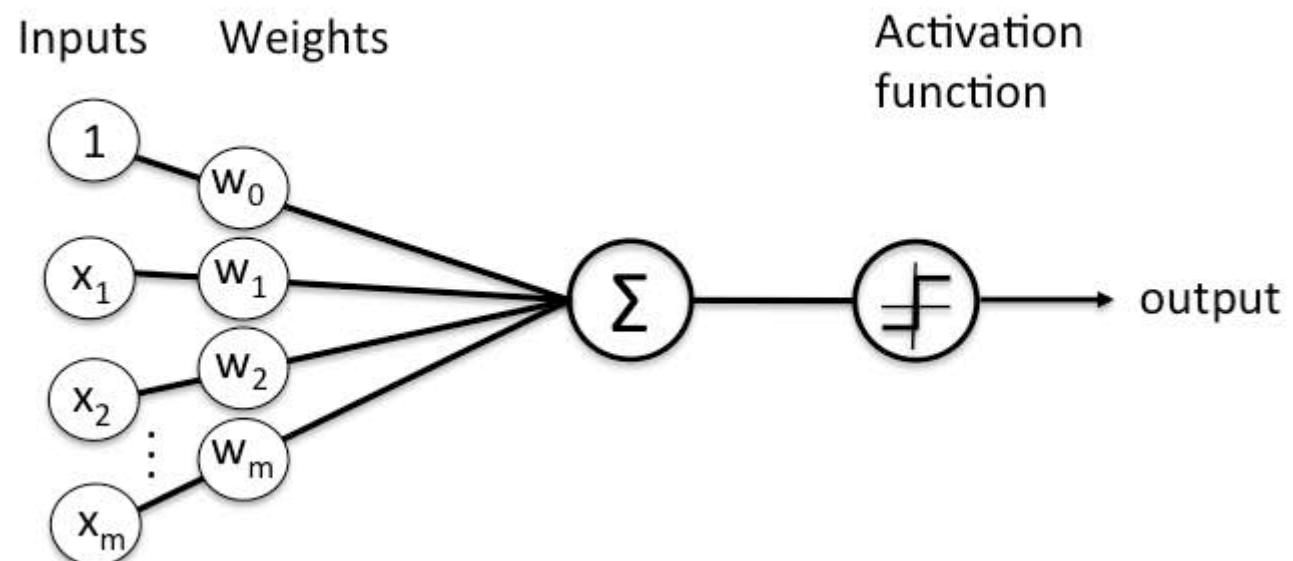
MACHINE LEARNING

DEEP LEARNING

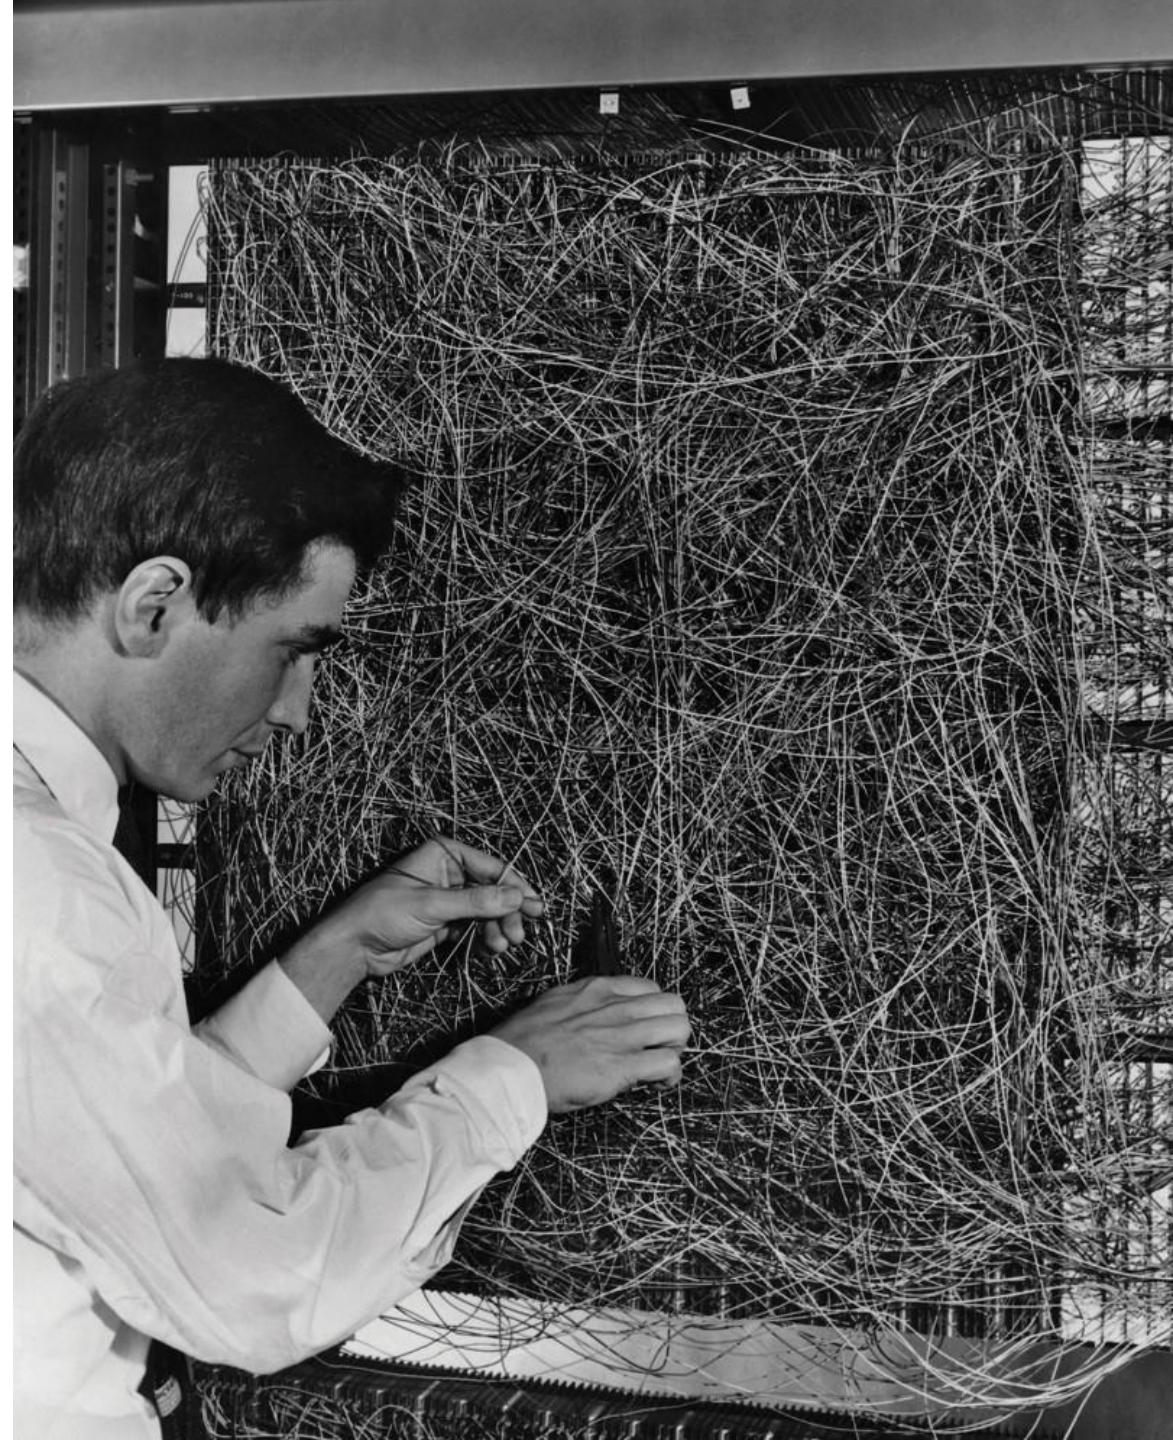
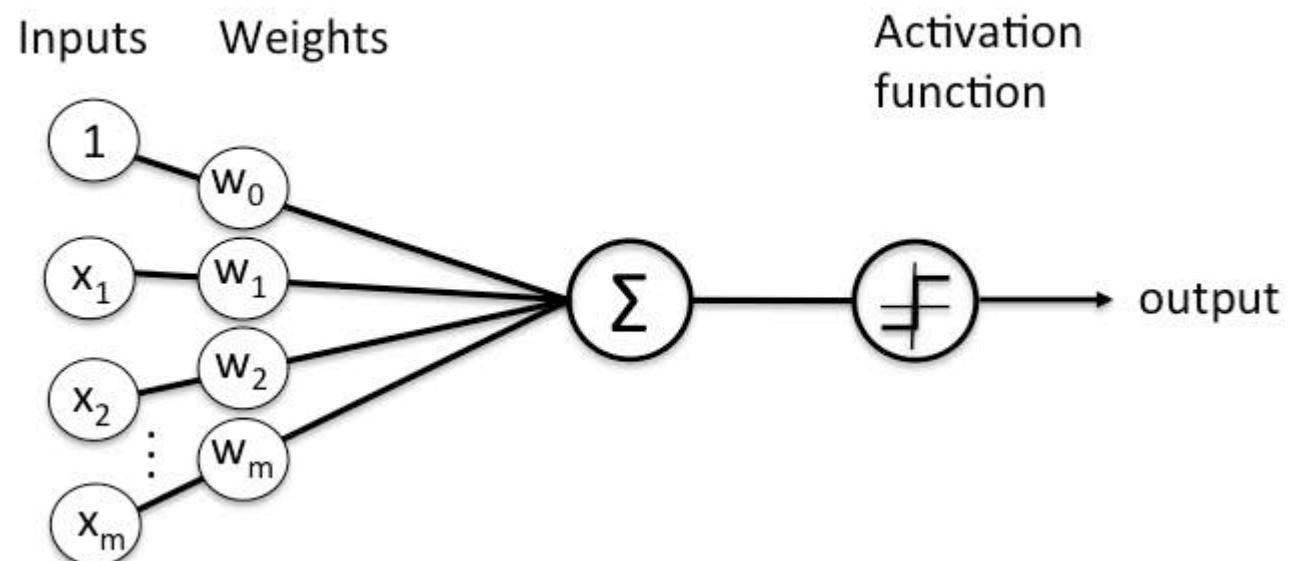




PERCEPTRON



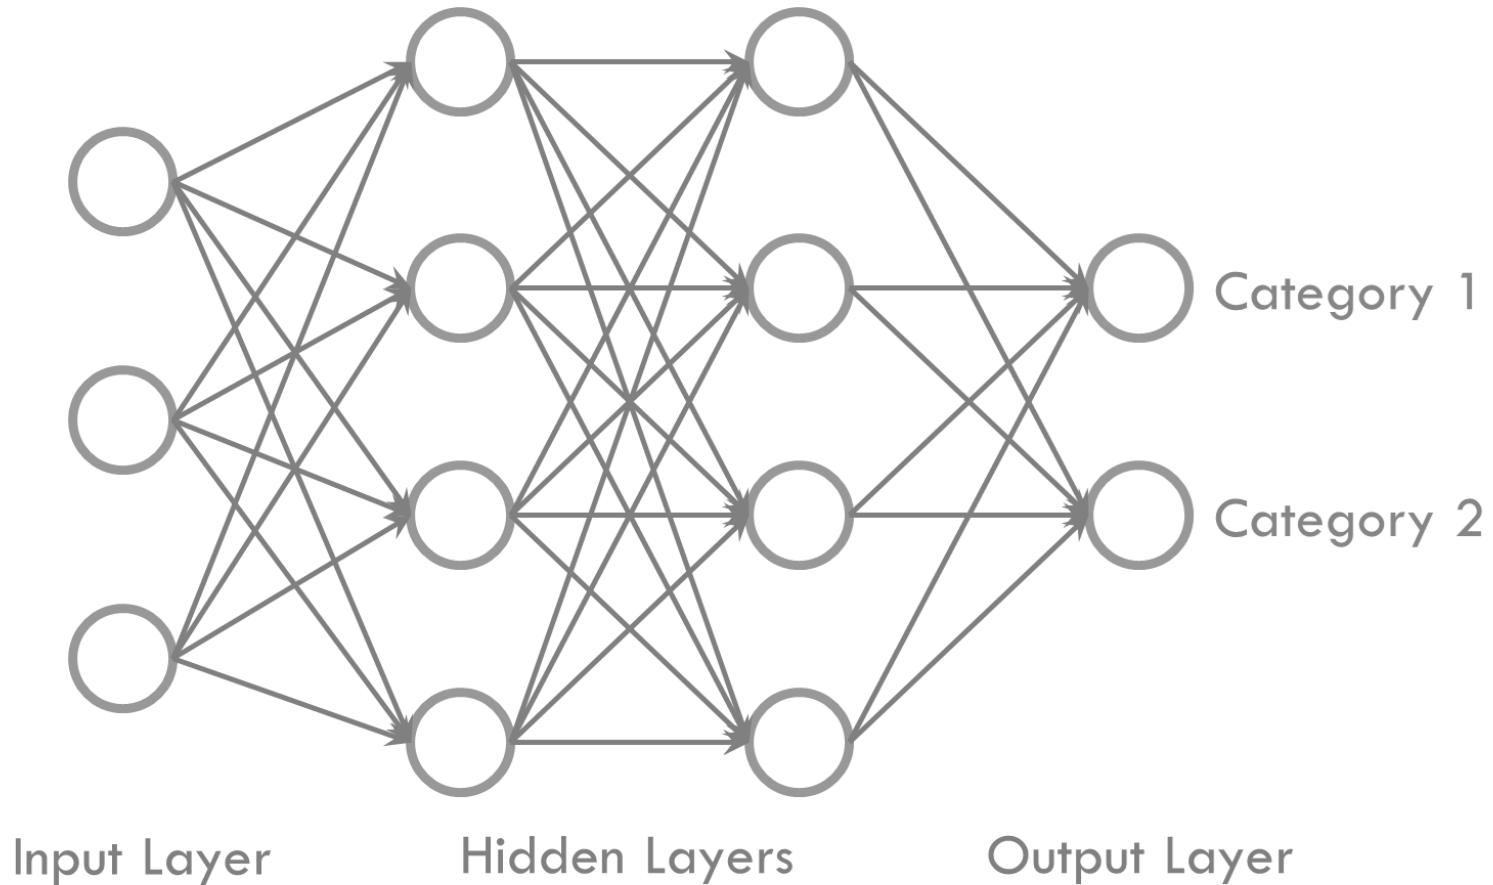
PERCEPTRON



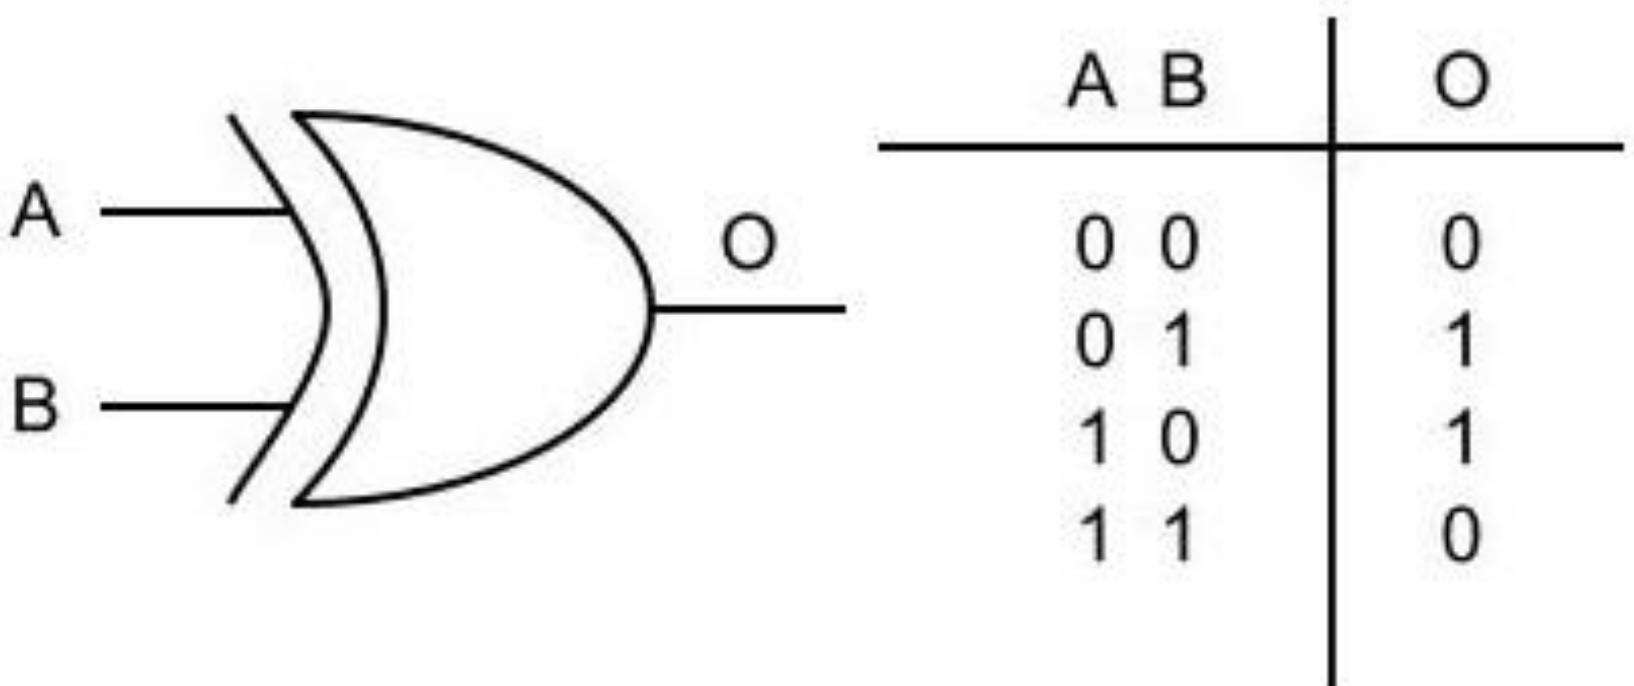
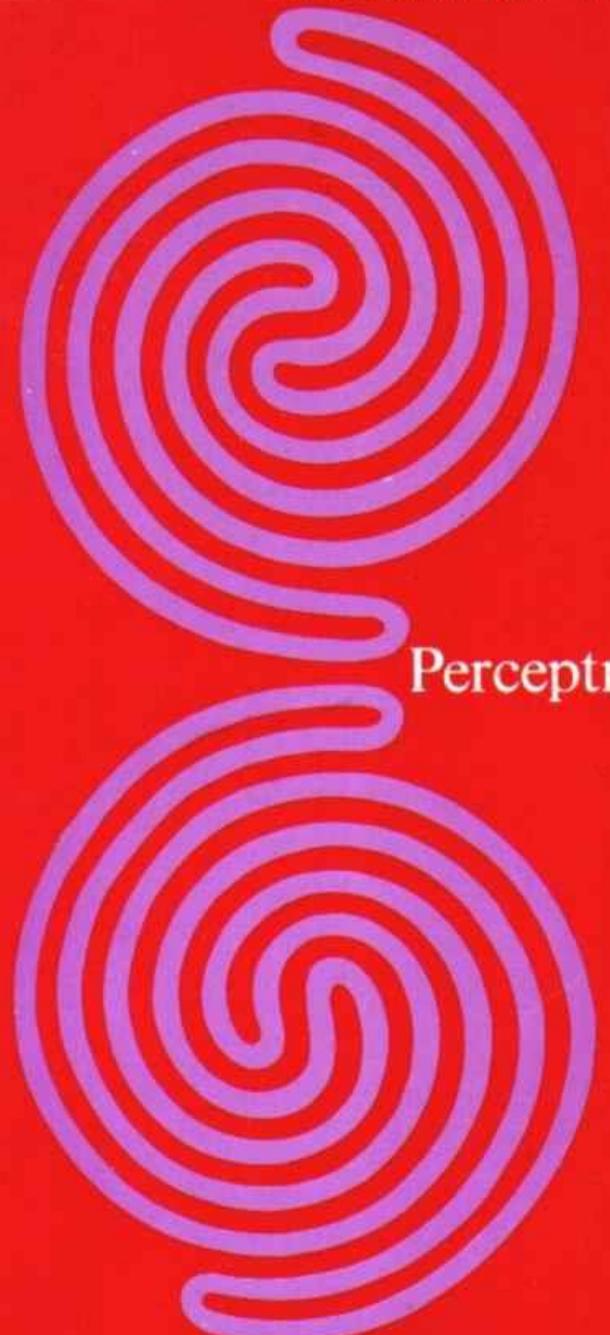
ARTIFICIAL INTELLIGENCE



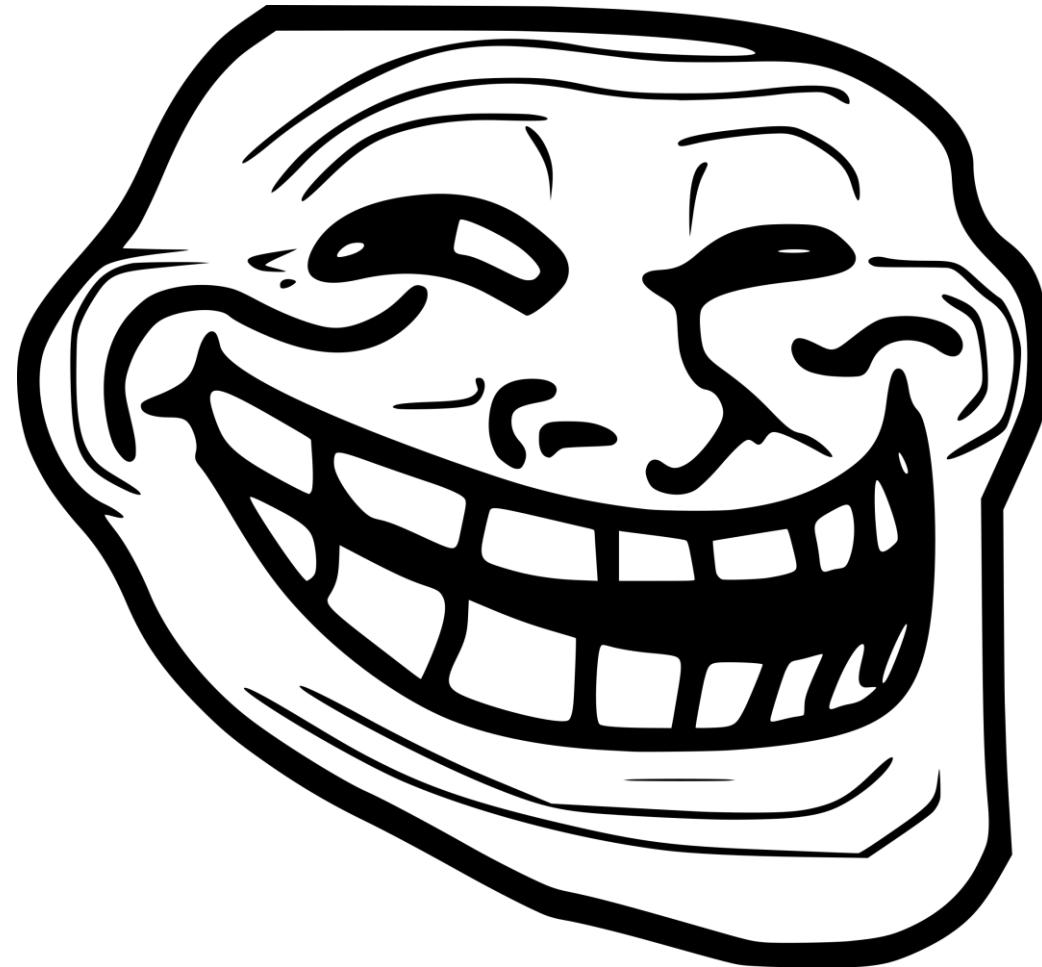
MULTILAYER PERCEPTRON

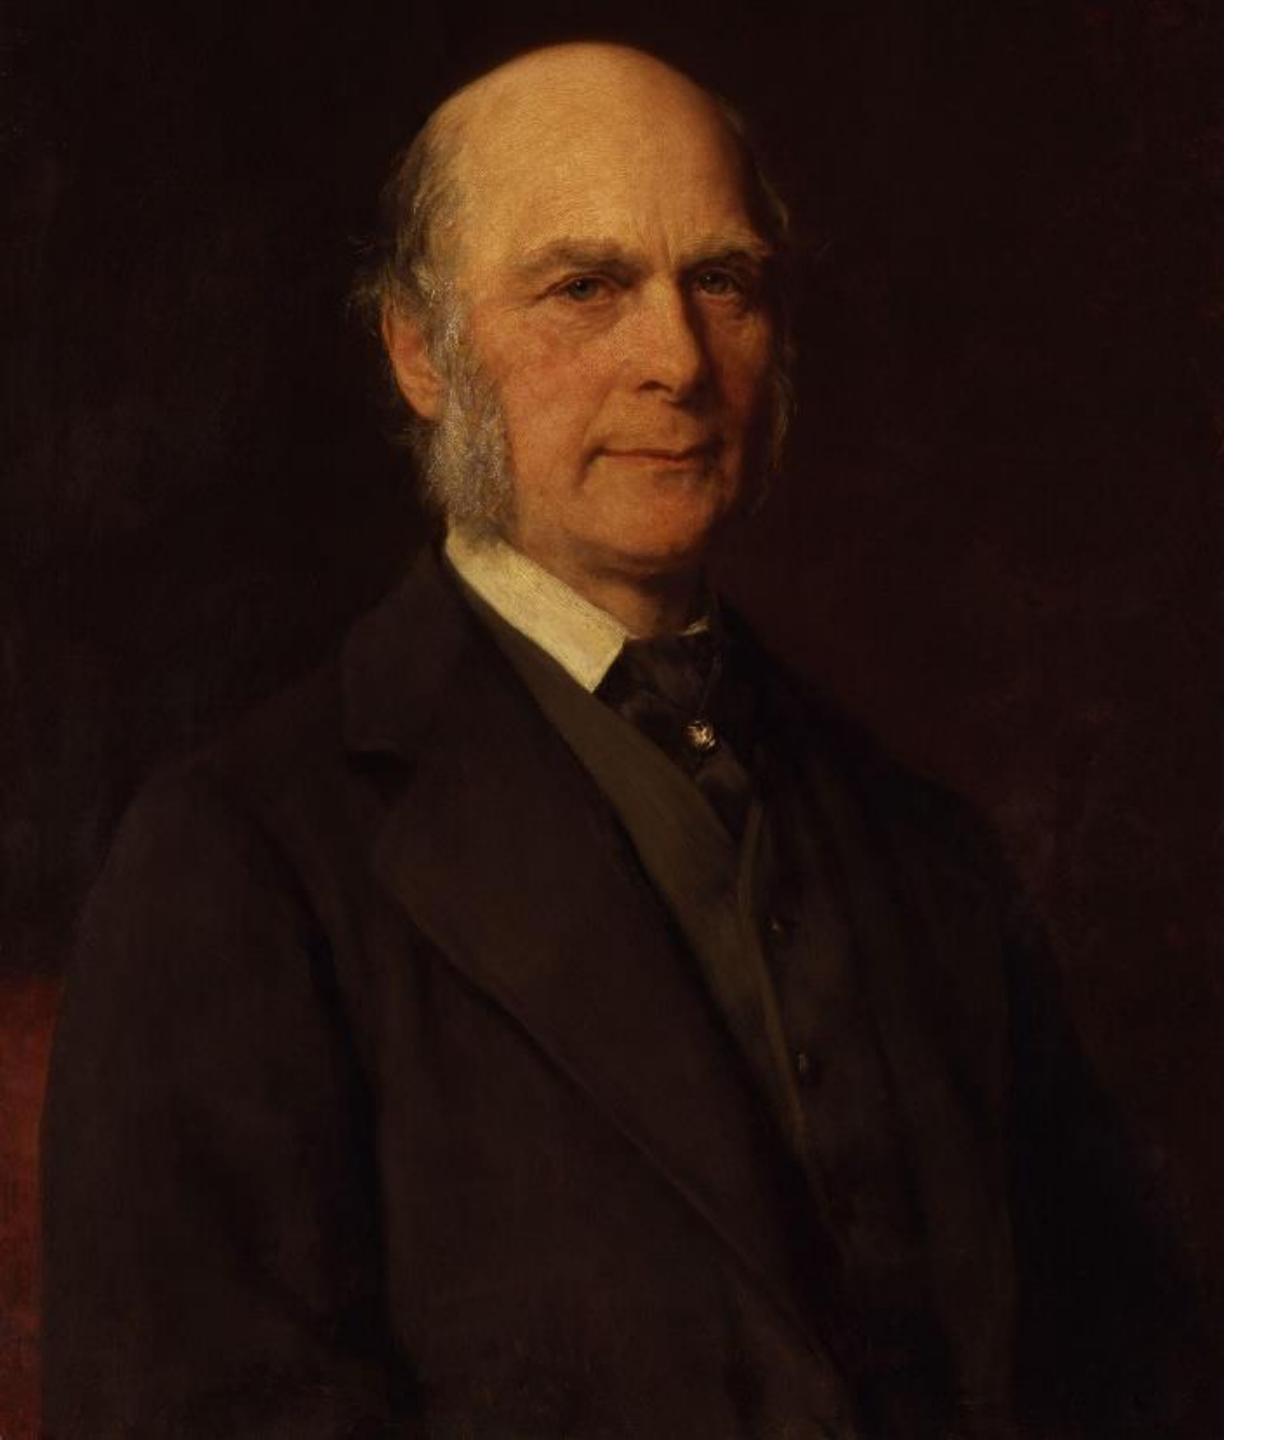


WINTER
IS COMING

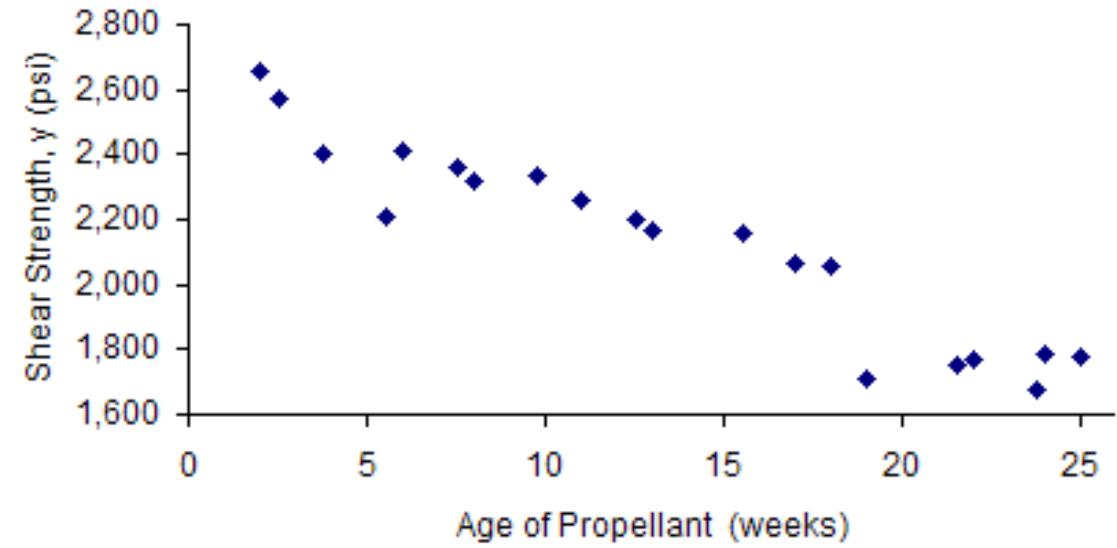


EXPERT SYSTEMS?

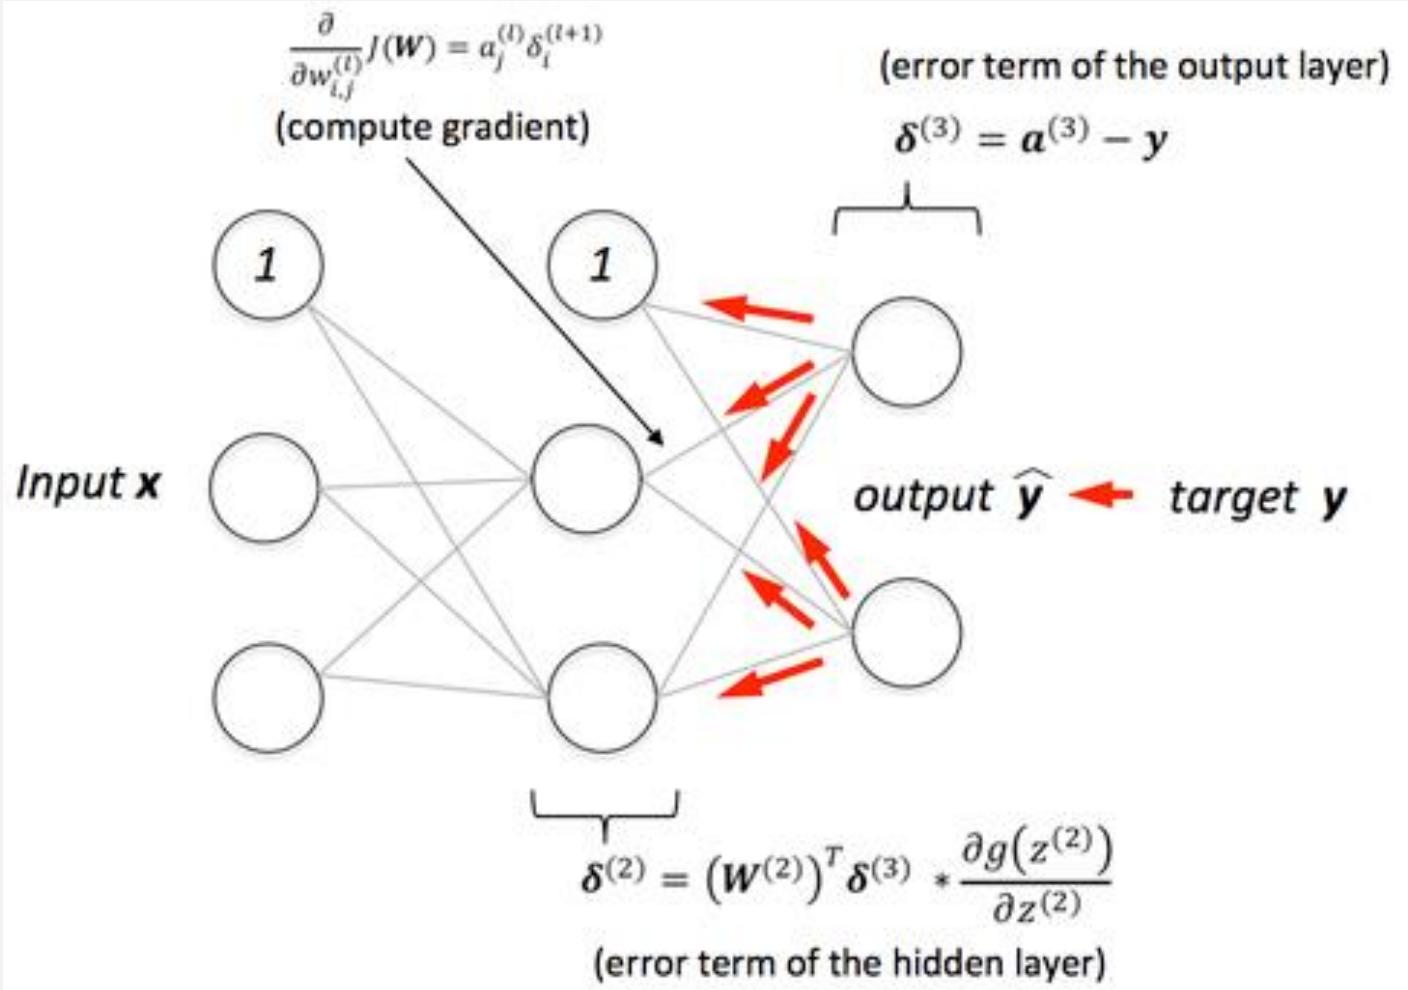


A portrait painting of a man with white hair, wearing a dark suit, white shirt, and patterned tie. He is looking slightly to his right.

MACHINE LEARNING

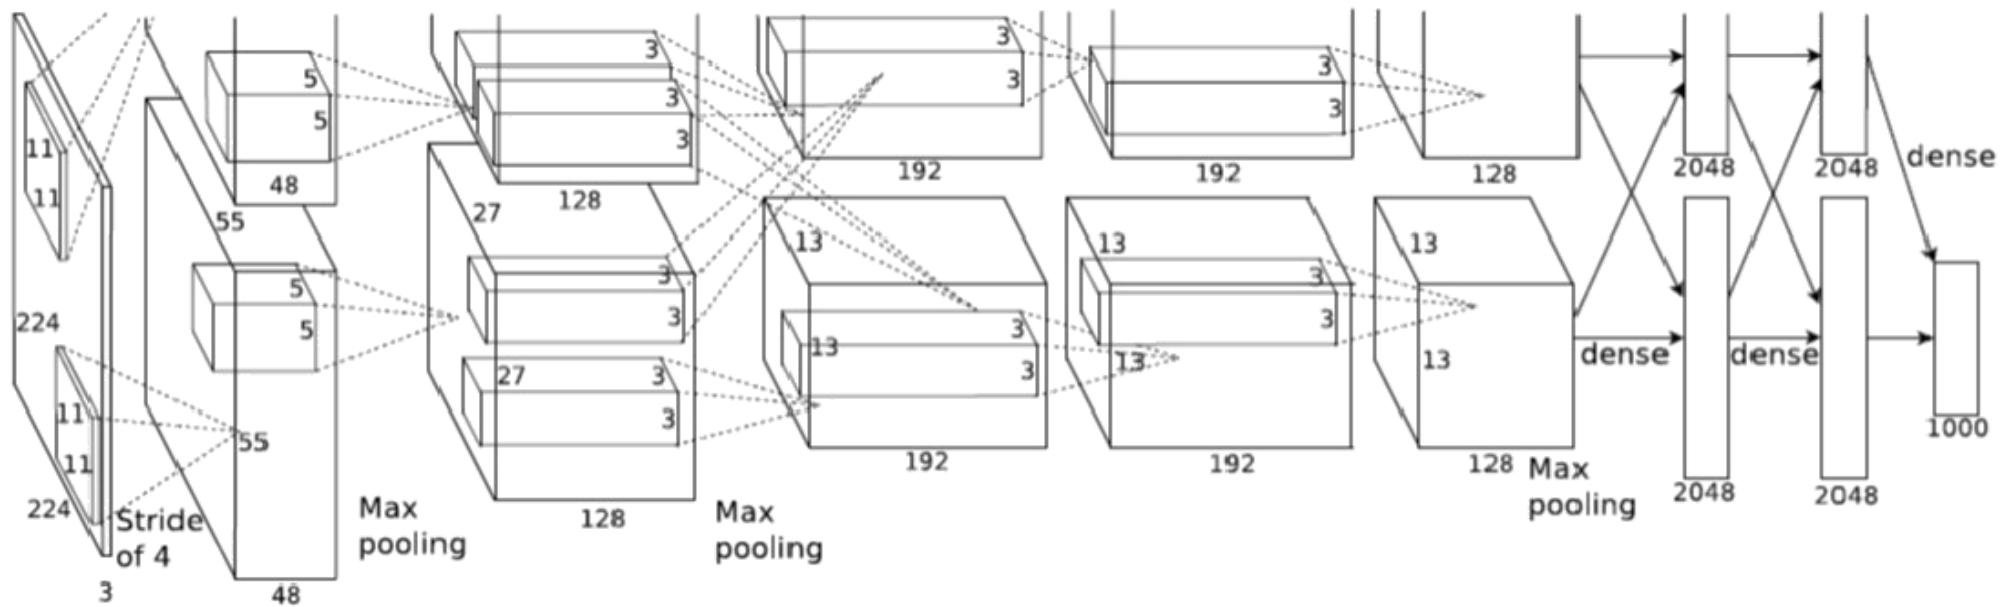


BACKPROPAGATION



DEEP LEARNING

A Machine Learning technique





MACHINE LEARNING



MACHINE LEARNING

Algorithms to find
patterns in data

Intelligent and
completely automatic
(the dream!)

WHERE CAN IT BE USEFUL?

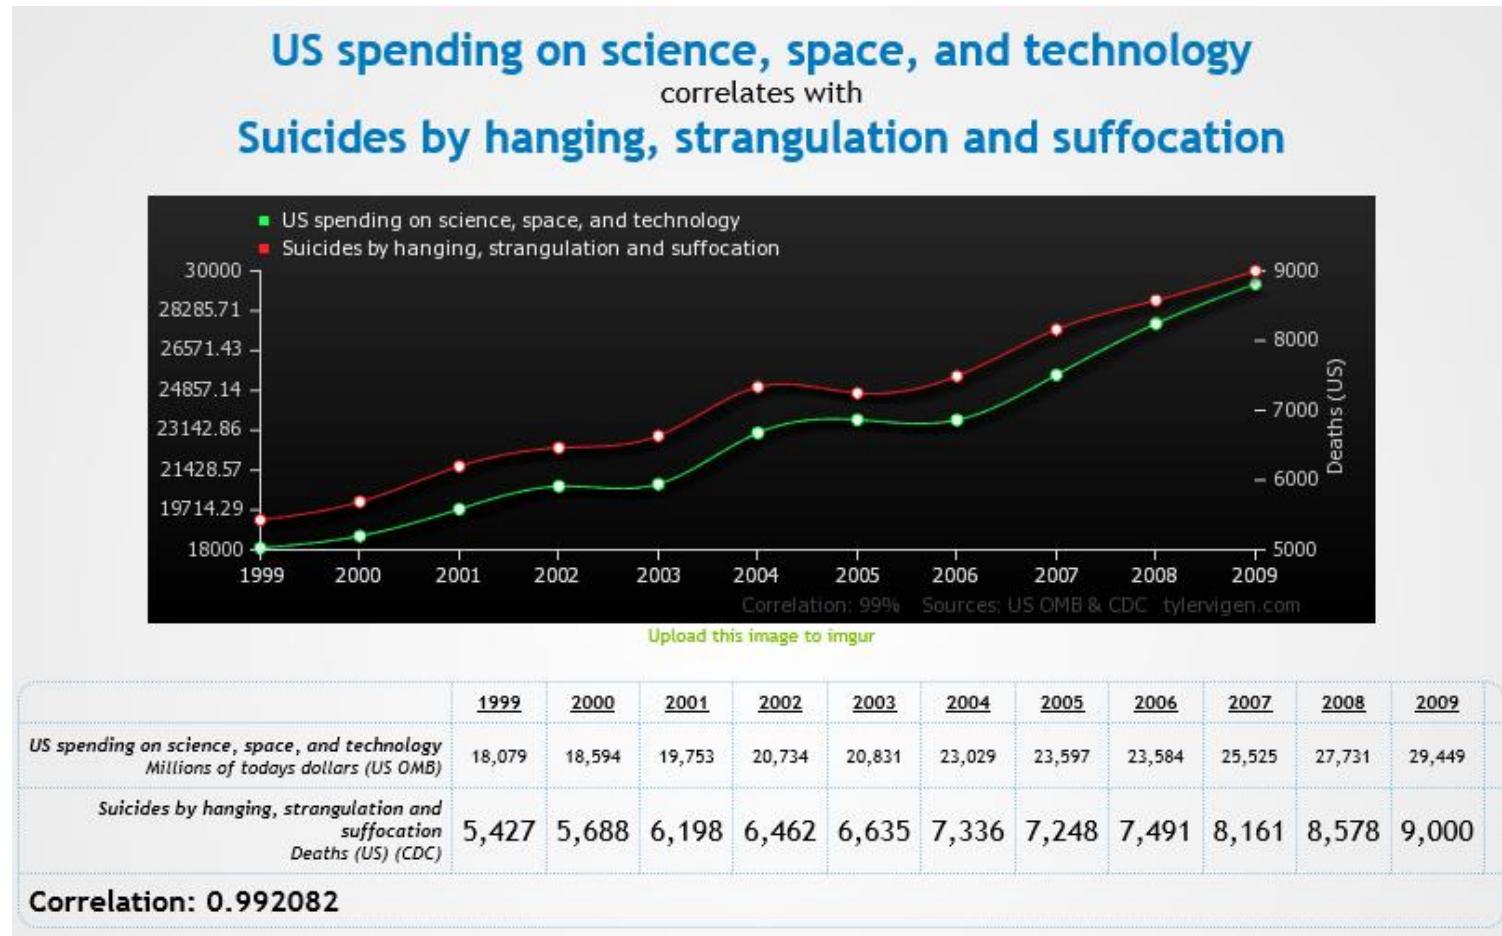
Where...

- There is no previous human experience
- Human experience exists, but it cannot be easily reduced to a set of rules (voice recognition)
- The scenario changes rapidly (spam email)
- Human training is expensive
- The amount of data is massive

COMMON SCENARIOS

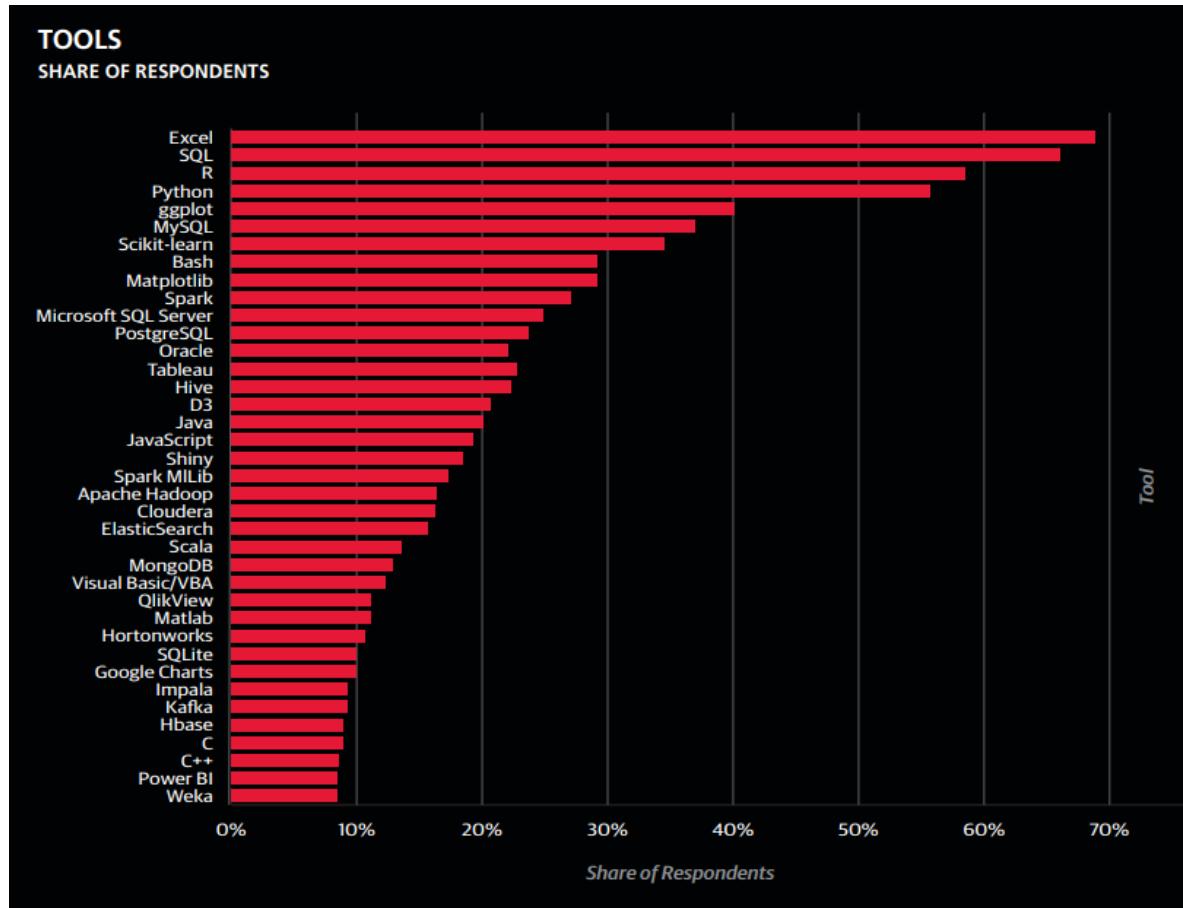
Domain	Scenario
Financial services	Modelling risks
	Threat and fraud detection
Media y entertainment	Segmented advertisement
	Recommendation engine
Commerce	Sentiment analysis
	Transaction analysis in point of sales (POS)
IT	Call Detail Records
Government	Environment monitorization
	Traffic clogging and re-routing
Healthcare	Research (genome, cancer, etc..)
	Pandemic early detection
Engineering	Predictive maintenance

CORRELATION != CAUSALITY



Source: <http://www.tylervigen.com>

TOOLS



"2017 European Data Science
Salary Survey"
[O'Reilly]

OUR TOOLS

Mainly

- SQL Server
- Python
- R (with Rstudio and Microsoft R)
- Excel + PowerBI
- HDInsight (Hive and Spark)
- CNTK

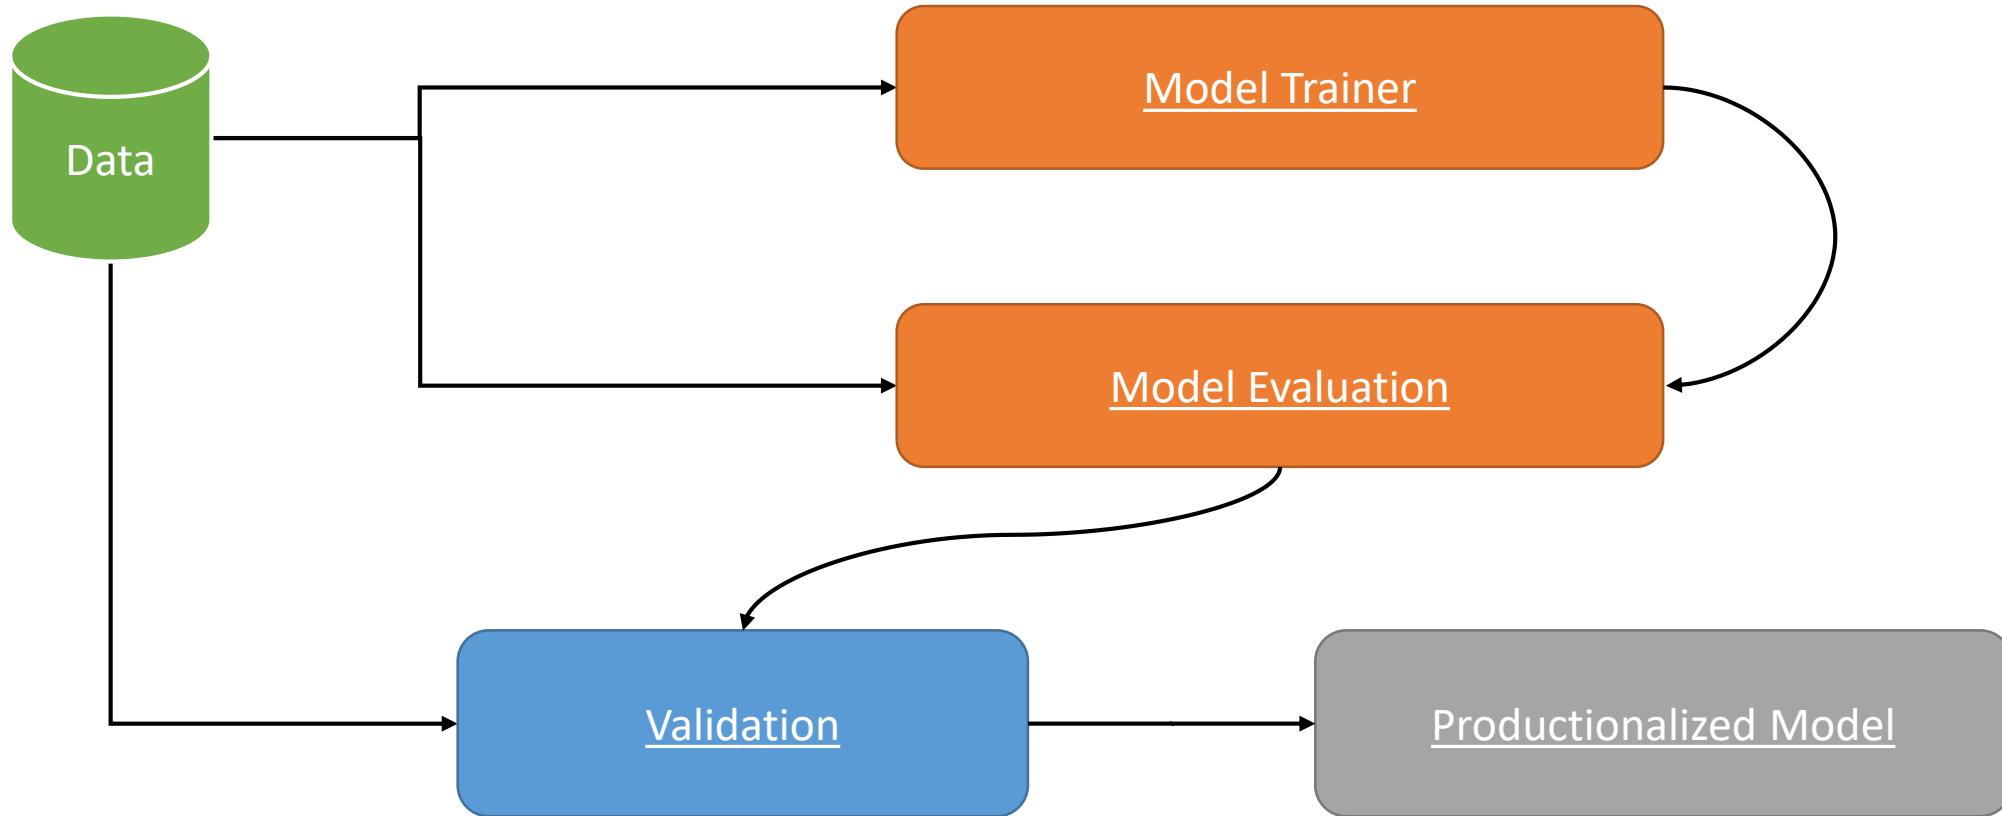
Secondly

- TensorFlow
- AzureML

We try to avoid ☺

- Mahout

BASIC CONCEPTS



MACHINE LEARNING CATEGORIES

Classification

- Assign a category
- i.e.: Restaurant (Chinese | Indian | Italian | Japanese)

Regression

- Predicting a real value for each element
- i.e.: Amount of a purchase, temperature, etc.

Ranking

- Ordering elements according to a criteria
- i.e.: web search

Clustering

- Partitioning of elements in heterogeneous groups
- i.e.: Tweets clustering by theme

Dimensionality reduction

- Transforming something complex to something simpler
- i.e.: Image pre-processing, voice recognition, signals

LEARNING / TRAINING

Task

- Applying a machine learning algorithm to domain data (training data)
- This domain represents a series of features that we want to model

Goal

- Minimize errors

Types of learning

- Supervised
- Not supervised
- Others (semi-supervised, active)

METHODS

Supervised

Non
Supervised

SUPERVISED LEARNING



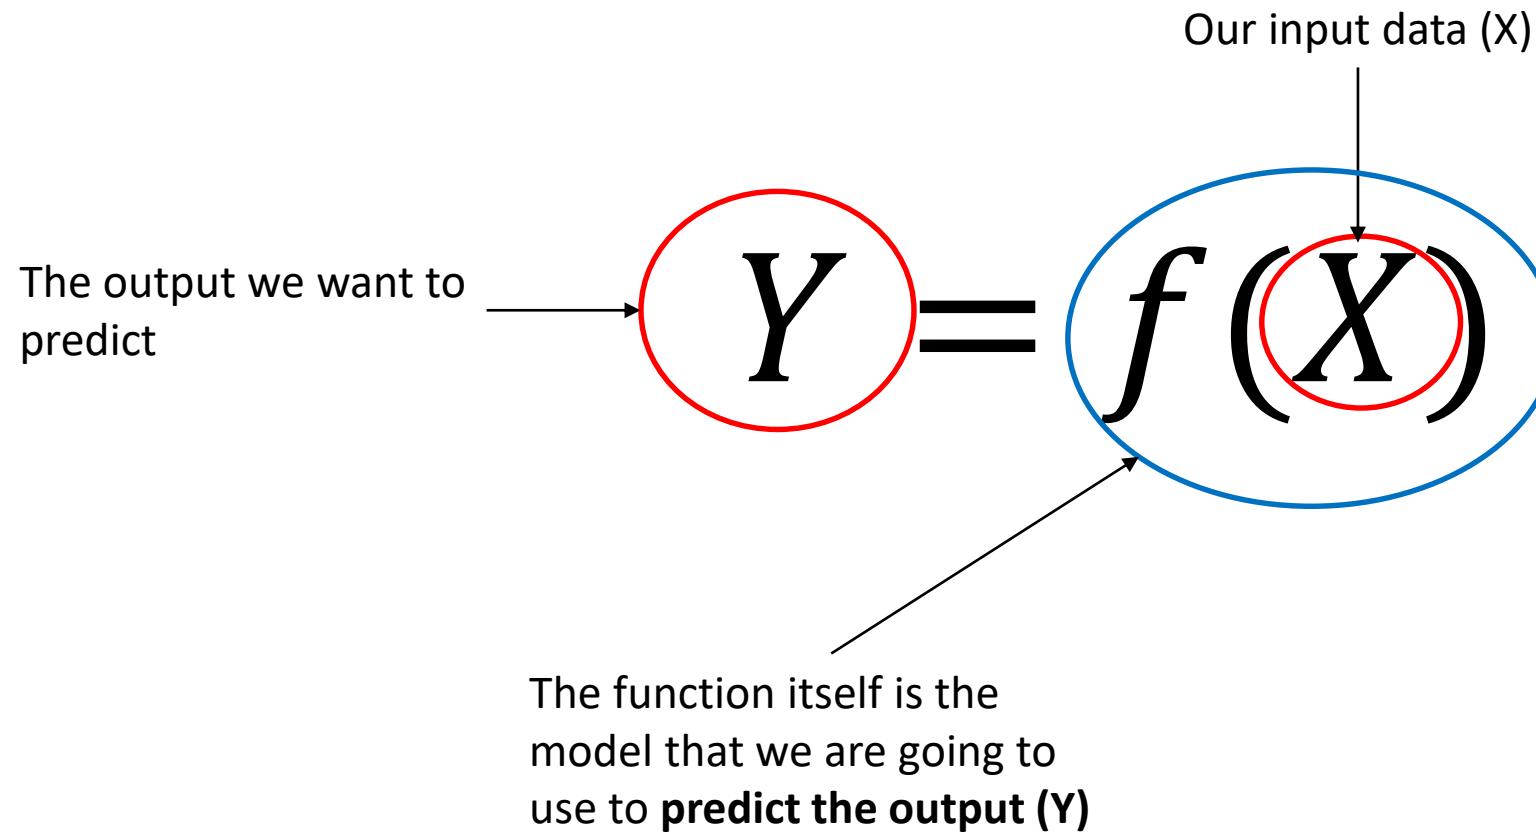
Training a model with data that has been previously labeled

SUPERVISED LEARNING

Labeled means that we have **historical data** about what we want to **predict**

SUPERVISED LEARNING

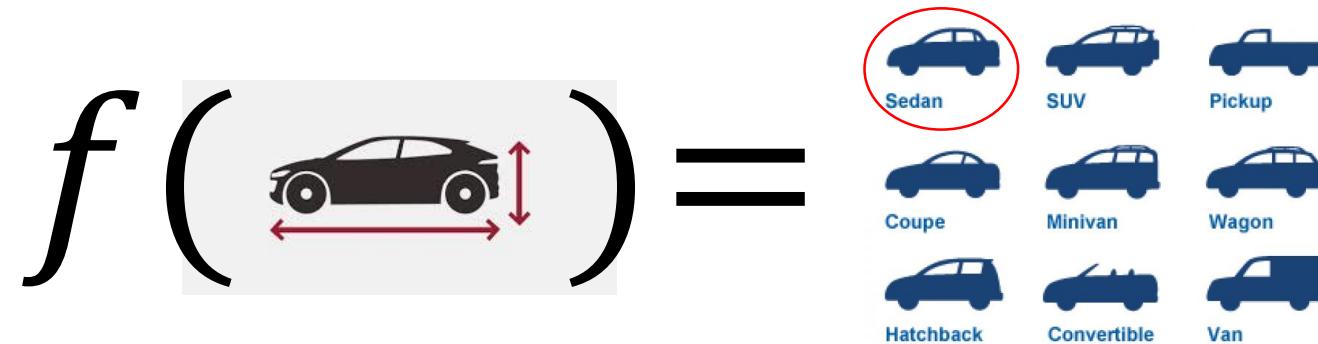
- Basically:



SO, WHERE IS THE LEARNING PROCESS?

The learning process consists on adjusting the model ($f(x)$) parameters in order to minimize the error when predicting

PREDICTING CAR TYPE USING DIMENSIONS



NON SUPERVISED LEARNING

- Unsupervised training is when you have input data but no output
- The goal of this kind training is not predicting an output but modelling the underlying structure or distribution in the input data
- It is called non supervised because there is no correct answers or teacher

TYPES OF NON SUPERVISED LEARNING

Clustering

- Discover groups in data
- Example: Grouping cars by horsepower

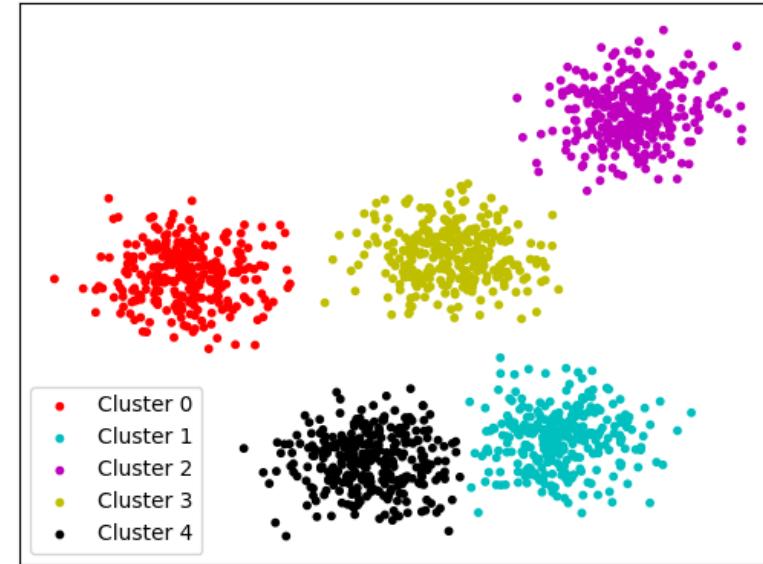
Association

- Search rules for large portion of your data
- Example: Amazon's "other customers also bought..."

EXAMPLE OF CLUSTERING

$f(horsepower,$
 $weight)$

=



WHEN IS IT USEFUL?

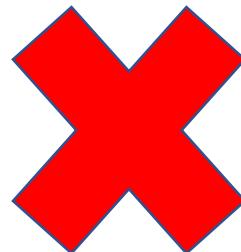
- When data cannot be labeled
 - Or when it is expensive to achieve
- When the amount of data is huge
 - Would anyone label/tag 10M of tweets?
- When there is no previous experience

SEMI SUPERVISED LEARNING

- Used combining labelled an unlabelled data to achieve better model results
- Consists on:
 1. Train a model using labelled data (supervised training)
 2. Predict using unlabelled data
 3. Add the best predictions (based in a threshold) to the labelled data set
 4. Repeat

EXAMPLE?

- Let's say we want to categorize websites
- Labelled data would force to
 1. Crawl data from the website
 2. Someone (human) would have to tag the site: "Forum", "News", "E-commerce", etc
- The bigger the sample the better the results



NOT AFFORDABLE

TAGGING WEBS

- But writing a crawler is a “simple” thing
- We could crawl the data, evaluate the website data using our “economic” model
- Use the best prediction as if they were labelled data
- Repeat



AFFORDABLE

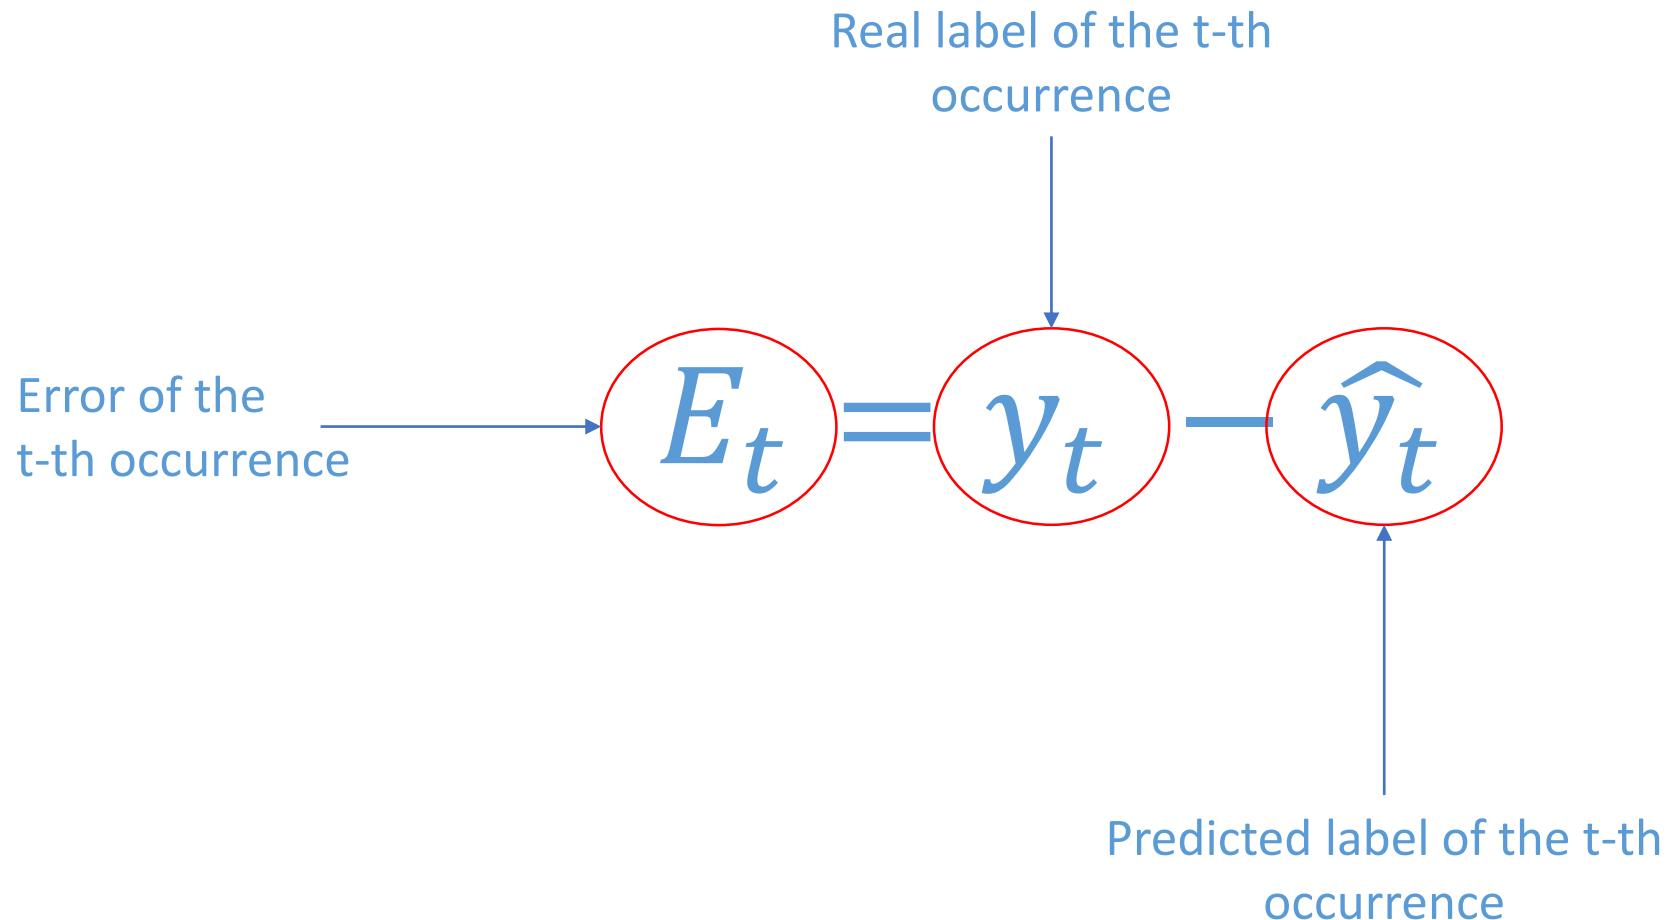
TYPES OF SEMI SUPERVISED LEARNING

- Self training (the simplest one)
- Generative Models
- S3VMs
- Graph-Bases Algorithms
- Multiview Algorithm

A QUICK INTRODUCTION TO ERROR MEASURING

- After the model is created and evaluated
- How can we know if it was a good model?
- We need to measure the **error** and be able to compare it

DEFINITION OF ERROR



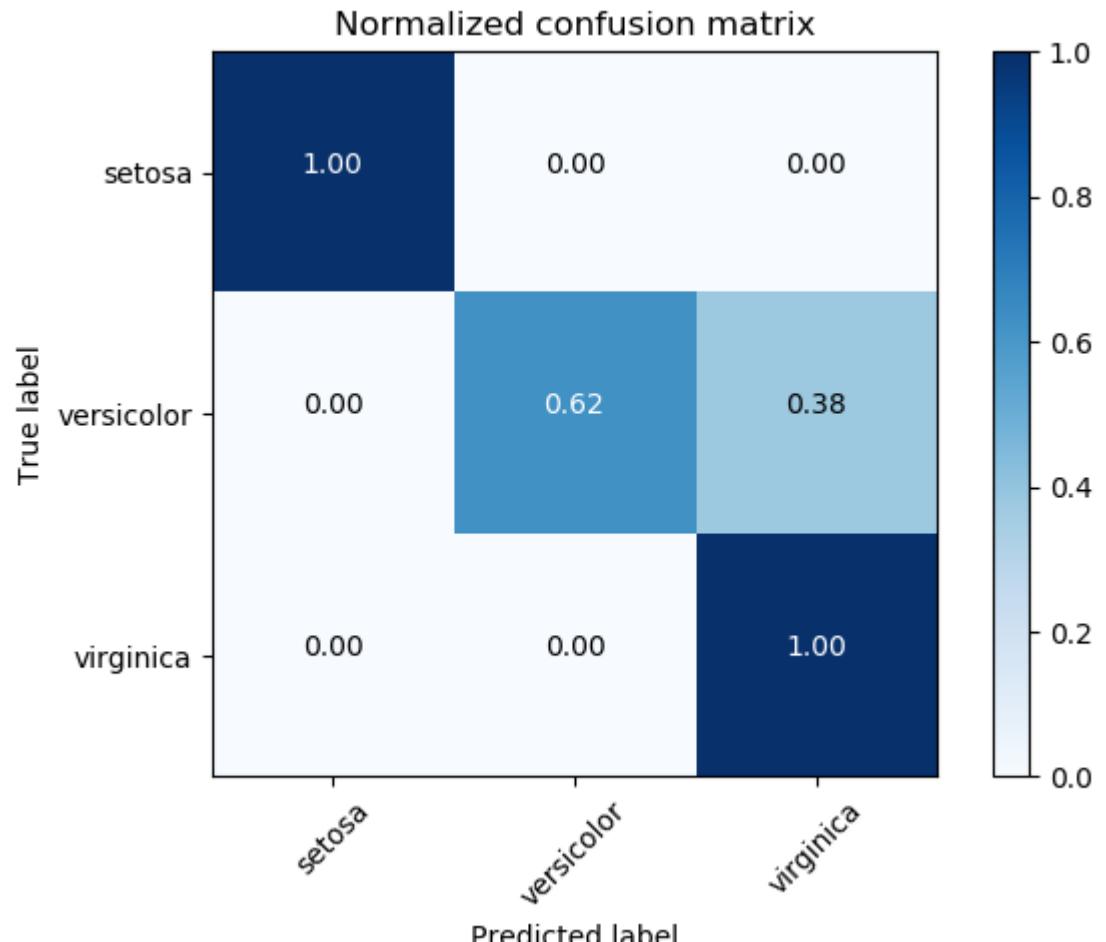
STATISTICS TOOLS

- Confusion matrix
- Recall and precision
- Error measurement:
 - Mean absolute error
 - Mean squared error
 - Root mean squared error
 - Mean absolute percentage error

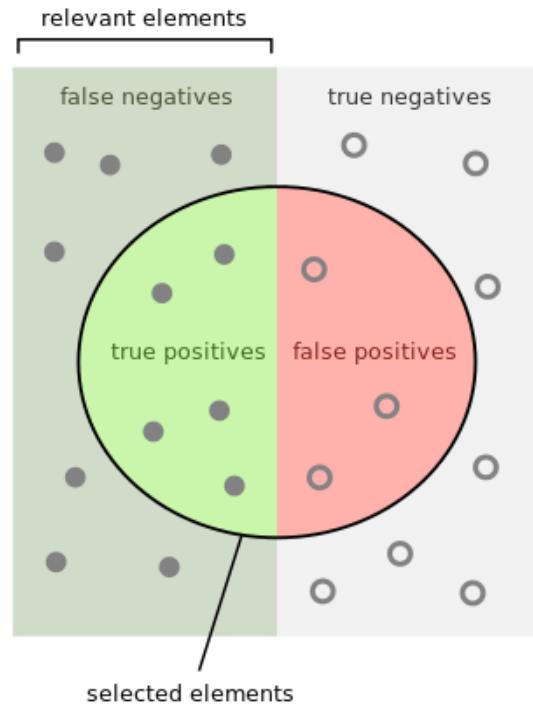
CONFUSION MATRIX

		Prediction	
		Positive	Negative
Actual	Positive	TP	FN
	Negative	FP	TN

CONFUSION MATRIX



RECALL AND PRECISION



How many selected items are relevant?

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

ERRORS MEASUREMENT

Mean squared error	$MSE = \frac{1}{n} \sum_{t=1}^n e_t^2$
Root mean squared error	$RMSE = \sqrt{\frac{1}{n} \sum_{t=1}^n e_t^2}$
Mean absolute error	$MAE = \frac{1}{n} \sum_{t=1}^n e_t $
Mean absolute percentage error	$MAPE = \frac{100\%}{n} \sum_{t=1}^n \left \frac{e_t}{y_t} \right $

plain concepts³

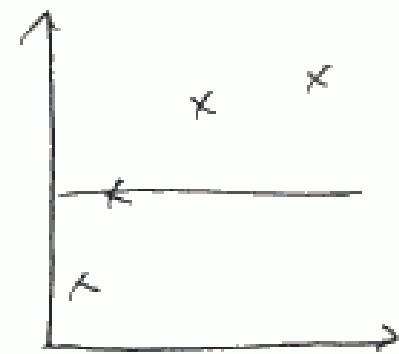
AZURE ML STUDIO



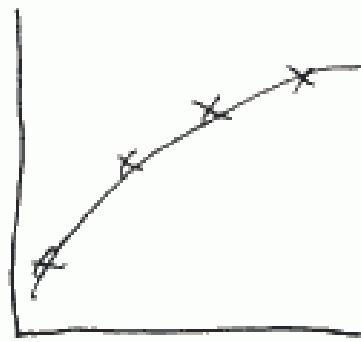
<https://studio.azureml.net/>

OVERFITTING AND UNDERFITTING

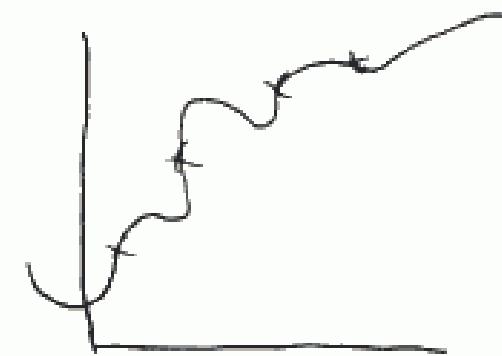
OVERFITTING AND UNDERFITTING



(a)



(b)

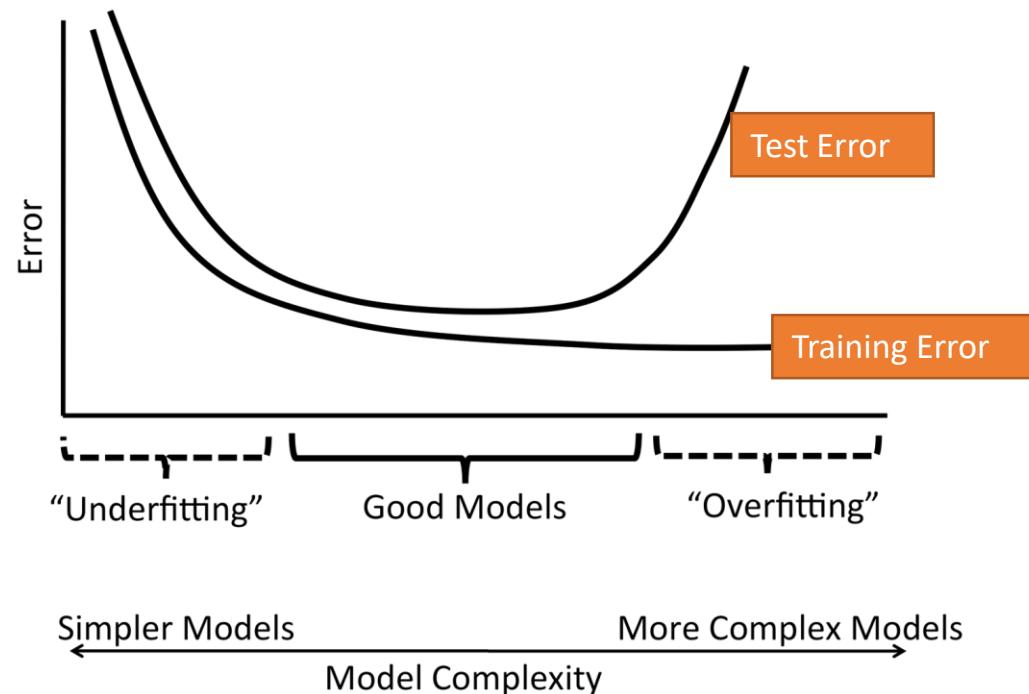


(c)

OVERFITTING AND UNDERFITTING

Occam Razor:

- Best models are the simplest which fit the data properly



CLASSIFICATION

CLASSIFICATION

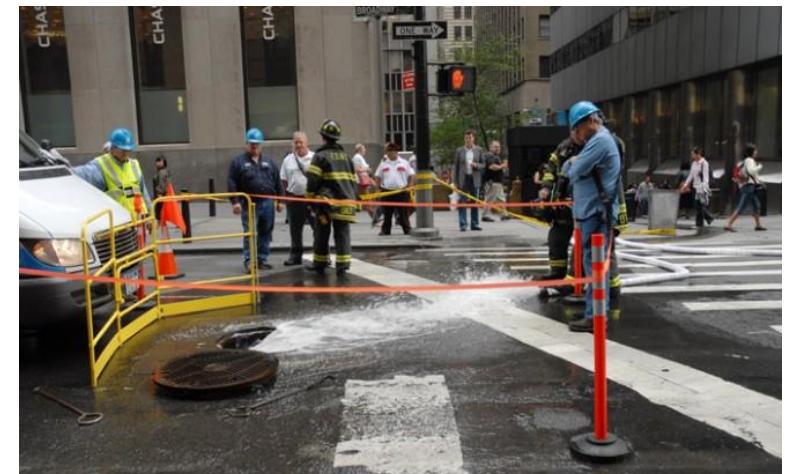
- There is something we want to predict (a class or a label)
- That class or label is categorical
 - That is, it's a value in a set of categories, not a number
 - True/False, A/B/C/D, ...
- There is a set of features associated to that class
- We want to know what features are able to predict that class

CLASSIFICATION

- Example: Manhole explosion analysis

Manhole model: [5 3 120 12 1]

Number of events last year
Number of serious incidents last year
Number of wires installed before 1930
Number of electric wires
Manhole with air circulation?
Inspected? 0



CLASSIFICATION

- Each observation is represented by a feature vector

Manhole model:

$$\begin{array}{ccccccccc} [& 5 & 3 & 120 & 12 & 1 & 0 & \dots &] & -1 \\ [& 0 & 0 & 89 & 5 & 1 & 1 & \dots &] & 1 \\ [& 1 & 0 & 20 & 0 & 0 & 1 & \dots &] & -1 \end{array}$$

Features

labels, Y

(predictors, covariates
or independent variables)

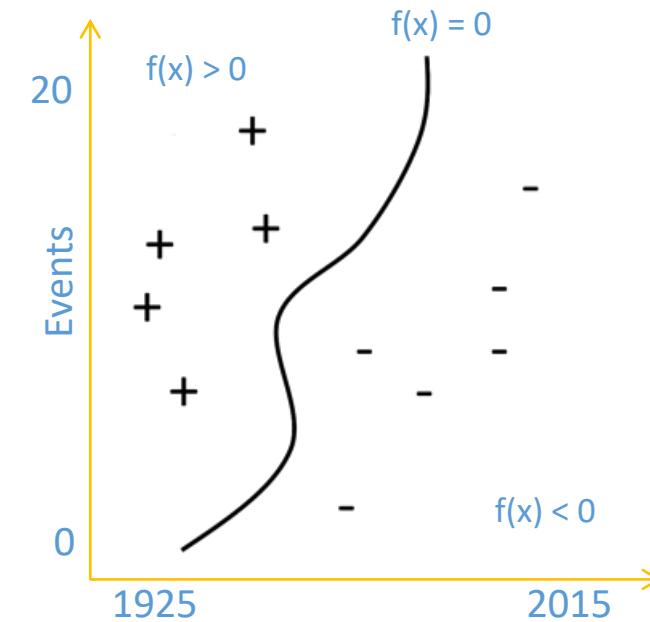
CLASSIFICATION

- Given a training set (x_i, y_i) for $i=1 \dots n$, we want to create a classification model which is able to predict a label y for the next value of x

Manhole model :

[1925 15]

Wire installation date
Number of events last year



CLASSIFICATION

Binary

Multivariate

Use cases

- Automatic handwritten recognition
- Spam detection
- Fraud detection
- Customer churn
- Voice recognition
- Image recognition
- Etc.



CLASSIFICATION



DECISION TREES

DECISION TREES

- Classification and prediction tools
- Can be represented using a tree structure, hierarchy and recursive
 - Easy to understand, even for not expert users
- Each node is divided in branches
- Each leaf represents a decision

TENNIS MATCH

Day	Weather	Temperature	Humidity	Wind	Match?
1	Sunny	Hot	High	Weak	No
2	Sunny	Hot	High	Strong	No
3	Cloudy	Hot	High	Weak	Yes
4	Rainy	Warm	High	Weak	Yes
5	Rainy	Cold	Normal	Weak	Yes
6	Rainy	Cold	Normal	Strong	No
7	Cloudy	Cold	Normal	Strong	Yes
8	Sunny	Warm	High	Weak	No
9	Sunny	Cold	Normal	Weak	Yes
10	Rainy	Warm	Normal	Weak	Yes
11	Sunny	Warm	Normal	Strong	Yes
12	Cloudy	Warm	High	Strong	Yes
13	Cloudy	Hot	Normal	Weak	Yes
14	Rainy	Warm	High	Strong	No

What will happen day 15?

Rain, High temperature, High humidity and weak wind.

TENNIS MATCH

Day	Weather	Temperature	Humidity	Wind	Match?
1	Sunny	Hot	High	Weak	No
2	Sunny	Hot	High	Strong	No
3	Cloudy	Hot	High	Weak	Yes
4	Rainy	Warm	High	Weak	Yes
5	Rainy	Cold	Normal	Weak	Yes
6	Rainy	Cold	Normal	Strong	No
7	Cloudy	Cold	Normal	Strong	Yes
8	Sunny	Warm	High	Weak	No
9	Sunny	Cold	Normal	Weak	Yes
10	Rainy	Warm	Normal	Weak	Yes
11	Sunny	Warm	Normal	Strong	Yes
12	Cloudy	Warm	High	Strong	Yes
13	Cloudy	Hot	Normal	Weak	Yes
14	Rainy	Warm	High	Strong	No

Training Data: 14 rows

TENNIS MATCH

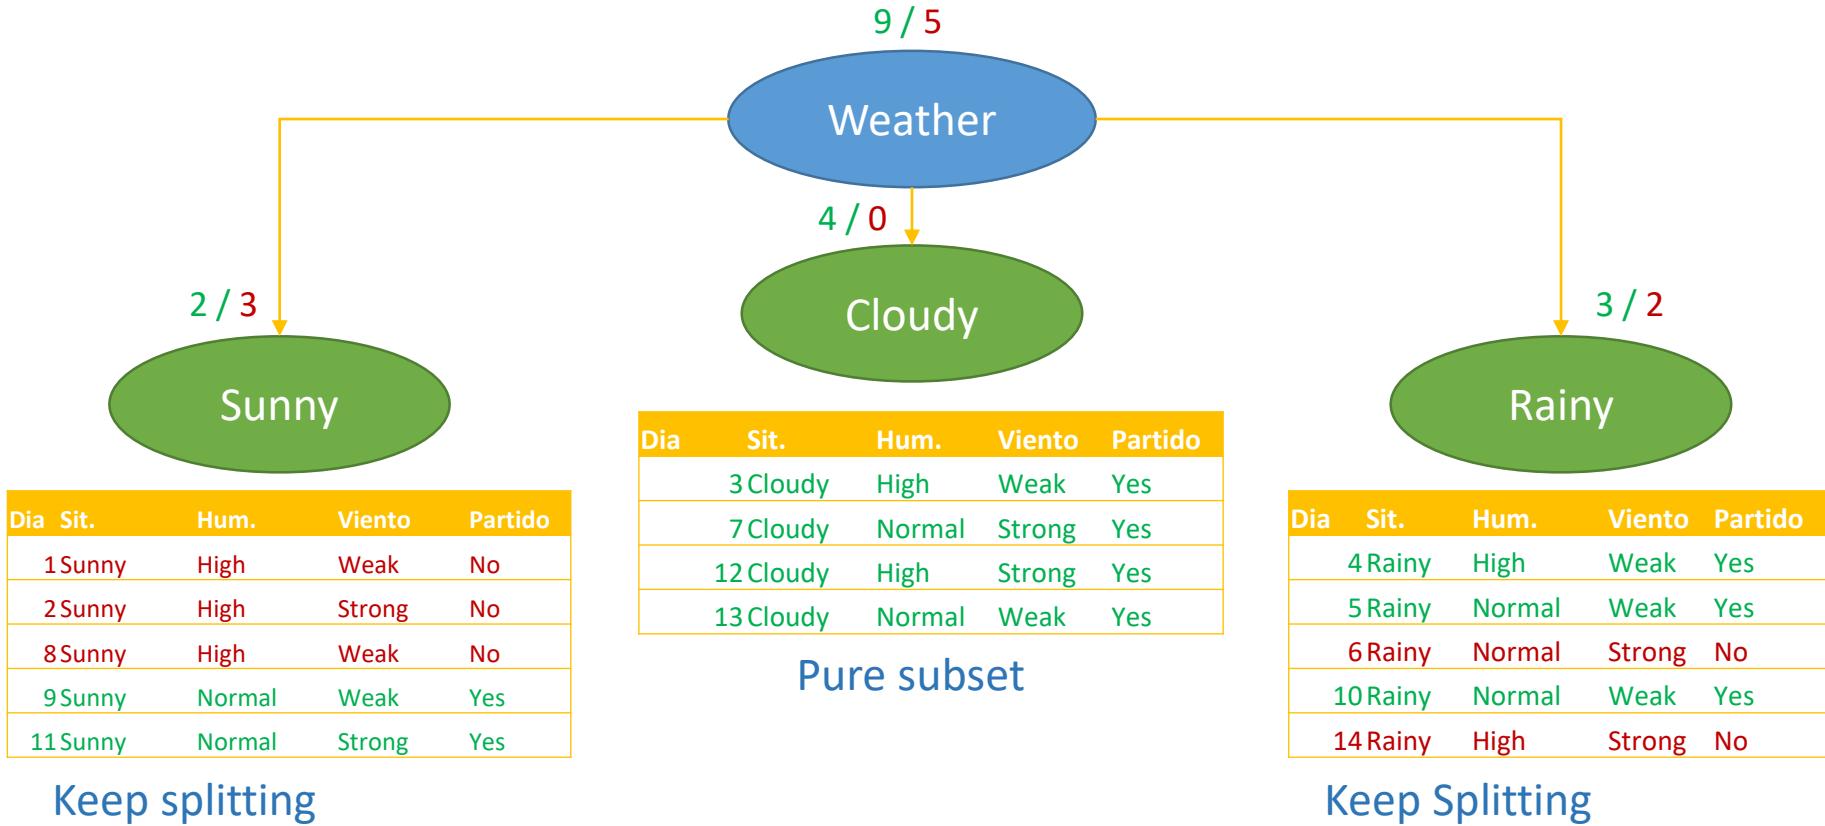
Day	Weather	Temperature	Humidity	Wind	Match?
1	Sunny	Hot	High	Weak	No
2	Sunny	Hot	High	Strong	No
3	Cloudy	Hot	High	Weak	Yes
4	Rainy	Warm	High	Weak	Yes
5	Rainy	Cold	Normal	Weak	Yes
6	Rainy	Cold	Normal	Strong	No
7	Cloudy	Cold	Normal	Strong	Yes
8	Sunny	Warm	High	Weak	No
9	Sunny	Cold	Normal	Weak	Yes
10	Rainy	Warm	Normal	Weak	Yes
11	Sunny	Warm	Normal	Strong	Yes
12	Cloudy	Warm	High	Strong	Yes
13	Cloudy	Hot	Normal	Weak	Yes
14	Rainy	Warm	High	Strong	No

Training Data:

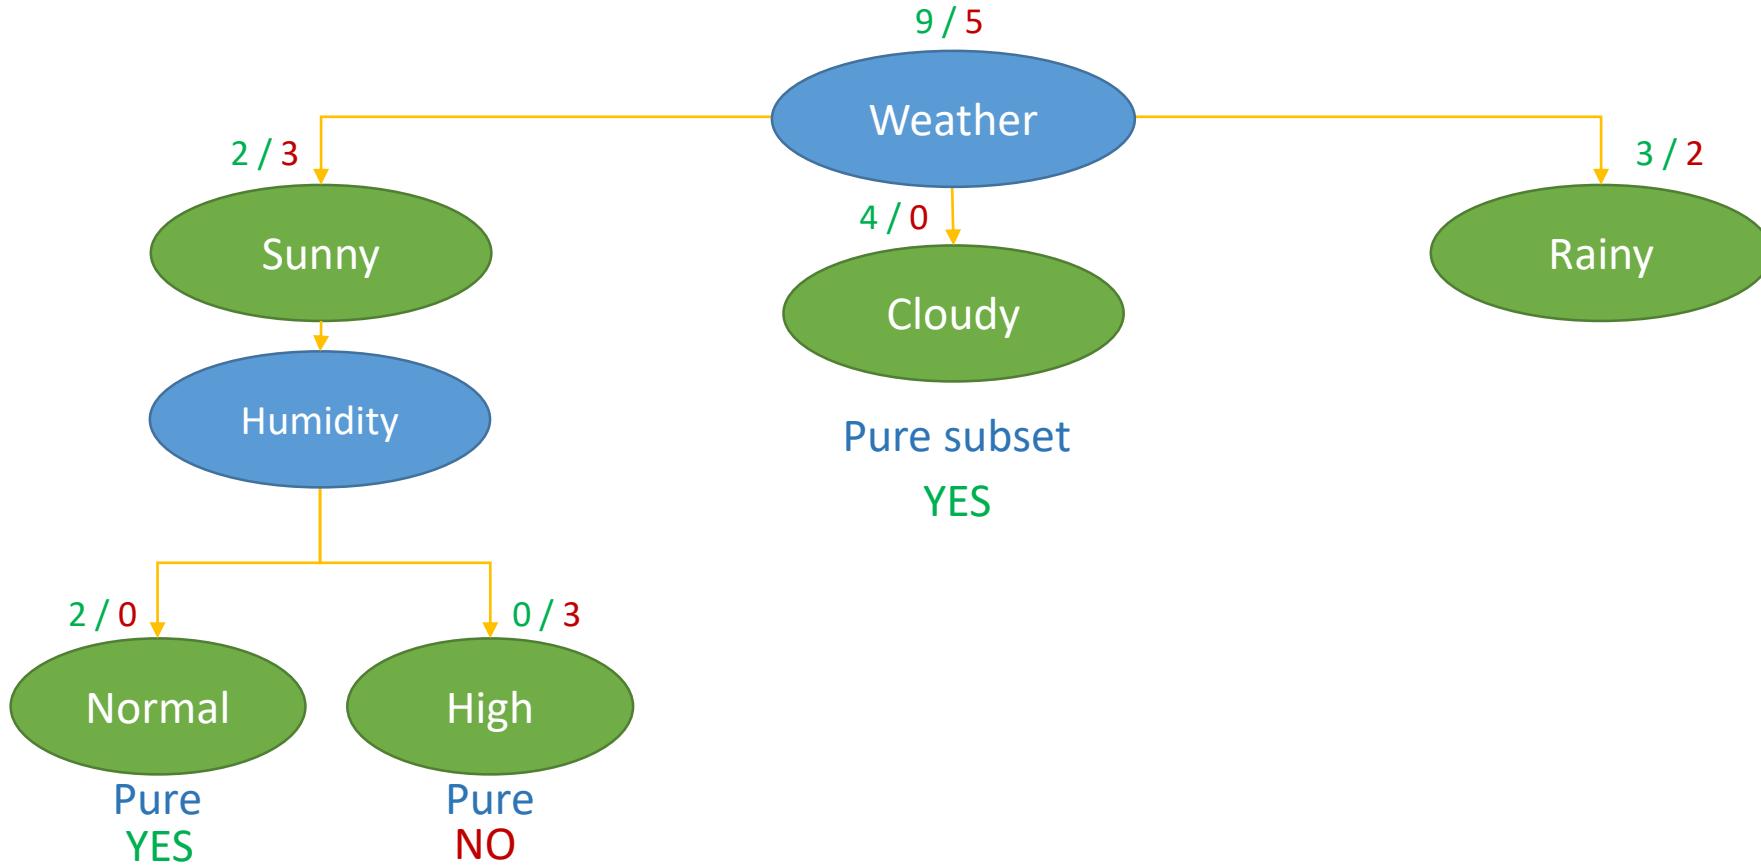
9 Yes

5 NO

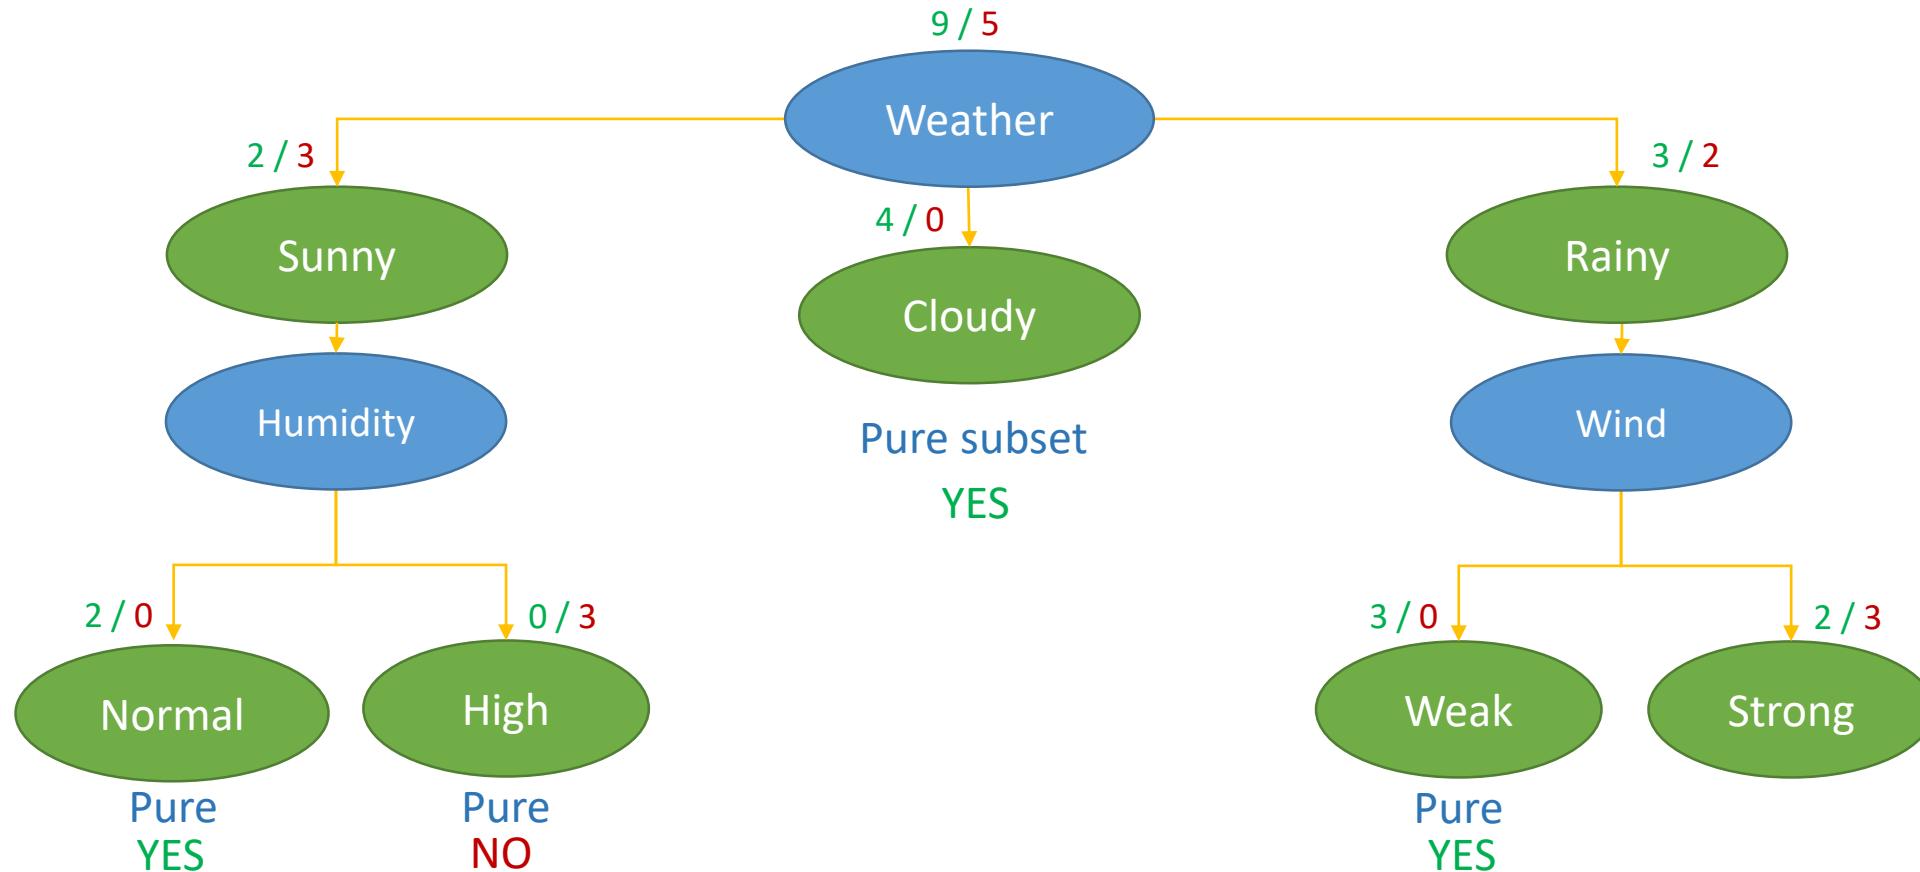
TENNIS MATCH



TENNIS MATCH



TENNIS MATCH



REGRESSION

REGRESSION

- There is something we want to predict (class or label)
- It's a numeric/real value
 - How many conversions are we going to get for this campaign in this week?
 - How many televisions will sell next year?
 - What is the salary of someone given his demographic information?
- There is a set of features associated to that class
- We want to know which features can predict that class

REGRESSION

- Each observation is represented by a feature vector

We model a person like this:

[5	3	120	12	1	0]	83
[0	0	89	5	1	1]	32
[1	0	20	0	0	1]	-10



Features, or X



labels, Y

(Predictors, covariables
or independent variables)

REGRESSION

- Each observation is represented like:

Modelling a person like this:

[5]
[0]
[1]



Only one feature

83
32
-10

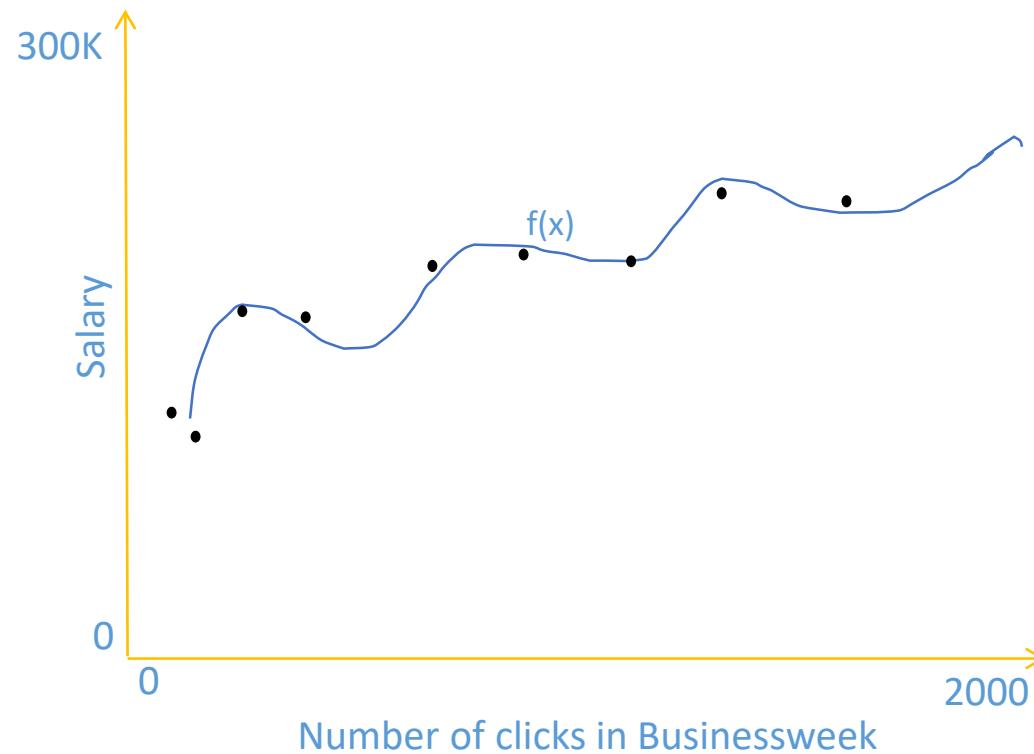


Y

REGRESSION

$f(x) = \text{function}(\text{Number of clicks in BusinessWeek})$

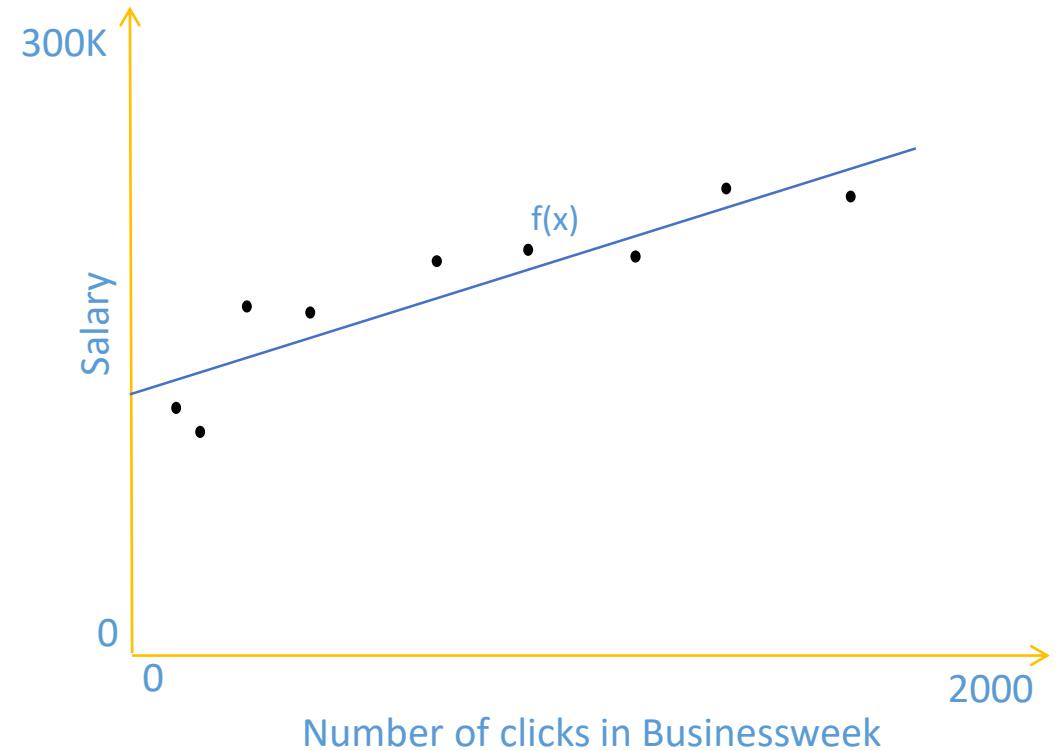
(Overfitting?)



REGRESSION

$$f(x) = \text{function}(\text{Number of clicks in BusinessWeek}) \\ = 5K * \text{Number of clicks in BusinessWeek} + 100K$$

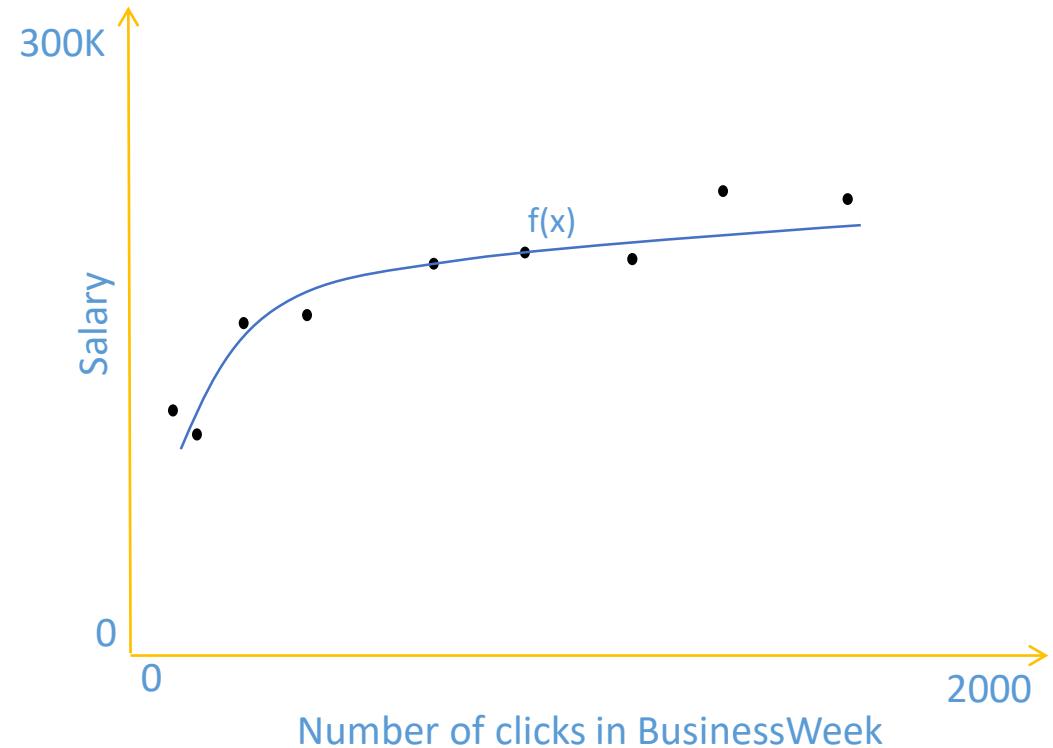
(Underfitting?)



REGRESSION

$f(x) = \text{function}(\text{Number of clicks in BusinessWeek})$

(Perfect?)



REGRESSION

Estimated salary

- $f(x) = \text{function}(\text{number of visits to a furniture shops}, \text{Number of clicks in BusinessWeek}, \text{number of different people who mails are sent per day}, \text{number of purchases above 5000€ within the last month}, \text{flights taken lately})$

Example:

$$\begin{aligned}f(x) = & 3 * \text{number of visits to a furniture shops} \\& + 10 * \text{Number of clicks in BusinessWeek} \\& + 100 * \text{number of different people who mails are sent per day} \\& + 2 * \text{number of purchases above 5000€ within the last month} \\& + 10 * \text{flights taken lately}\end{aligned}$$



LINEAR REGRESSION



CLUSTERING

K-MEANS

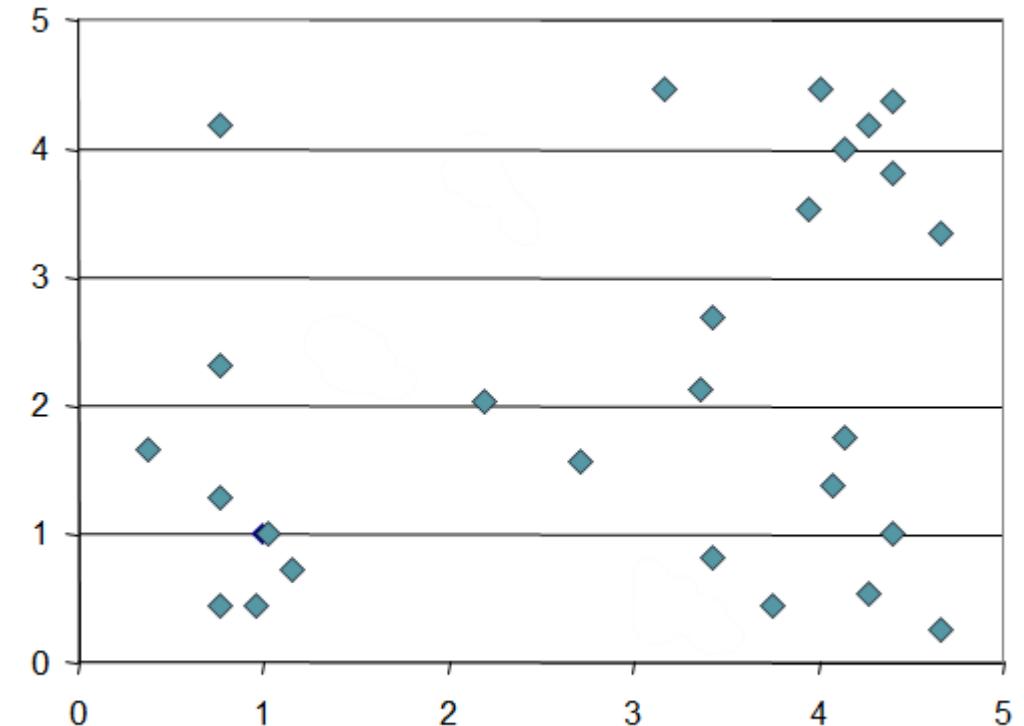
One of the most used clustering algorithms

Groups data depending on its features

Popular

Fast

Conceptually simple



K-MEANS

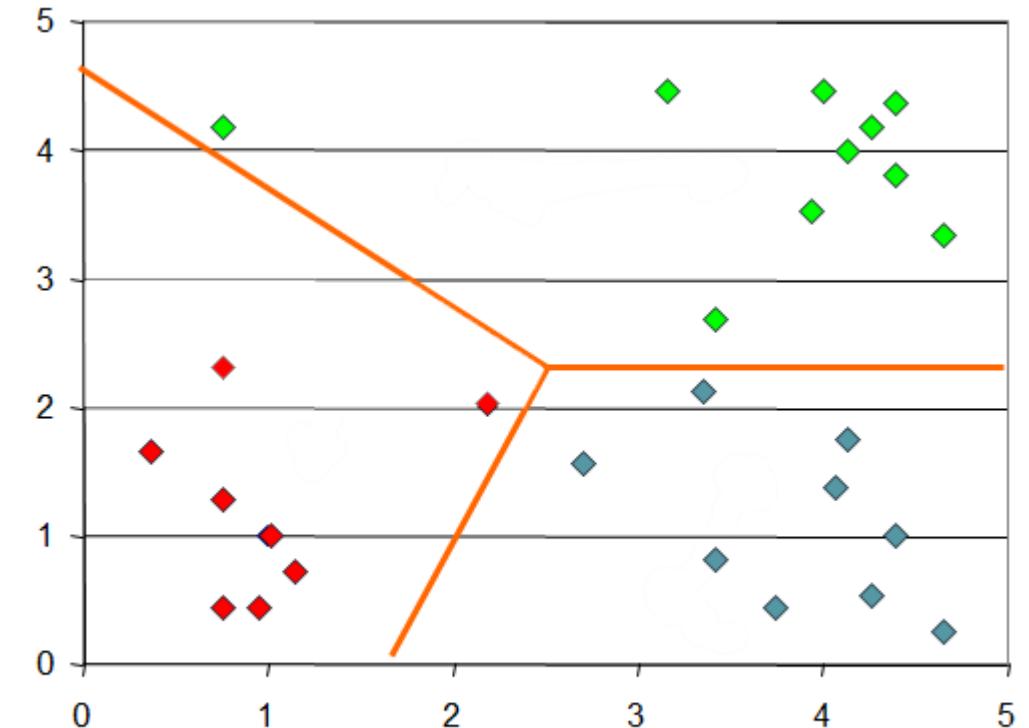
One of the most used clustering algorithms

Groups data depending on its features

Popular

Fast

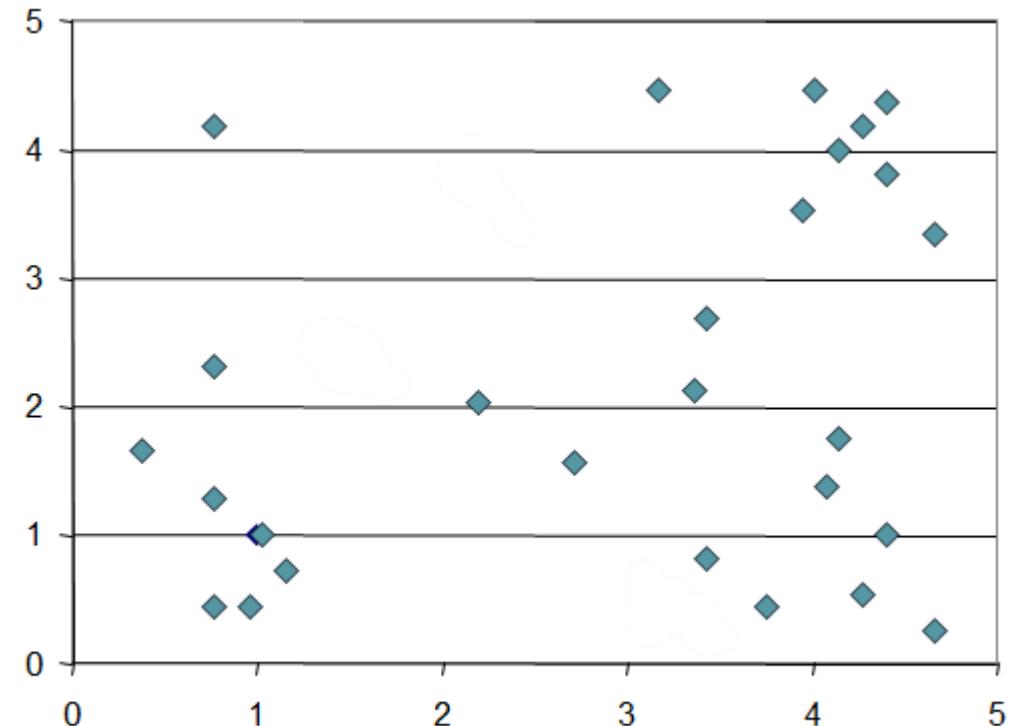
Conceptually simple



K-MEANS

Given a collection of values

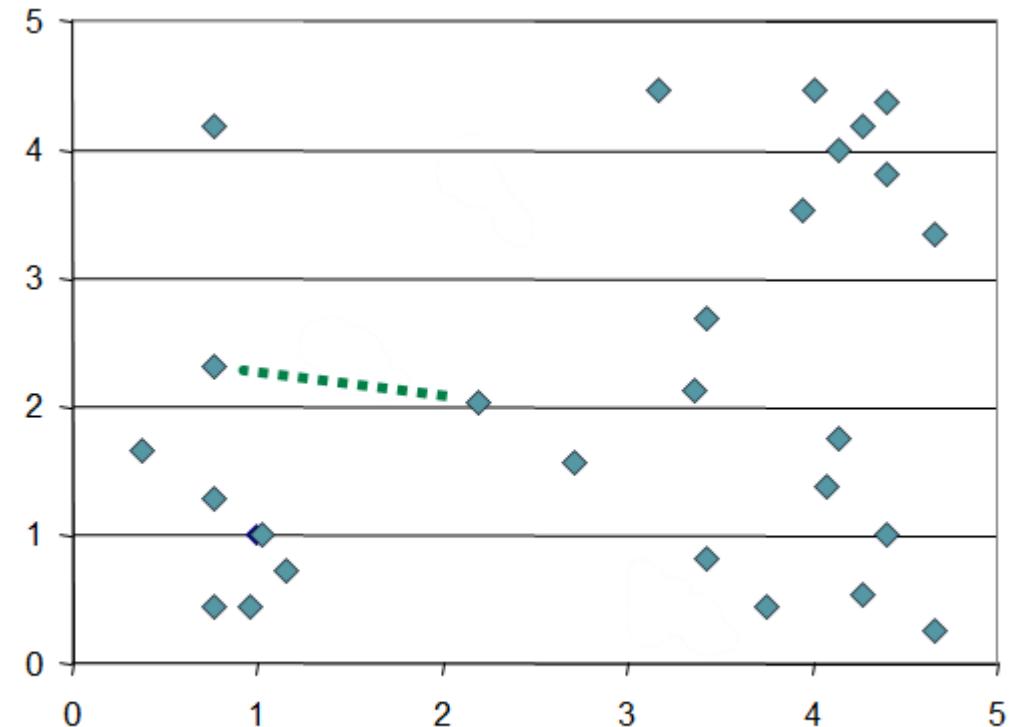
```
data = lines.map(line =>  
    parseVector(line))
```



K-MEANS

We define the difference between points
as the square of the euclidean distance

dist = p.squaredDist(q)

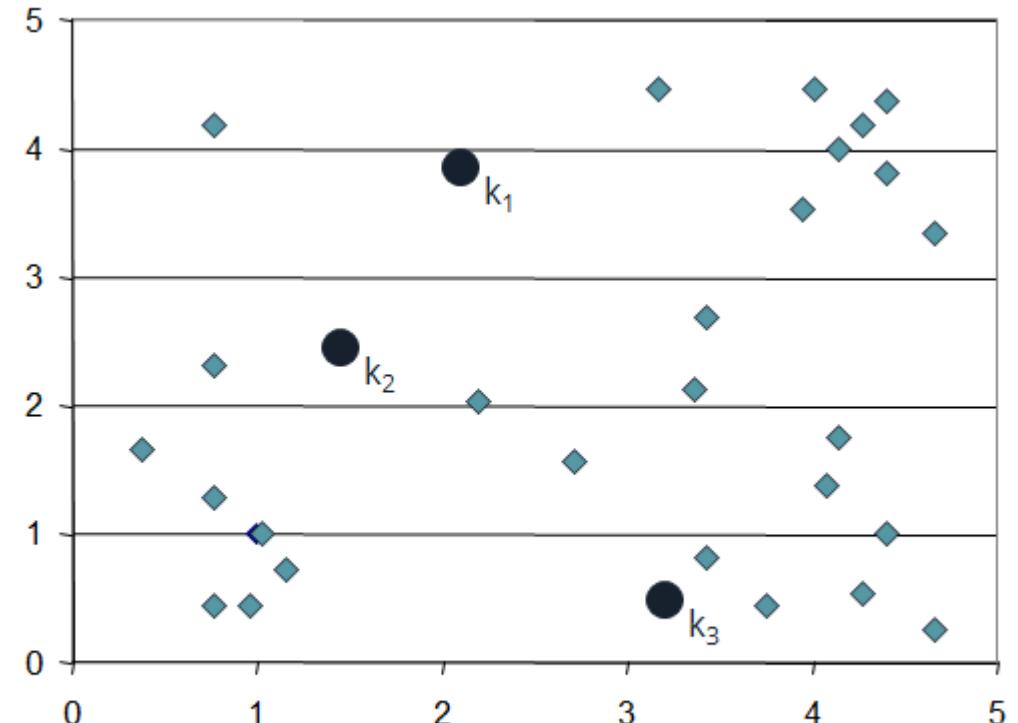


K-MEANS

With n clusters: $k_1, k_2, k_3, \dots, k_n$

We initialize the cluster centers

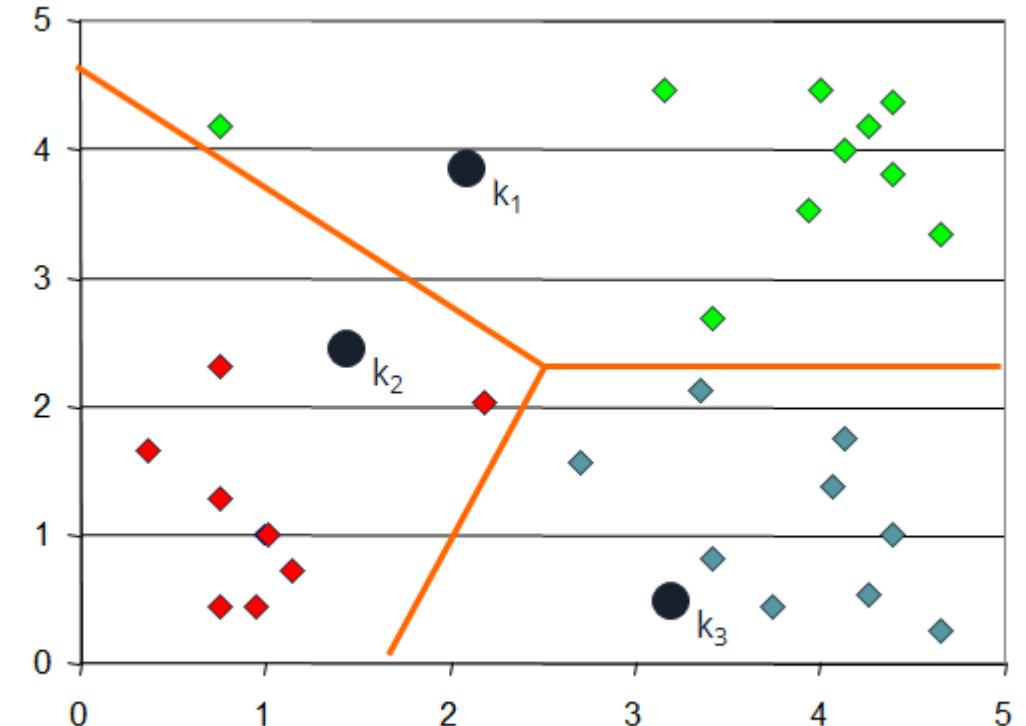
```
centers = data.takeSample(false,  
K, seed)
```



K-MEANS

Assign each point with the nearest cluster centroid

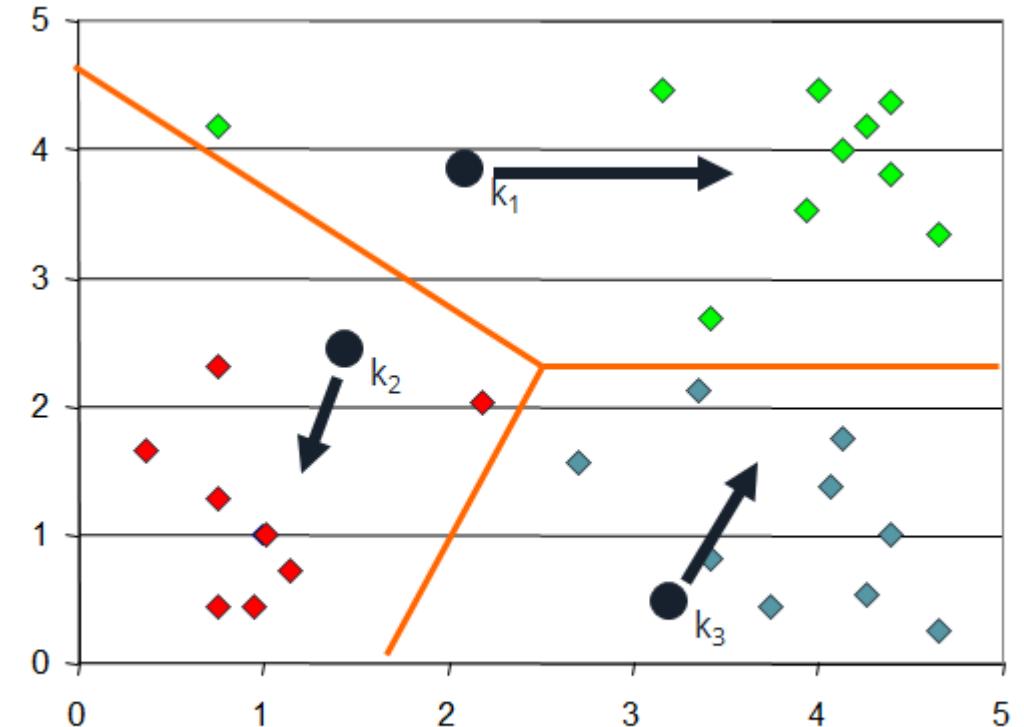
```
closest = data.map(p =>  
    (closestPoint(p, centers), p))
```



K-MEANS

We move the centroid to the average
distance of its points

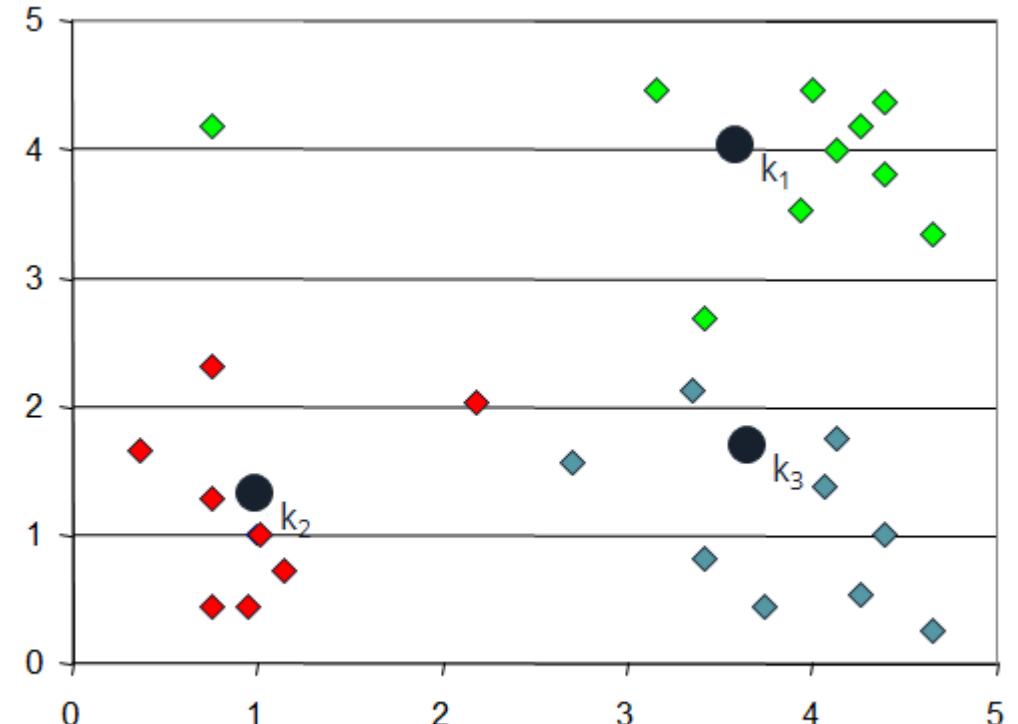
```
pointsGroup = closest.groupByKey()  
  
newCenters = pointsGroup.mapValues(ps  
          => average(ps))
```



K-MEANS

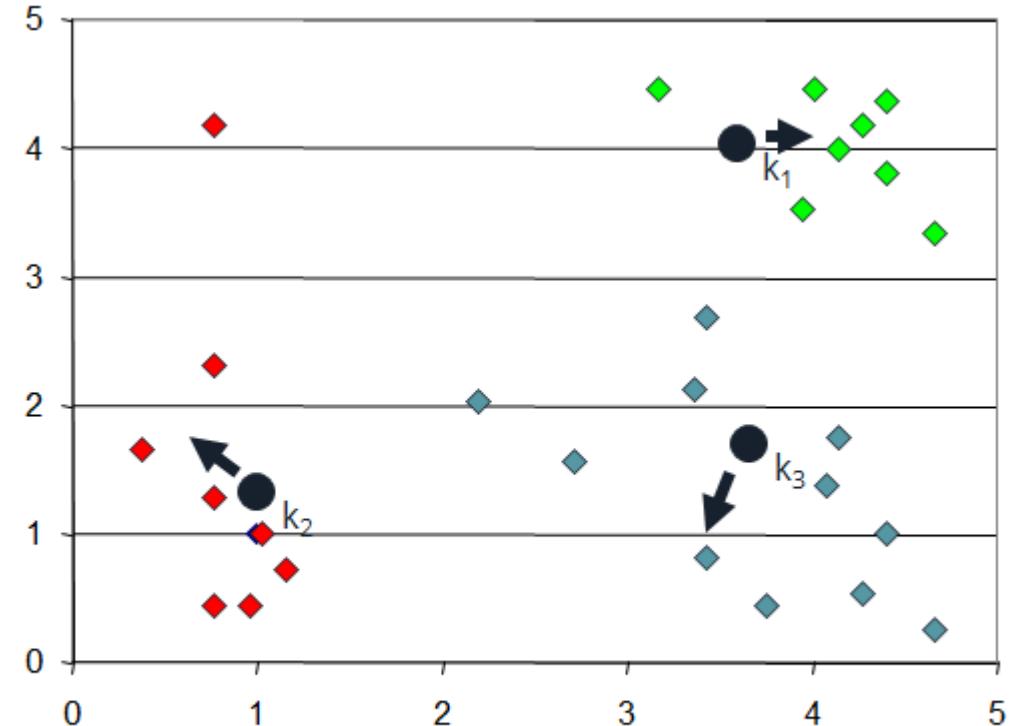
We move the centroid to the average
distance of its points

```
pointsGroup = closest.groupByKey()  
  
newCenters = pointsGroup.mapValues(ps  
        => average(ps))
```



K-MEANS

Repeat until convergence



K-MEANS

```
centers = data.takeSample(false, K, seed)

while (d > ε)
{
    closest = data.map(p => (closestPoint(p, centers), p))

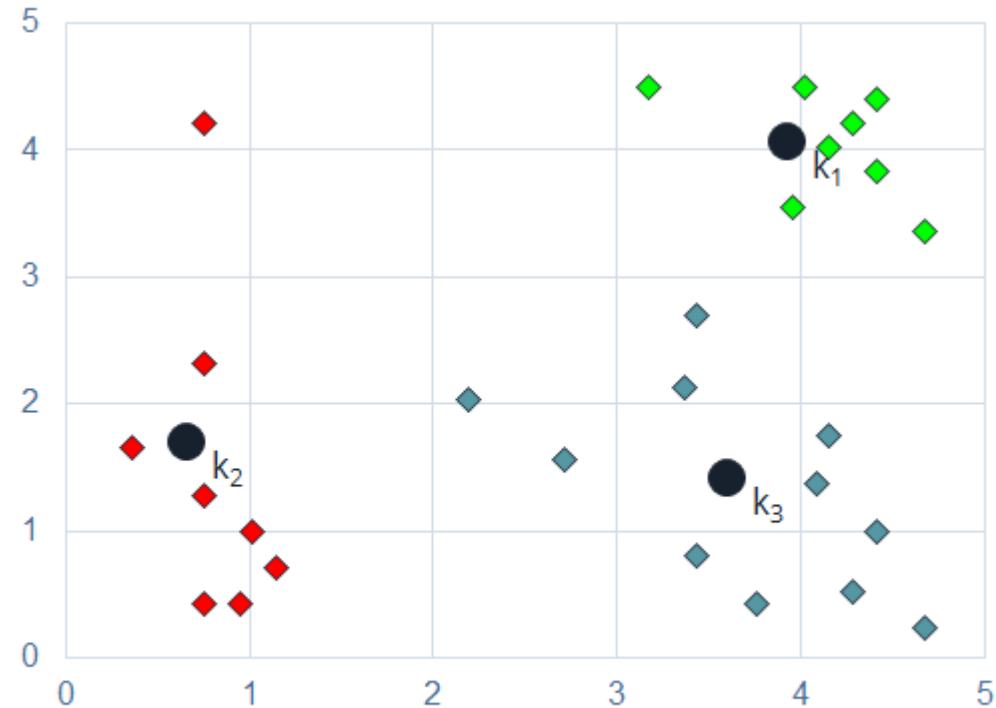
    pointsGroup = closest.groupByKey()

    newCenters = pointsGroup.mapValues(ps => average(ps))

    d = distance(centers, newCenters)

    centers = newCenters.map(_)
}
```

K-MEANS





K-MEANS



SUPPORT VECTOR MACHINES

SVM

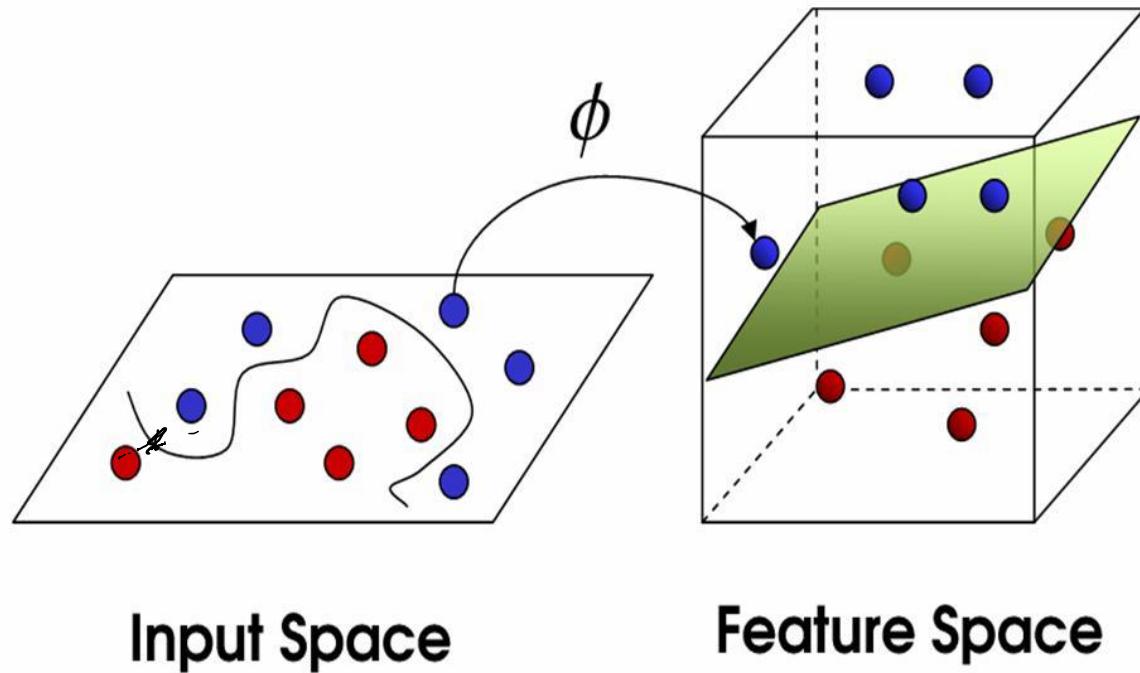
- Given a point set:
 - Subset of a bigger set (space)
 - Each point belongs to a two possible categories
- SVM builds a model able to predict the category of a new point

SVM

- The goal is look for the plane which divides both classes and, at the same time, maximizes the distance to the nearest point on each side of the plane
- Once the perfect plane is achieved, it is called “perceptron of optimal stability”
- Reminder, this scenario is being simplified (2 classes, 1 plane)
- In fact, when we say plane, we mean hyperplane

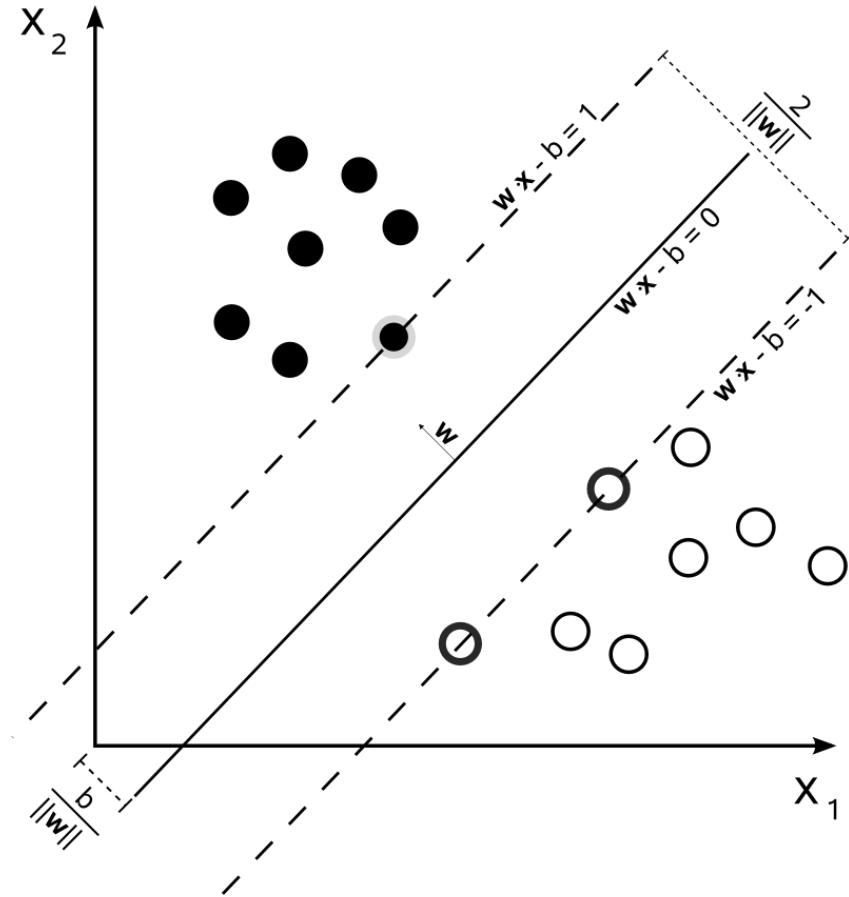
SVM

Principle of Support Vector Machines (SVM)

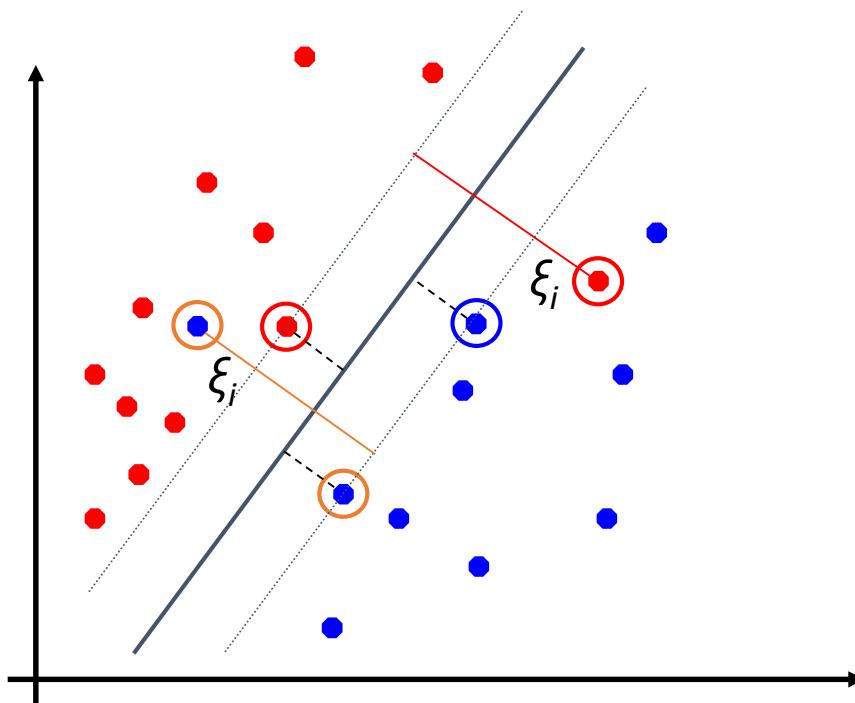


CLASSIFICATION MARGIN

- Hyperplane divides de subsets
- Support vectors are defined by the nearest points to the hyperplane
- The classification margin is the distance between hyperplane and support vectors

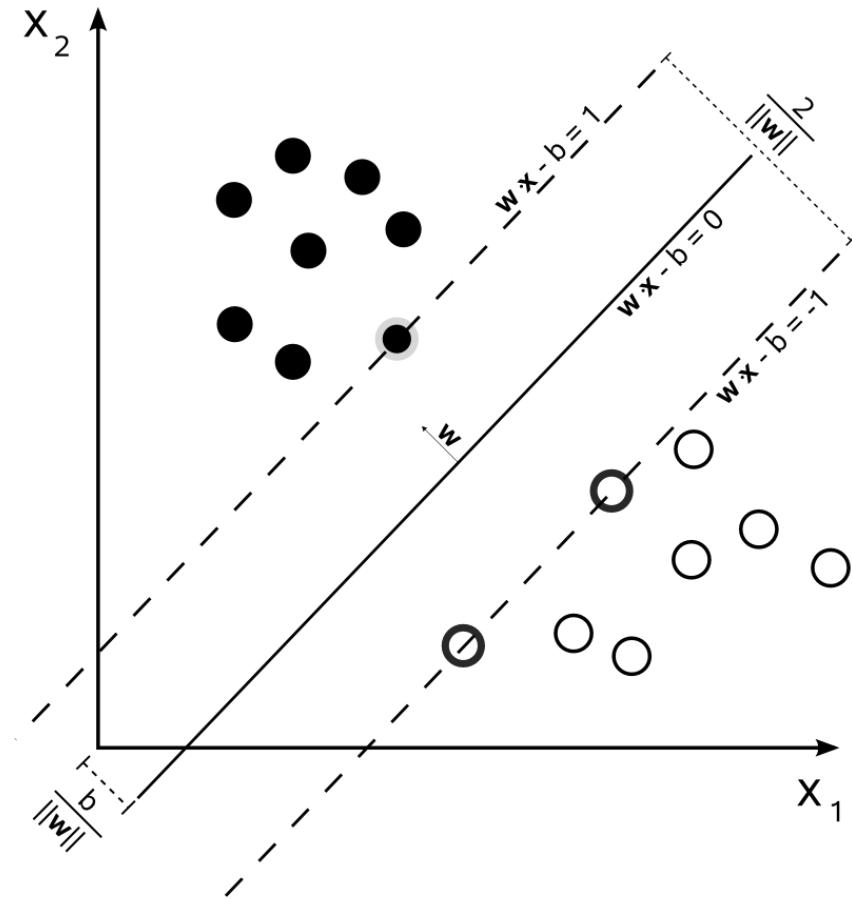


ERROR



WHY MAXIMIZE THE MARGIN?

- Points near the decision surface are inaccurate when predicting
- The bigger the margin, the more accurate the predictions



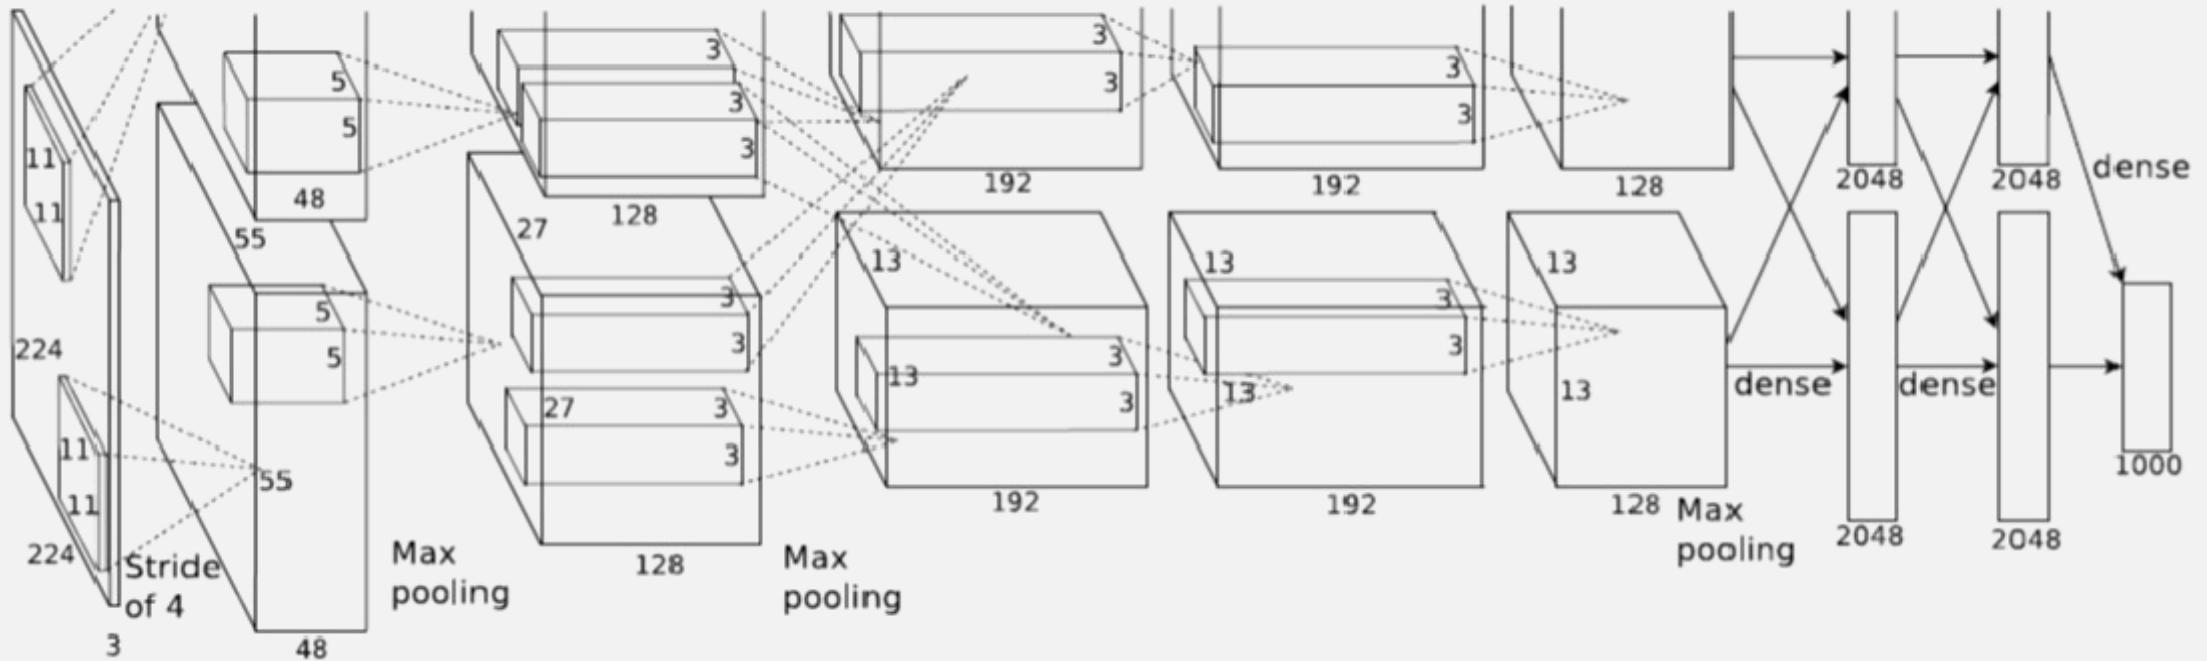
SVM ADVANTAGES

- Maximizing the margin it reduces overfitting
- Efficient: $O(n^3 * m)$
- Used for linear and non linear scenarios

DEEP LEARNING

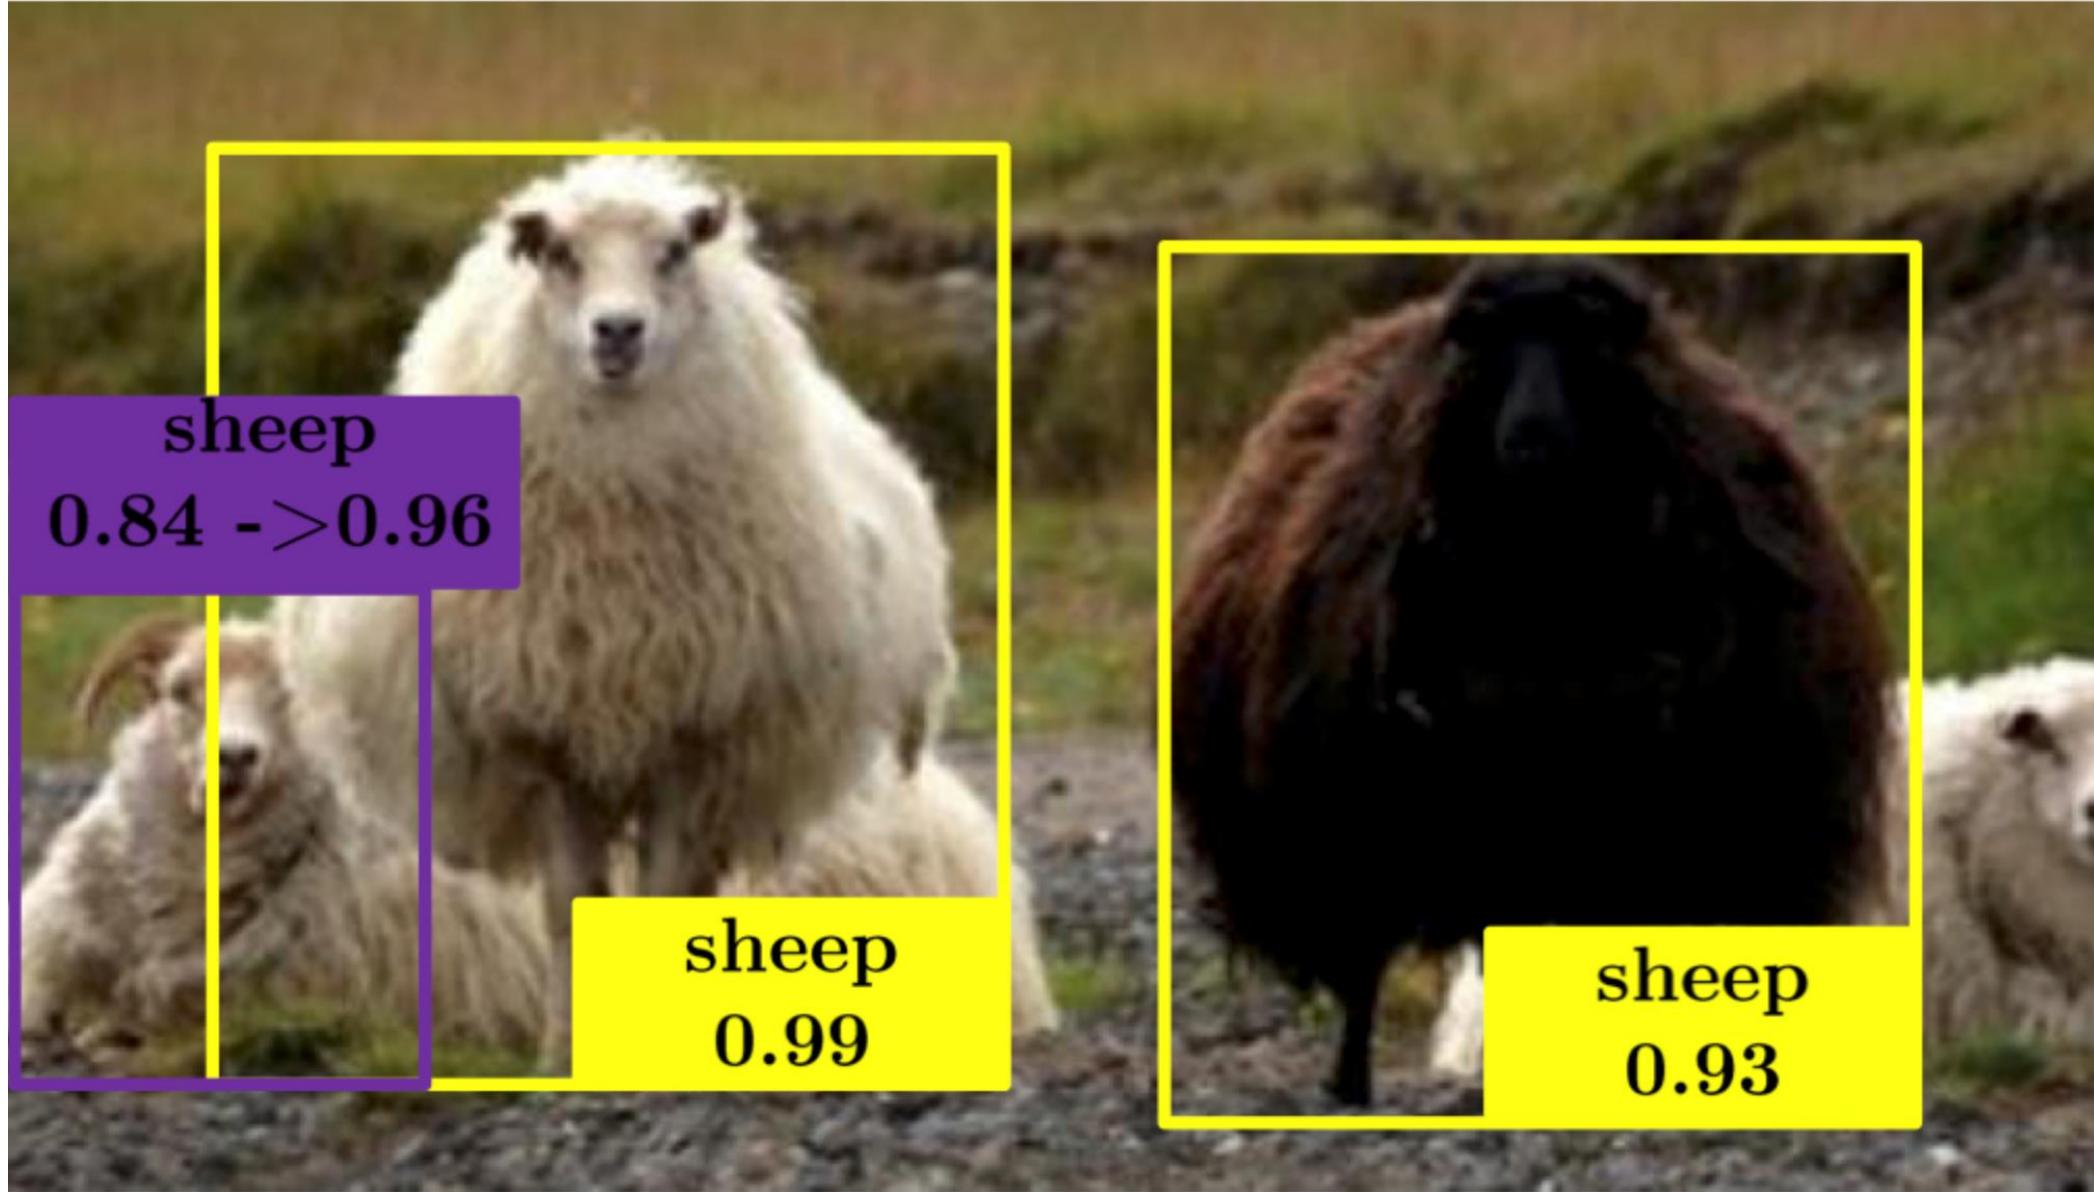
Who are **you**?

WHAT IS THIS THING?





WHAT IS THIS USEFUL FOR?

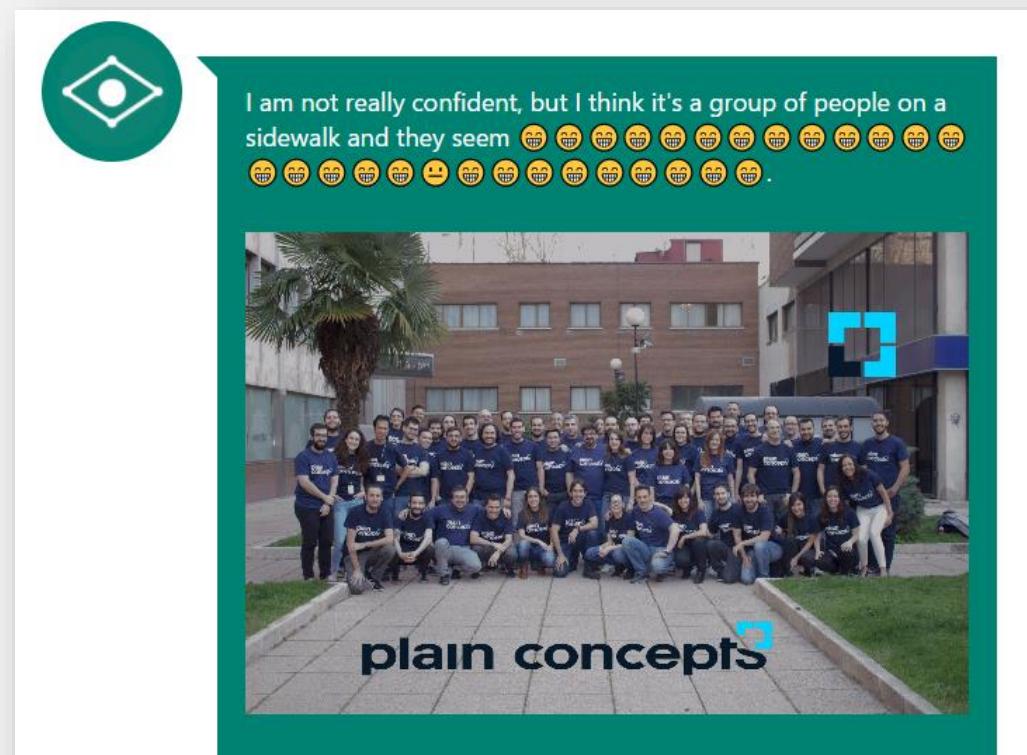








WHEN NOT TO USE THE COGNITIVE SERVICES?



How did I do?



plain concepts

Vision

Speech

Language

Knowledge

Search



How did I do?



Vision

- Computer Vision
 - Content Moderator
 - Emotion API
 - Face API

Speech

Language

Knowledge

Search



How did I do?



Vision

Speech

- Bing Speech API
 - Custom Speech Service
 - Speaker Recognition Service

Language

Knowledge

Search



How did I do?



Vision

Speech

Language

- Spell Check API
 - LUIS
 - Linguistic Analysis
 - Translator API

Knowledge

Search



How did I do?



Vision

Speech

Language

Knowledge

- Entity Linking Service
 - Academic Knowledge Service
 - Q&A Maker
 -

Search



How did I do?



Vision

Speech

Language

Knowledge

Search

- Autosuggest API
 - Image Search API
 - News Search API
 - ...



How did I do?



WHEN NOT TO USE THE COGNITIVE SERVICES?



WHEN NOT TO USE THE COGNITIVE SERVICES?



WHEN NOT TO USE THE COGNITIVE SERVICES?



WHEN NOT TO USE THE COGNITIVE SERVICES?

- Vienen pre-entrenados



DEEP LEARNING WITH CNTK

DEEP LEARNING

Machine Learning

Based on Neural Networks

More than one hidden layer (*Deep Models*)

CNTK

- Don't ask me for the acronym. Really. Don't.
-

Why?

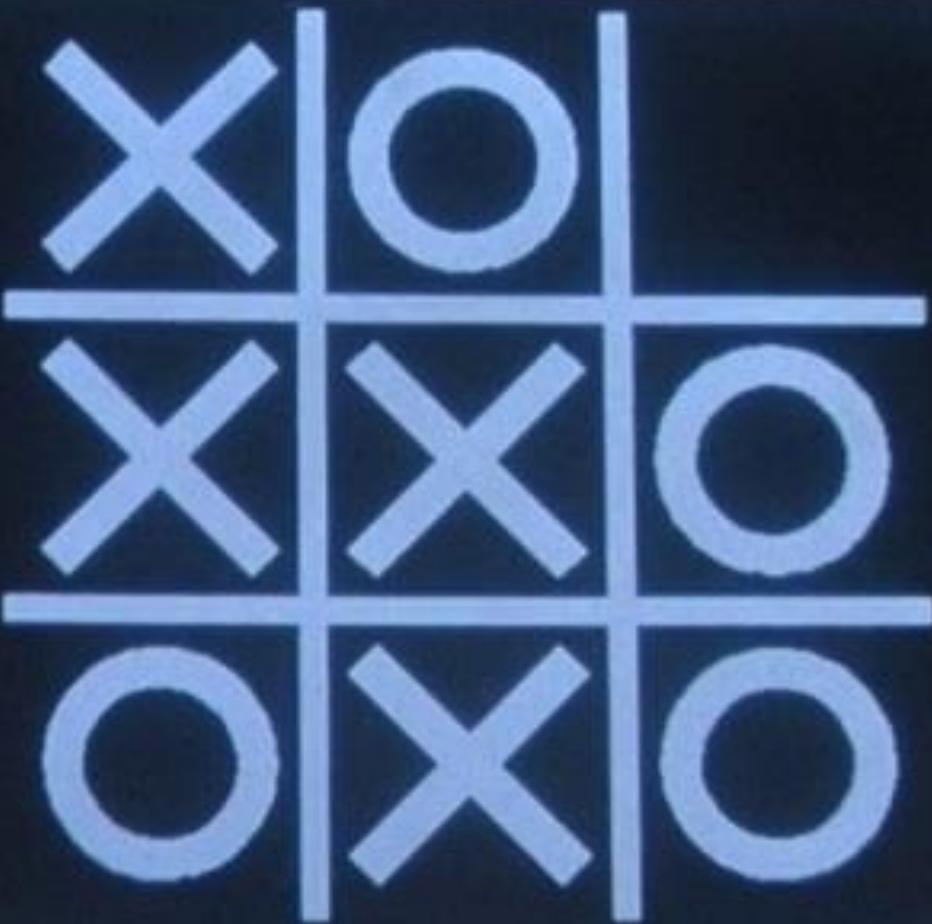
- Faster model development
- Handles multiple CPUs and GPUs
- Multi-platform
- Internally used at Microsoft. For real.

CNTK

- Don't try. I still don't want to talk about the meaning of the acronym.
-

- It is an acronym of something 😊
- Deep Learning framework, open-source with external contributions by MIT and Stanford, among others
- Allows for the expression of complex neural networks in a straightforward way, and performs all the steps in the lifecycle: from training to evaluation.
- Development in BrainScript, Python (.NET in progress)

SHALL WE PLAY A LITTLE GAME?



COM STS
8365

GPM STS
3603

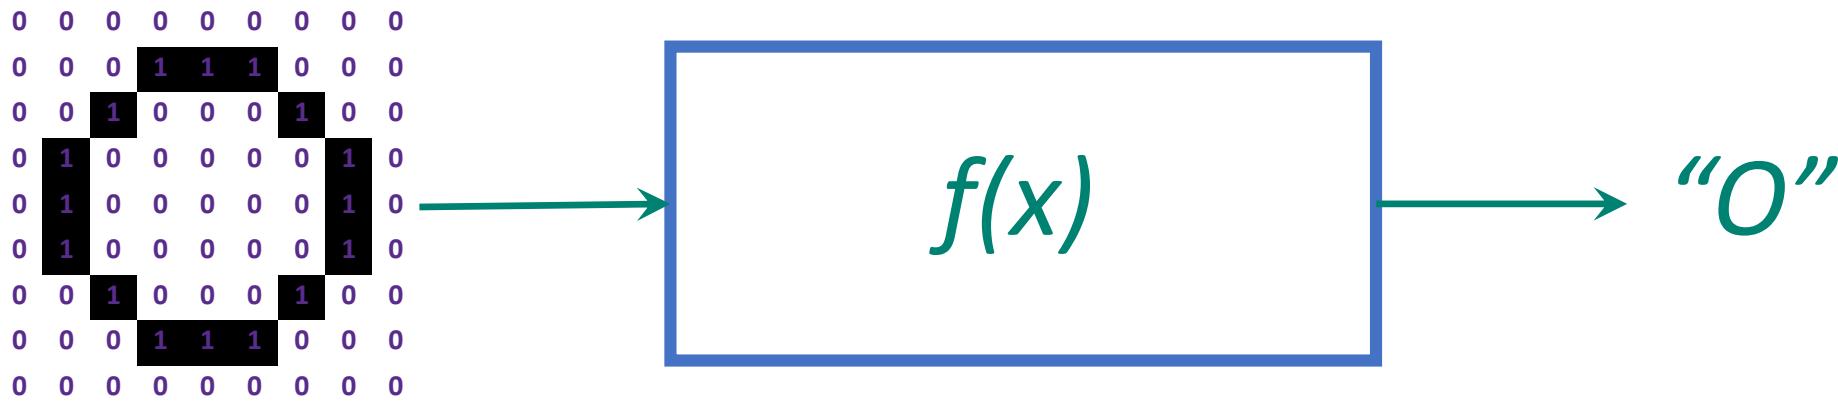
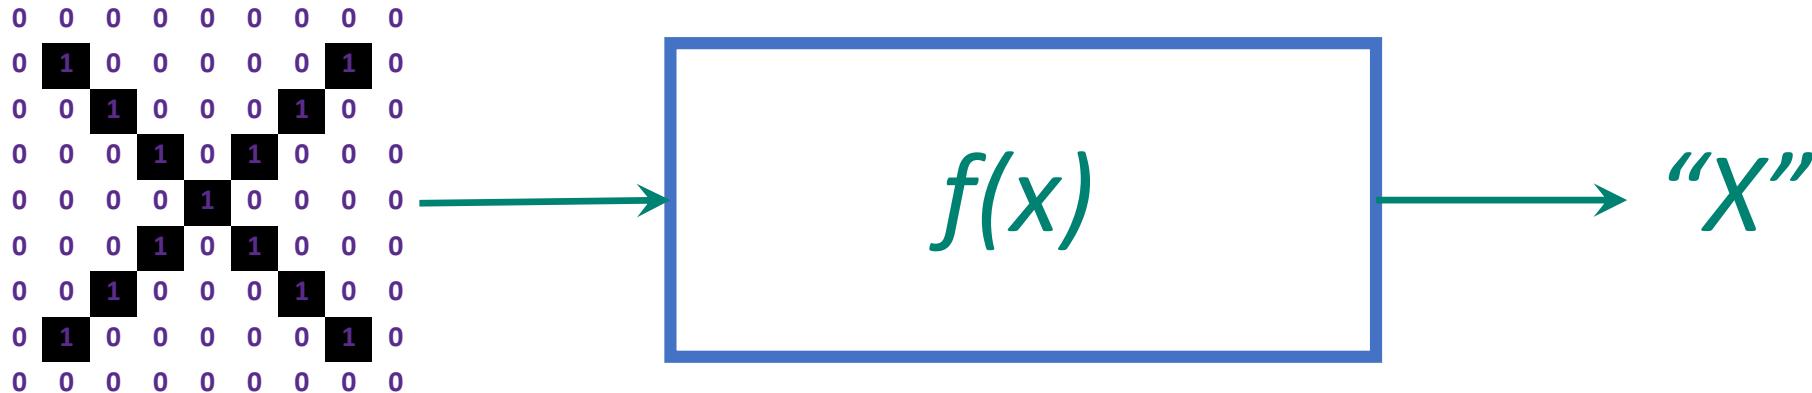
RPL STS
9071

EOT STS
4531

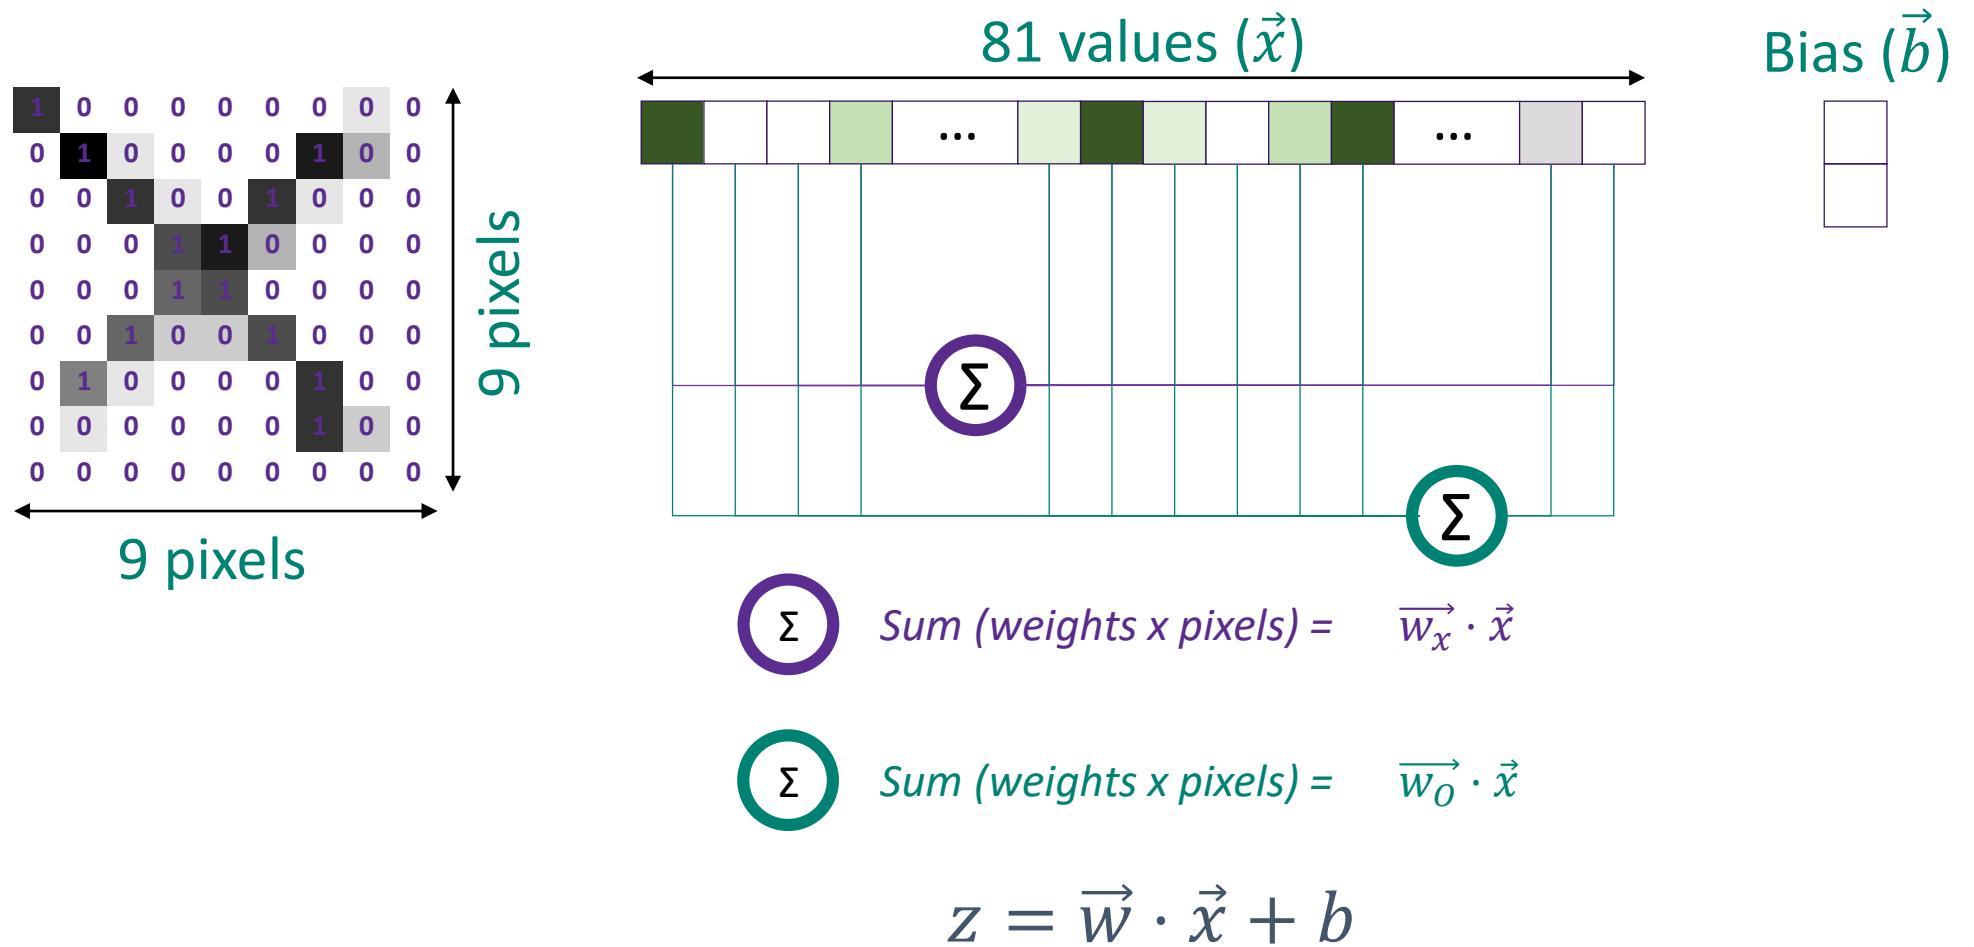
PAC STS
8838

ADL STS
6632

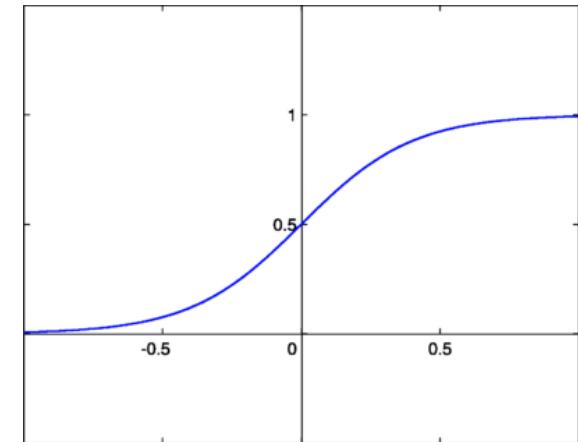
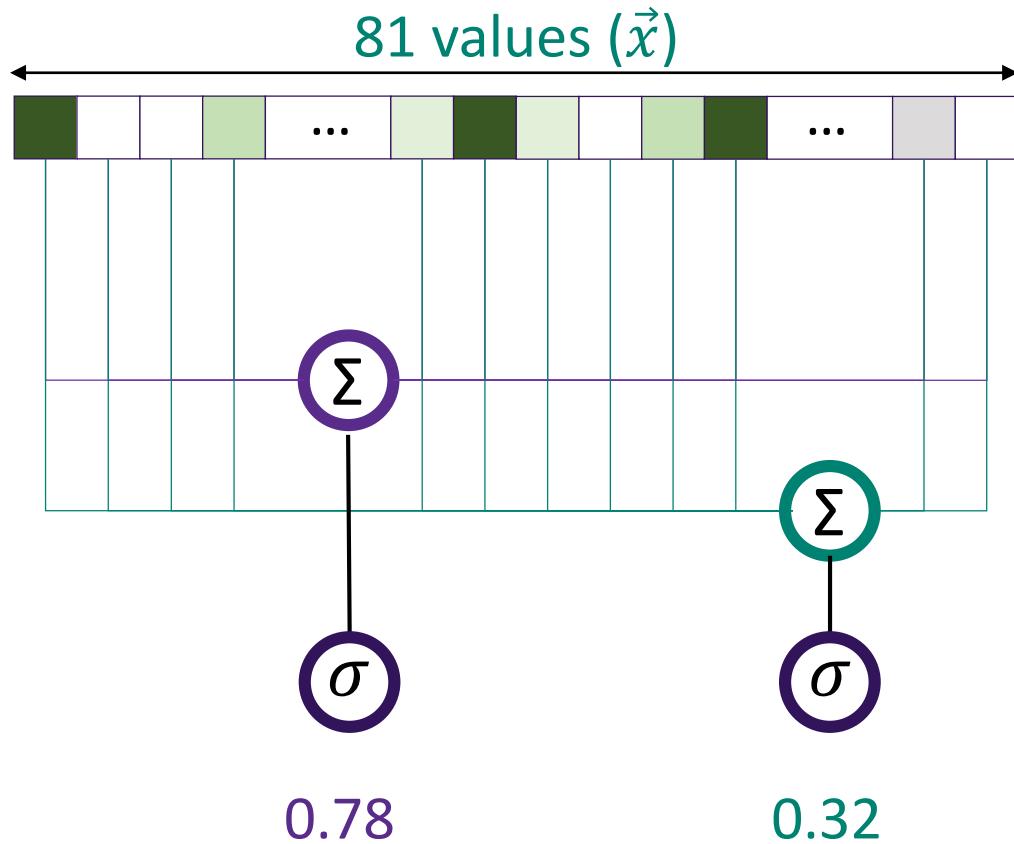
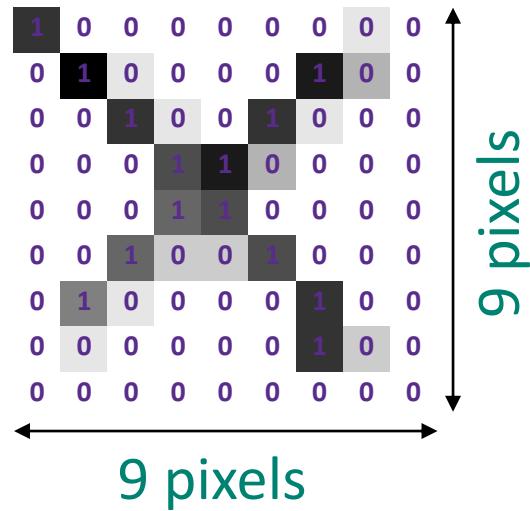
A SIMPLE EXAMPLE



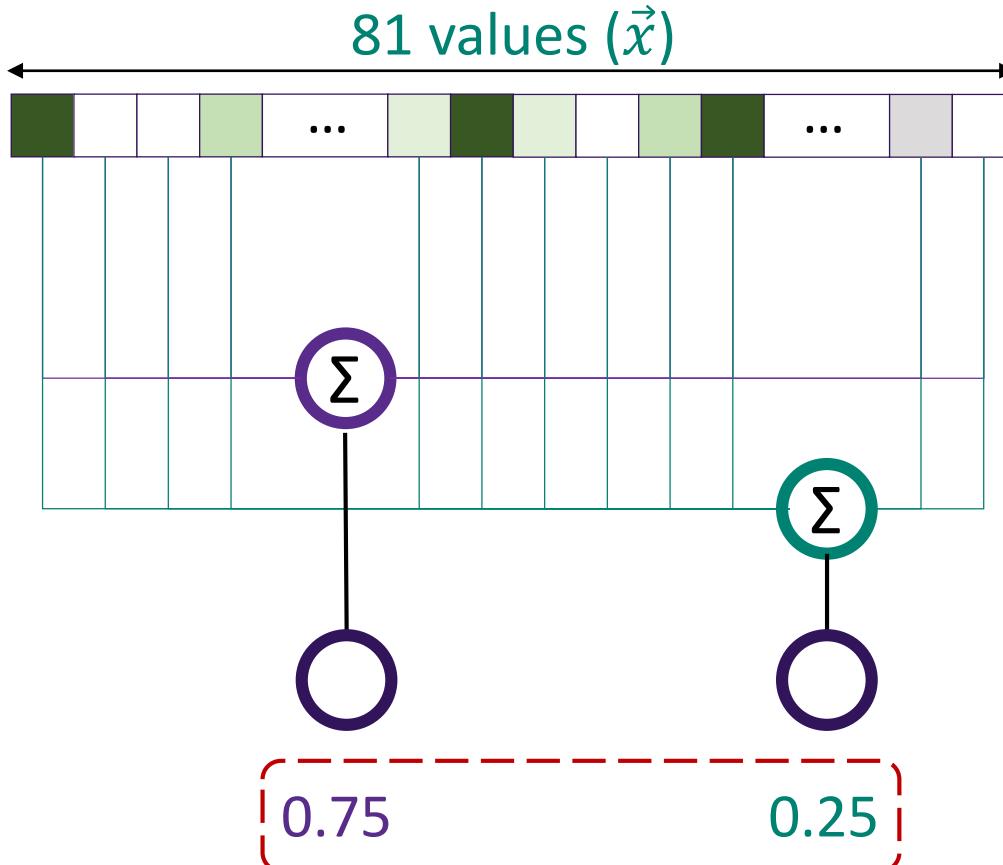
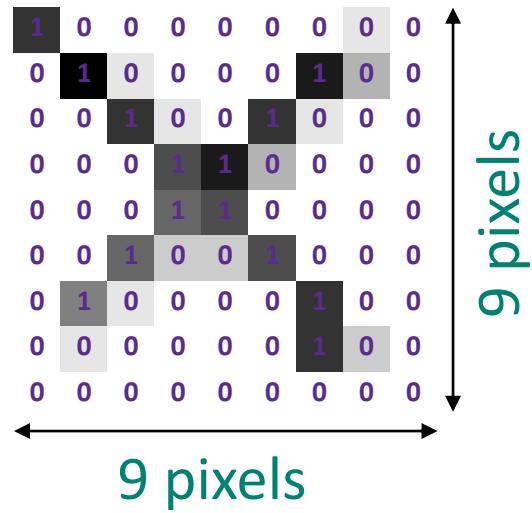
A SIMPLE EXAMPLE



SIGMOID

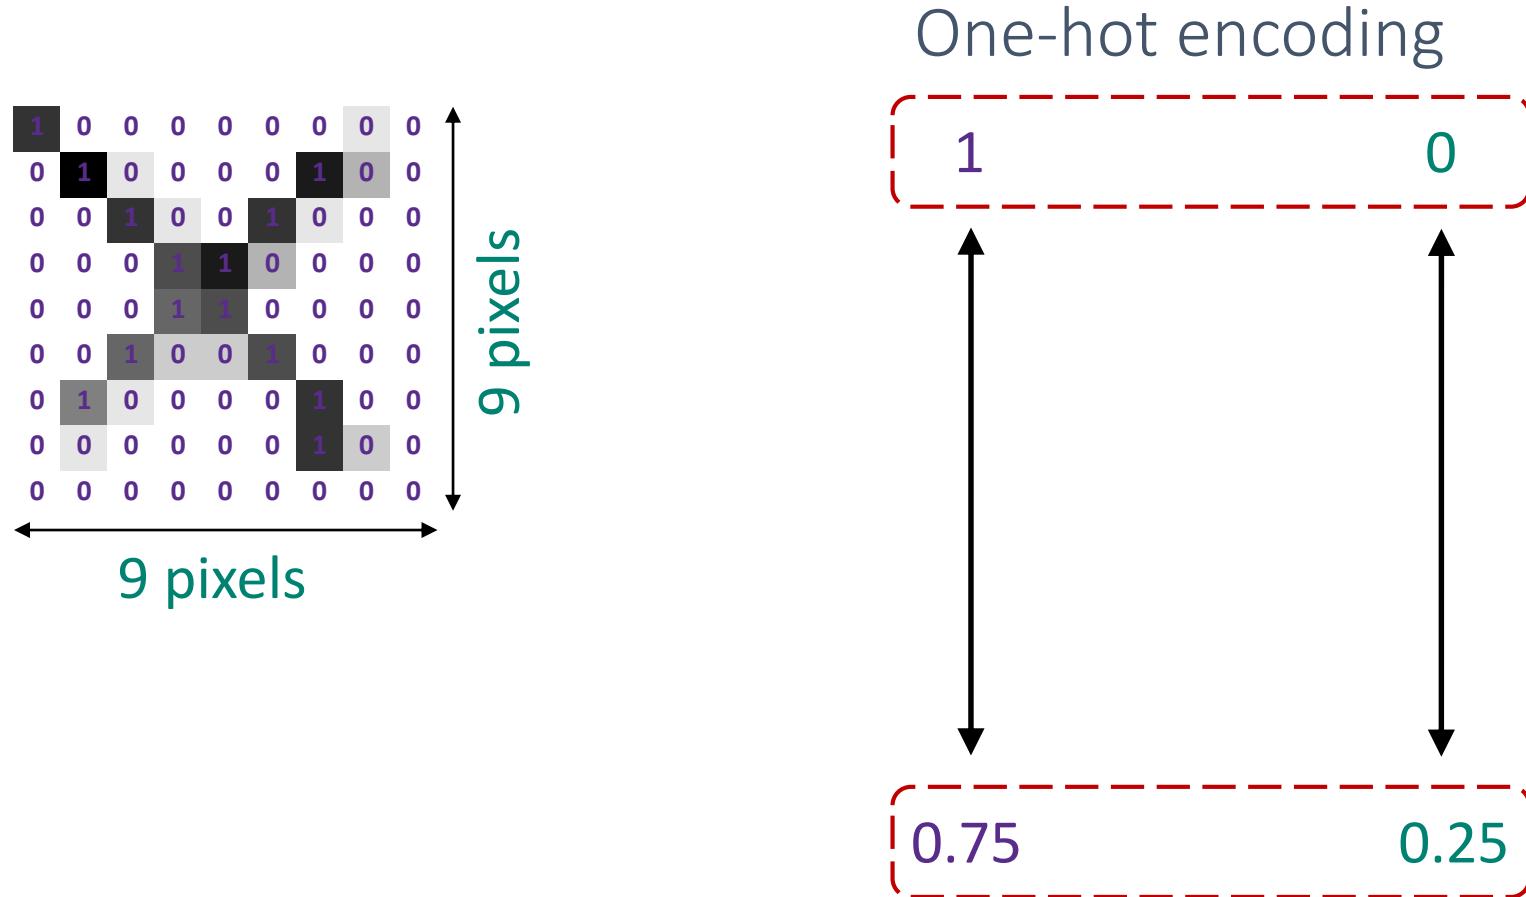


SOFTMAX

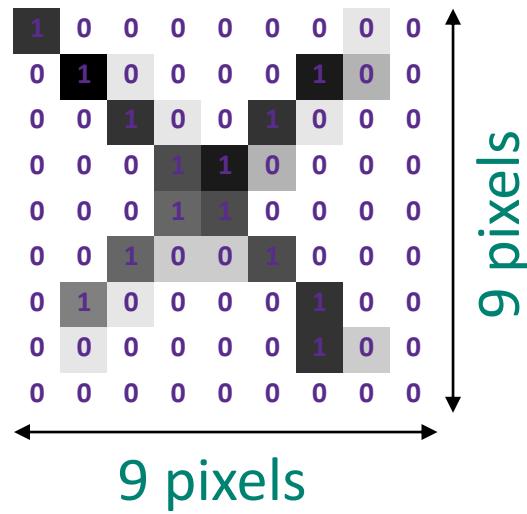


$$s(y_i) = \frac{e^{y_i}}{\sum_j e^{y_i}}$$

LOSS FUNCTION



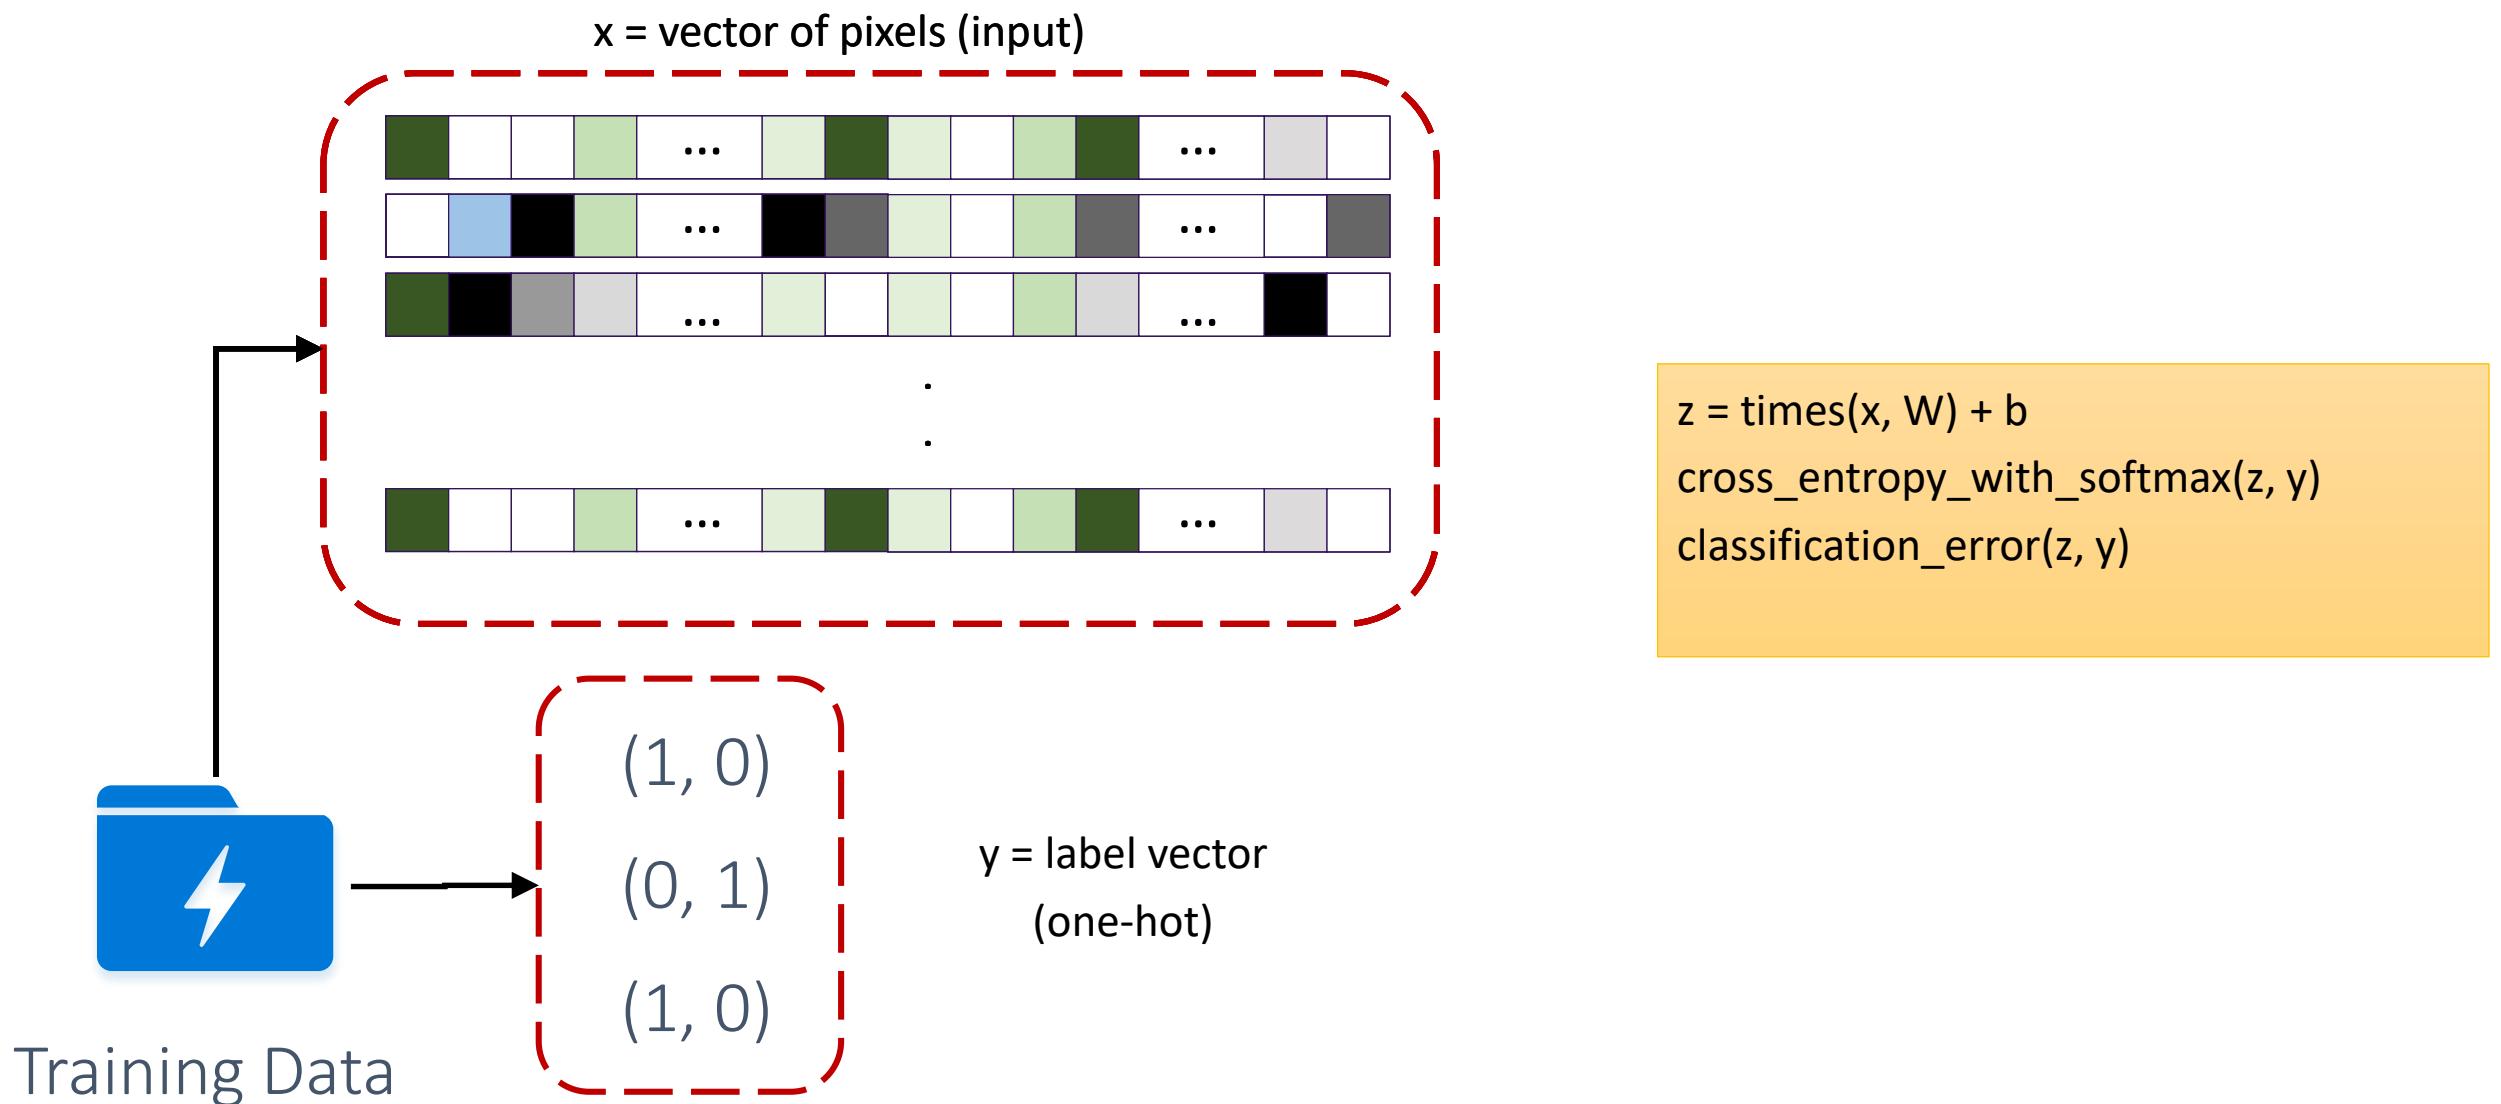
LOSS FUNCTION



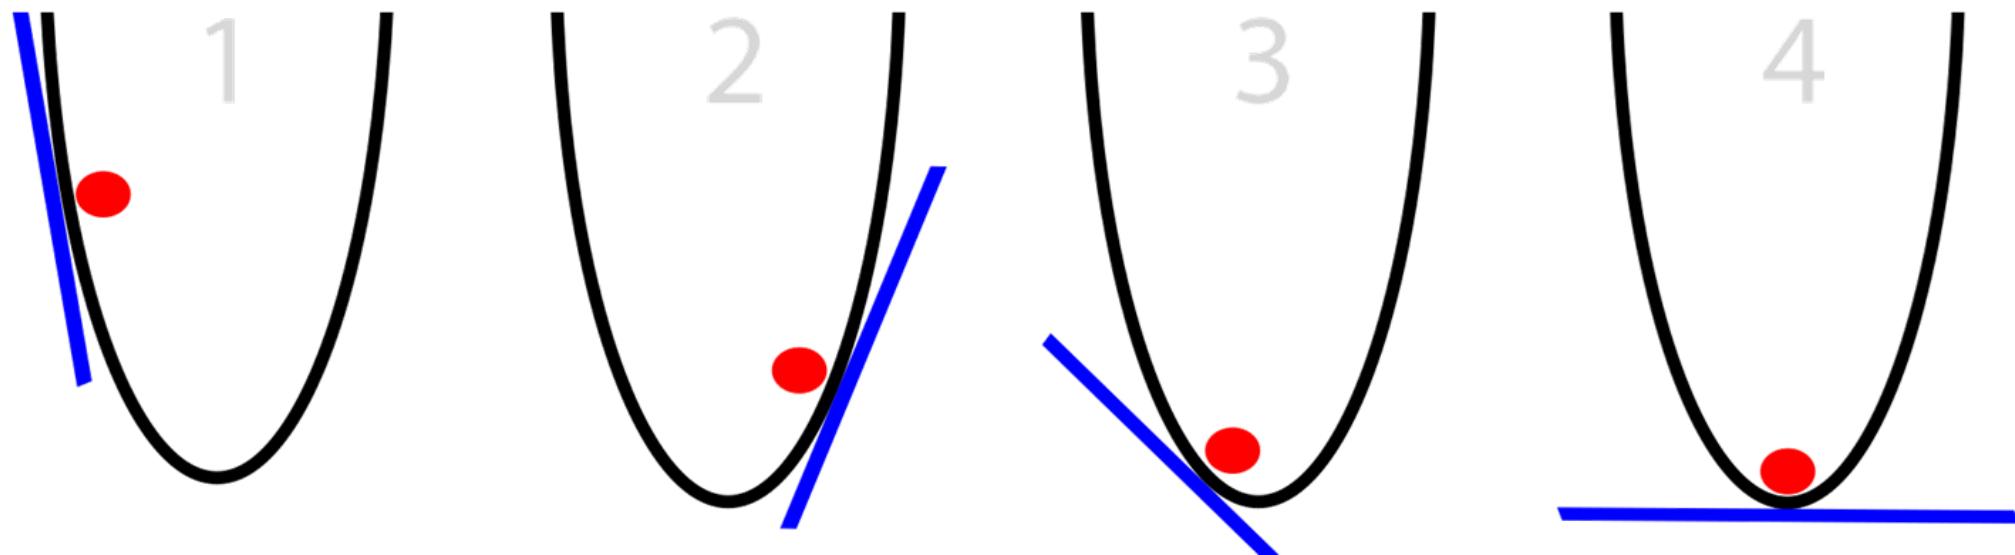
$$\text{error} = \frac{1}{n} \sum_{i=1}^n (y_j - p_j)^2$$



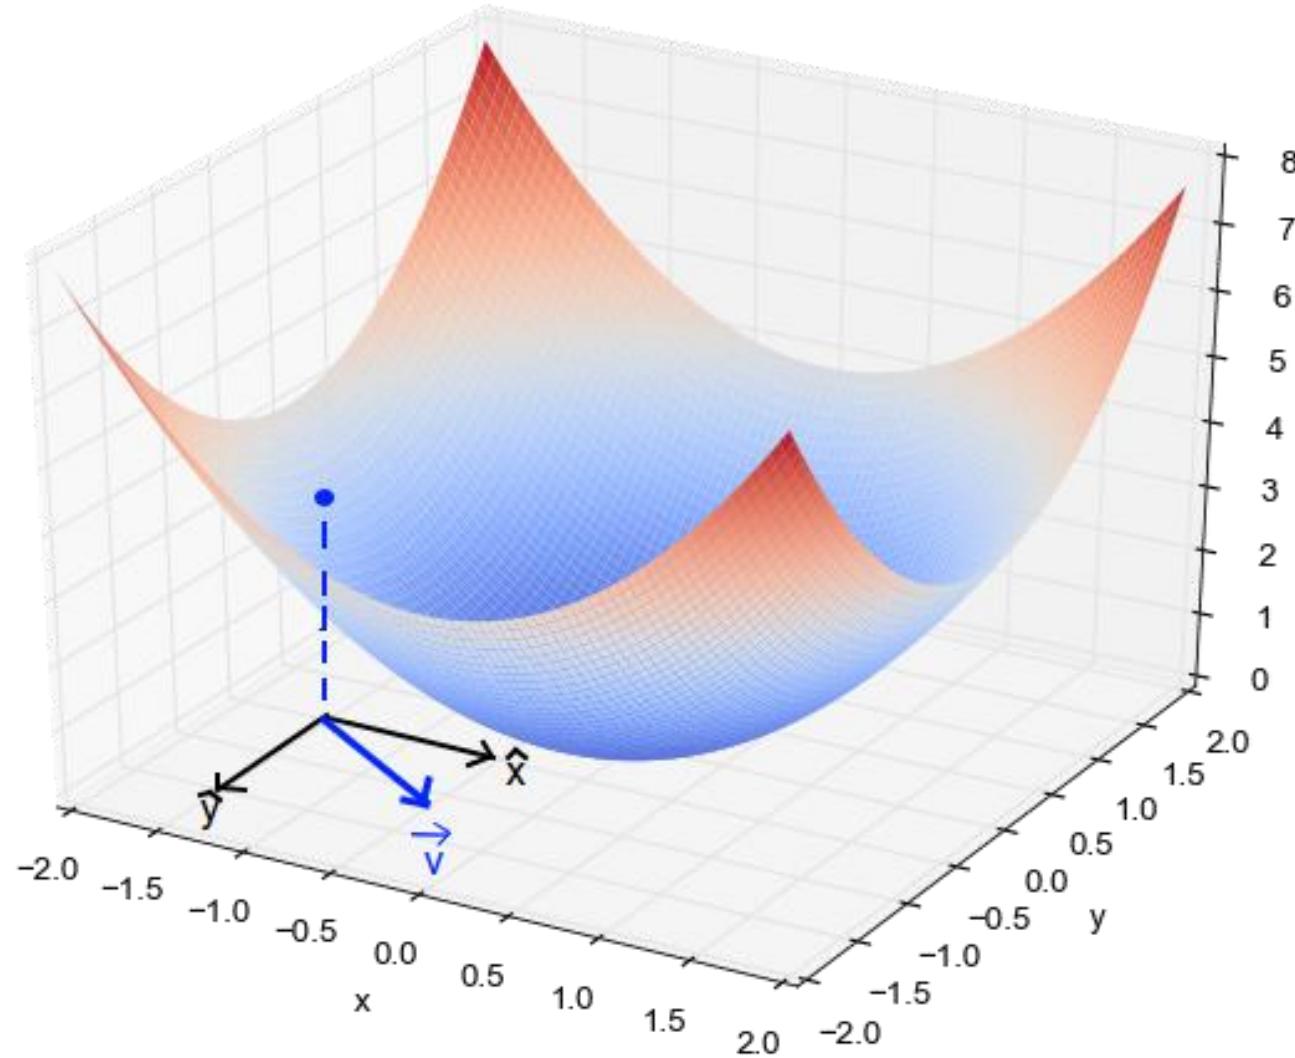
TRAINING



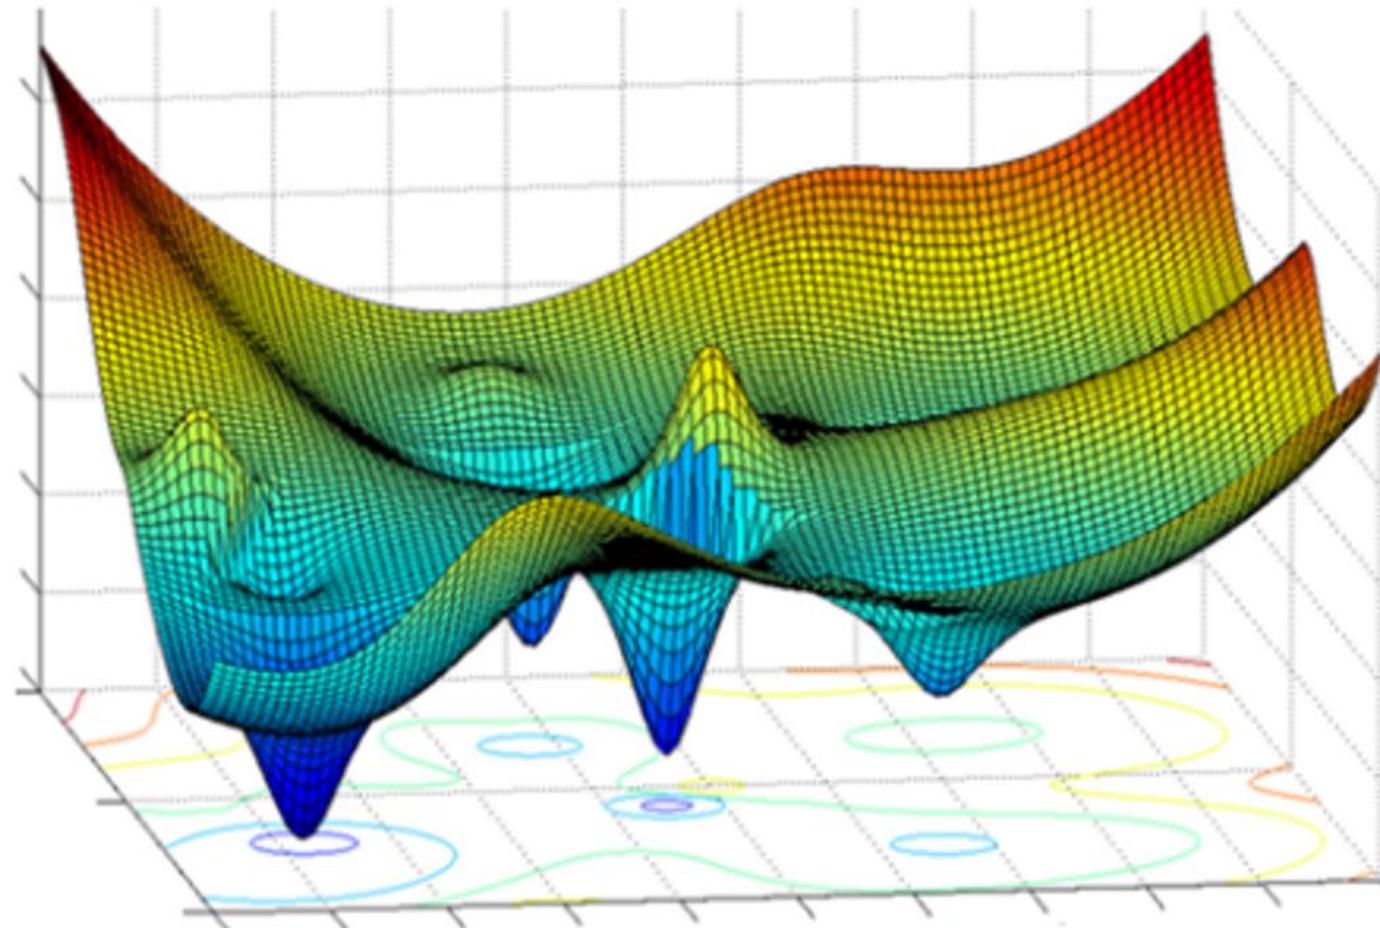
GRADIENT DESCENT



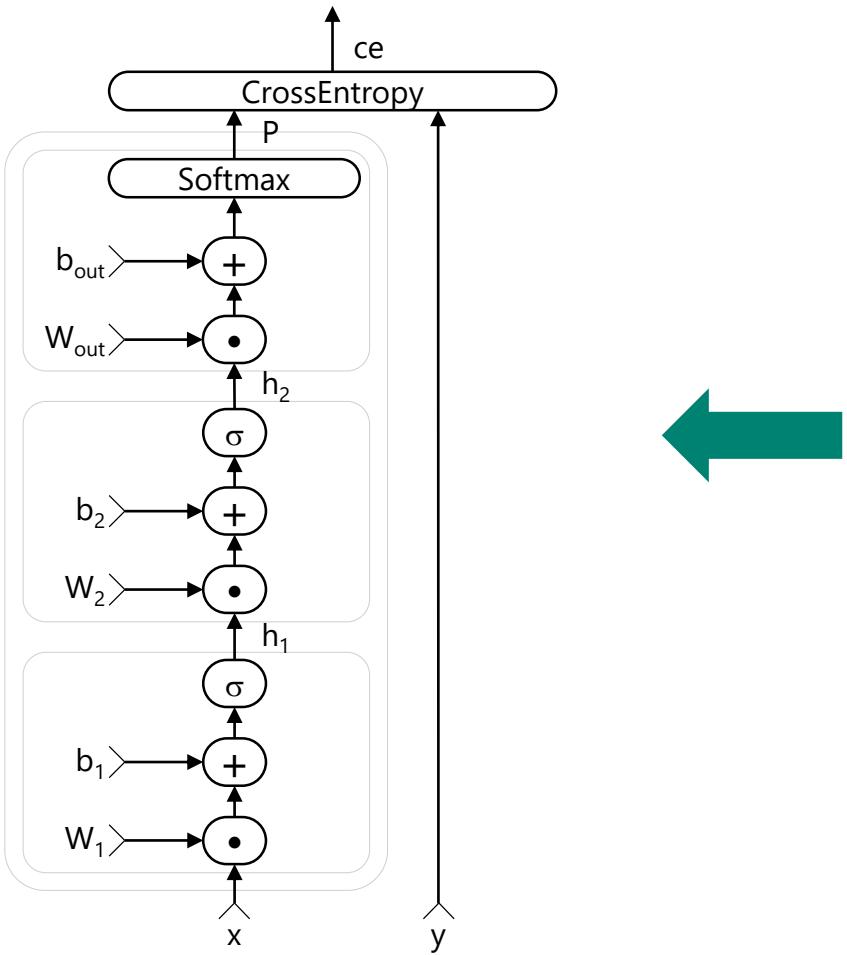
GRADIENT DESCENT



GRADIENT DESCENT



HOW DO WE DESIGN THIS IN CNTK?

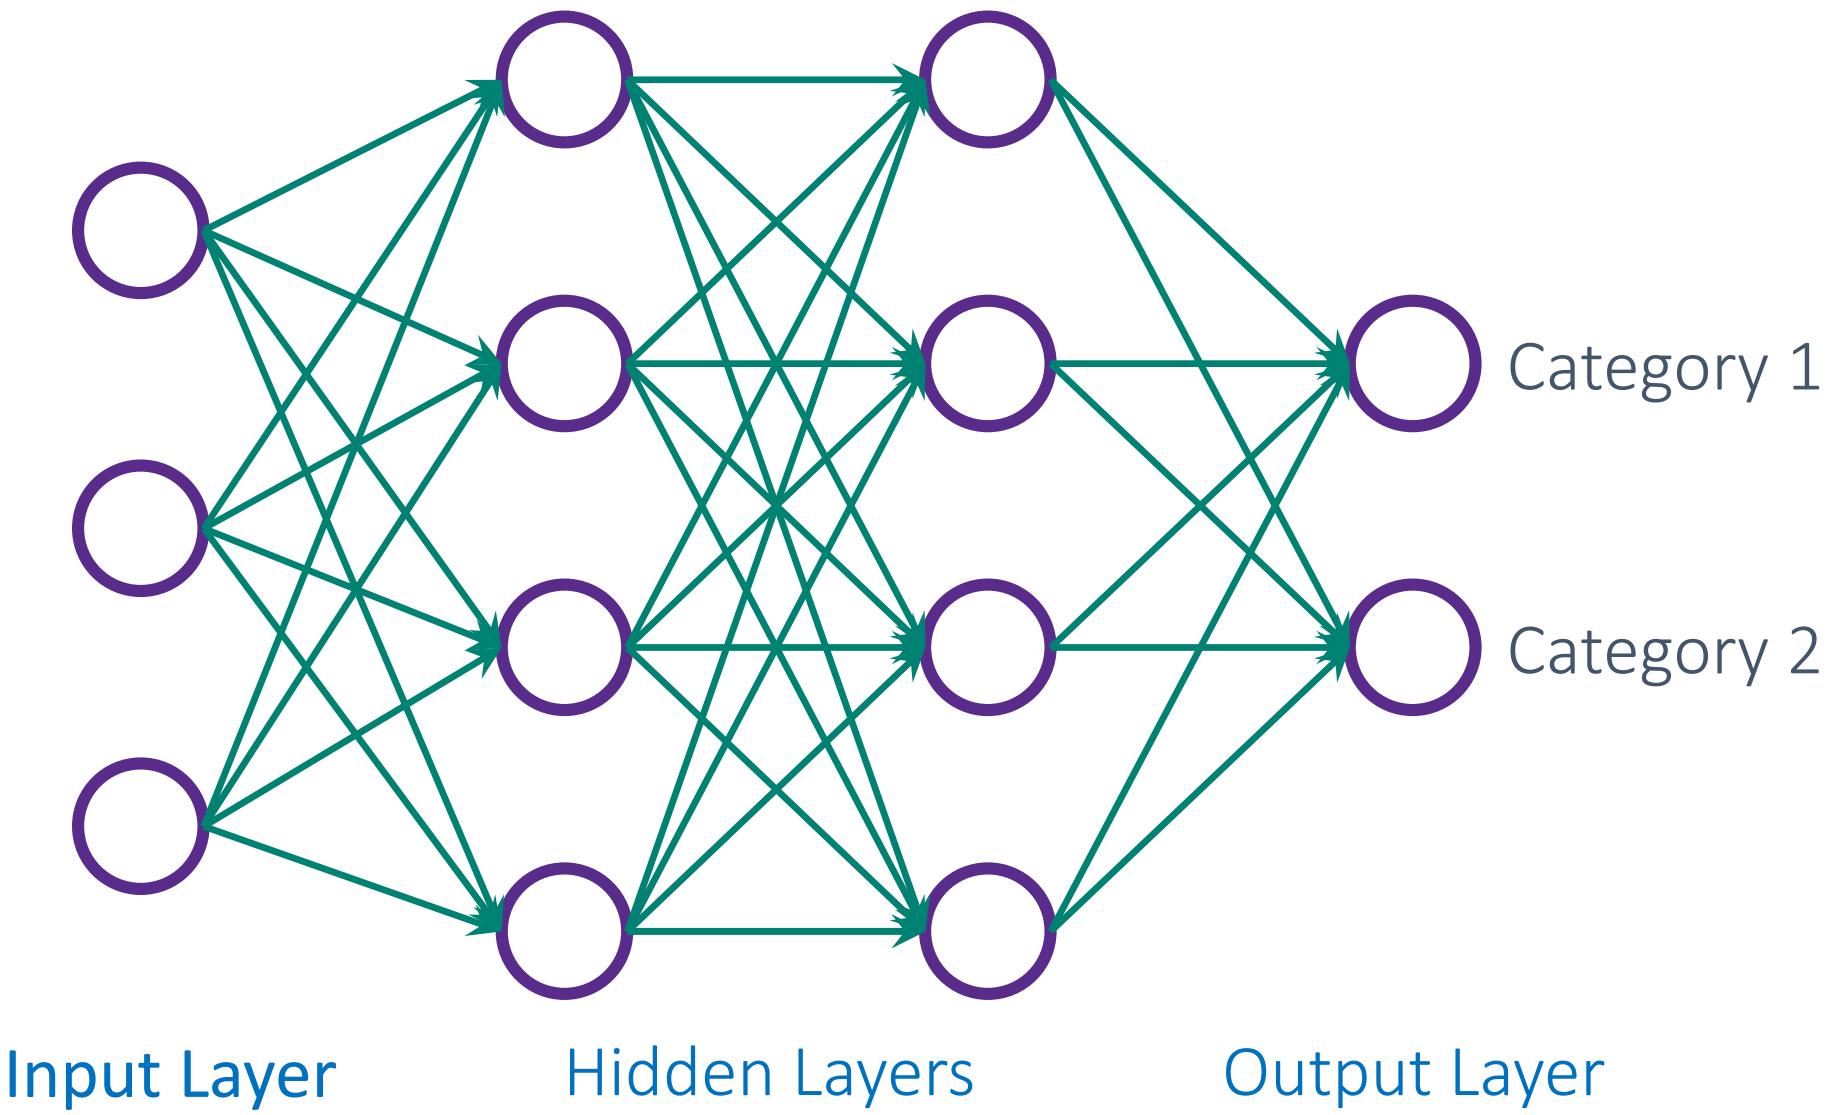


h1 = Sigmoid ($w_1 * x + b_1$)
h2 = Sigmoid ($w_2 * h_1 + b_2$)
P = Softmax ($w_{out} * h_2 + b_{out}$)
ce = CrossEntropy (y, P)

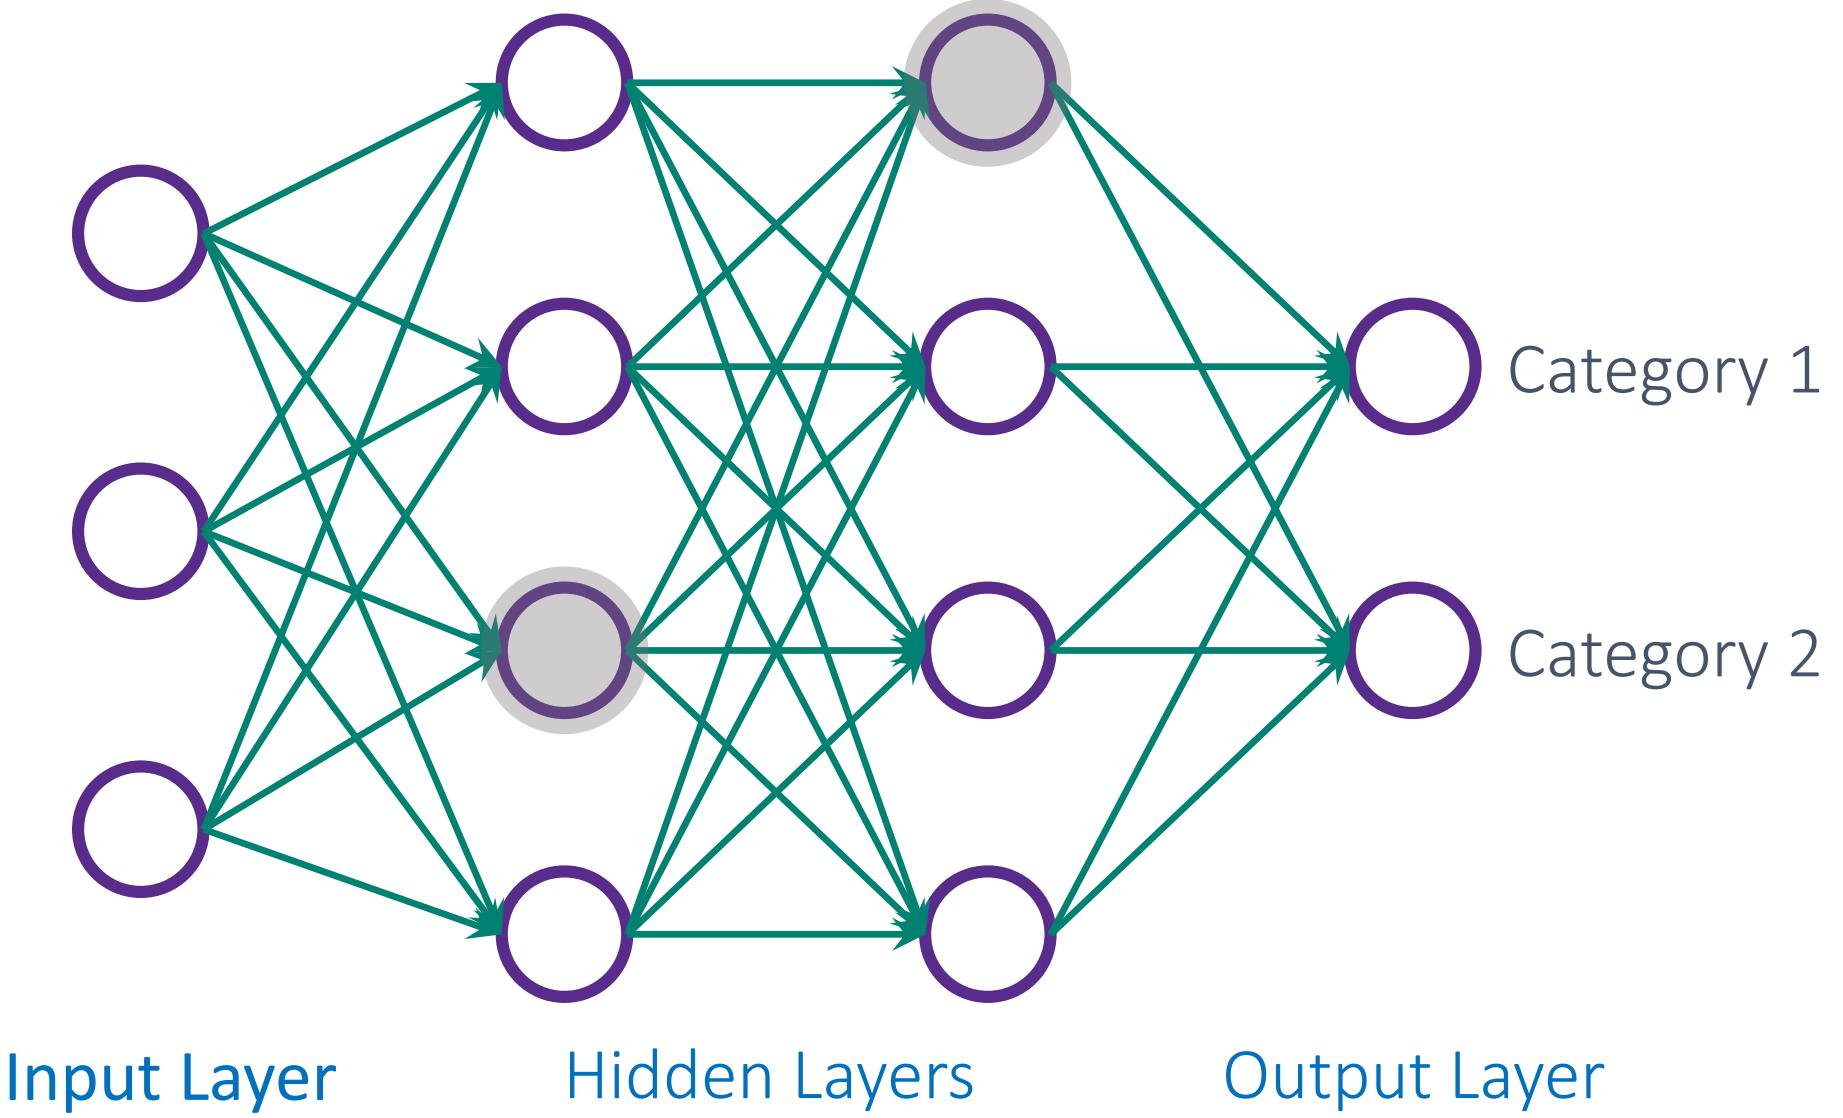
LAB 1: LOGISTIC REGRESSION

NEURAL NETWORKS

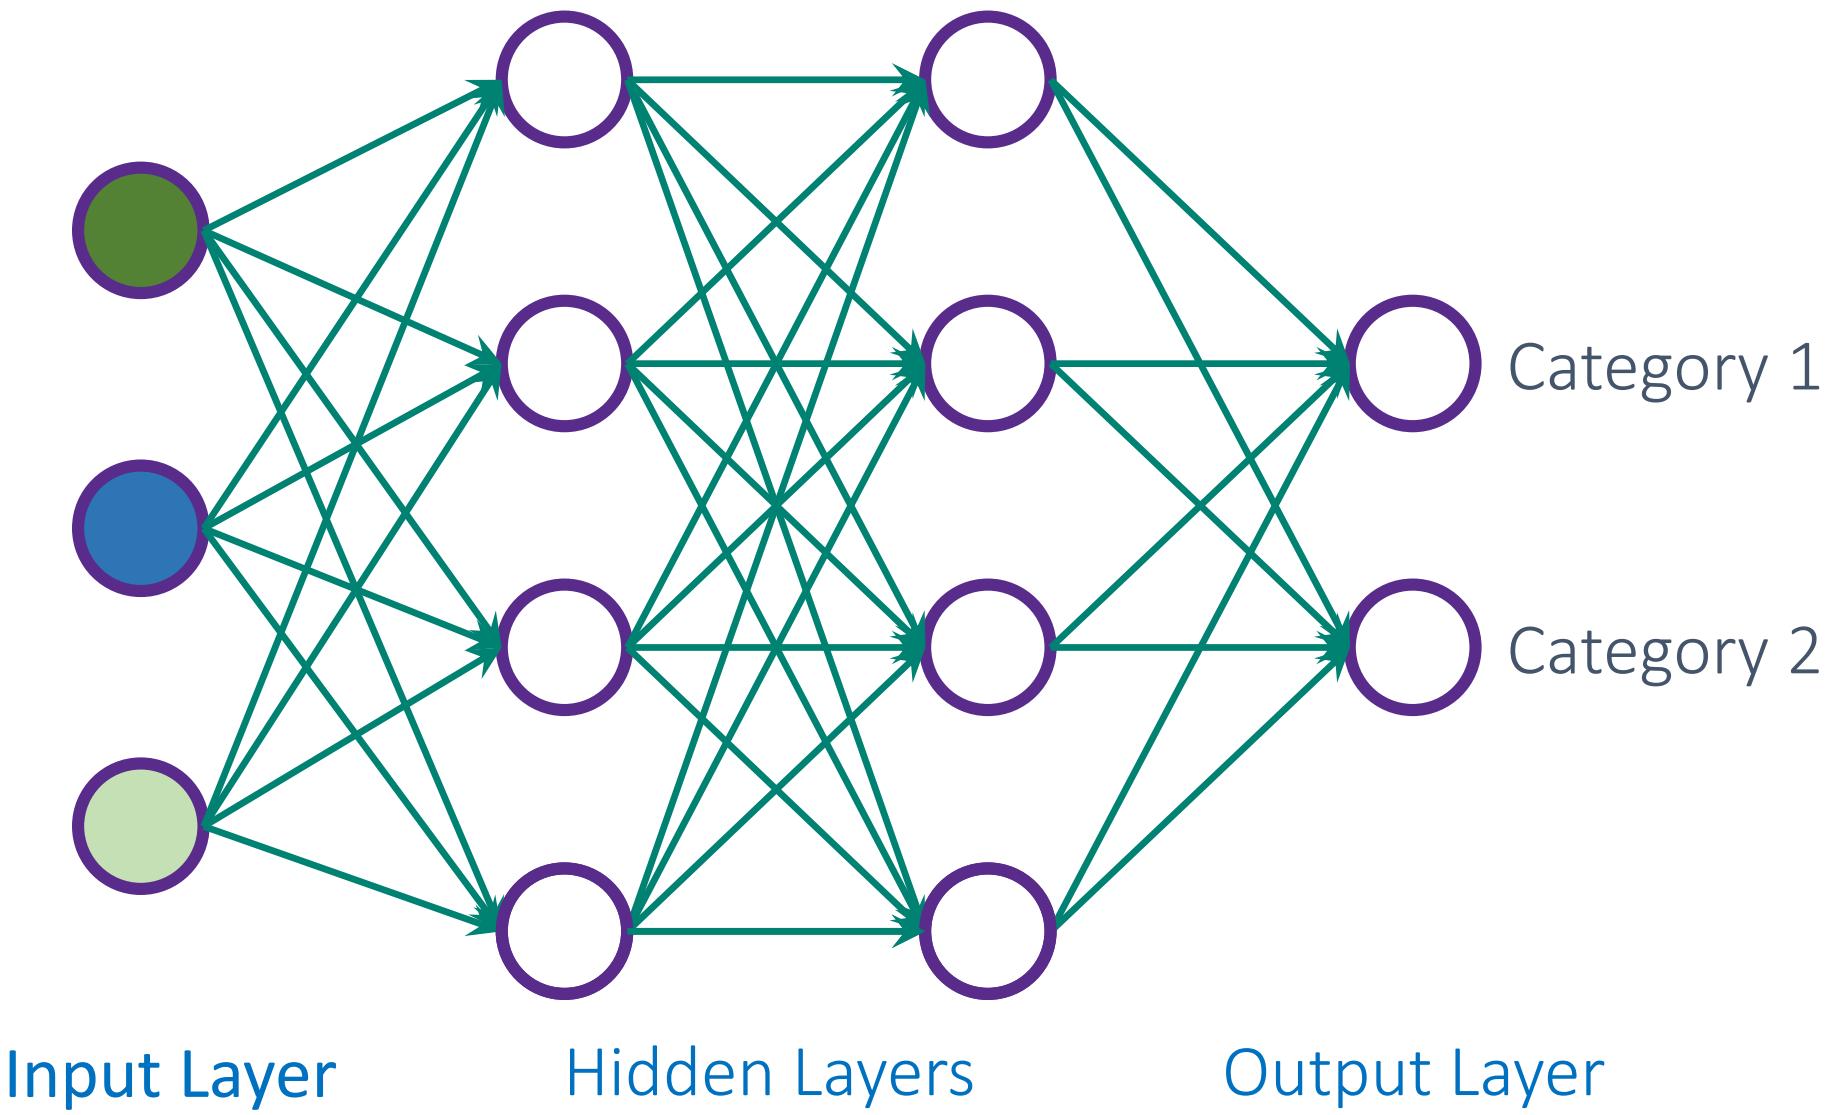
NEURAL NETWORK (CLASSIFIER)



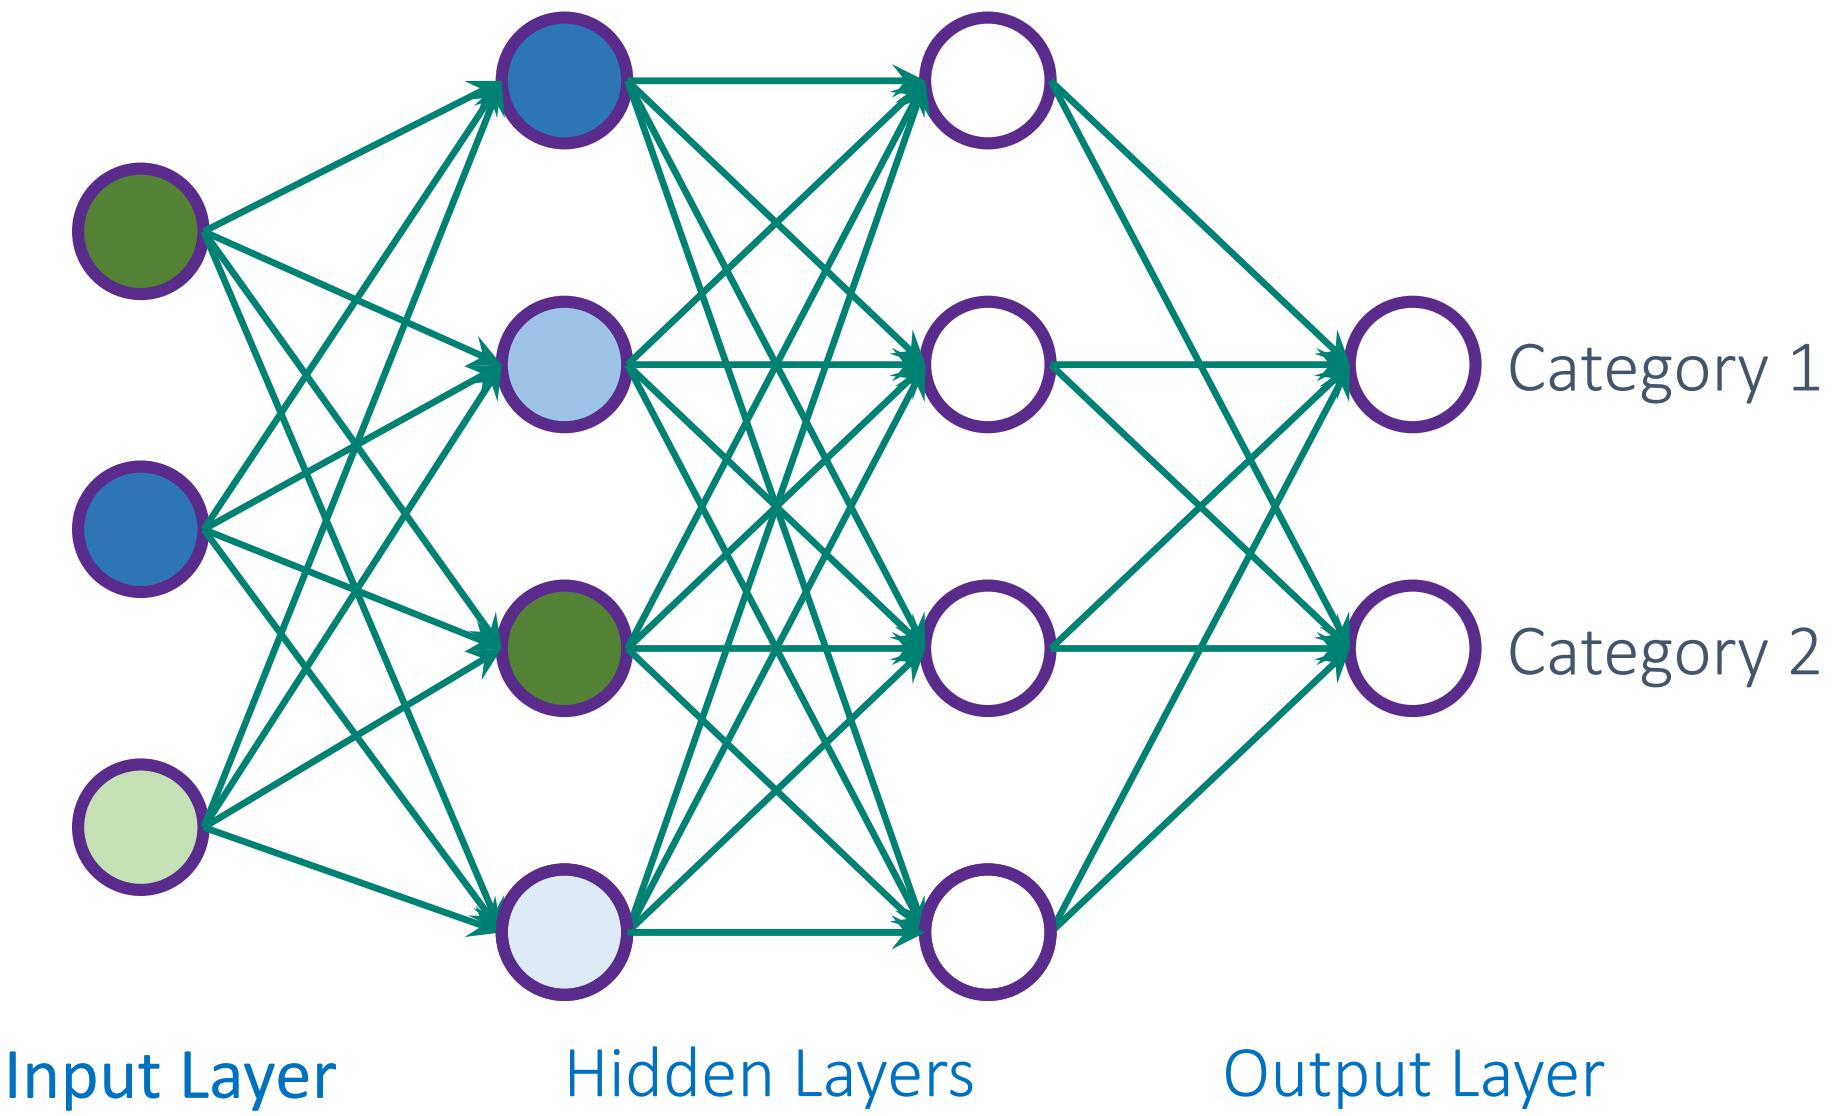
NEURAL NETWORK (CLASSIFIER)



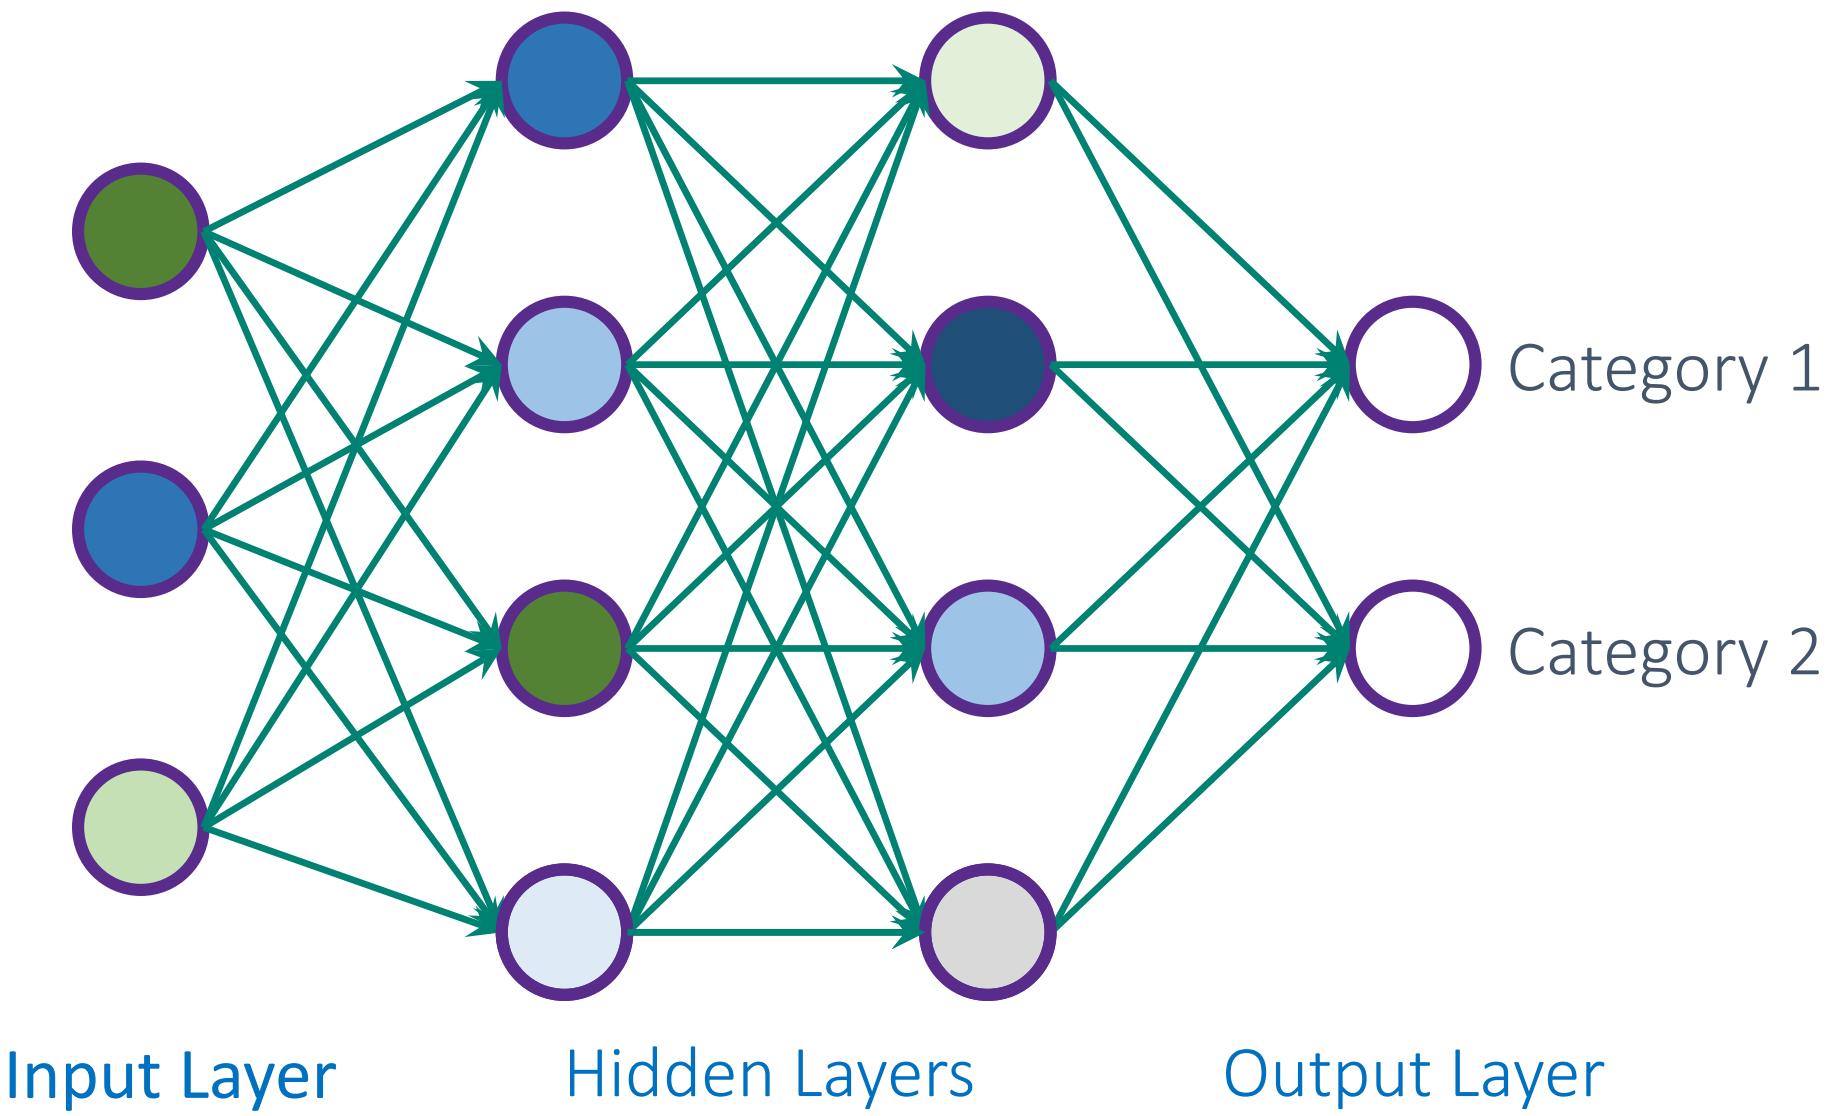
FORWARD PROPAGATION



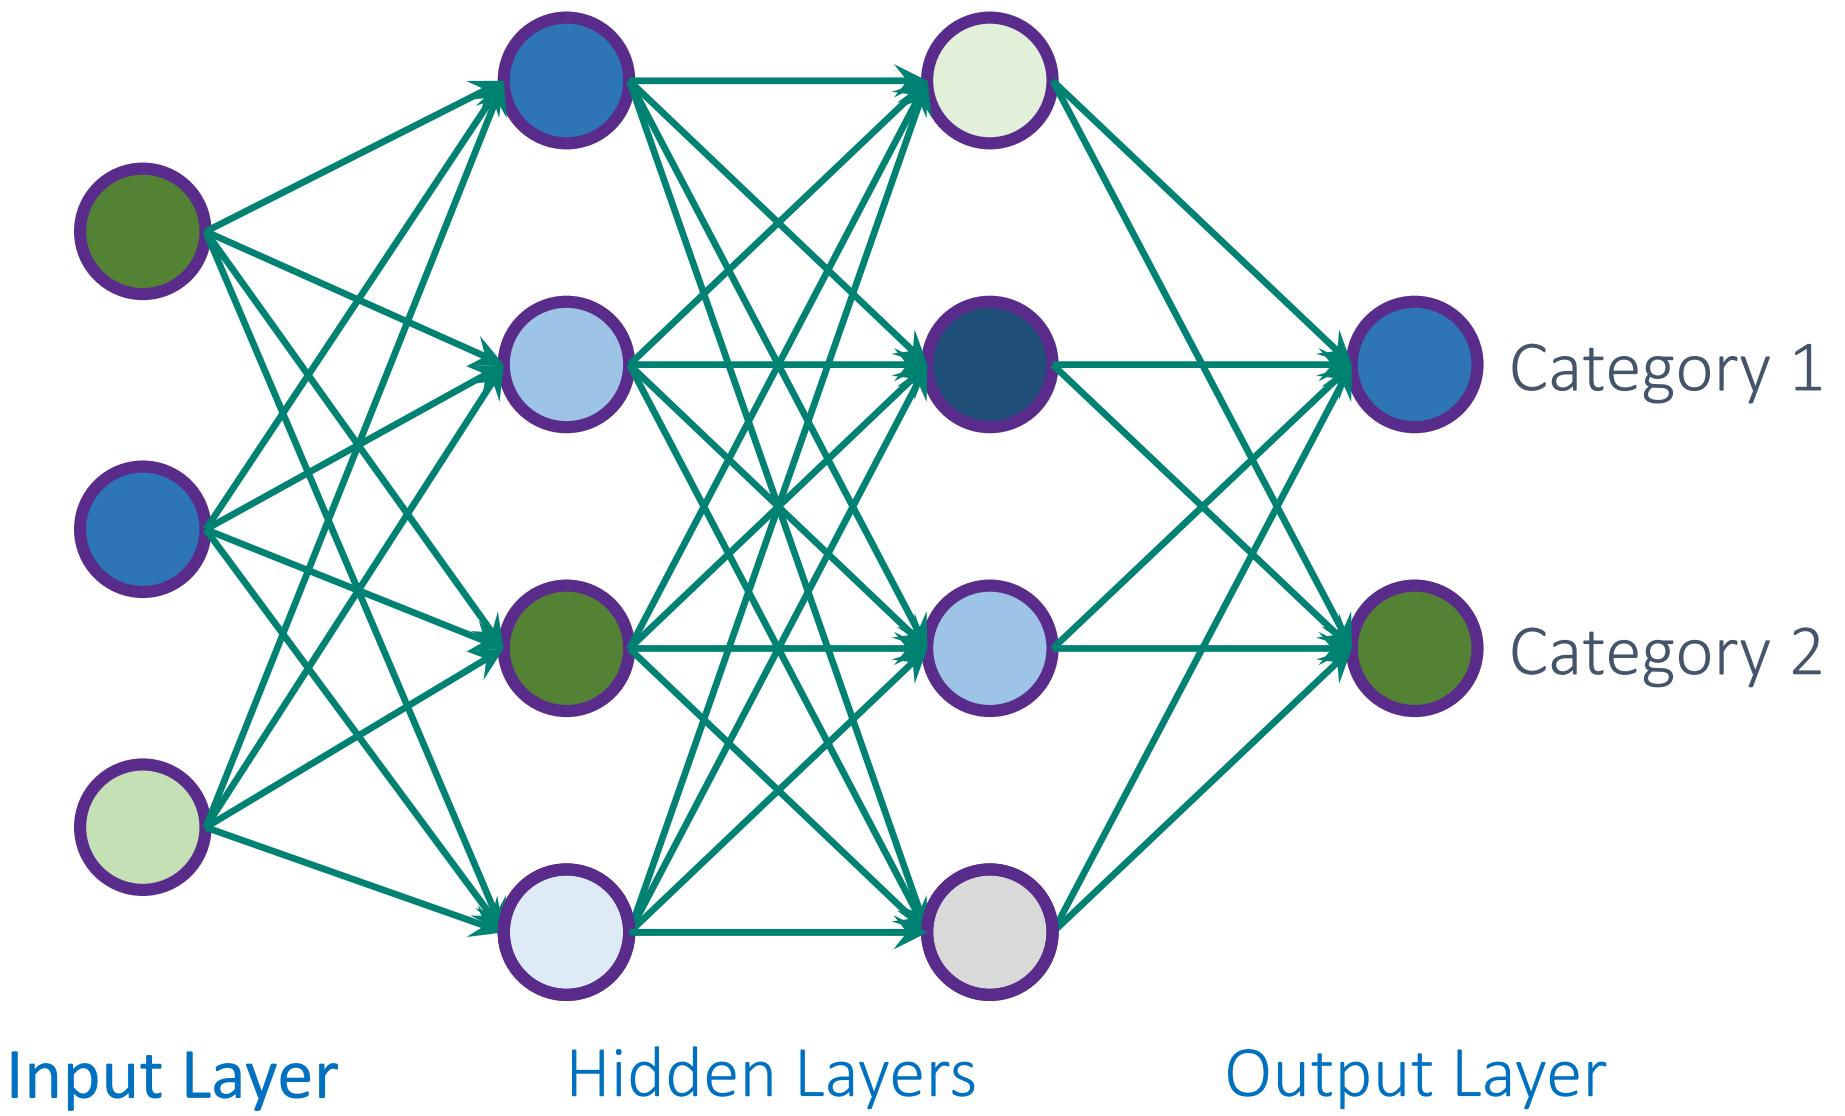
FORWARD PROPAGATION



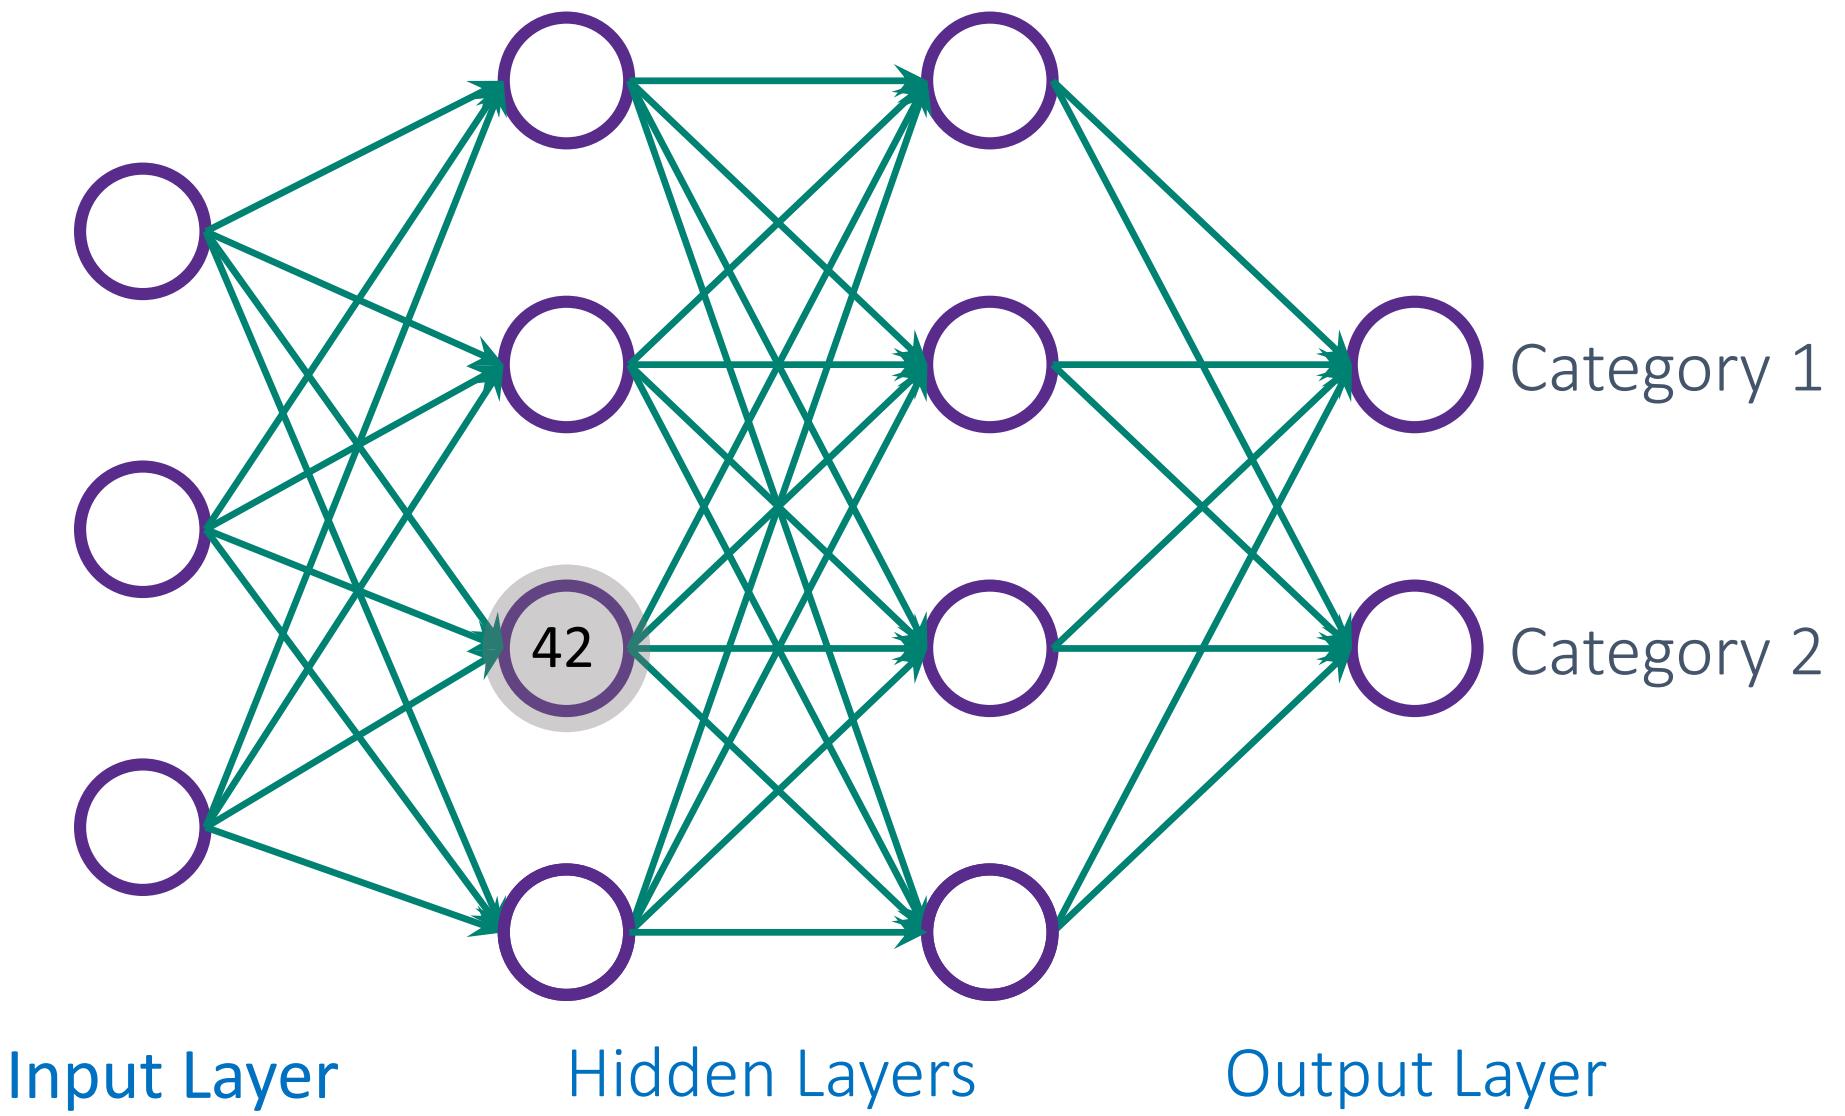
FORWARD PROPAGATION



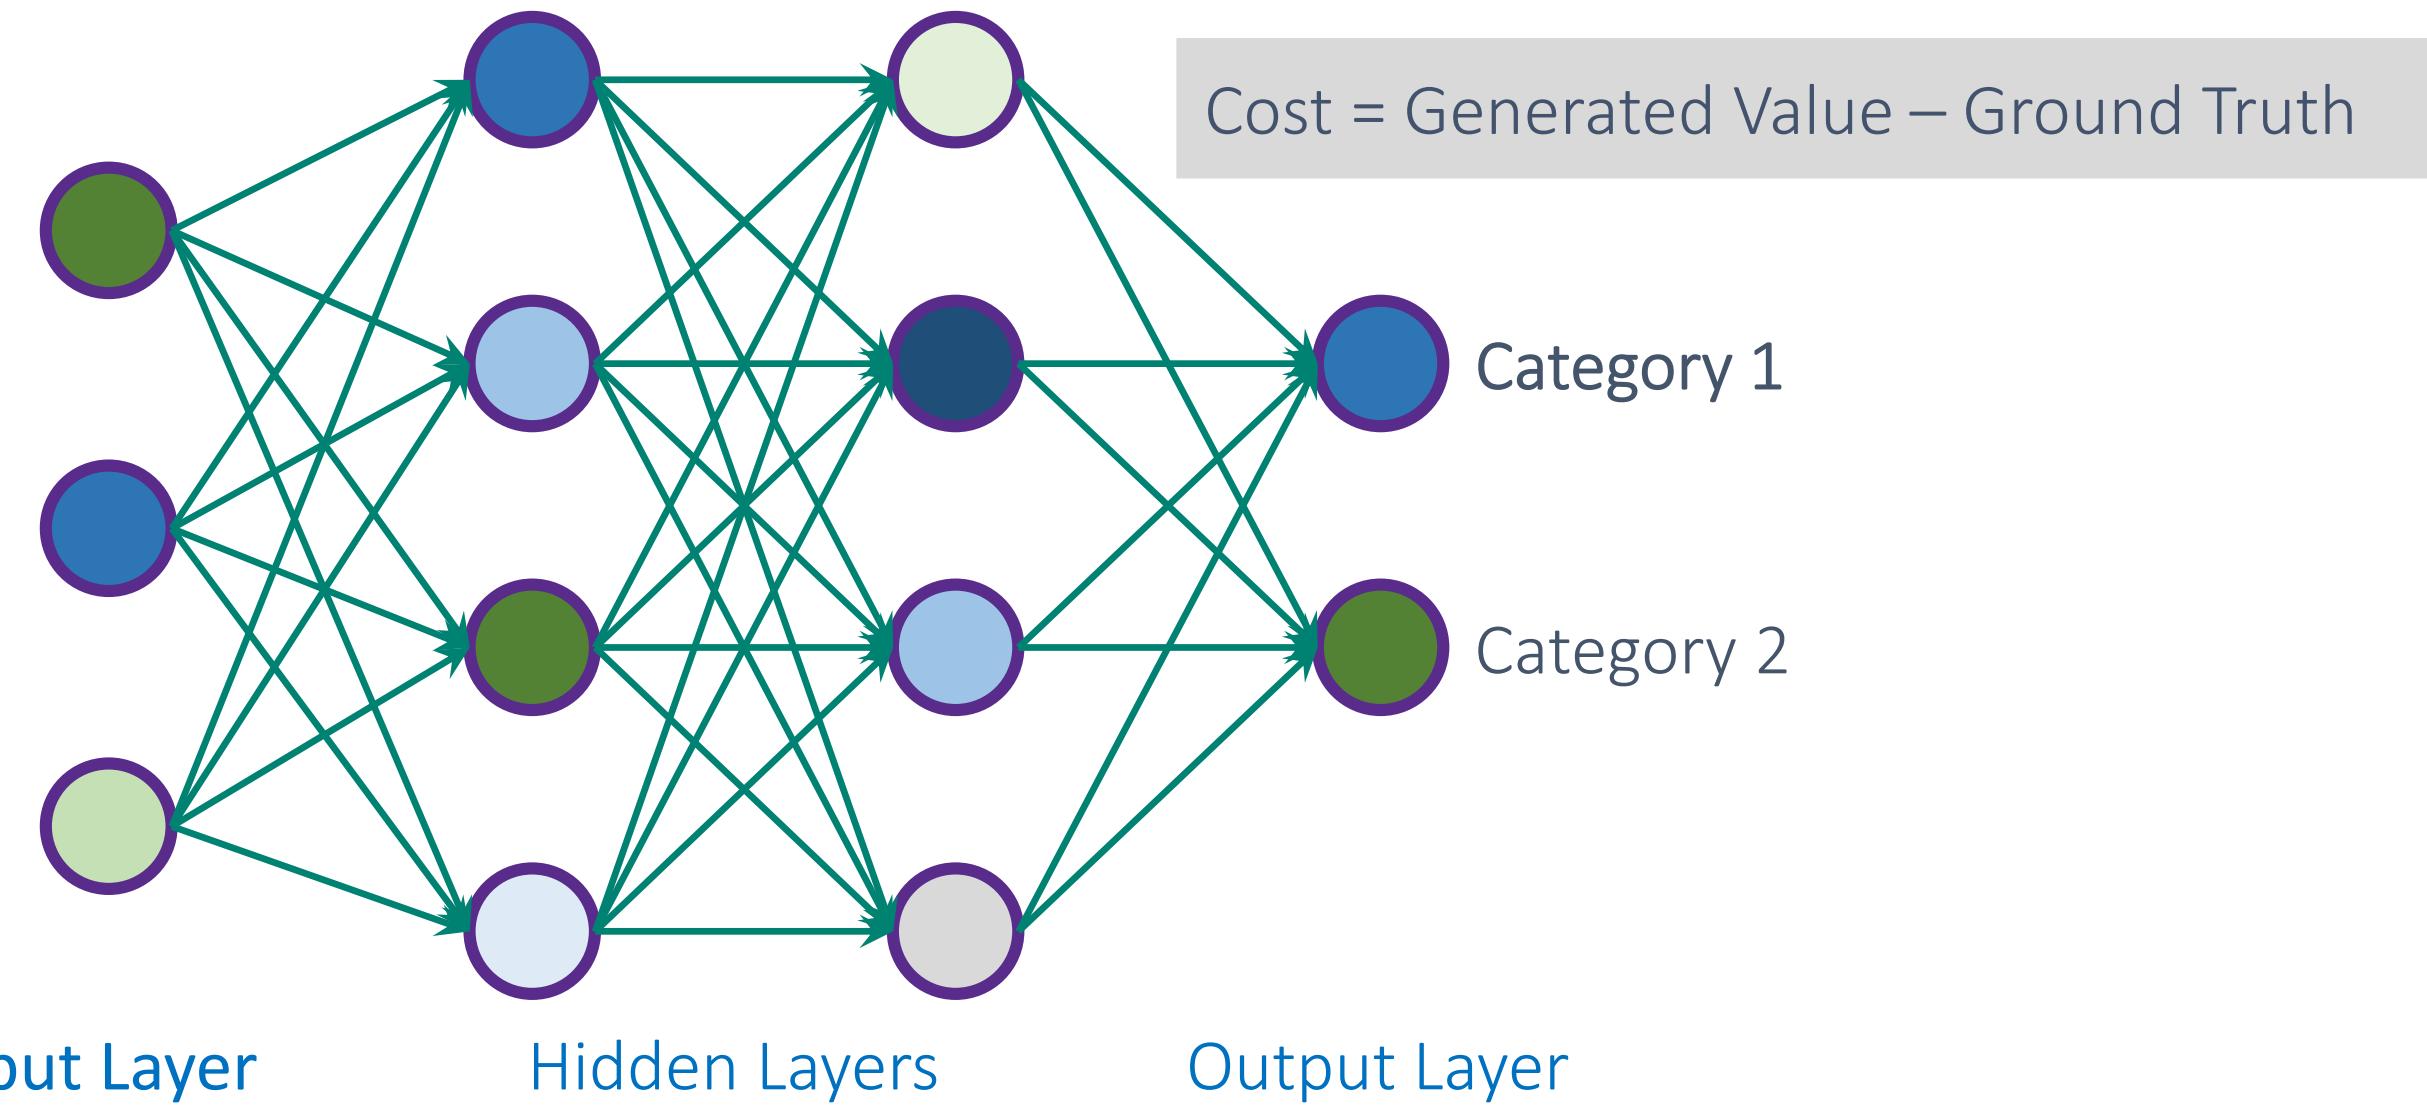
FORWARD PROPAGATION



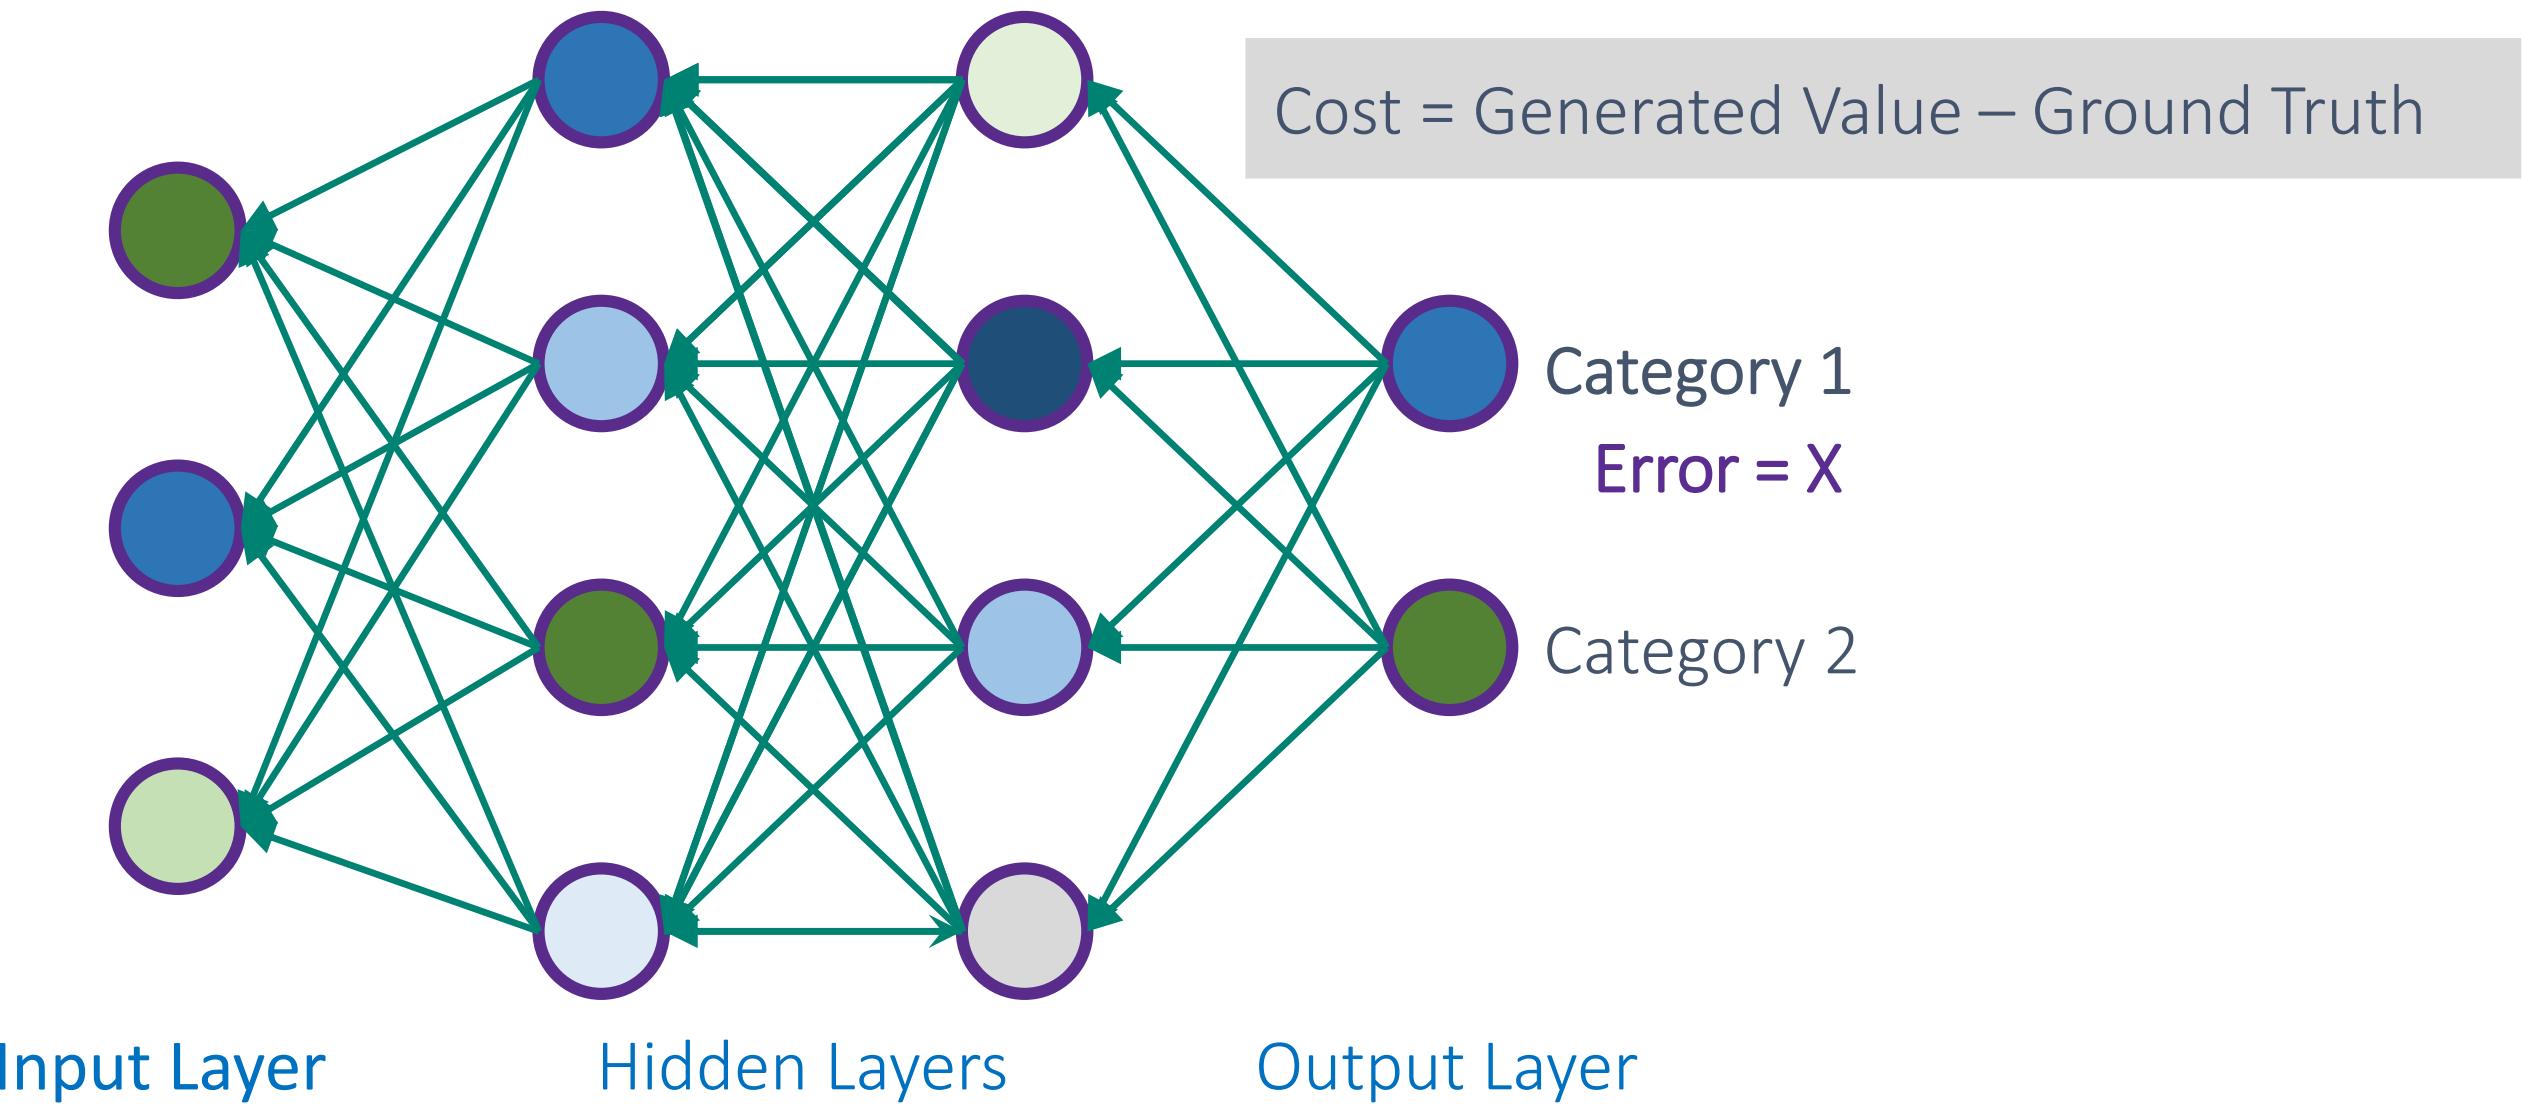
WEIGHTS AND BIAS



TRAINING



BACKPROPAGATION

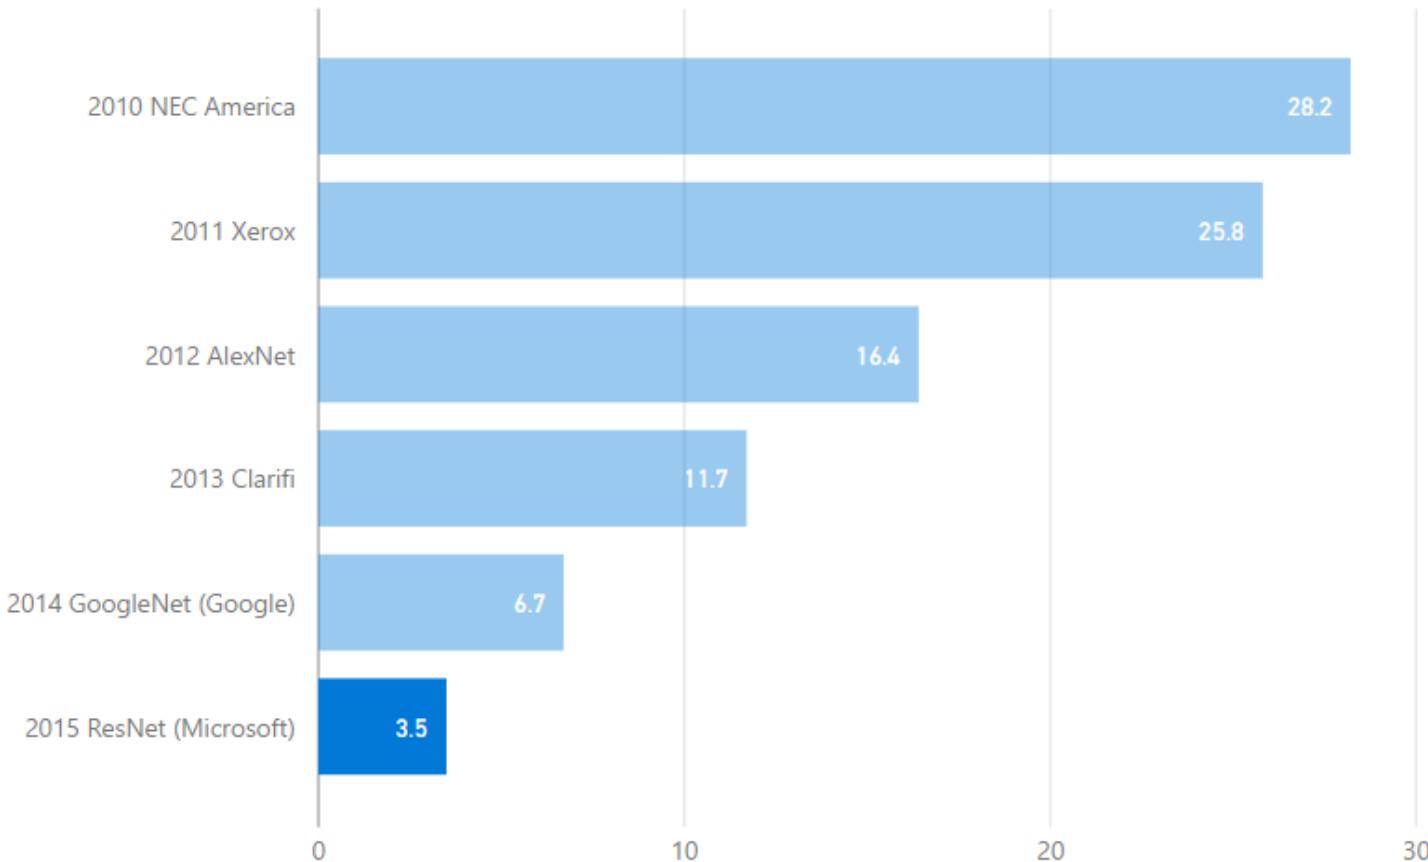


LAB 2: MULTI-LAYER PERCEPTRON

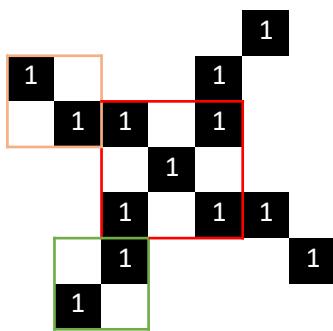
DL: CONVOLUTIONAL NETWORKS

IMAGENET CHALLENGE

% de Error por Año y Equipo



FILTERING



FILTERING

...looking for common characteristics

- Filter 1 Filter 2 Filter 3

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

1	-1	-1
-1	1	-1
-1	-1	1

-1	-1	1
-1	1	-1
1	-1	-1

1	-1	1
-1	1	-1
1	-1	1

1	1	1
1	1	1
1	1	1

-1	1	-1
1	1	1
-1	1	-1

1	1	-1
1	1	1
-1	1	1

1.0

0.11

0.55

CONVOLUTION

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

• Filter

1	-1	-1
-1	1	-1
-1	-1	1

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

CONVOLUTION

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

CONVOLUTION

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

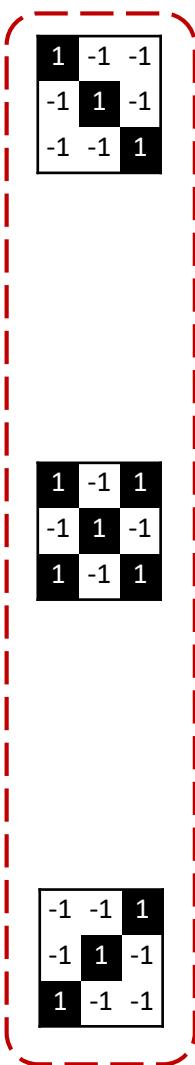


-1	-1	1
-1	1	-1
1	-1	-1

=

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



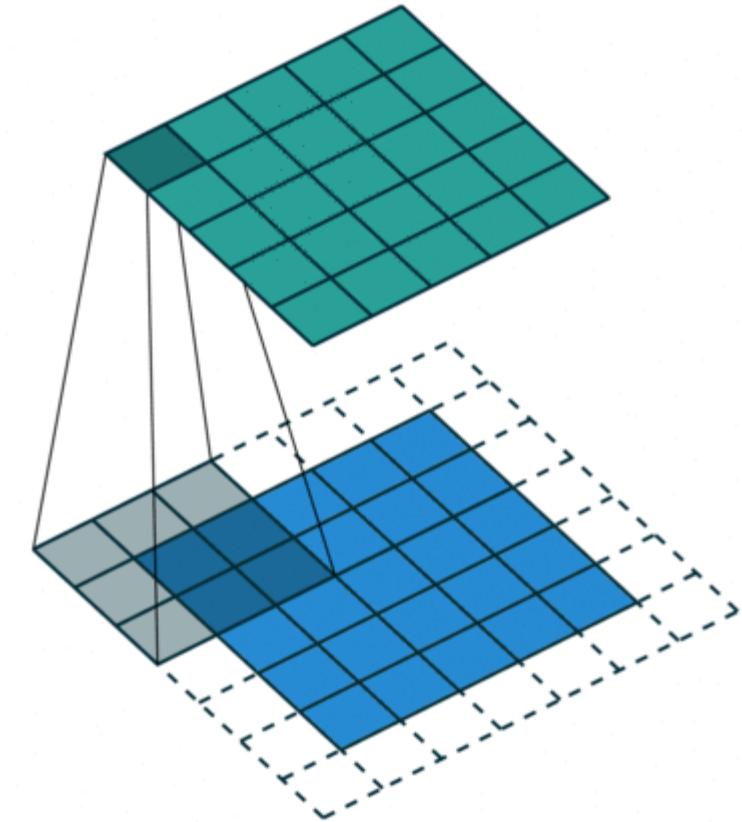
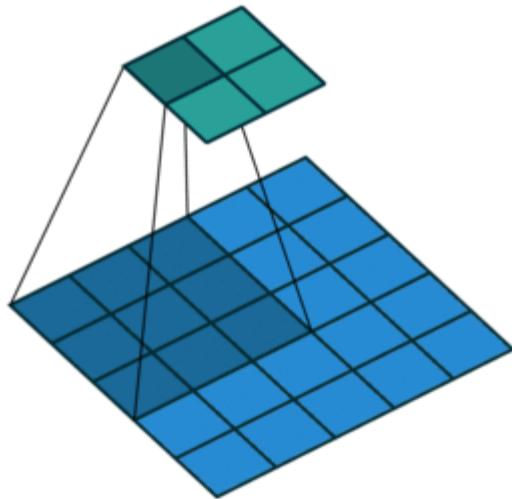
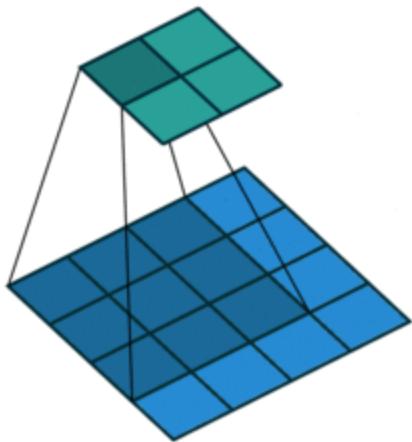
=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

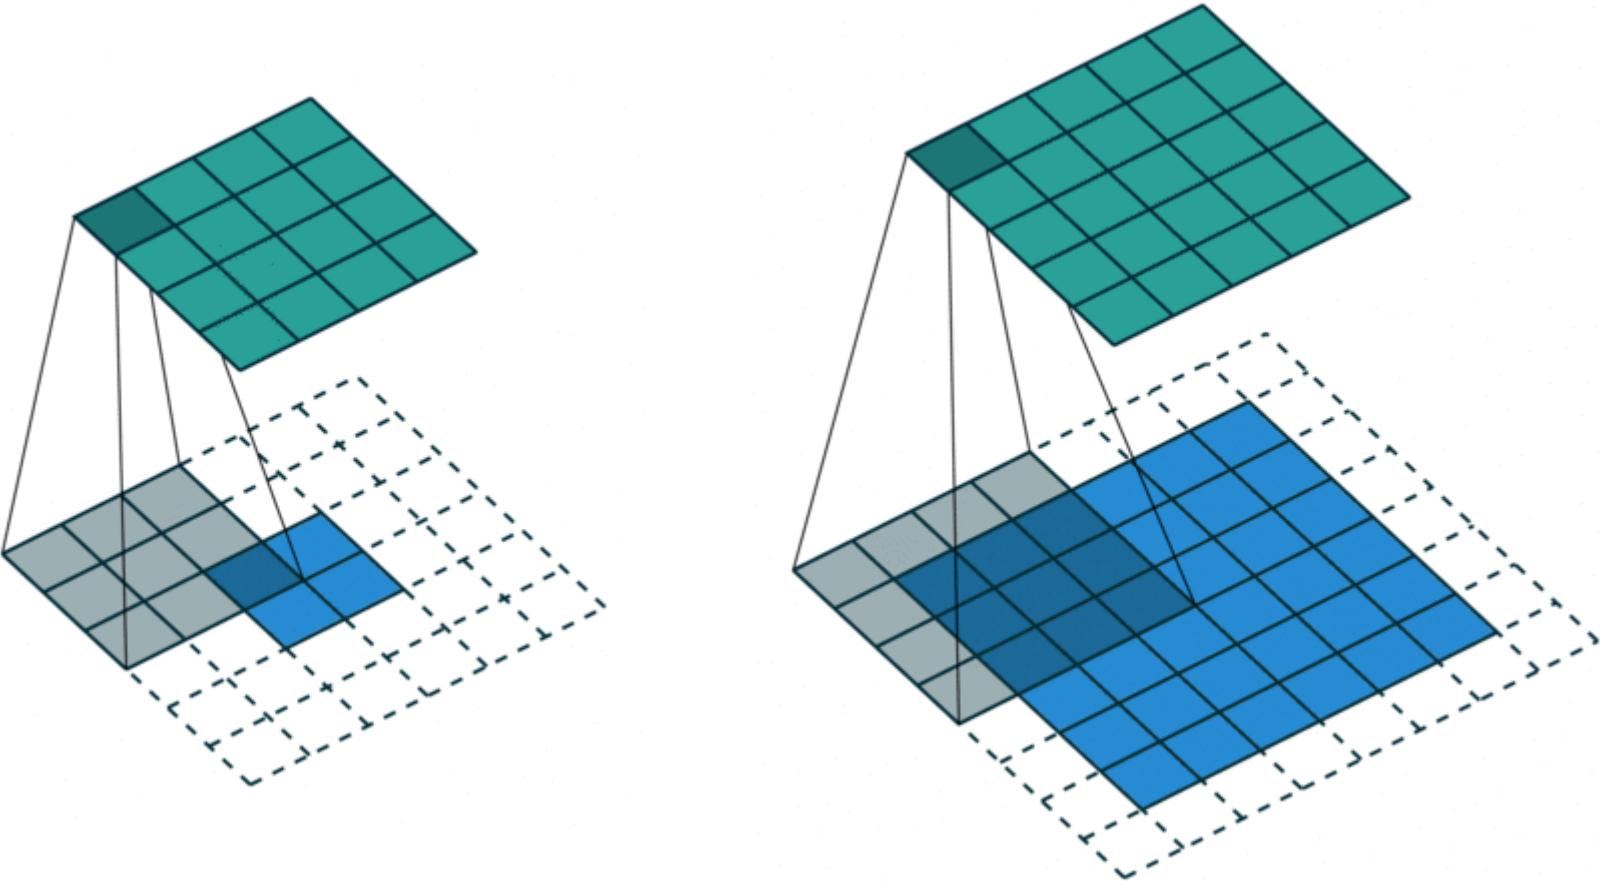
0.33	-0.55	-0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

STRIDES AND PADDING



DECONVOLUTION – TRANSPOSED CONVOLUTION



POOLING

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

POOLING

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

AVERAGE POOLING – STRIDE 1

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.39	0.22	0.17	0.28	0.11	0.06
0.22	0.45	0.22	0.00	0.00	0.11
0.17	0.22	0.22	0.00	0.00	0.28
0.28	0.00	0.00	0.22	0.22	0.17
0.11	0.00	0.00	0.22	0.45	0.22
0.06	0.11	0.28	0.17	0.22	0.39

MAX POOLING – STRIDE 1

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

1.00	1.00	0.33	0.55	0.55	0.33
1.00	1.00	1.00	0.33	0.11	0.55
0.33	1.00	1.00	0.55	0.33	0.55
0.55	0.33	0.55	1.00	1.00	0.33
0.55	0.11	0.33	1.00	1.00	1.00
0.33	0.55	0.55	0.33	1.00	1.00

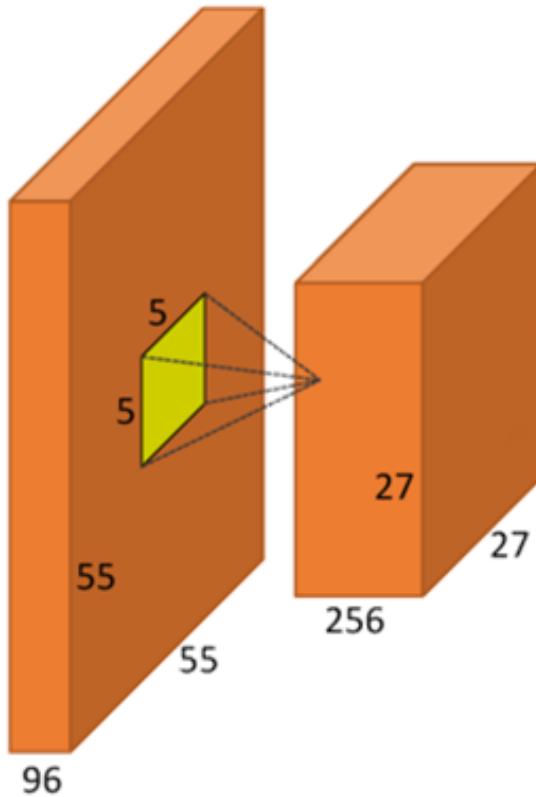
AVERAGE POOLING – STRIDE 2

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

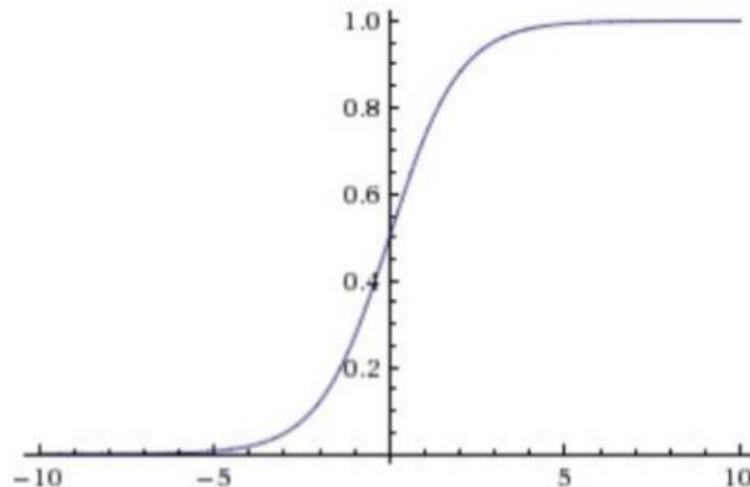
WHY POOLING?

ACTIVATION FUNCTIONS?



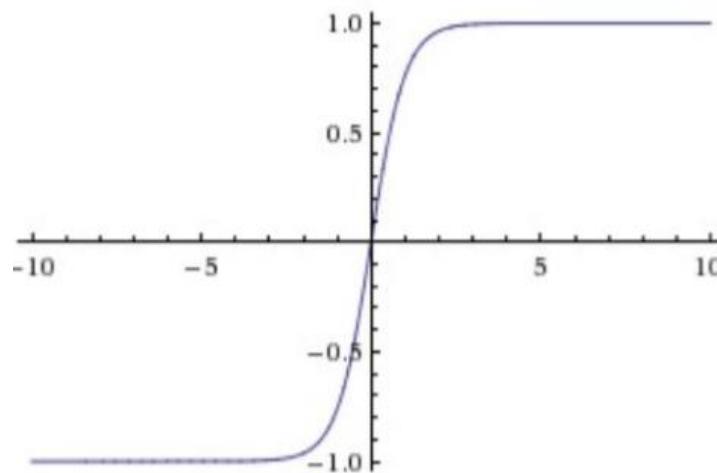
ACTIVATION FUNCTIONS?

LOGISTIC FUNCTION (“sigmoid”):



- Simple
- Drawbacks:
 - Saturates
 - Non-centered in zero
- We don't use it anymore, but we thank her for the services provided to the cause :)

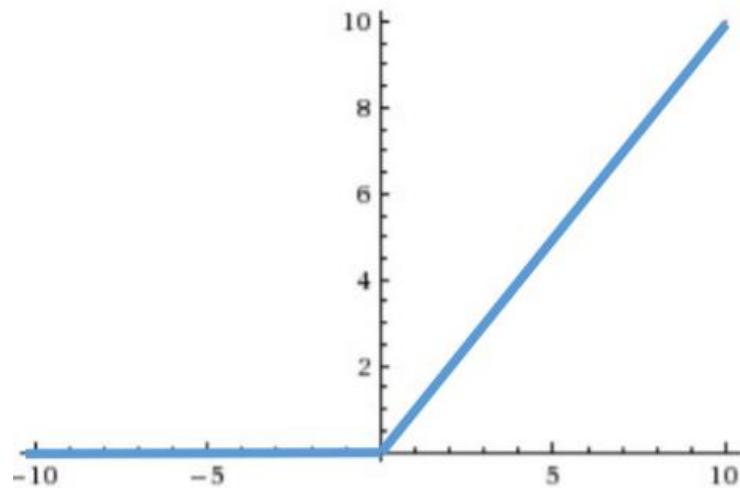
ACTIVATION FUNCTIONS?



HYPERBOLIC TANGENT (\tanh):

- Solves the issue with not being centered in zero
- Still ‘killing gradients’ due to saturating behaviour
- It is always better than the logistic function

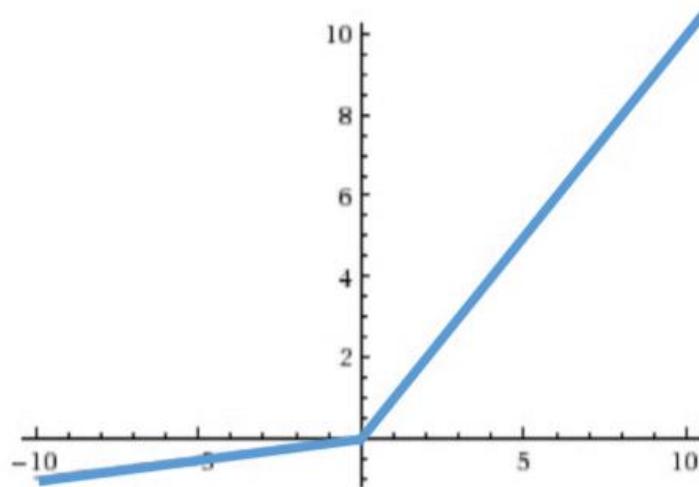
ACTIVATION FUNCTIONS?



RECTIFIED LINEAR UNITS (ReLU):

- Optimization for the SGD convergence
- More simple from a computational point of view
- Fragile

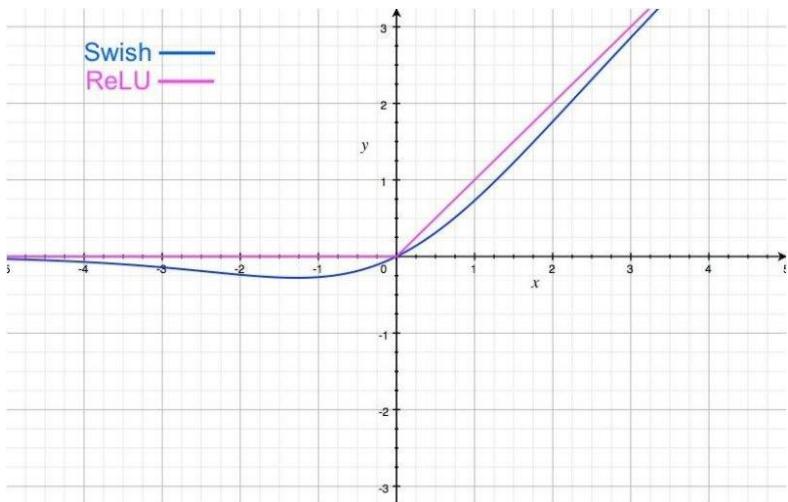
ACTIVATION FUNCTIONS?



LEAKY ReLU:

- More stable than traditional ReLU
- Particular case of MaxOut

ACTIVATION FUNCTIONS?



SWISH:

- Self-Gating activation function
- Better performance than ReLU

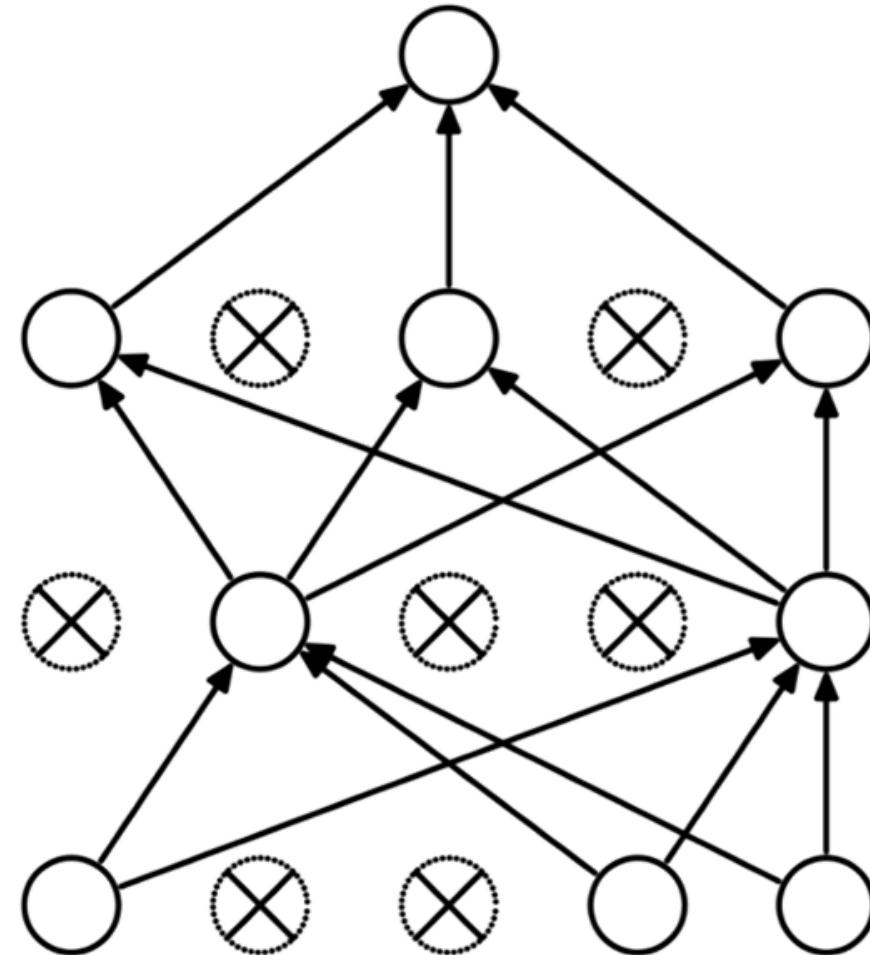
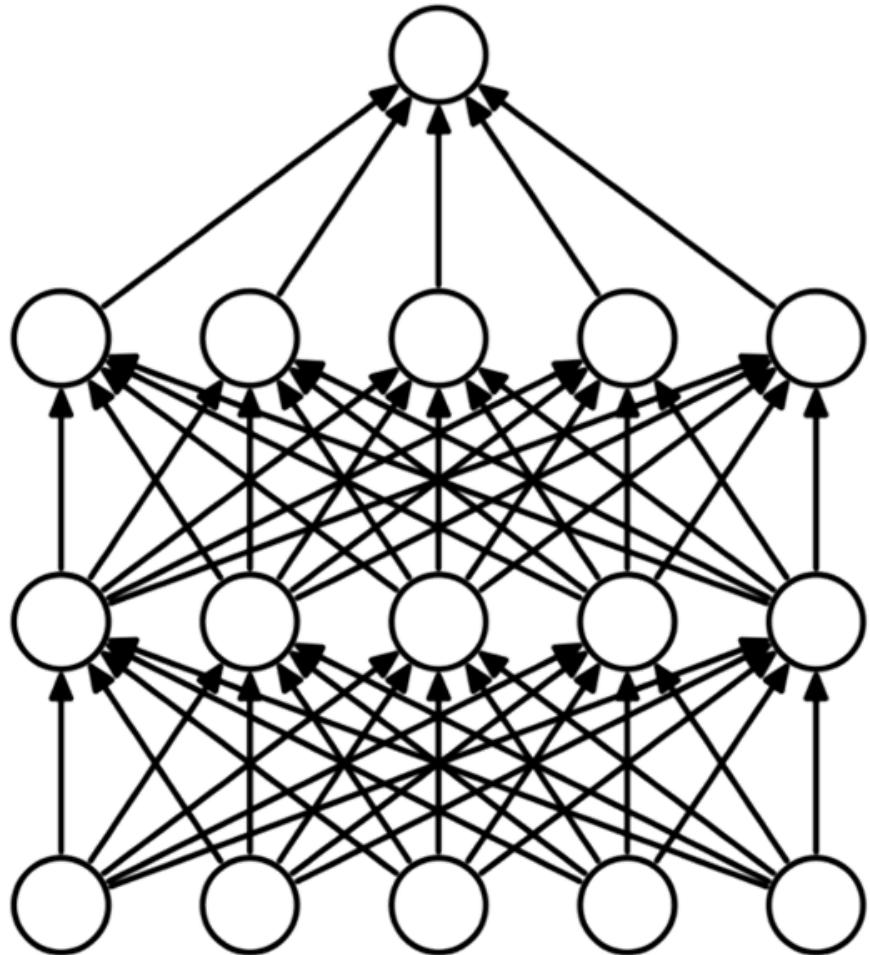
RECTIFIED LINEAR UNITS (ReLU)

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

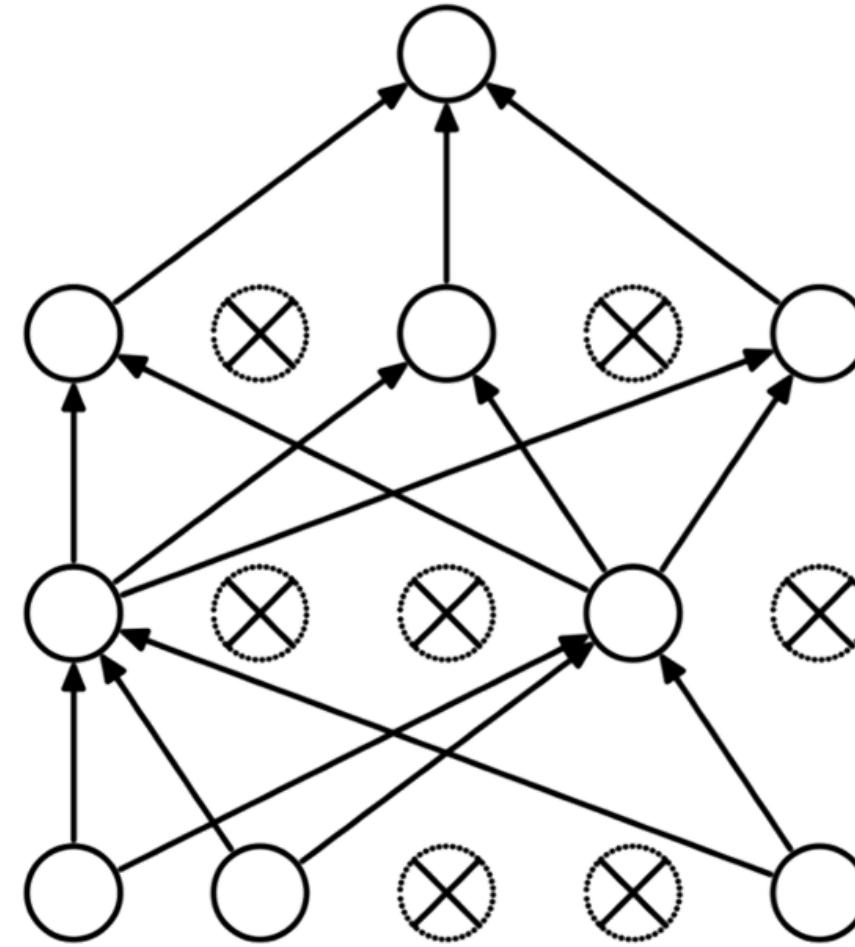
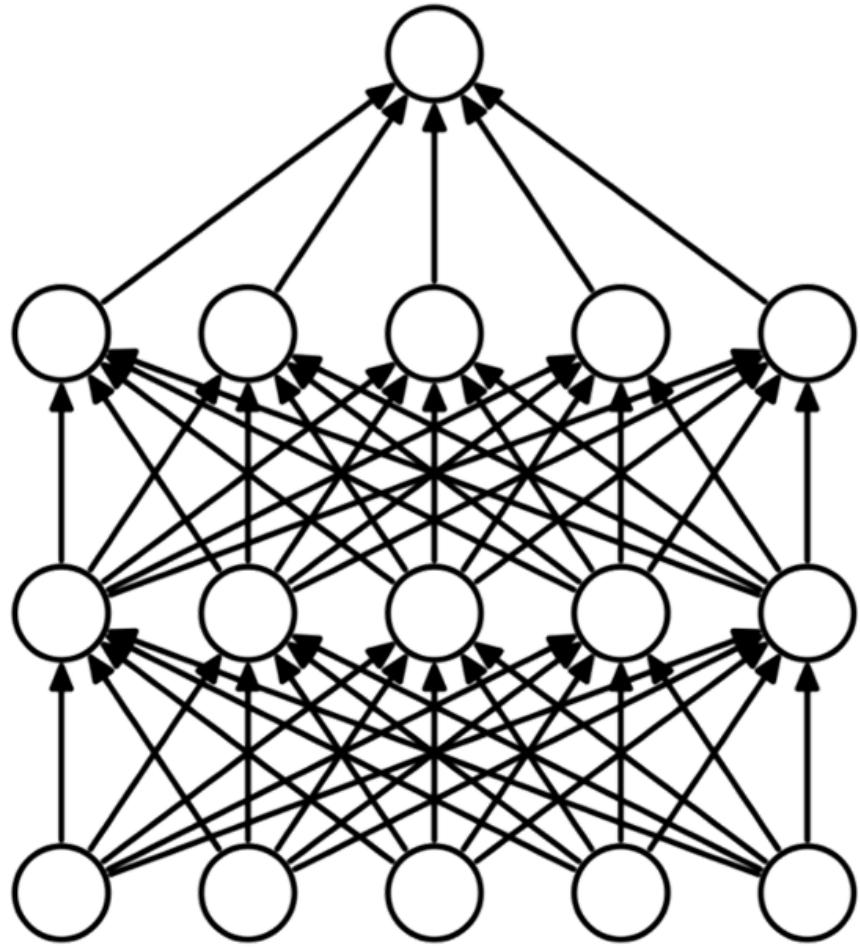


0.77	0.00	0.11	0.33	0.55	0.00	0.33
0.00	1.00	0.00	0.33	0.00	0.11	0.00
0.11	0.00	1.00	0.00	0.11	0.00	0.55
0.33	0.33	0.00	0.55	0.00	0.33	0.33
0.55	0.00	0.11	0.00	1.00	0.00	0.11
0.00	0.11	0.00	0.33	0.00	1.00	0.00
0.33	0.00	0.55	0.33	0.11	0.00	0.77

DROPOUT



DROPOUT



WHAT IS MISSING?

- Filter 1 Filter 2 Filter 3

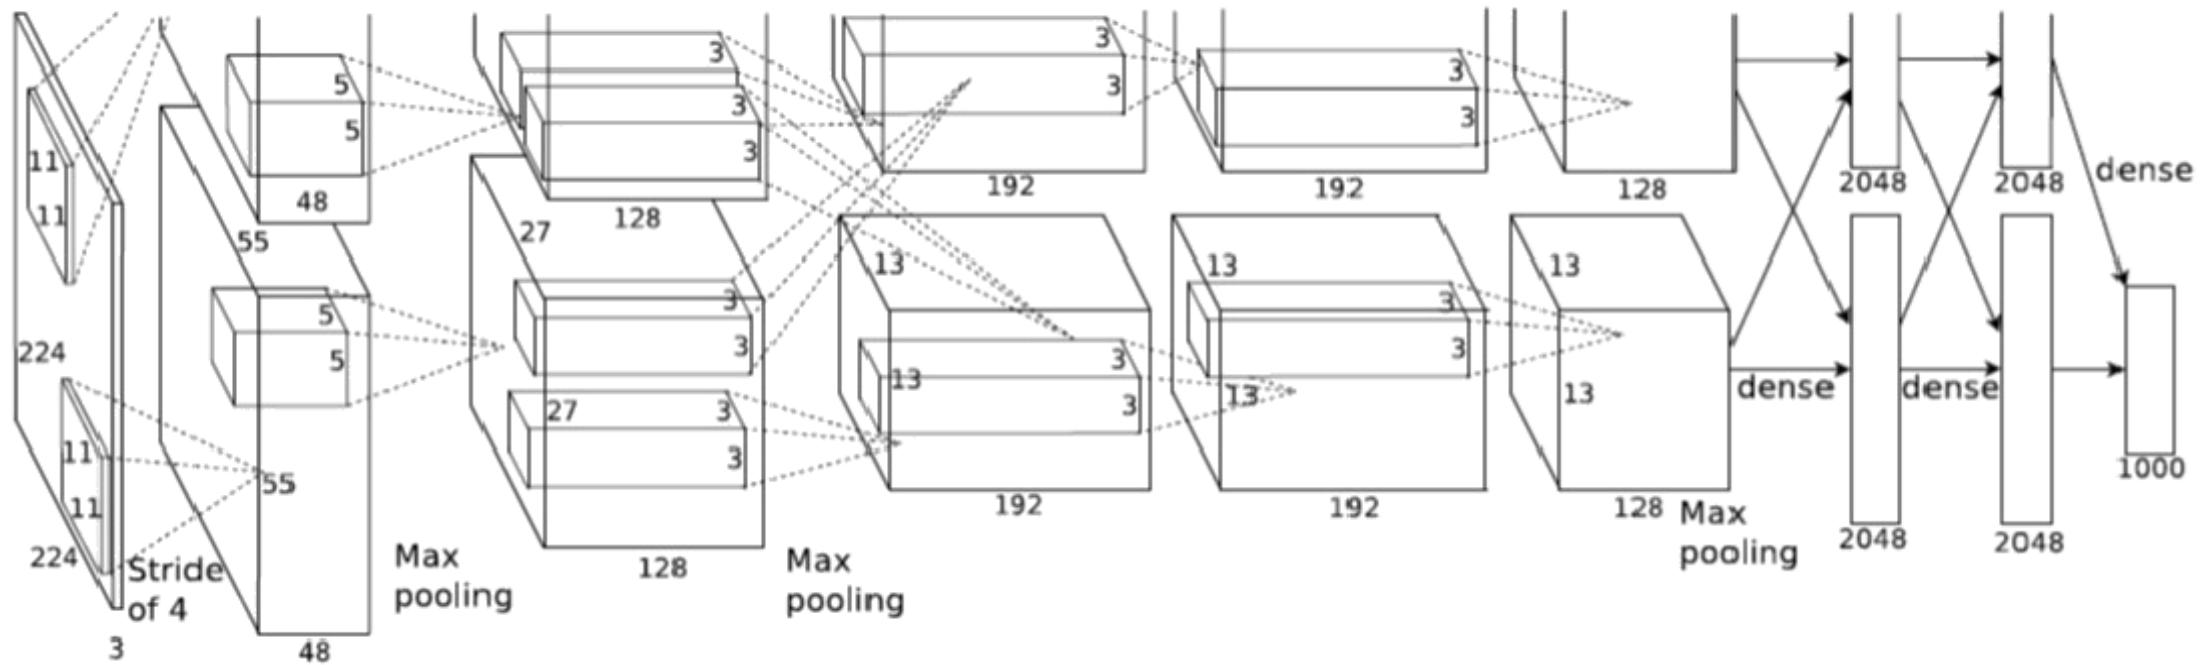
$$\begin{bmatrix} 1 & -1 & -1 \\ -1 & \boxed{1} & -1 \\ -1 & -1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} -1 & -1 & 1 \\ -1 & \boxed{1} & -1 \\ 1 & -1 & -1 \end{bmatrix}$$

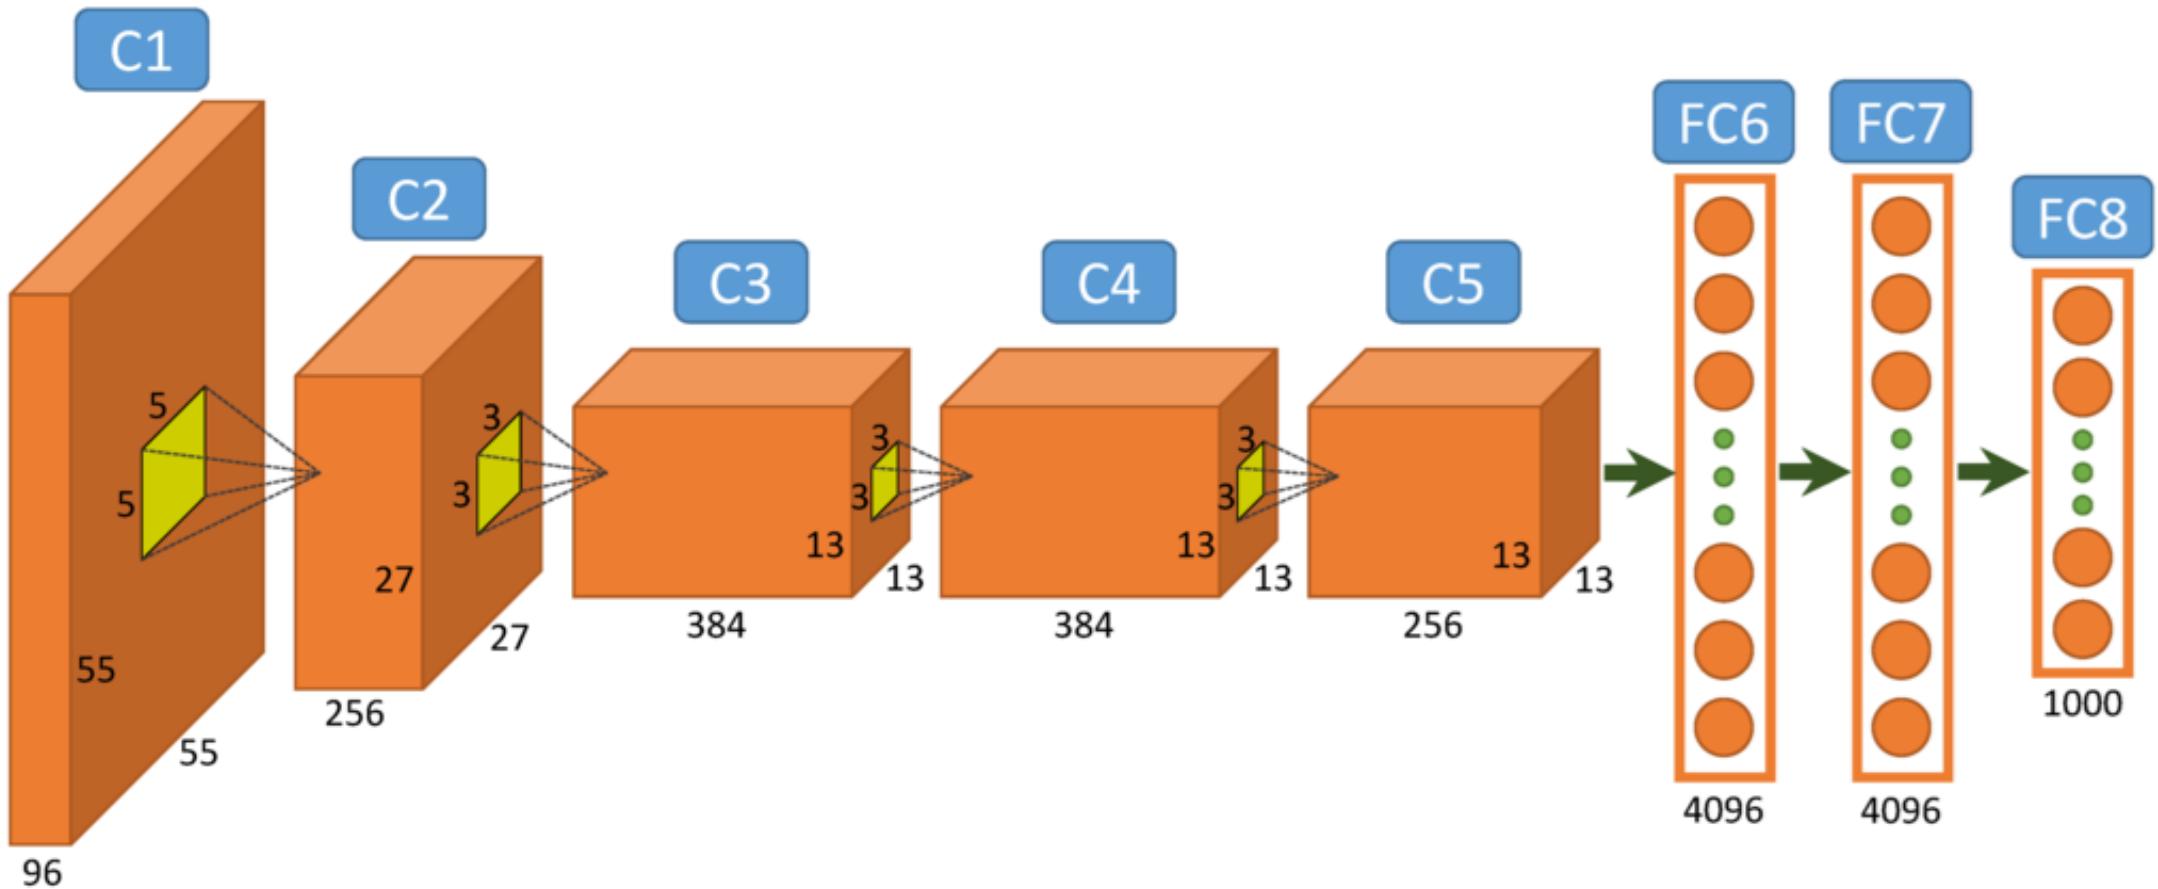
$$\begin{bmatrix} 1 & -1 & 1 \\ -1 & \boxed{1} & -1 \\ 1 & -1 & 1 \end{bmatrix}$$

¿De dónde han salido estos?

AlexNet



AlexNet



HOW A CNN SEES



Layer 1

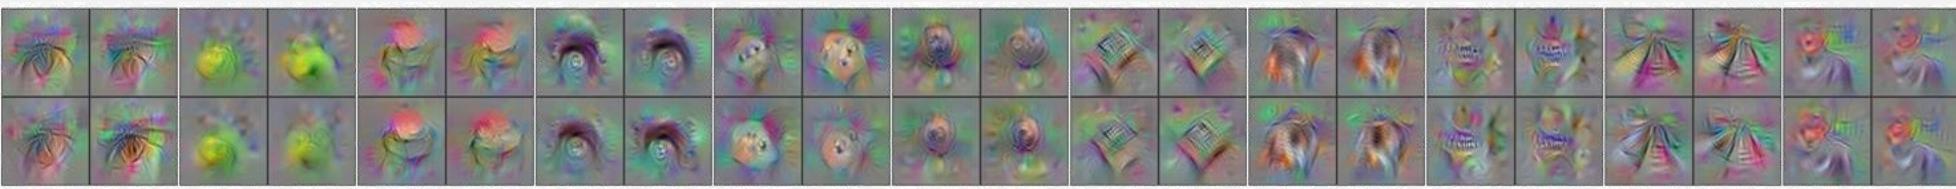
HOW A CNN SEES

Layer 2



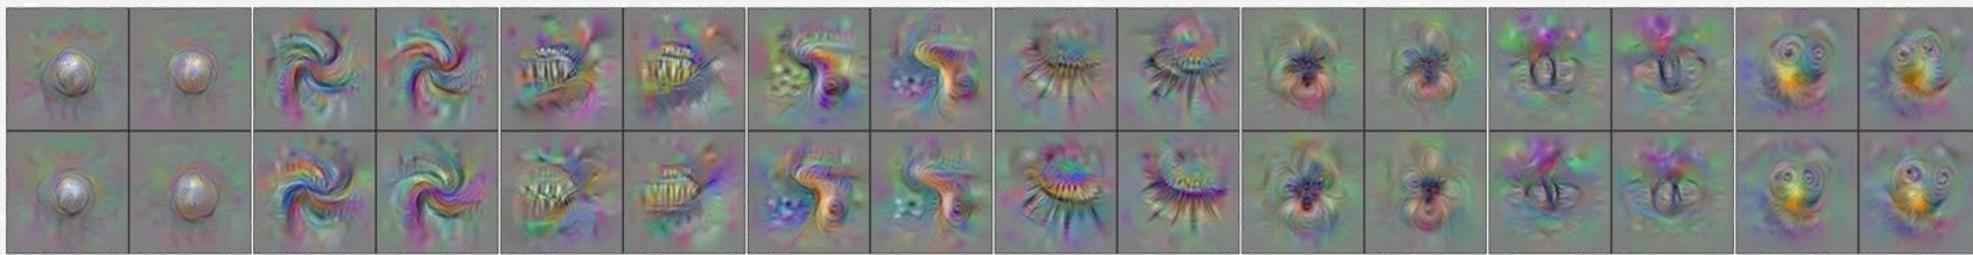
HOW A CNN SEES

Layer 3



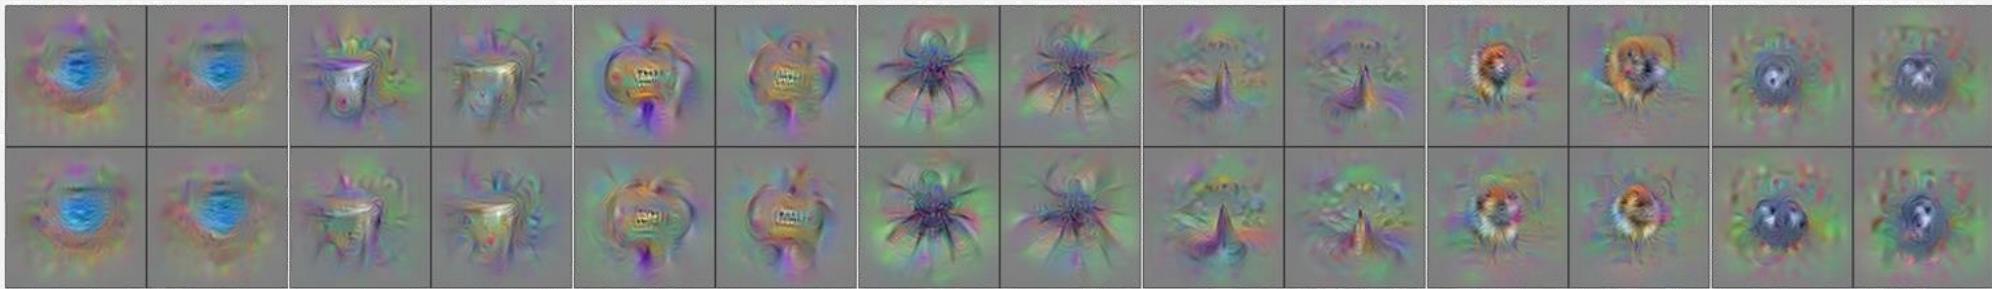
HOW A CNN SEES

Layer 4

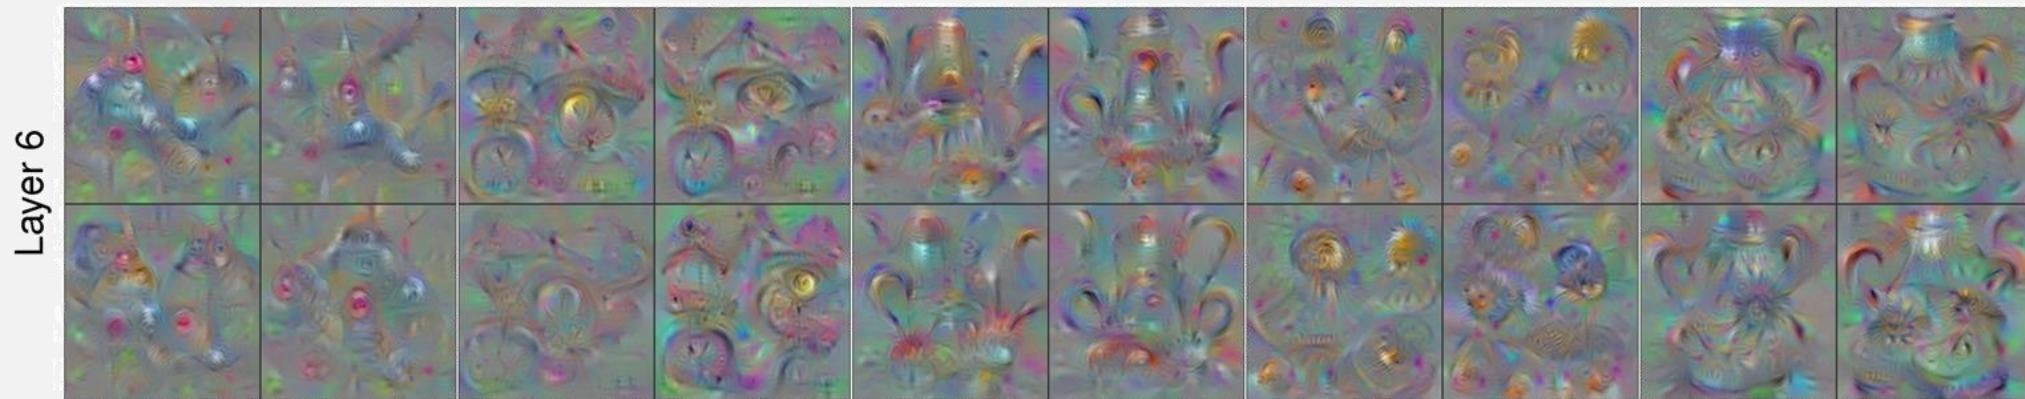


HOW A CNN SEES

Layer 5

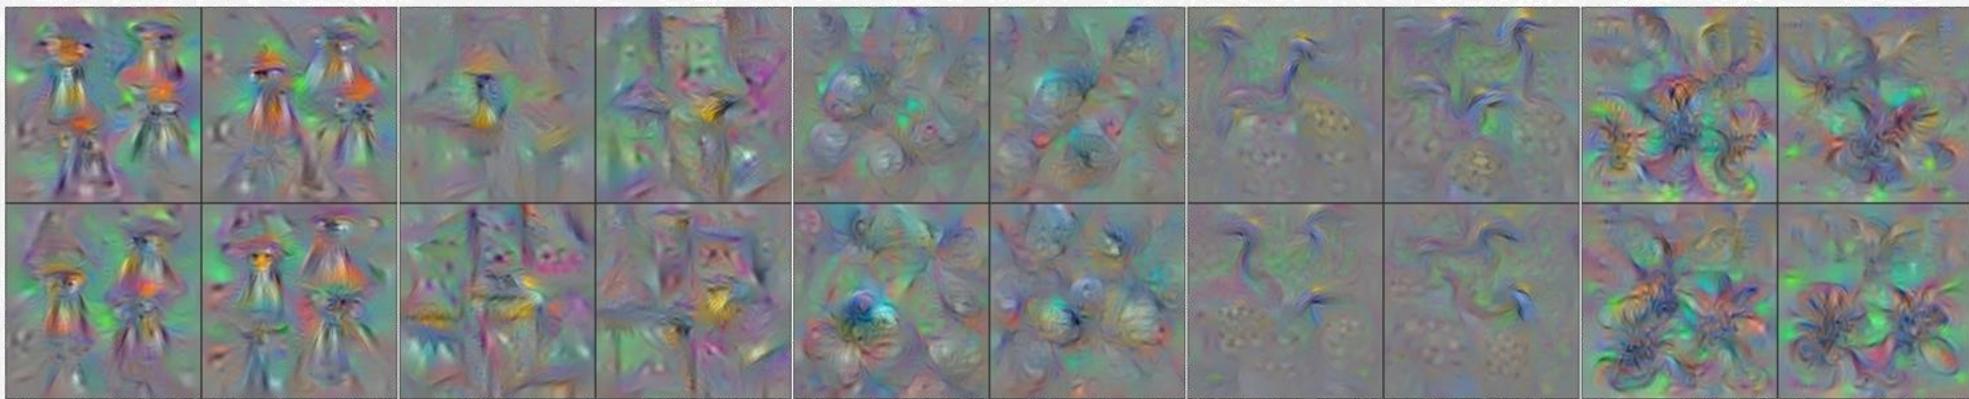


HOW A CNN SEES



HOW A CNN SEES

Layer 7



HOW A CNN SEES

Layer 8

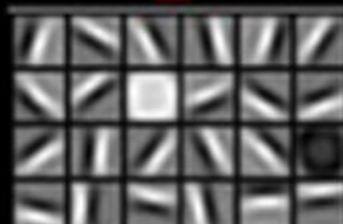
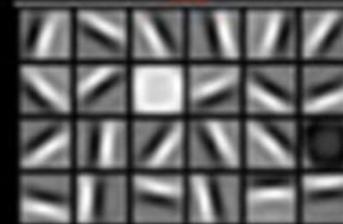
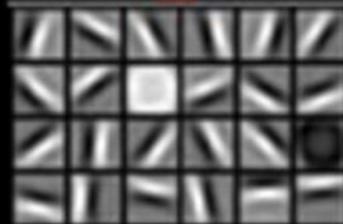
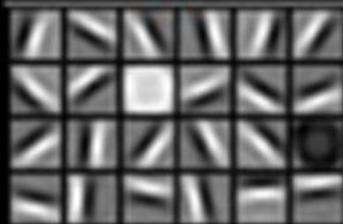
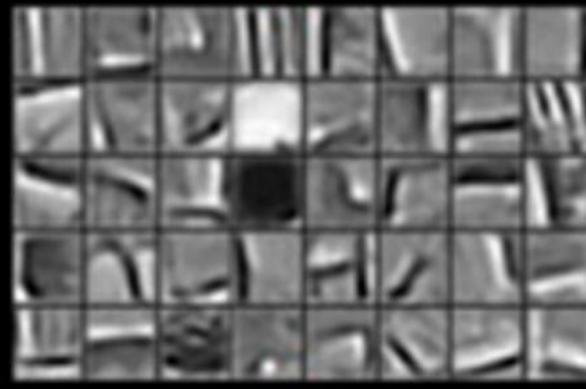
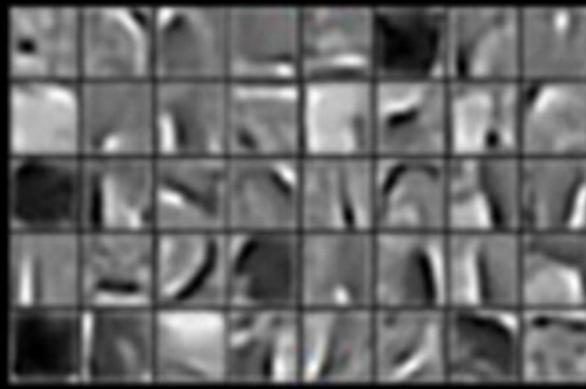
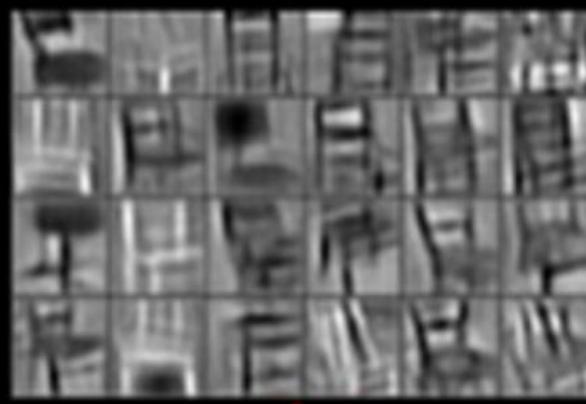


Faces

Cars

Elephants

Chairs

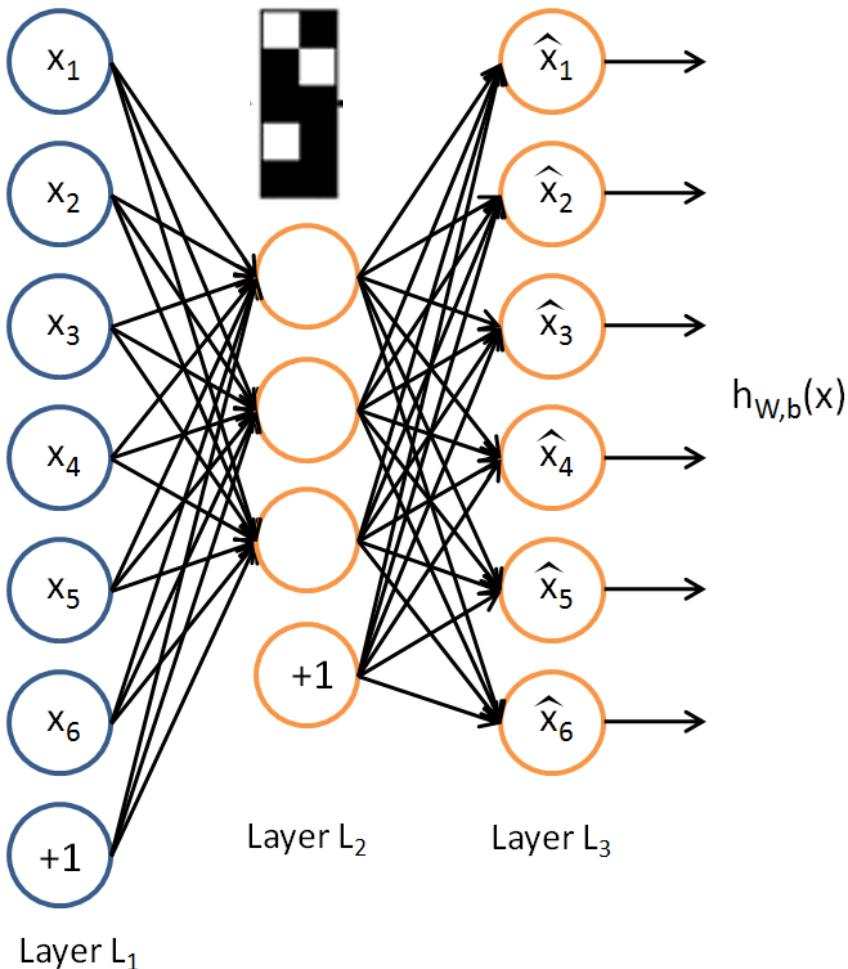


LAB 3: CONVOLUTIONAL NETWORKS



AUTOENCODERS

BASIC AUTOENCODER

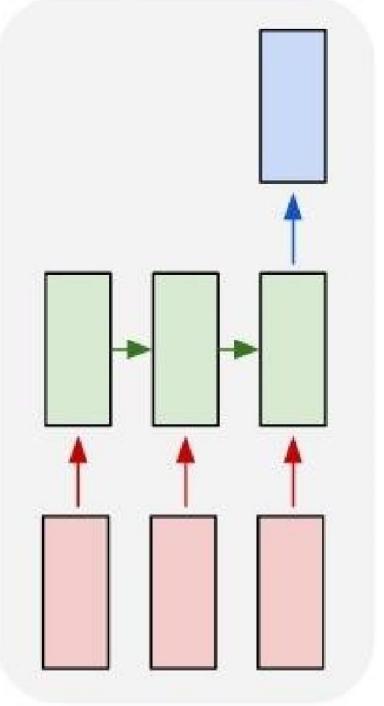


$$\hat{x} = h_{W,b} \approx x$$

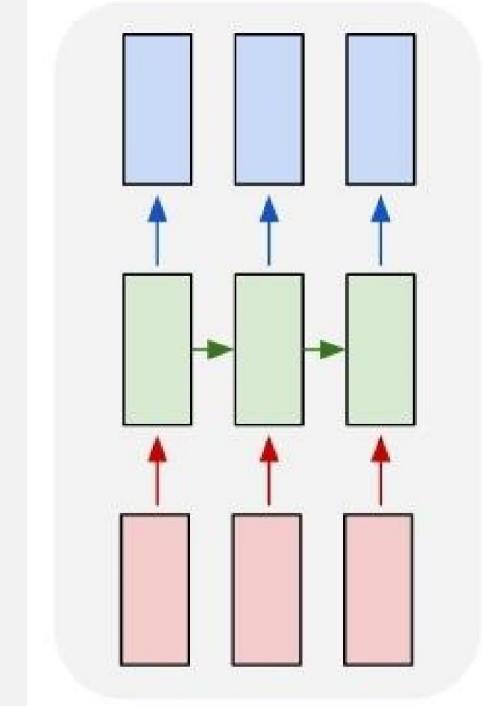
RECURRENT NEURAL NETWORKS

HANDLING SEQUENCES

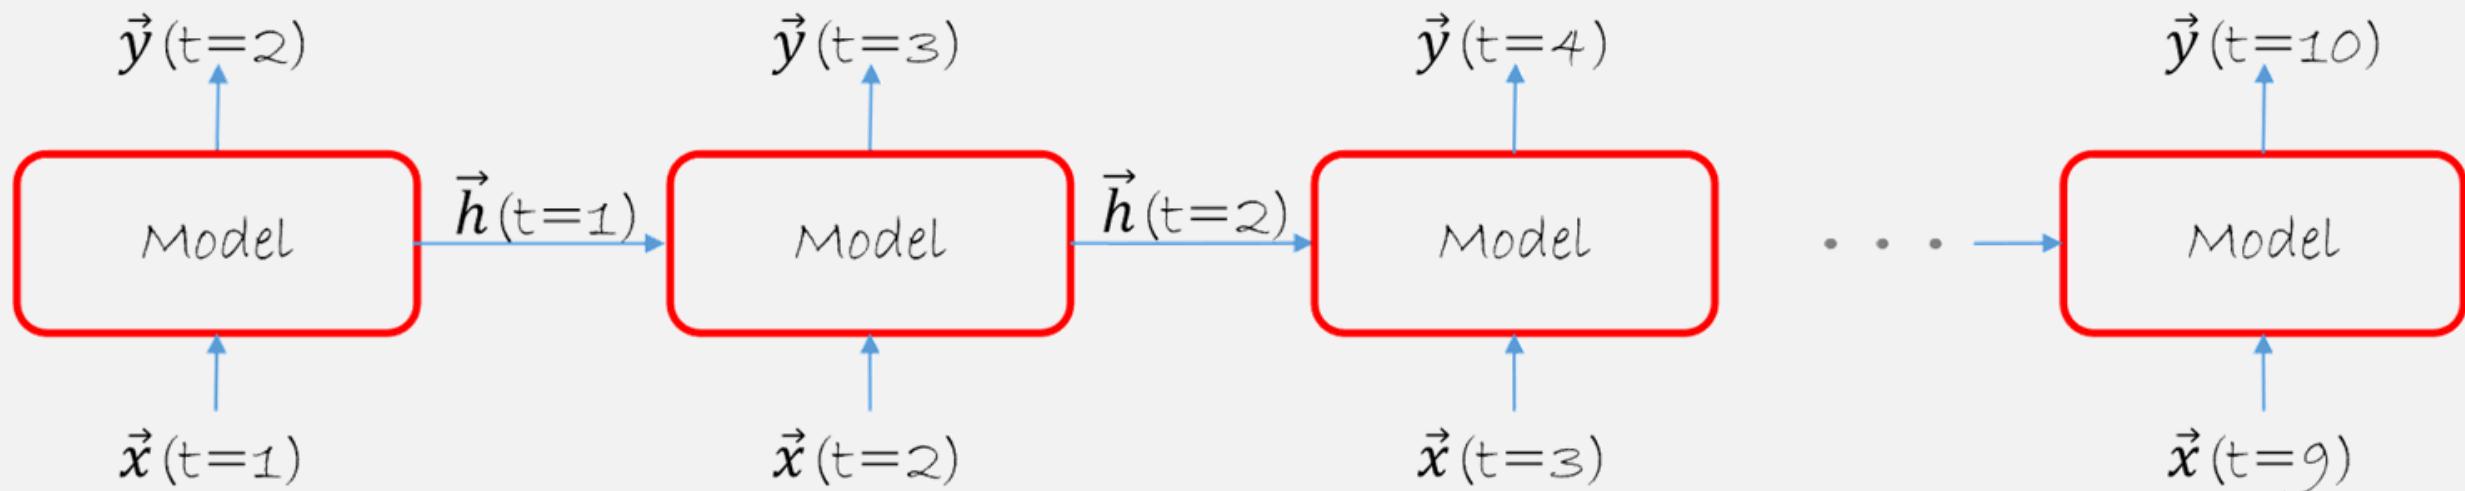
many to one



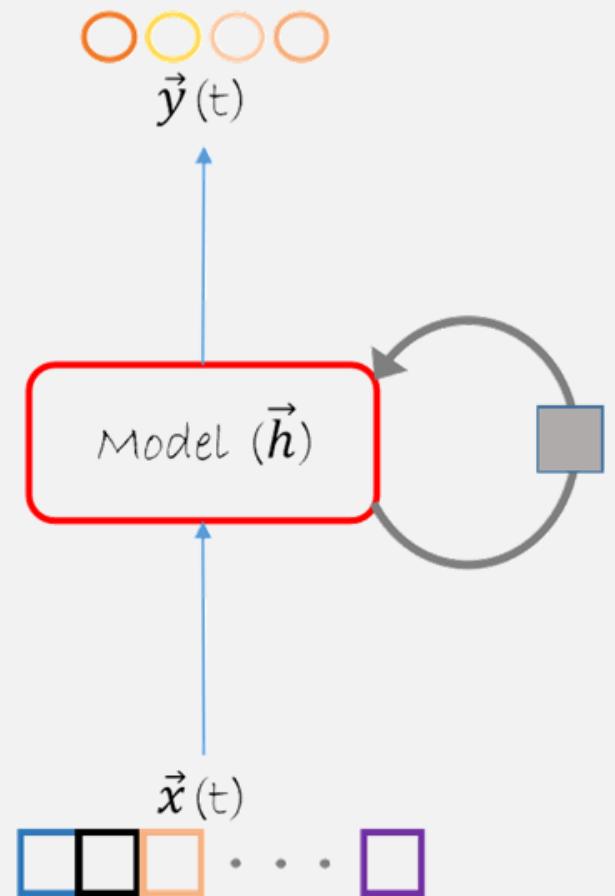
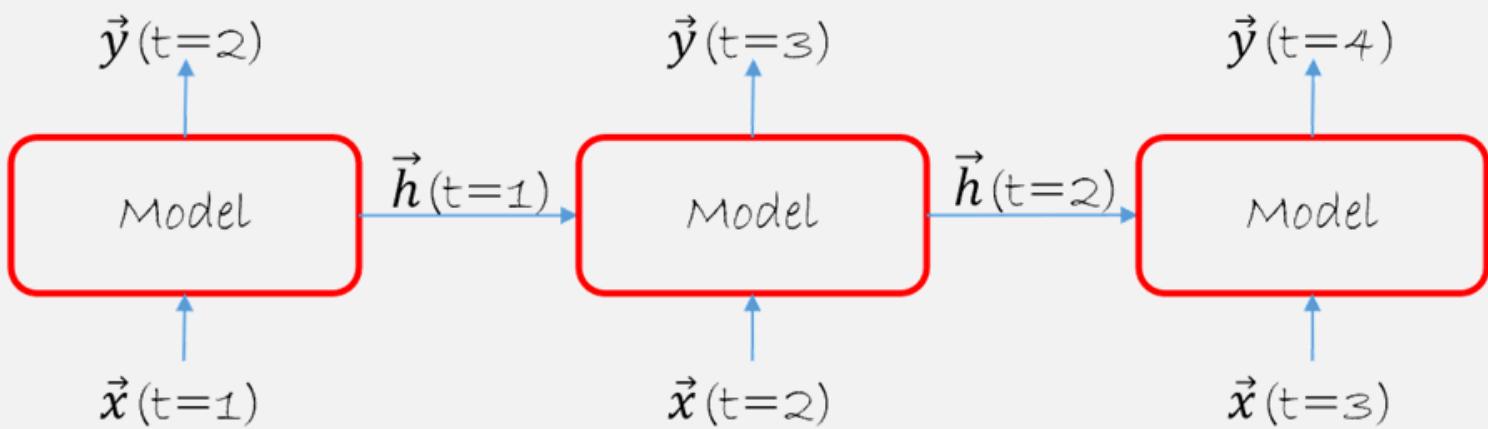
many to many

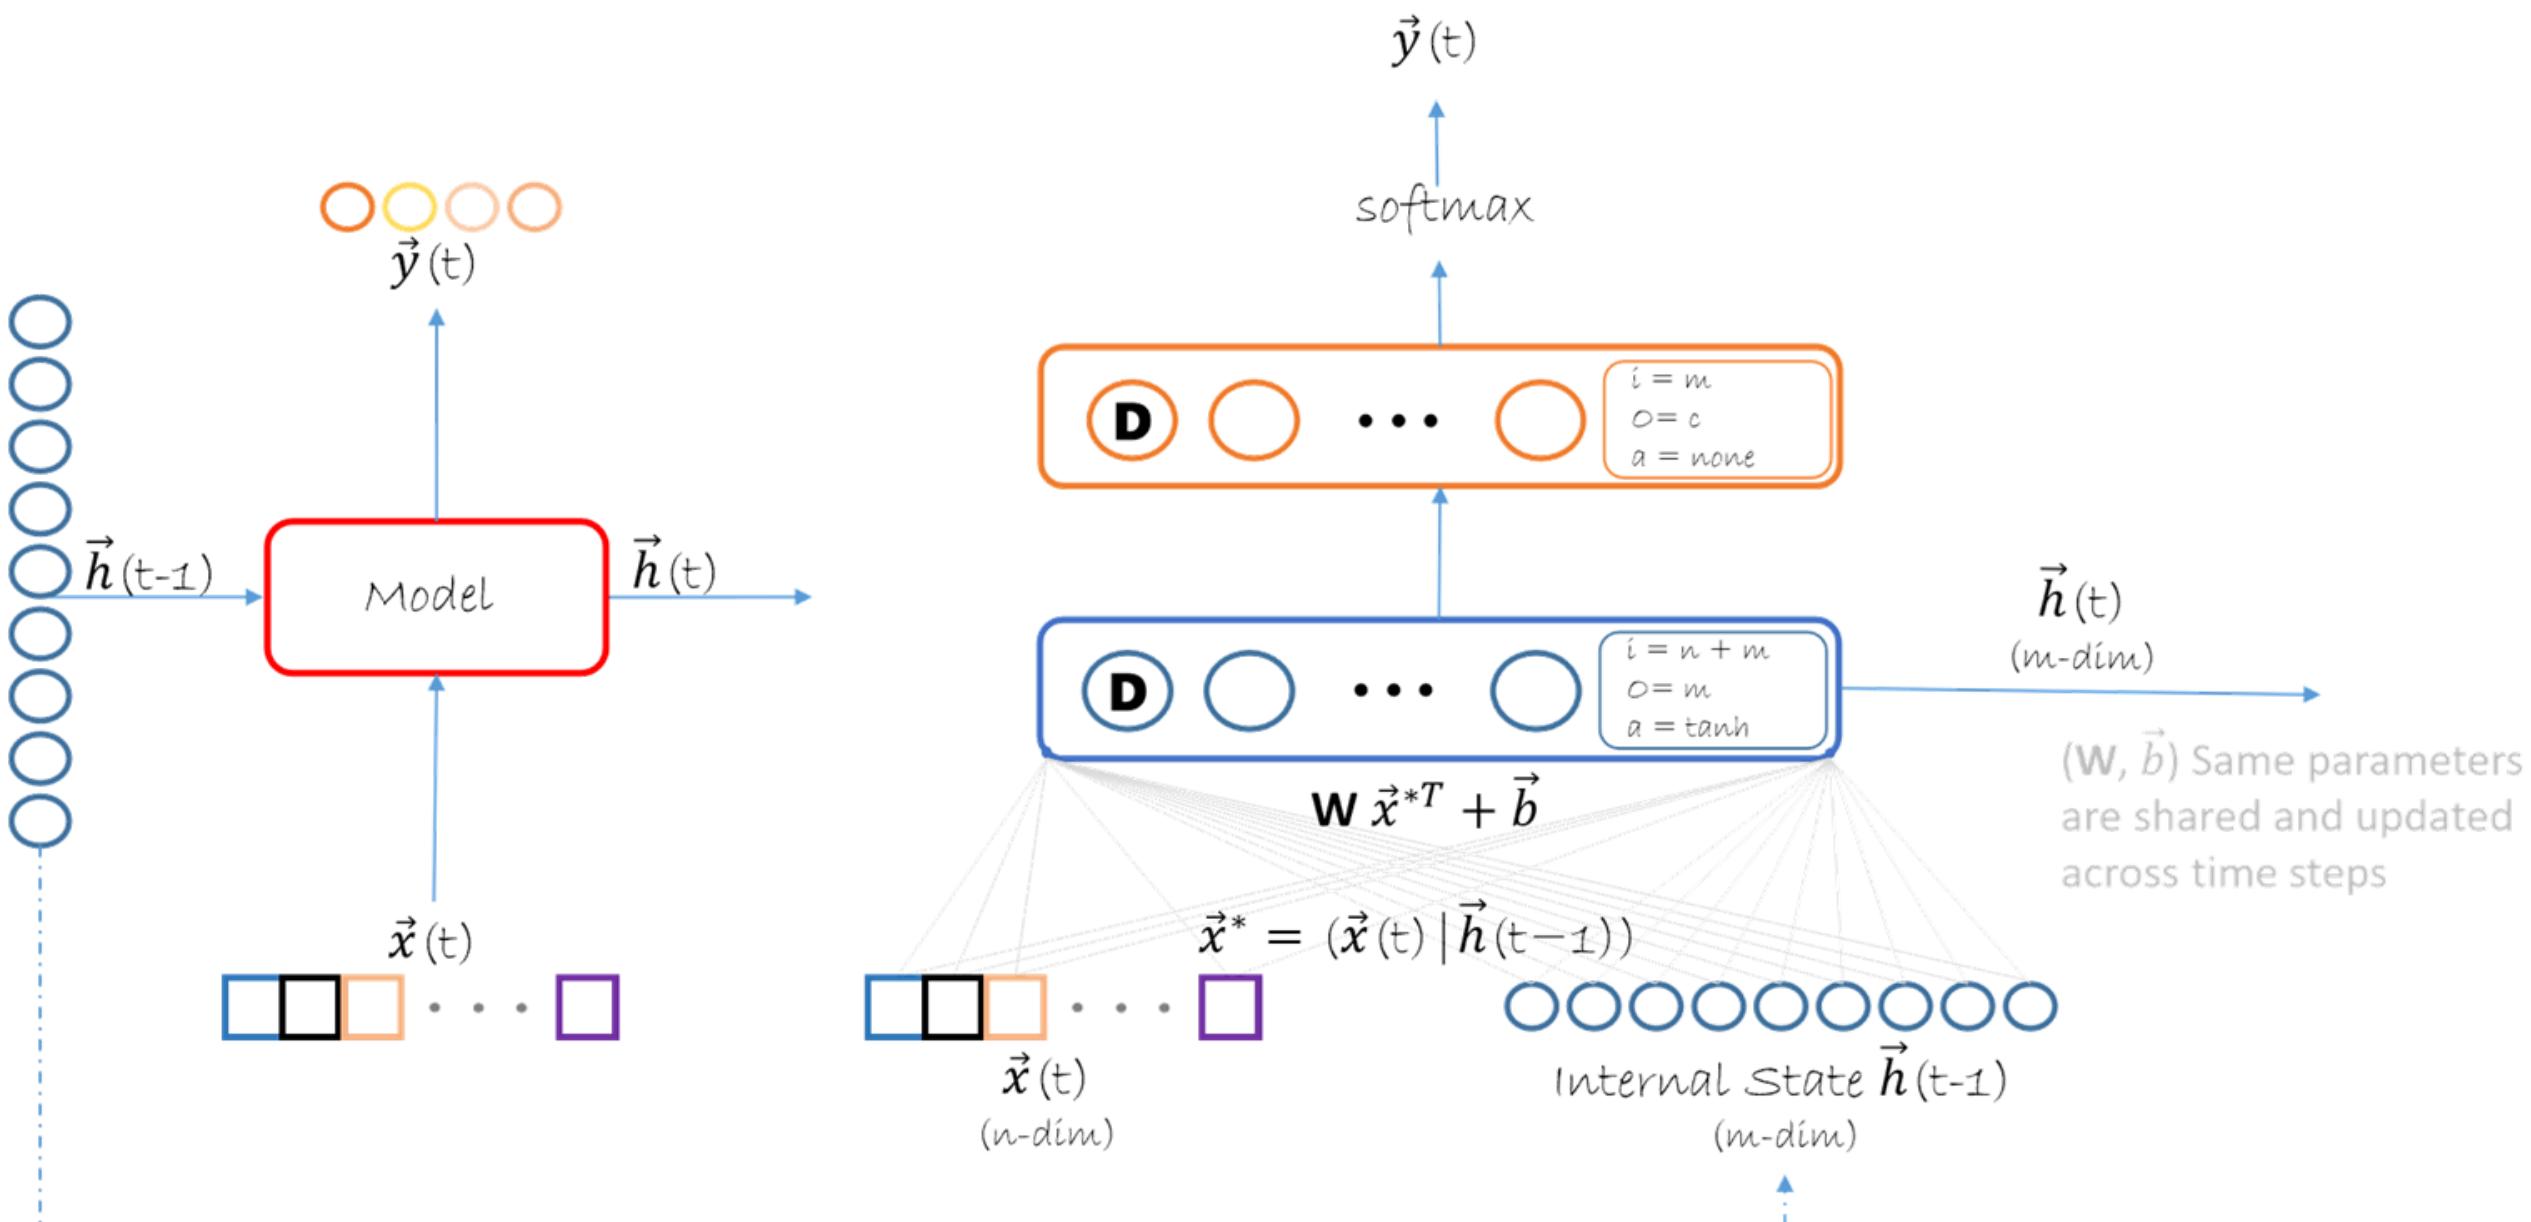


TRACKING THE HISTORY

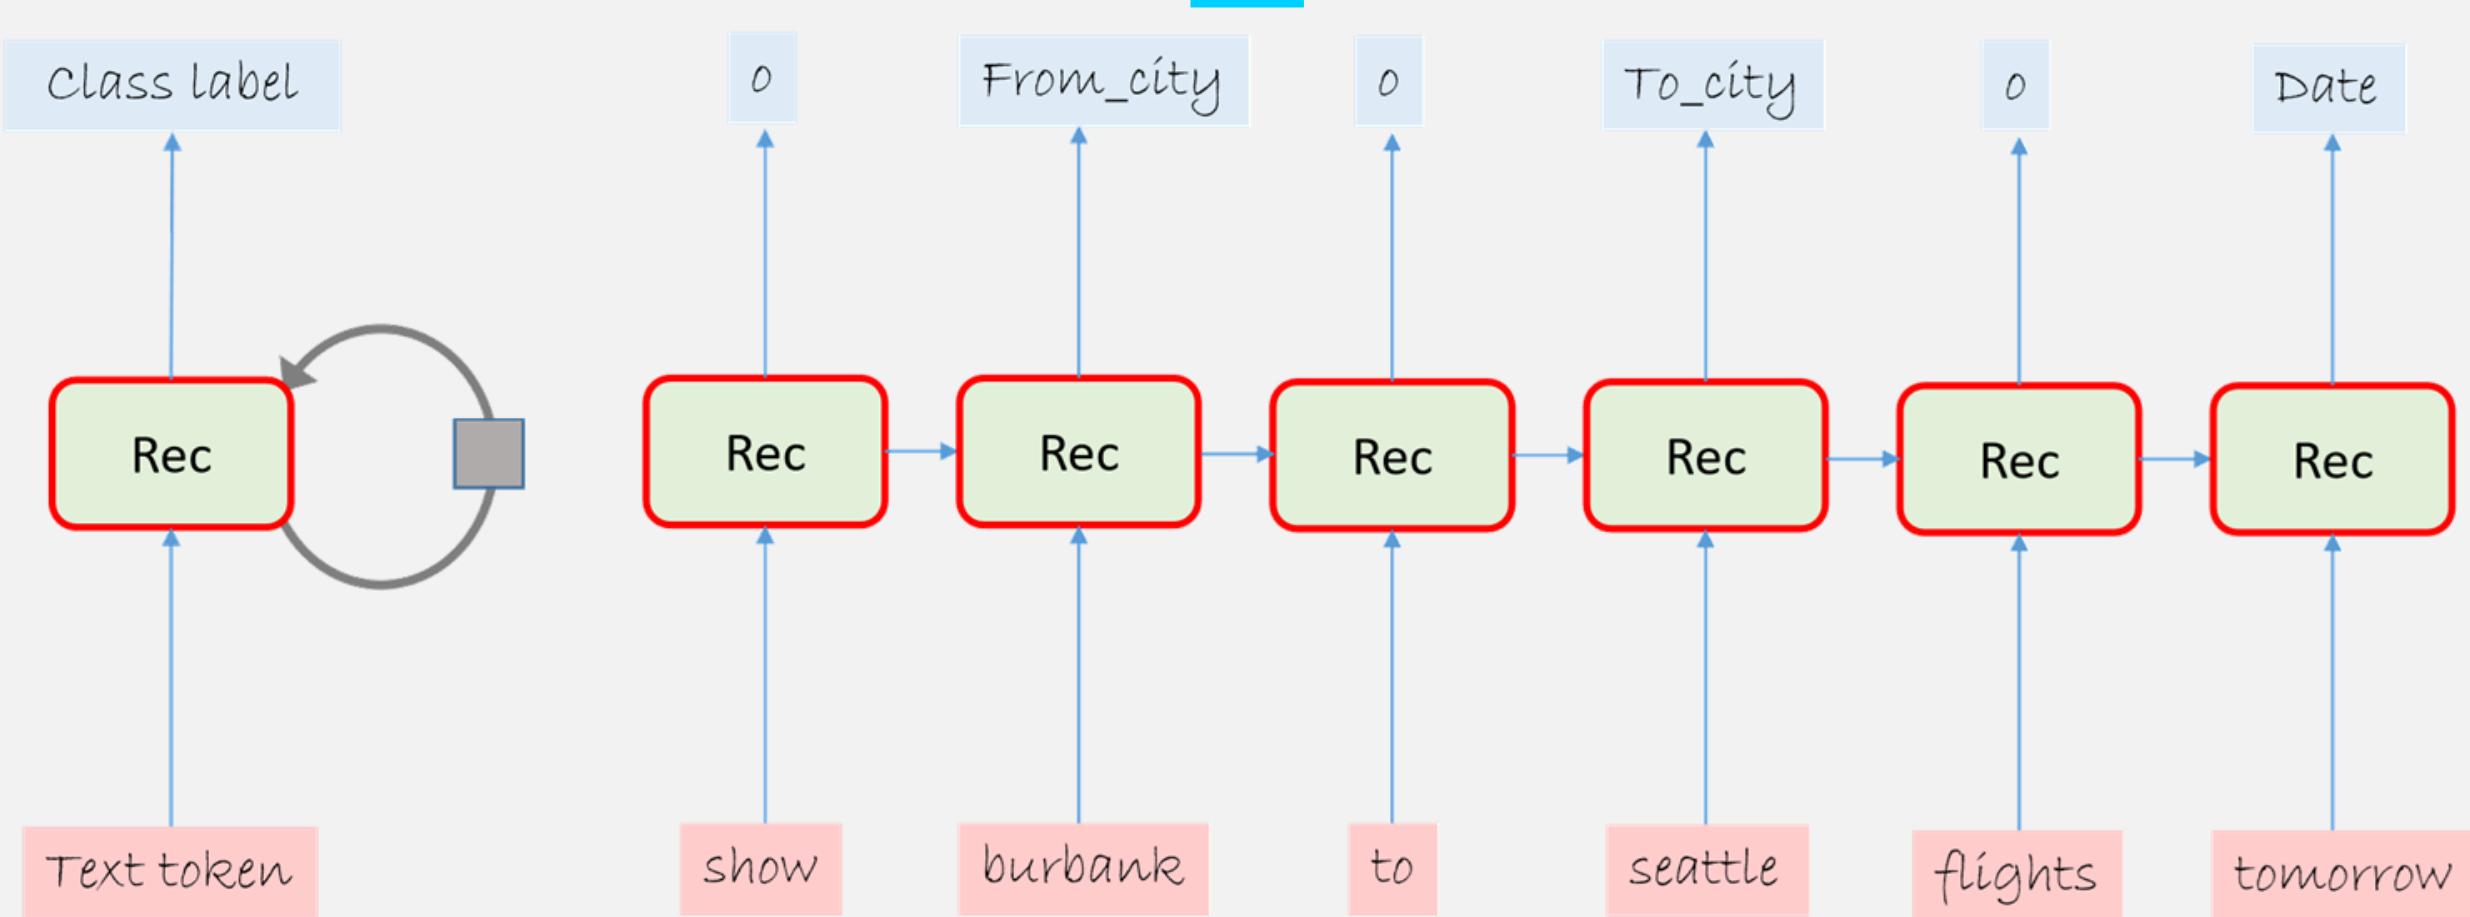


RECURRENCE

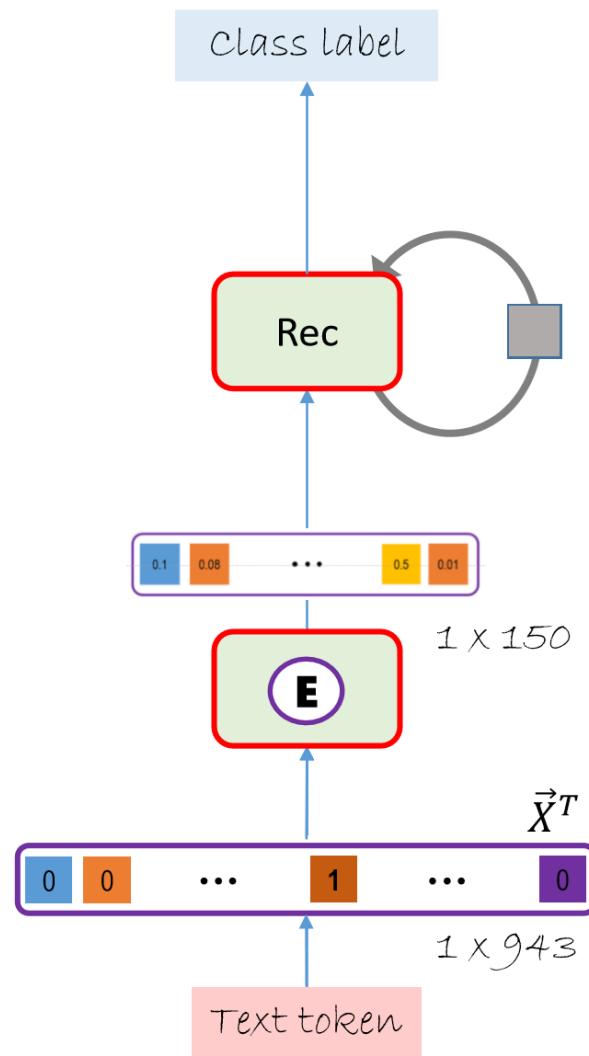


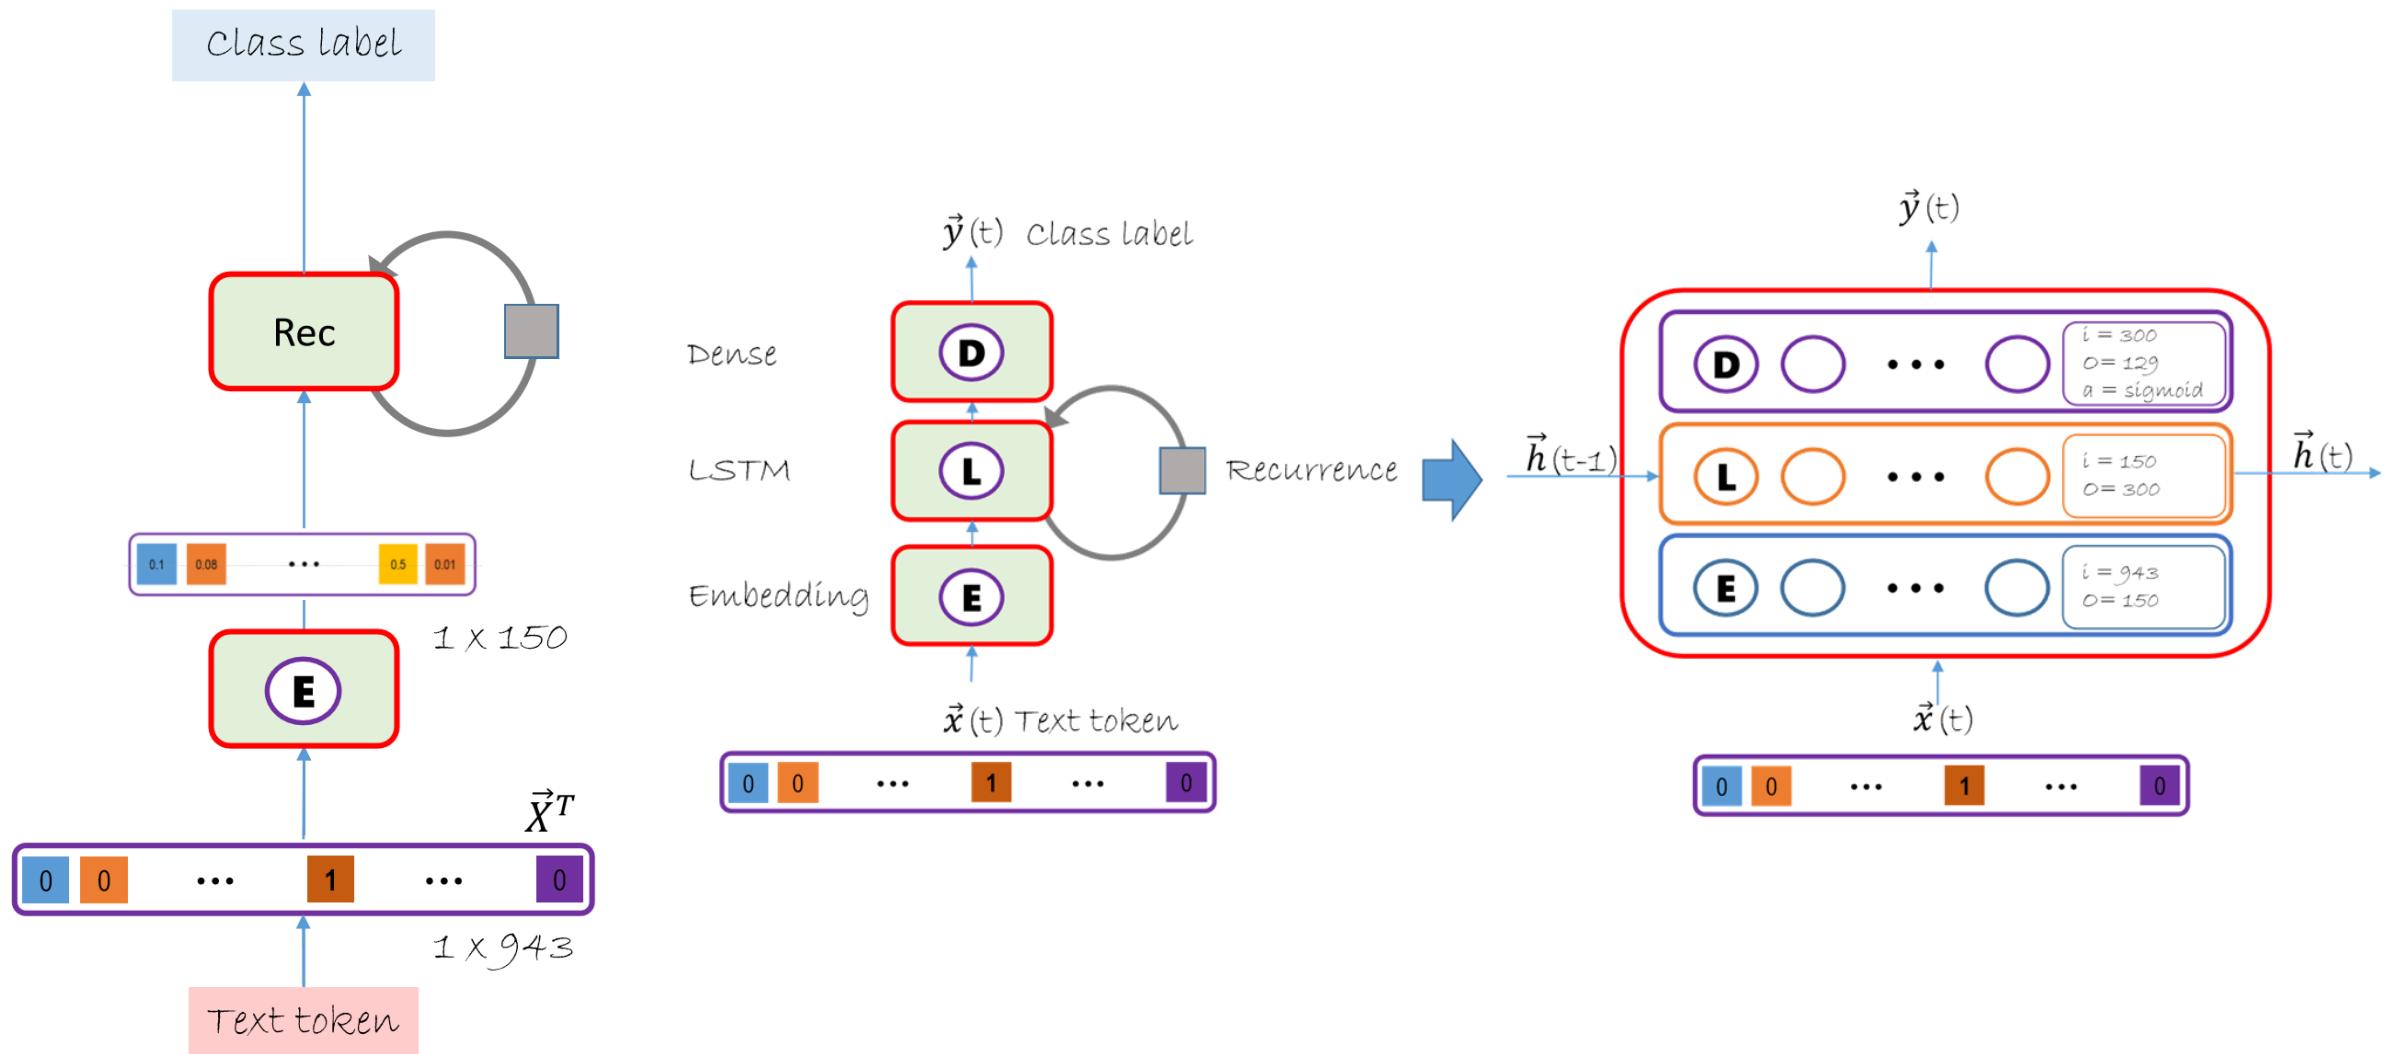


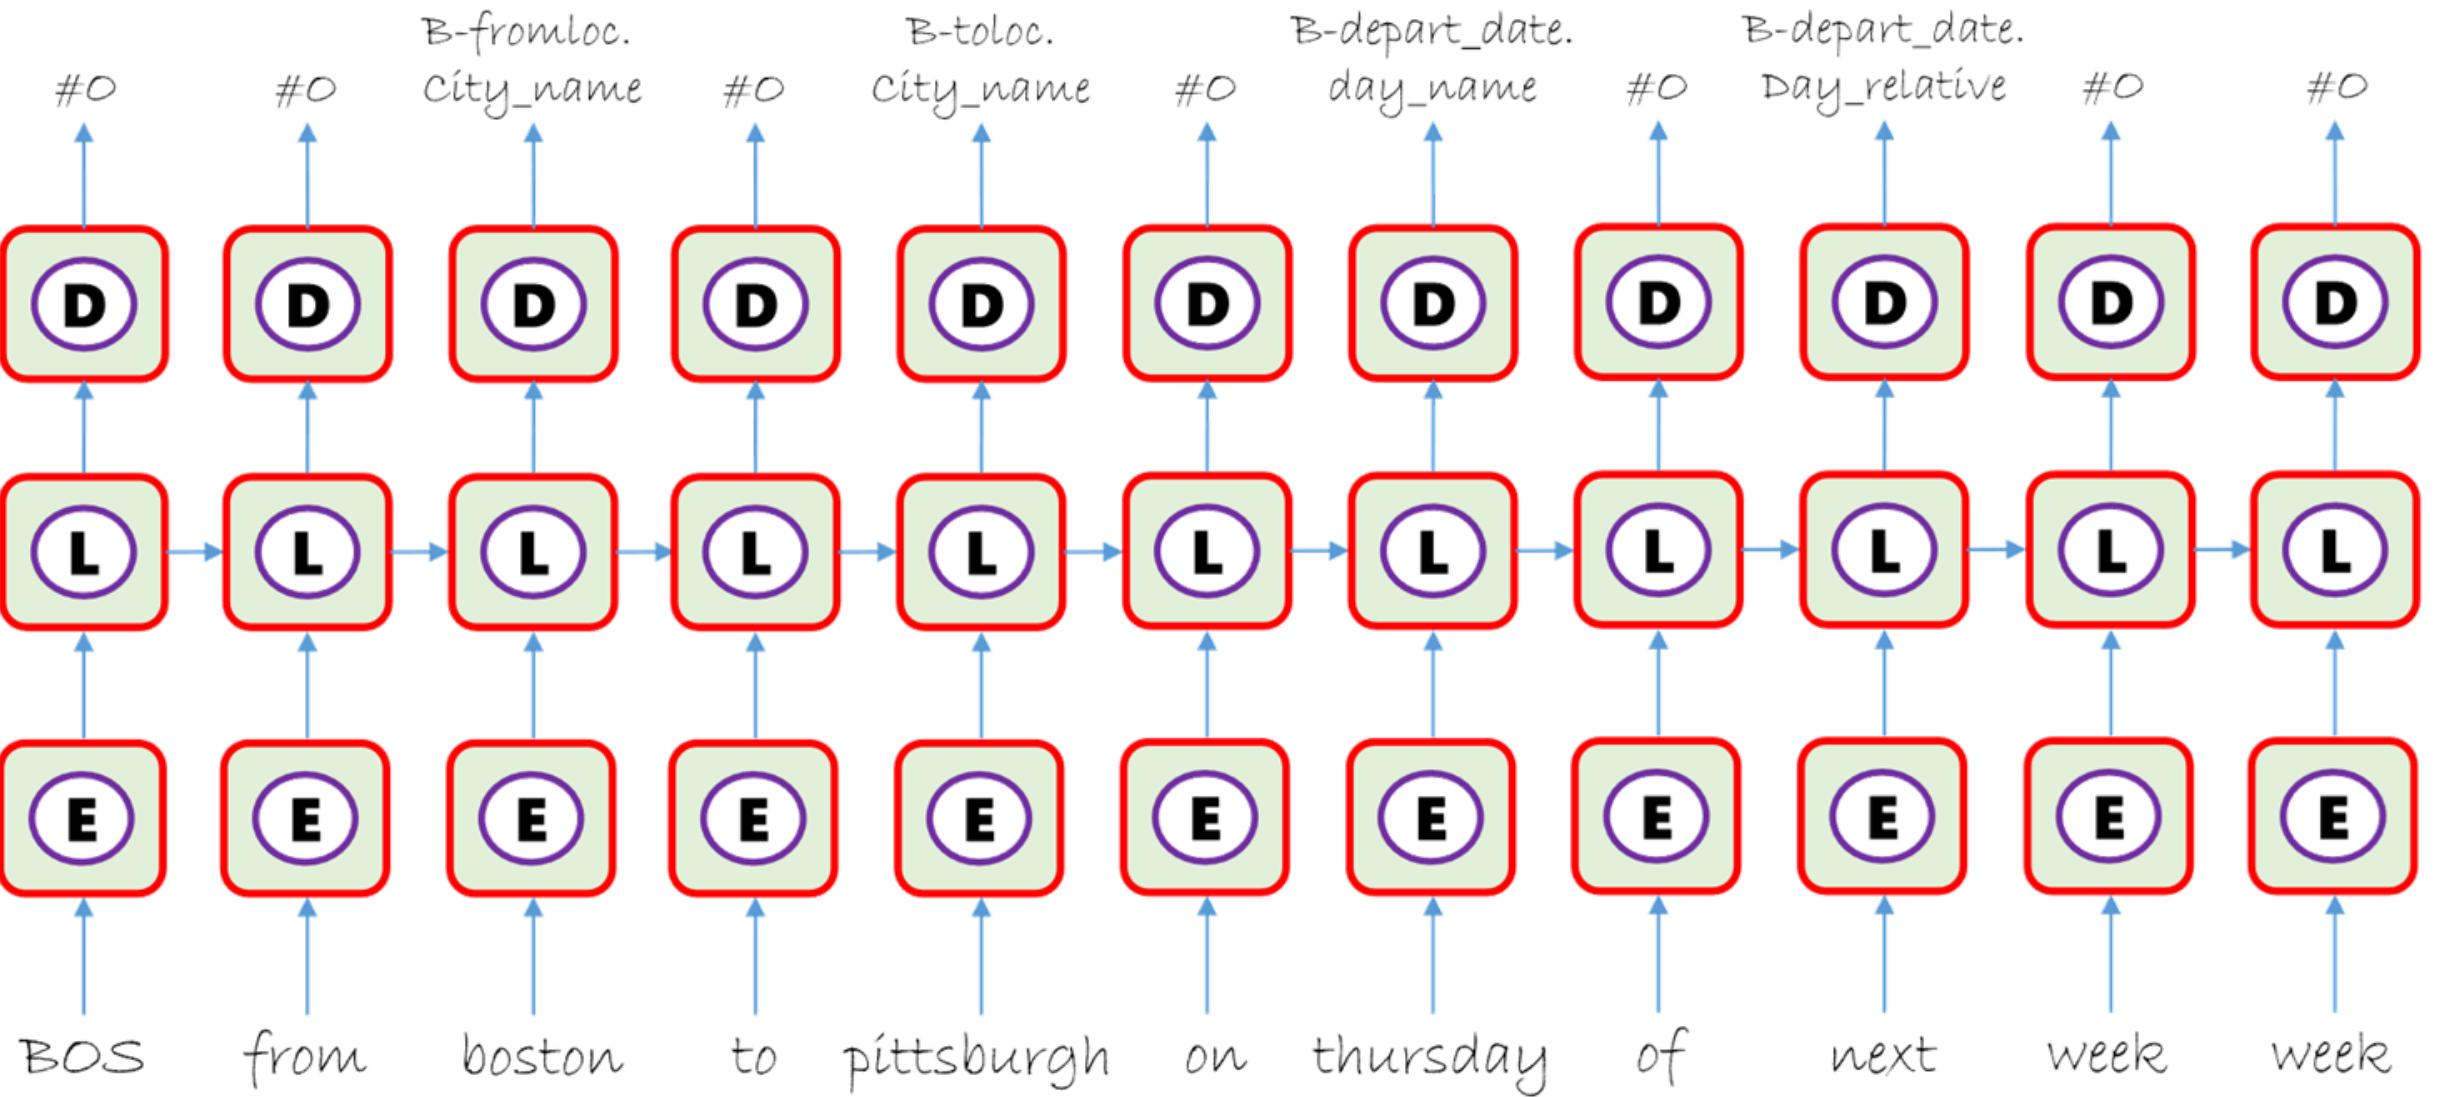
LONG SHORT TERM MEMORY



Sequence Id	Input Word (sample)	Word Index (in vocabulary) S0	Word Label	Label Index (S2)
19	# BOS	178:1	# O	128:1
19	# please	688:1	# O	128:1
19	# give	449:1	# O	128:1
19	# me	581:1	# O	128:1
19	# the	827:1	# O	128:1
19	# flights	429:1	# O	128:1
19	# from	444:1	# O	128:1
19	# boston	266:1	# B-fromloc.city_name	48:1
19	# to	851:1	# O	128:1
19	# pittsburgh	682:1	# B-toloc.city_name	78:1
19	# on	654:1	# O	128:1
19	# thursday	845:1	# B-depart_date.day_name	26:1
19	# of	646:1	# O	128:1
19	# next	621:1	# B-depart_date.date_relative	25:1
19	# week	910:1	# O	128:1
19	# EOS	179:1	# O	128:1







EMBEDINGS

WHAT IS WORD2VEC?

Two distinct models

- CBoW
- Skip-Gram

Various training methods

- Negative Sampling
- Hierarchical Softmax

A rich preprocessing pipeline

- Dynamic Context Windows
- Subsampling
- Deleting Rare Words

SKIP-GRAMS WITH NEGATIVE SAMPLING (SGNS)

Marco saw a furry little wampimuk hiding in the tree.

SKIP-GRAMS WITH NEGATIVE SAMPLING (SGNS)

Marco saw a furry little wampimuk hiding in the tree.

SKIP-GRAMS WITH NEGATIVE SAMPLING (SGNS)

Marco saw a **furry** little wampimuk **hiding** in the tree.

words

wampimuk

wampimuk

wampimuk

wampimuk

...

contexts

furry

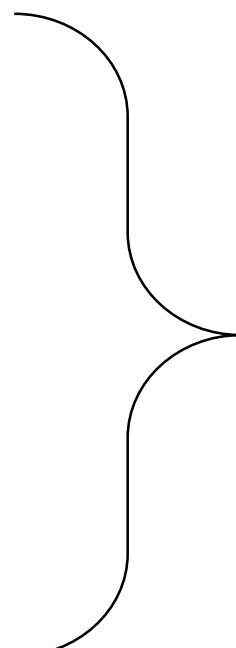
little

hiding

in

...

D (data)



SKIP-GRAMS WITH NEGATIVE SAMPLING (SGNS)

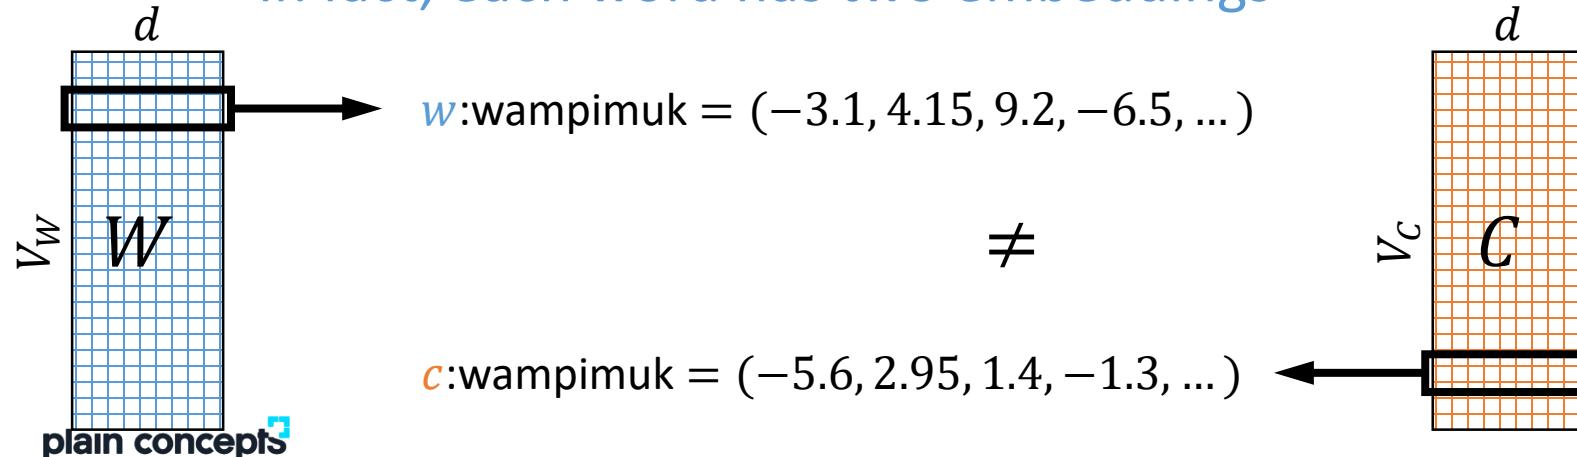
SGNS finds a vector \vec{w} for each word w in our vocabulary V_W

Each such vector has d latent dimensions (e.g. $d = 100$)

Effectively, it learns a matrix W whose rows represent V_W

Key point: it also learns a similar auxiliary matrix C of context vectors

In fact, each word has two embeddings



SKIP-GRAMS WITH NEGATIVE SAMPLING (SGNS)

SKIP-GRAMS WITH NEGATIVE SAMPLING (SGNS)

- **Maximize:** $\sigma(\vec{w} \cdot \vec{c})$
 - c was observed with w

<u>words</u>	<u>contexts</u>
wampimuk	furry
wampimuk	little
wampimuk	hiding
wampimuk	in

SKIP-GRAMS WITH NEGATIVE SAMPLING (SGNS)

- Maximize: $\sigma(\vec{w} \cdot \vec{c})$
 - c was observed with w

<u>words</u>	<u>contexts</u>
wampimuk	furry
wampimuk	little
wampimuk	hiding
wampimuk	in

- Minimize: $\sigma(\vec{w} \cdot \vec{c}')$
 - c' was **hallucinated** with w

<u>words</u>	<u>contexts</u>
wampimuk	Australia
wampimuk	cyber
wampimuk	the
wampimuk	1985

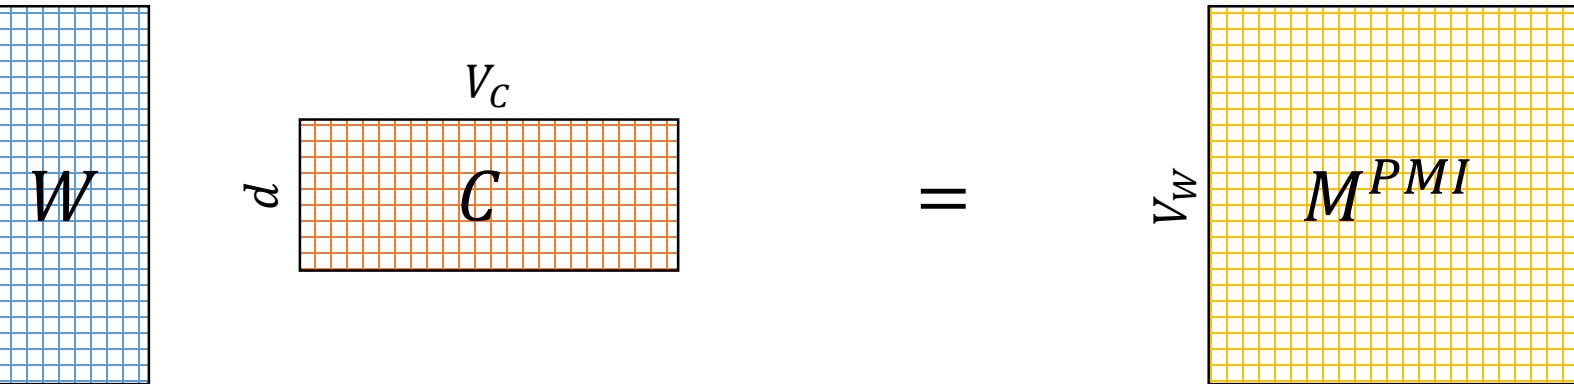
SKIP-GRAMS WITH NEGATIVE SAMPLING (SGNS)

- “Negative Sampling”
- SGNS samples k contexts c' at random as **negative examples**
- “Random” = unigram distribution

$$P(c) = \frac{\#c}{|D|}$$

- **Spoiler:** Changing this distribution has a significant effect

WHAT IS SGNS LEARNING?

$$\begin{matrix} d \\ V_w \end{matrix} \begin{matrix} W \\ \times \\ d \end{matrix} \begin{matrix} V_c \\ C \end{matrix} = \begin{matrix} V_c \\ V_w \end{matrix} \begin{matrix} M^{PMI} \\ - \log k \end{matrix}$$


“Neural Word Embeddings as Implicit Matrix Factorization”
Levy & Goldberg, NIPS 2014

SOME OTHER ADVANCED IDEAS

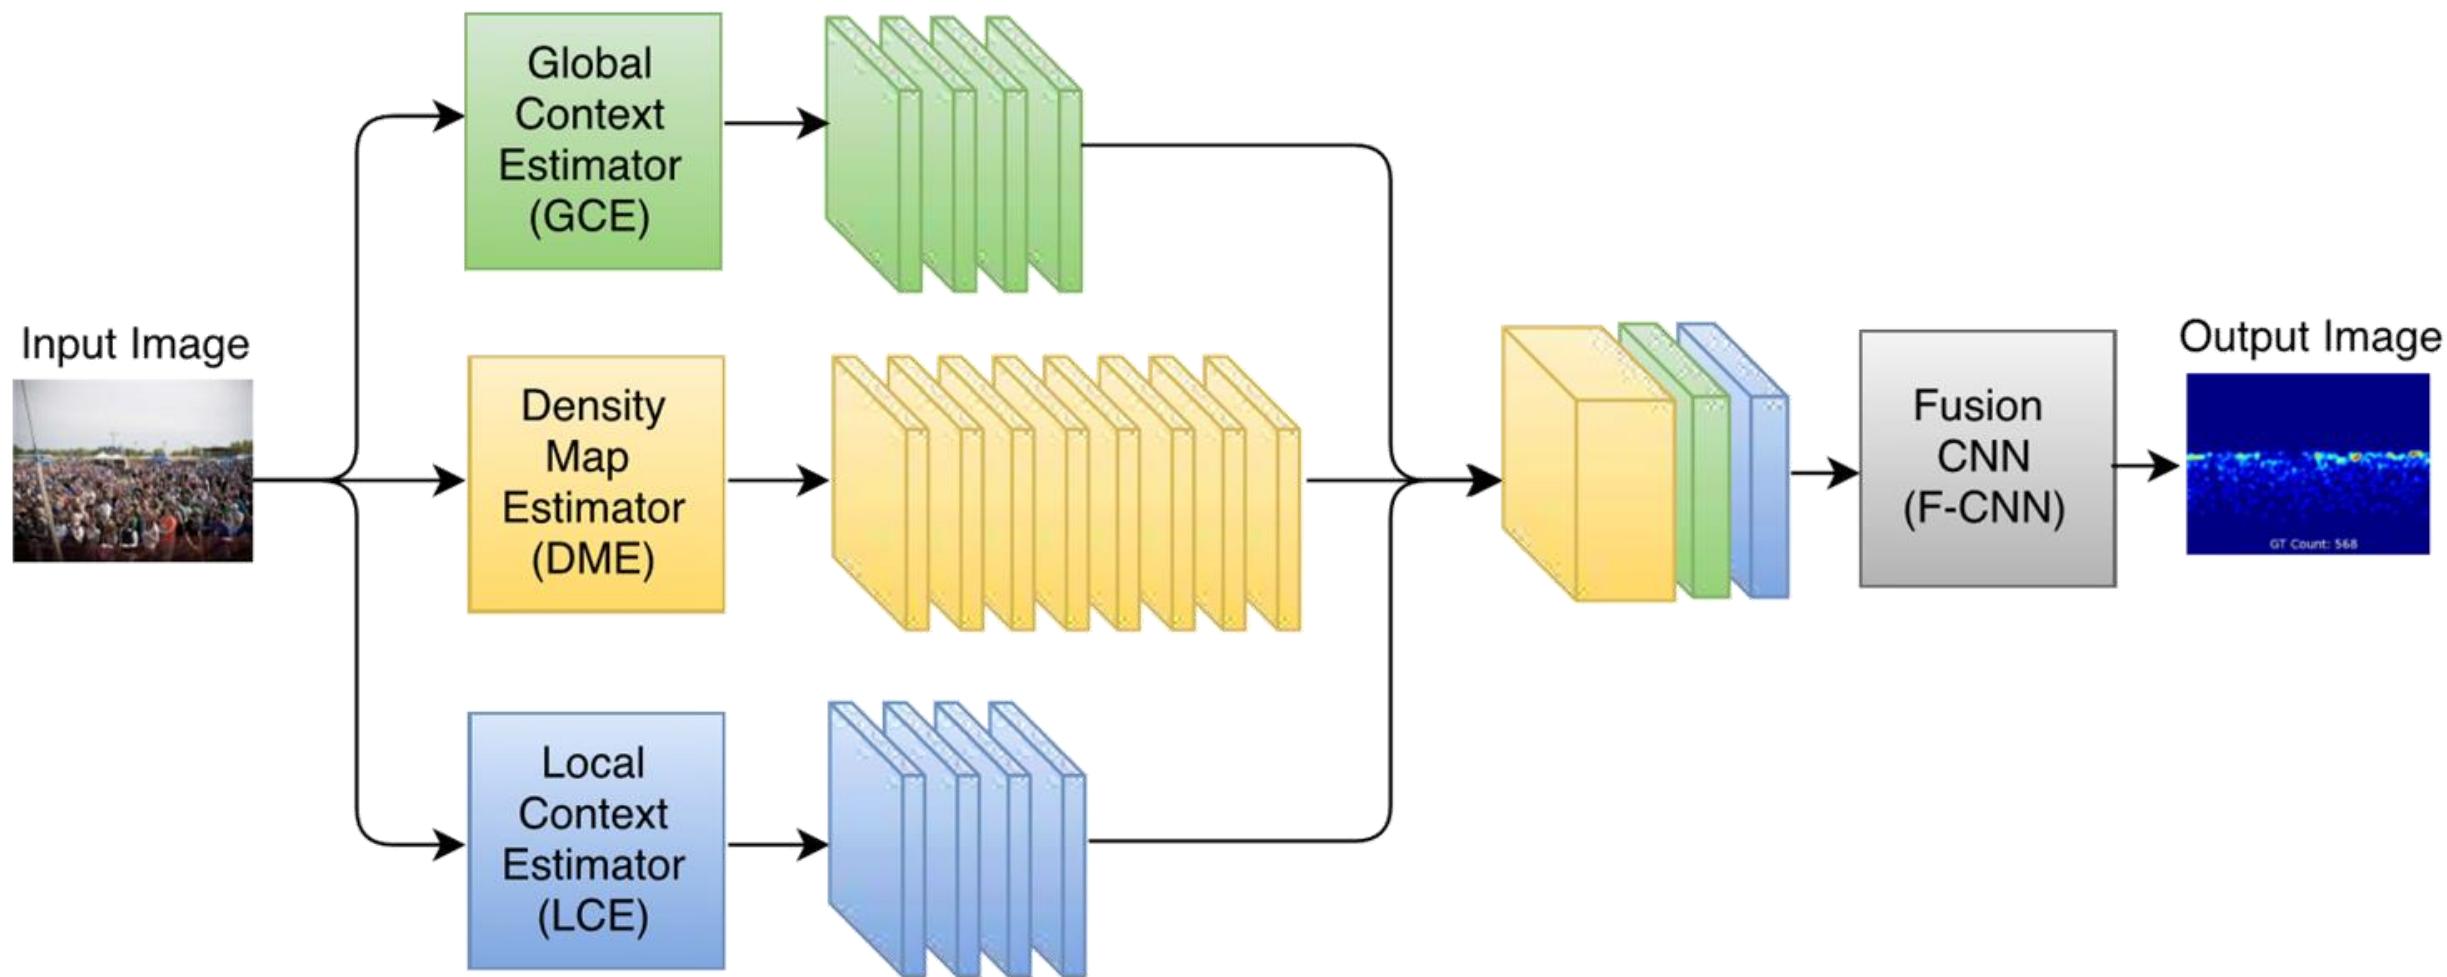
Counting people in a multitude

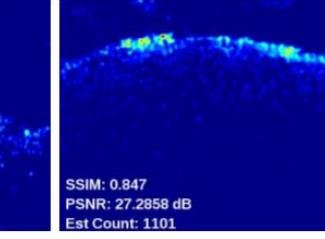
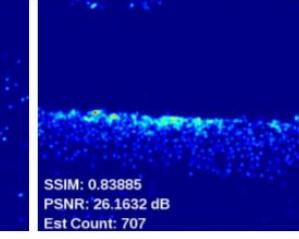
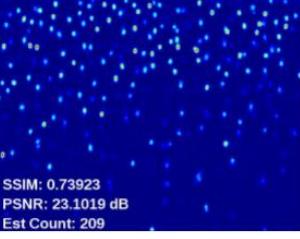
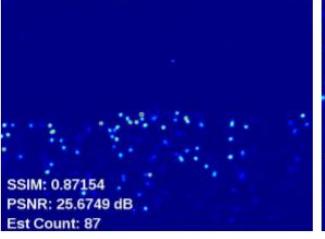
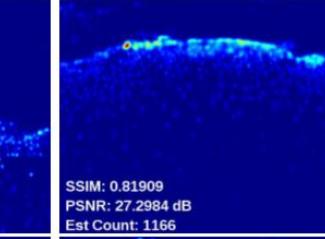
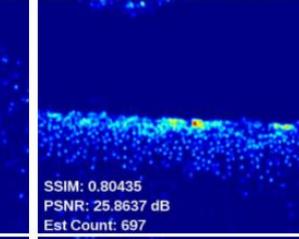
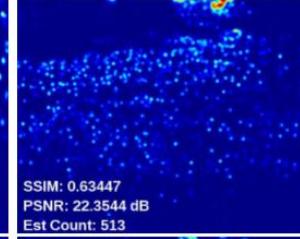
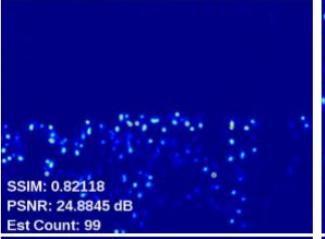
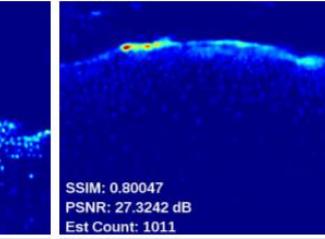
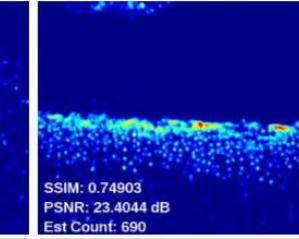
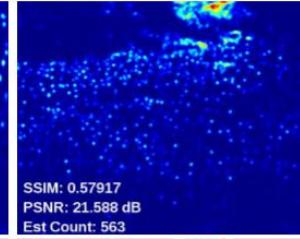
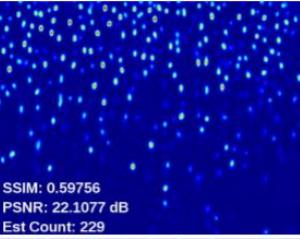
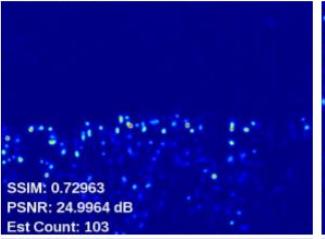
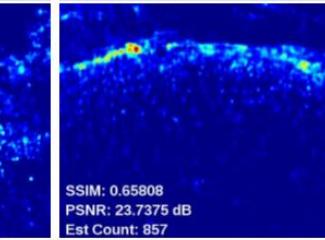
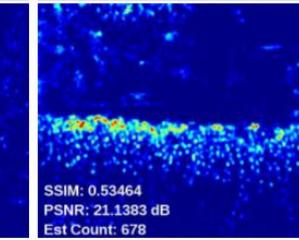
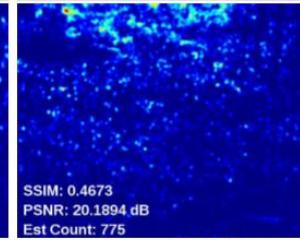
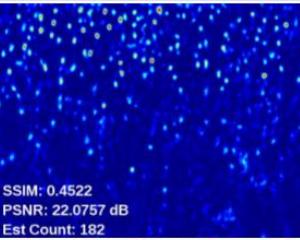
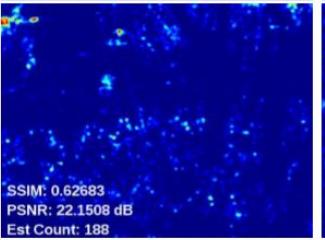
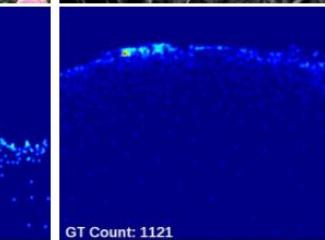
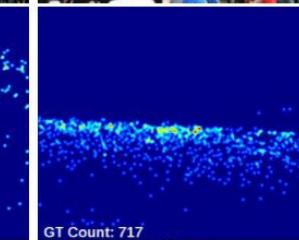
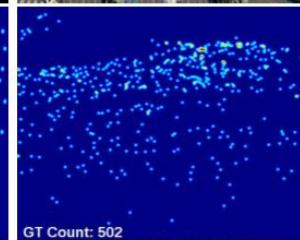
- CONTEXTUAL PYRAMID CNN (CP-CNN)

A new method named Contextual Pyramid CNN (CP-CNN) is proposed here to generate density maps and influx estimations, by explicitly incorporating global and local context information. Composed of four modules: Global Context Estimator (GCE), Local Context Estimator (LCE), Density Map Estimator (DME) and a Fusion-CNN (F-CNN) convolutional network.

Vishwanath A. Sindagi, Vishal M. Patel; The IEEE International Conference on Computer Vision (ICCV), 2017, pp. 1861-1870









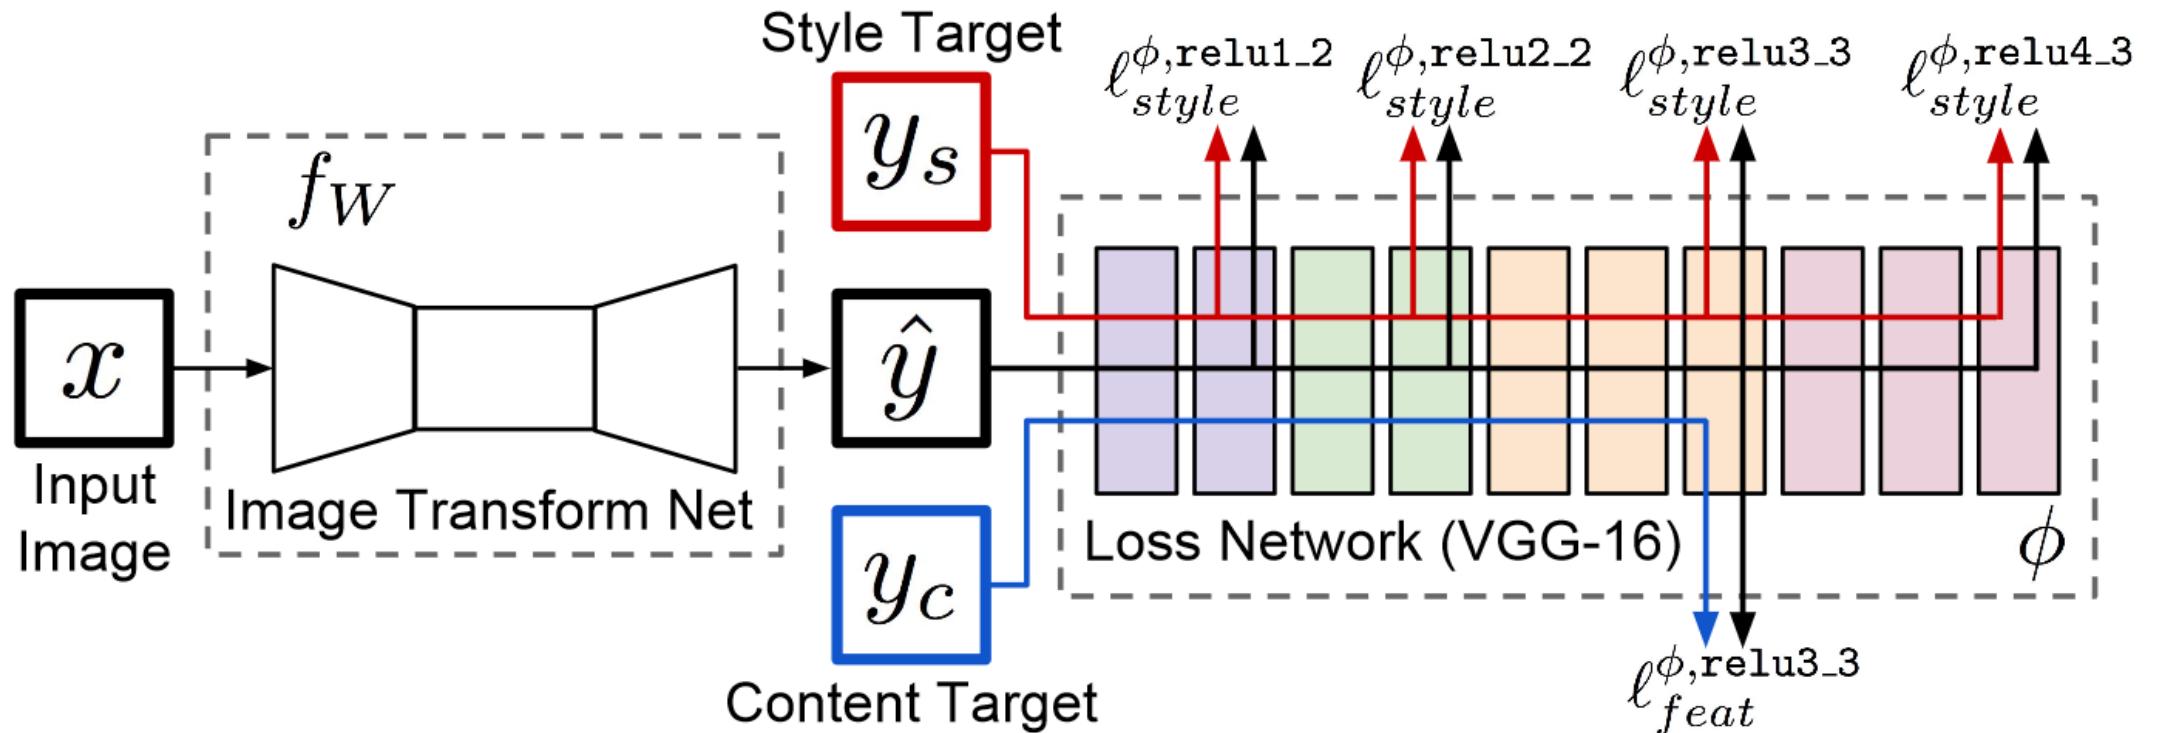
Transferring style between images

- CONVOLUTIONAL NEURAL NETWORKS

By using a perceptual loss functions based on high-level features extracted from pretrained networks, networks for image transformation tasks can be trained, and by fine tuning the loss function different features can be kept for the source image and the style image.

Justin Johnson, Alexandre Alahi, Li Fei-Fei; Perceptual Losses for Real-Time Style Transfer and Super-Resolution, 2016









Using GANs to drive design decisions

- GENERATIVE ADVERSARIAL NETWORKS

By taking advantage of Generational Adversarial Networks, synthetic images based on the training data can be generated. Including an external array of features, the generated images can be tailored to a specific set of requirements.

Jaime Deverall, Jiwoo Lee, Miguel Ayala; Using Generative Adversarial Networks to Design Shoes

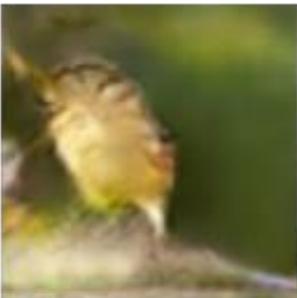


Text
description

This bird is blue with white and has a very short beak



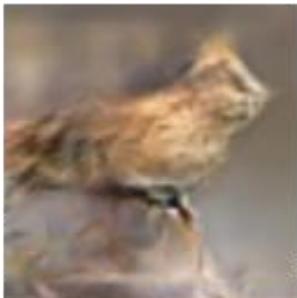
This bird has wings that are brown and has a yellow belly



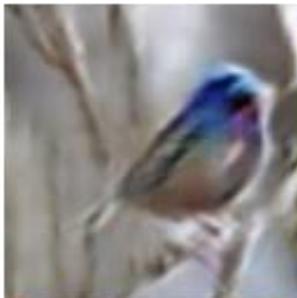
A white bird with a black crown and yellow beak



This bird is white, black, and brown in color, with a brown beak



The bird has small beak, with reddish brown crown and gray belly



This is a small, black bird with a white breast and white on the wingbars.



This bird is white black and yellow in color, with a short black beak



Stage-I
images

Stage-II
images

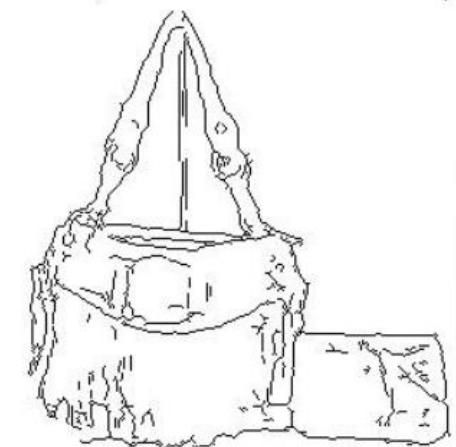
Input



Ground truth

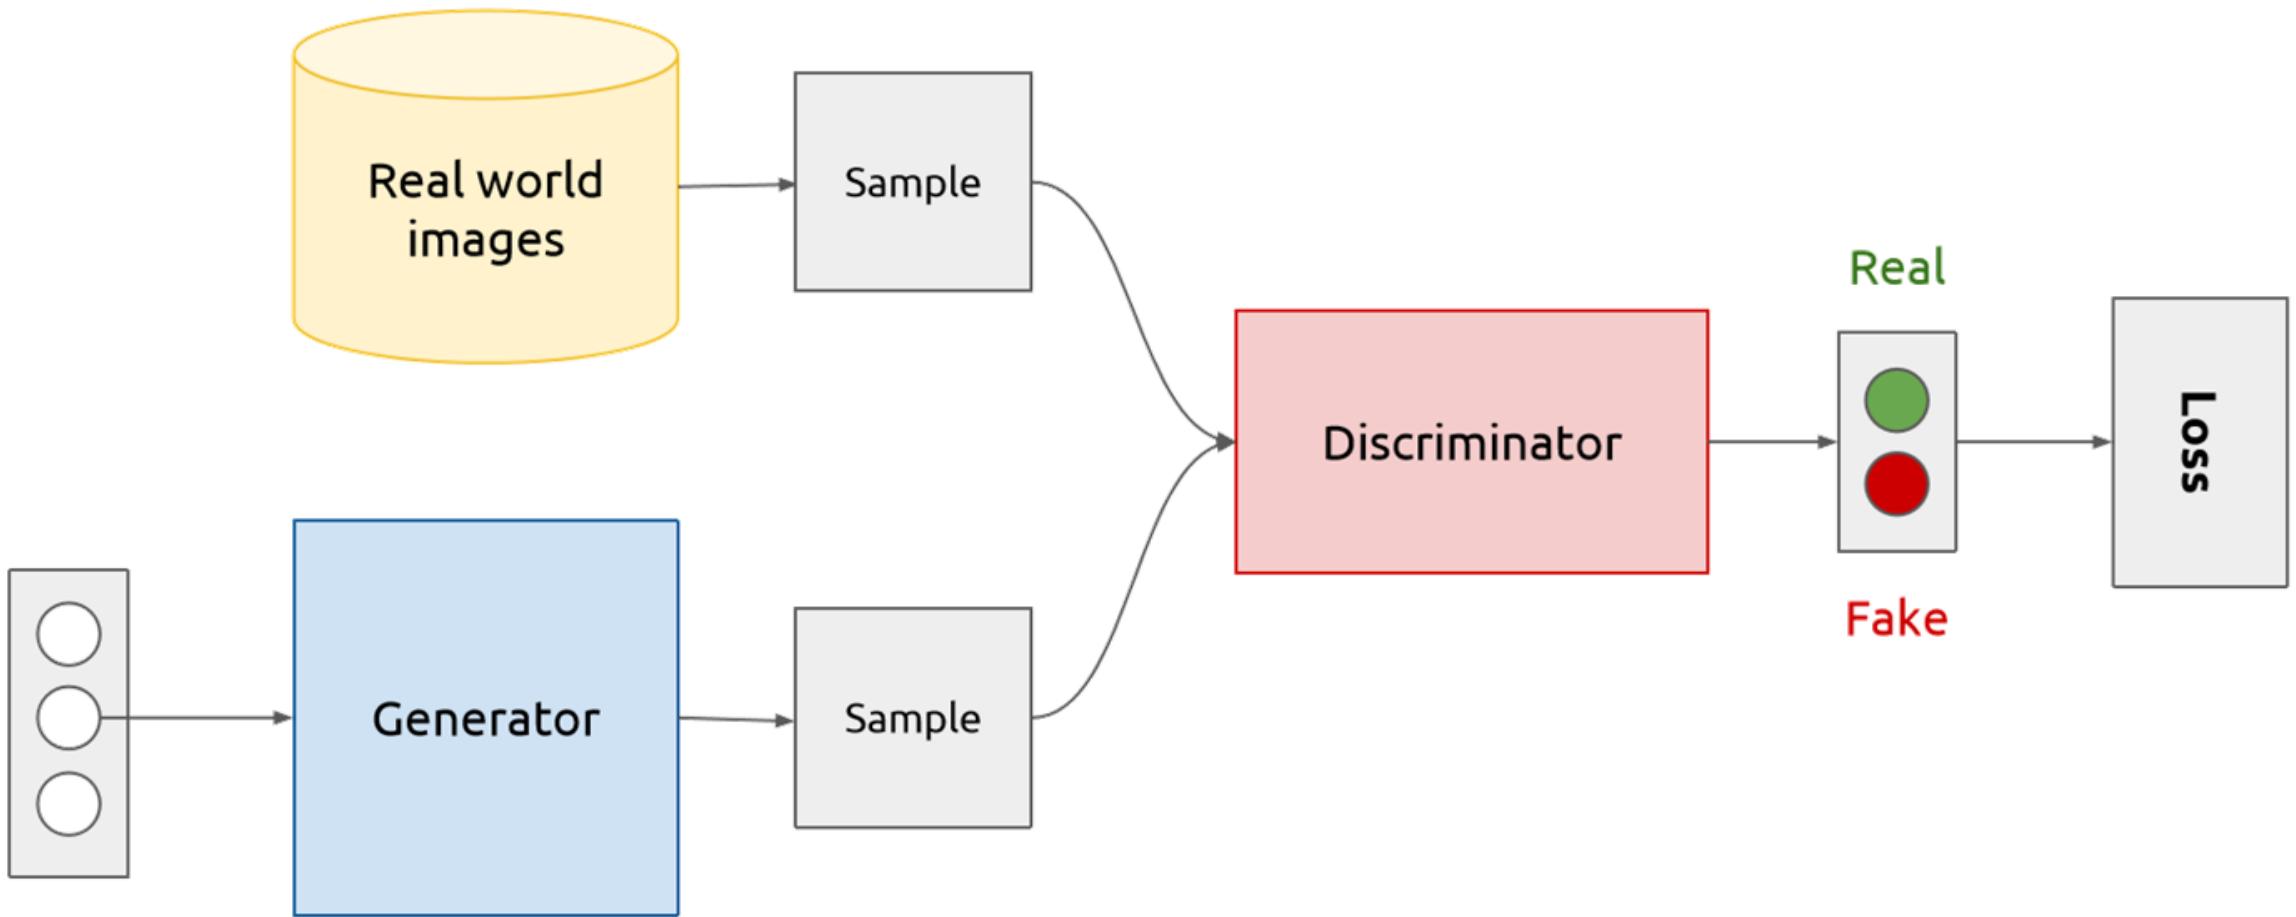


Output





Latent random variable



AZURE MACHINE LEARNING

DATA SCIENCE & AI

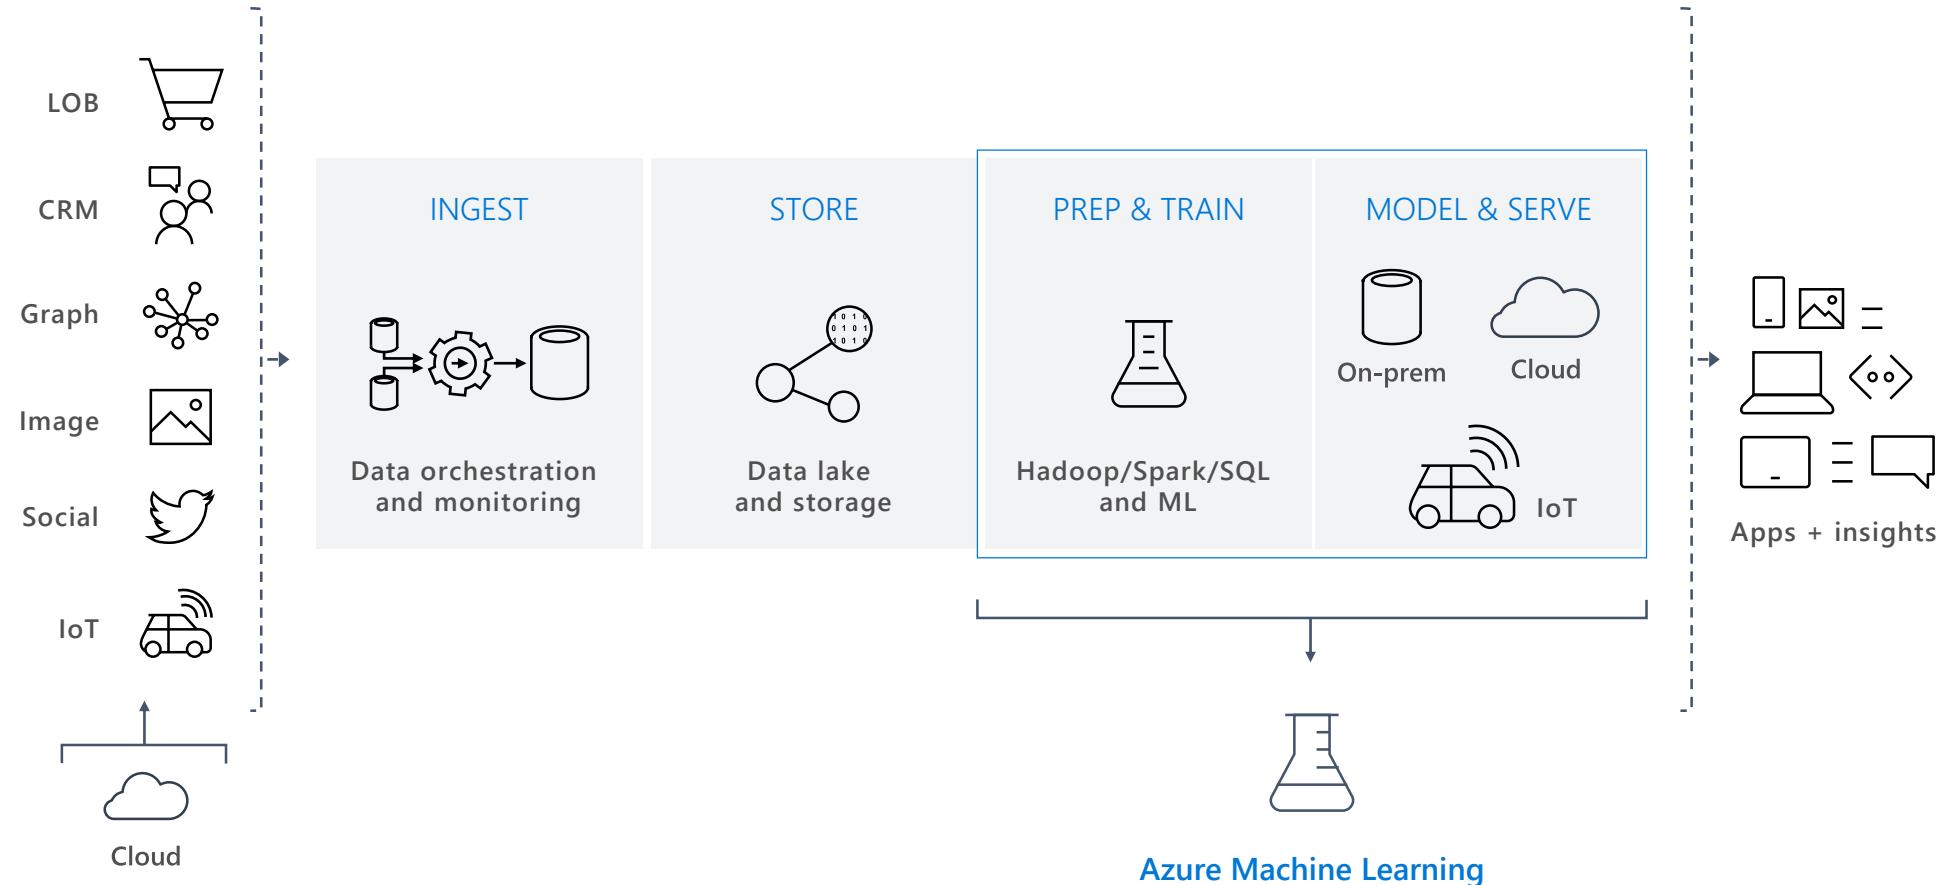
KEY TRENDS

- > Accelerating adoption of AI by developers (consuming models)
- > Rise of hybrid training and scoring scenarios
- > Push scoring/inference to the event (edge, cloud, on-prem)
- > Some developers moving into deep learning as non-traditional path to DS / AI dev
- > Growth of diverse hardware arms race across all form factors (CPU / GPU / FPGA / ASIC / device)

CHALLENGES

- ⚠ Data prep
- ⚠ Model deployment & management
- ⚠ Model lineage & auditing
- ⚠ Explain-ability

THE AI DEVELOPMENT LIFECYCLE



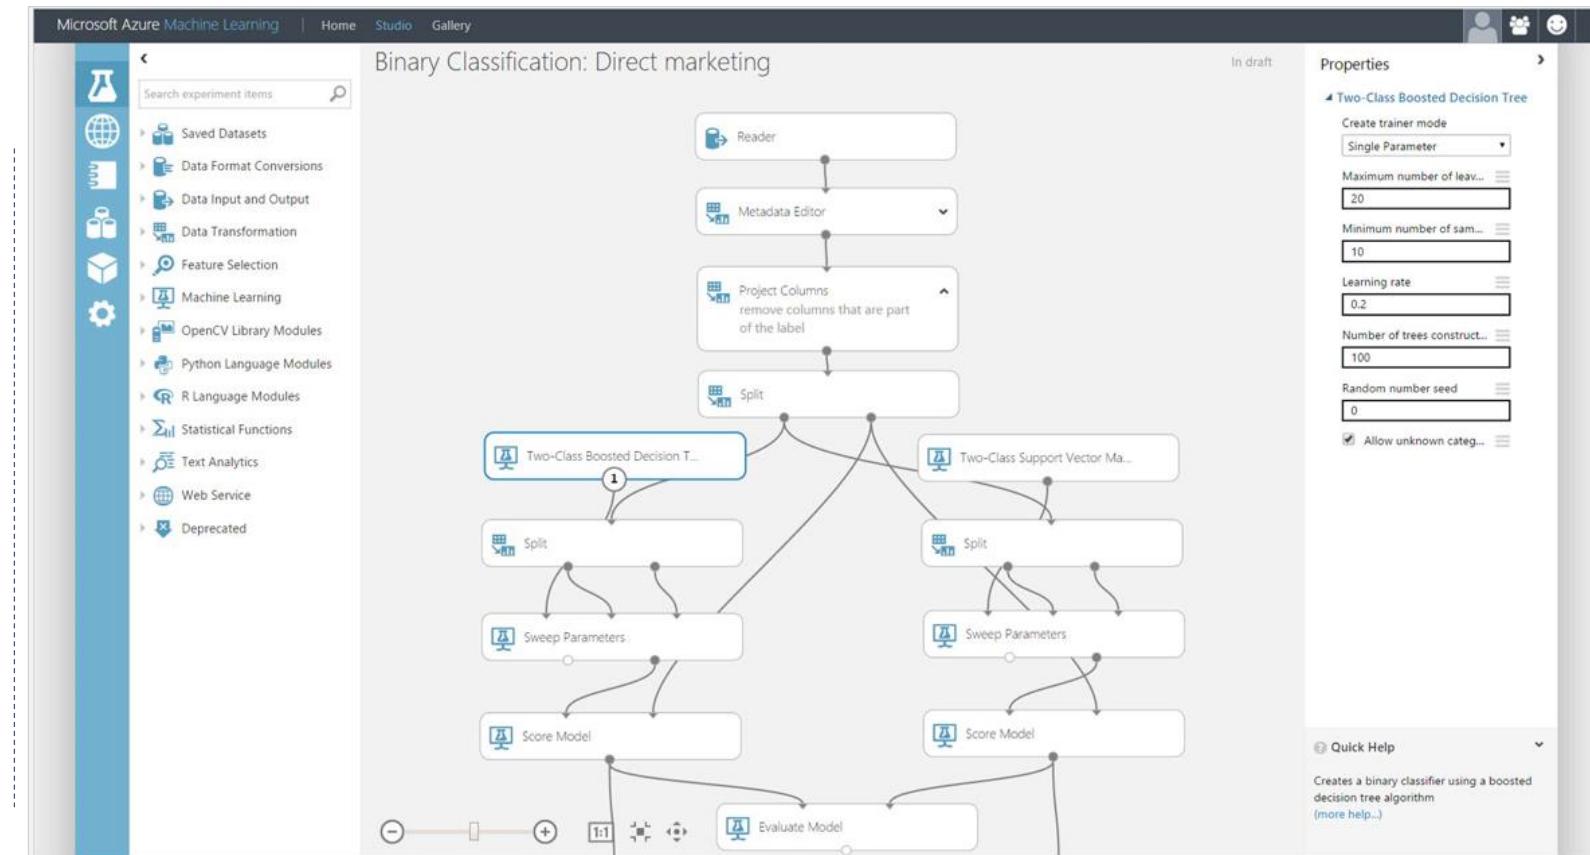
AZURE MACHINE LEARNING STUDIO

Platform for emerging data scientists to graphically build and deploy experiments

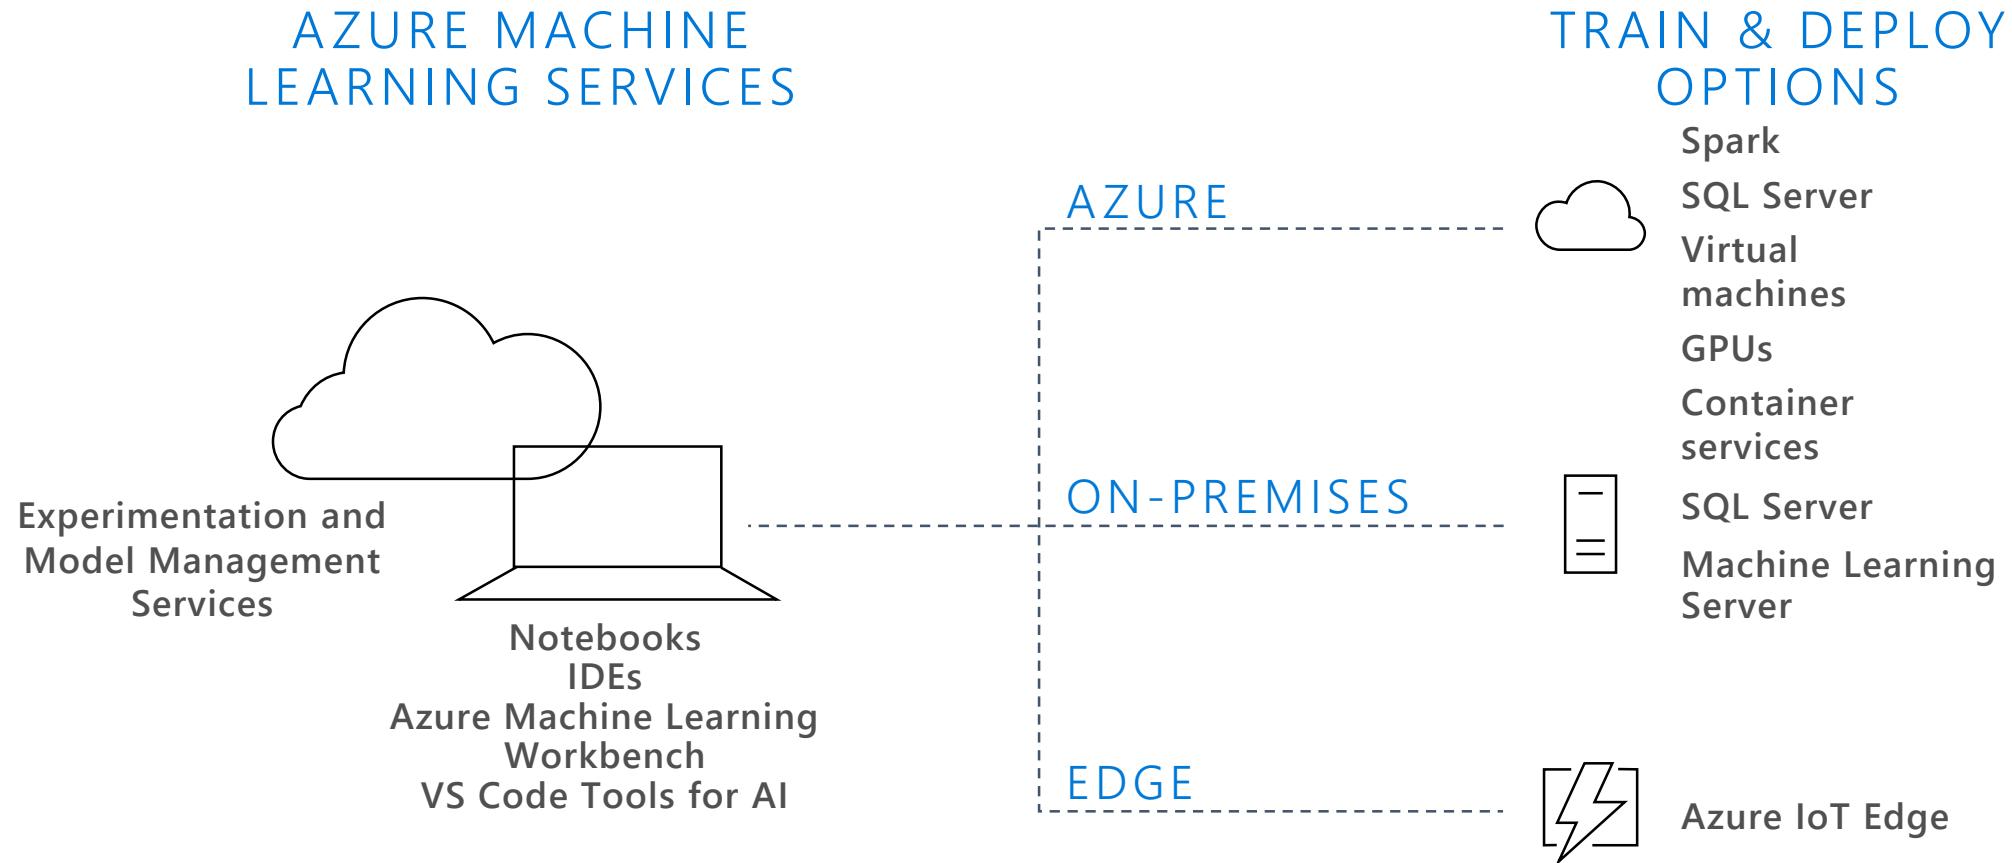
- Rapid experiment composition
- > 100 easily configured modules for data prep, training, evaluation
- Extensibility through R & Python
- Serverless training and deployment

Some numbers:

- 100's of thousands of deployed models serving billions of requests



NEW CAPABILITIES





New

Dashboard

Data factories

SQL databases

HDInsight clusters

Resource groups

Storage accounts

Data Lake Analytics

Storage accounts (class...)

All resources

Recent

Data Lake Store

App Services

Virtual machines

Subscriptions

SQL servers

Virtual machines (classic)

Virtual networks

Cloud services (classic)

More services >

Marketplace



AI + Cognitive Services



Filter

Search AI + Cognitive Services



Machine Learning Experimenta...

Microsoft

Create an Experimentation Account to start preparing your data, engineer features, author models using existing tools and manage reproducibility through tracking history of runs of all your machine learning models.

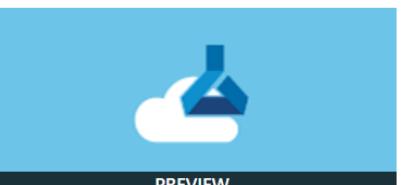
Create

Machine Learning

More

Machine Learning Experimentation
(preview)

Microsoft

Machine Learning Model Management
(preview)

Microsoft

Data Science Virtual Machine - Windows
2016

Microsoft

Bot Service

More



PREVIEW



New

Dashboard

Data factories

SQL databases

HDInsight clusters

Resource groups

Storage accounts

Data Lake Analytics

Storage accounts (class...)

All resources

Recent

Data Lake Store

App Services

Virtual machines

Subscriptions

SQL servers

Virtual machines (classic)

Virtual networks

Cloud services (classic)

More services >

ML Experimentation



Machine Learning Experimentation

* Experimentation account name

Enter the experimentation account name

* Subscription

DX Hackathon

* Resource group

 Create new Use existing

* Location

Australia East

* Number of seats. First two are free. i

2

* Storage account i Create new Use existing

Enter the storage account name

* Workspace for Experimentation account i

MattWorkspace

* Assign owner for the workspace

Matt Winkler (AI&R)

 Create Model Management account i

* Account name

Enter the Model Management account name

 Pin to dashboard**Create**

Automation options



DemoExp1

Machine Learning Experimentation - PREVIEW



New

Dashboard

Data factories

SQL databases

HDInsight clusters

Resource groups

Storage accounts

Data Lake Analytics

Storage accounts (class...

All resources

Recent

Data Lake Store

App Services

Virtual machines

Subscriptions

SQL servers

Virtual machines (classic)

Virtual networks

Cloud services (classic)

More services >

Search (Ctrl+J)

Delete

Add workspace

Update seats

Resource group

projectviennademos

Status

Enabled

Location

East US 2

Subscription

Project Vienna Demo 1

Subscription ID

Storage

projectviennademostorage

Number of seats

2

Number of workspaces

4

Getting Started



Download Azure Machine Learning Workbench for Windows

Azure Machine Learning Workbench enables you to explore data, build experiments, and compare model performance.



Download Azure Machine Learning Workbench for Mac OS X

Azure Machine Learning Workbench enables you to explore data, build experiments, and compare model performance.



Create an account for Machine Learning Model Management

Machine Learning Model Management accounts allow you to easily deploy your models and manage them in Azure.



Learn how to configure Machine Learning Compute for Model Management

To take advantage of deploying and managing your models in production, you must first configure the environment where the models will be deployed. Learn how to set up your environment to deploy and manage your models in Azure.

WHAT DID I GET?

Experimentation Account

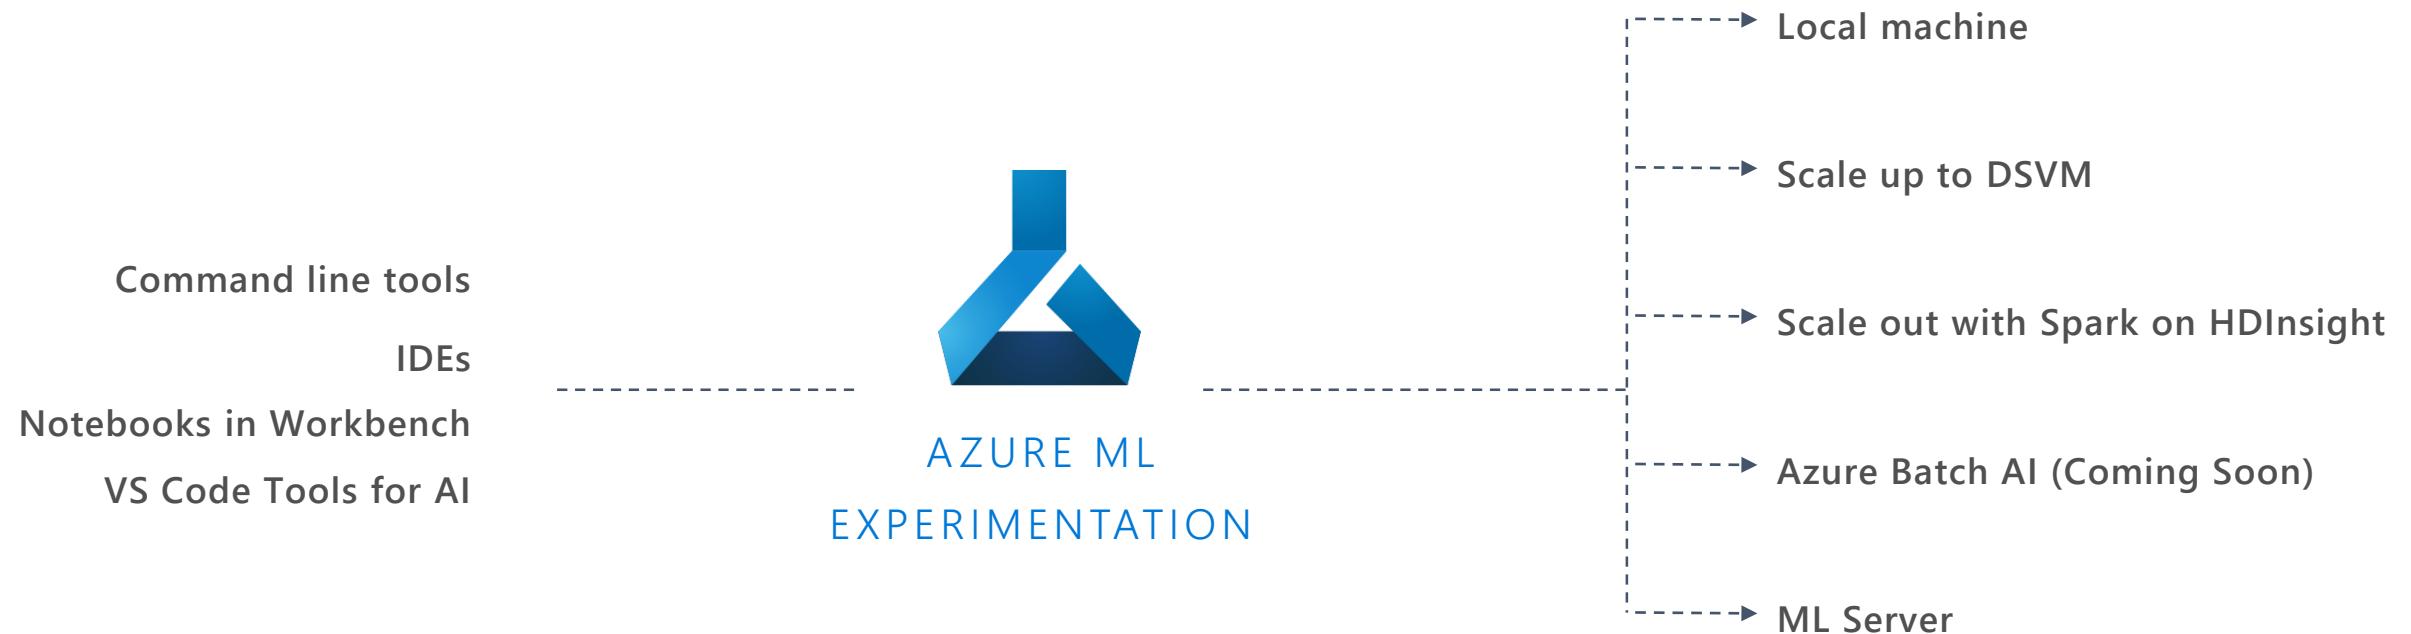
- Keep track of projects
- Local, scale-up, and scale-out training
- Run history tracking

Model Management Account

- Create containers for models
- Manage and monitor deployed models

EXPERIMENTATION SERVICE

Experiment Everywhere



EXPERIMENTATION SERVICE

Manage project dependencies

USE ANY FRAMEWORK OR LIBRARY



Manage training jobs locally, scaled-up or scaled-out

Git based checkpointing and version control

Service side capture of run metrics, output logs and models

Use your favorite IDE, and any framework

USE ANY TOOL



USE THE MOST POPULAR INNOVATIONS



MODEL MANAGEMENT SERVICE

DEPLOY EVERYWHERE



AZURE ML

MODEL MANAGEMENT



DOCKER

- Single node deployment (cloud/on-prem)
- Azure Container Service
- Azure IoT Edge
- Microsoft ML Server
- Spark clusters
- SQL Server

MANAGE MODELS

Deployment and management of models as HTTP services

Container-based hosting of real time and batch processing

Management and monitoring through Azure Application Insights

First class support for SparkML, Python, Cognitive Toolkit, TF, R, extensible to support others (Caffe, MXnet)

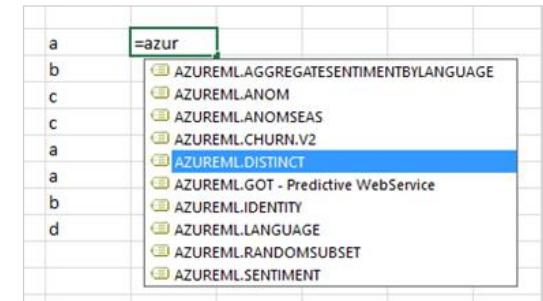
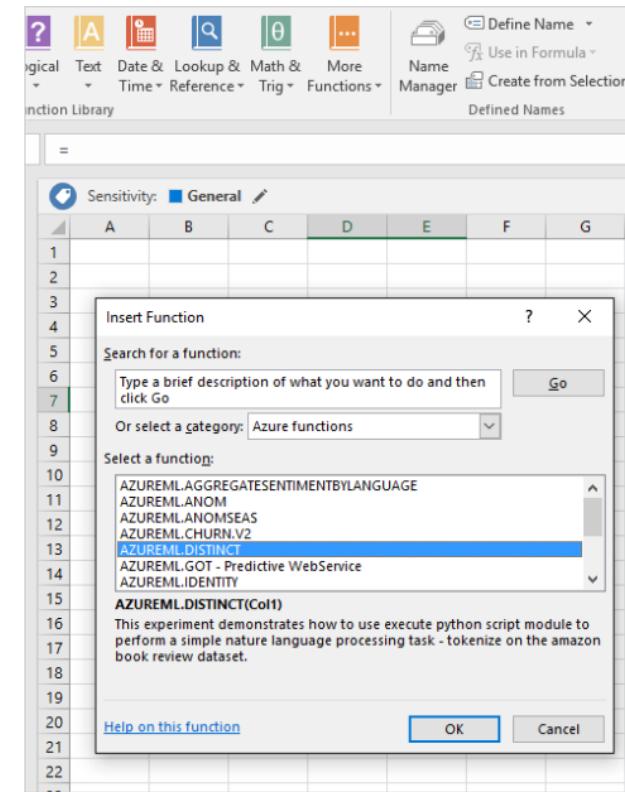
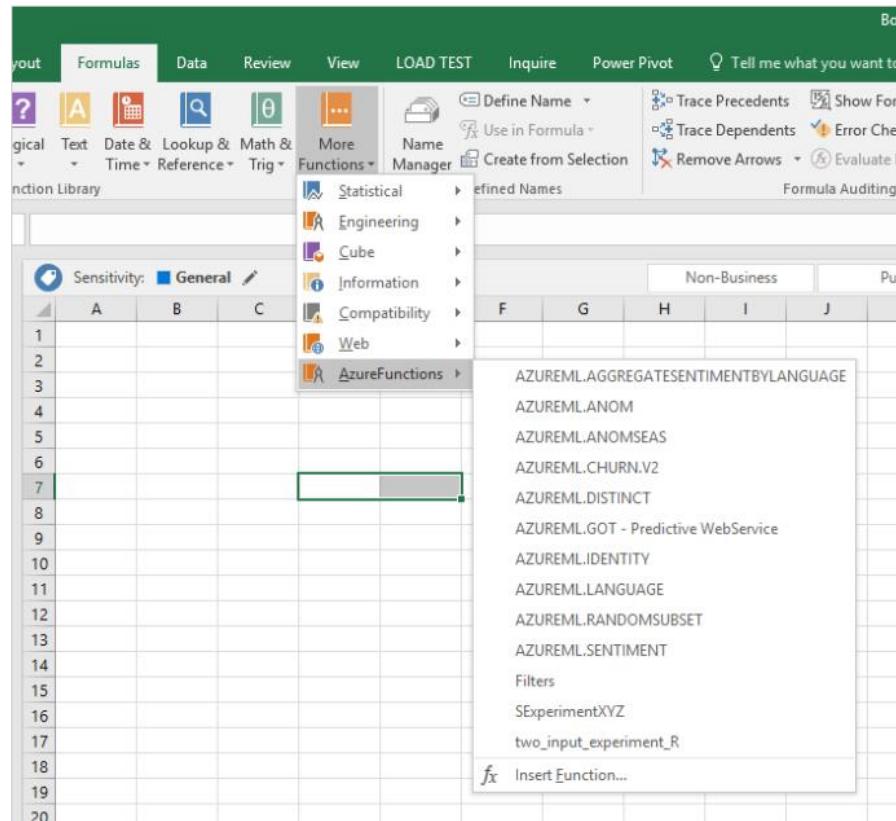
Service authoring in Python

```
# -n app name  
# -f scoring script file name  
# -m dependencies, in this case it is the pickled model file  
# -r type of model, in this case it is the scikit-learn model  
$ az ml service create realtime -n irisapp -f iris_score.py -m model.pkl -r scikit-py
```

```
$ az ml service run realtime -n irisapp -d '{"input": "[[1.0, 2.2, 2.3, 2.7]]"}'  
2
```

AI POWERED SPREADSHEETS

Any model published through Model Management can be called from Excel



VS CODE TOOLS FOR AI

VS CODE TOOLS FOR AI

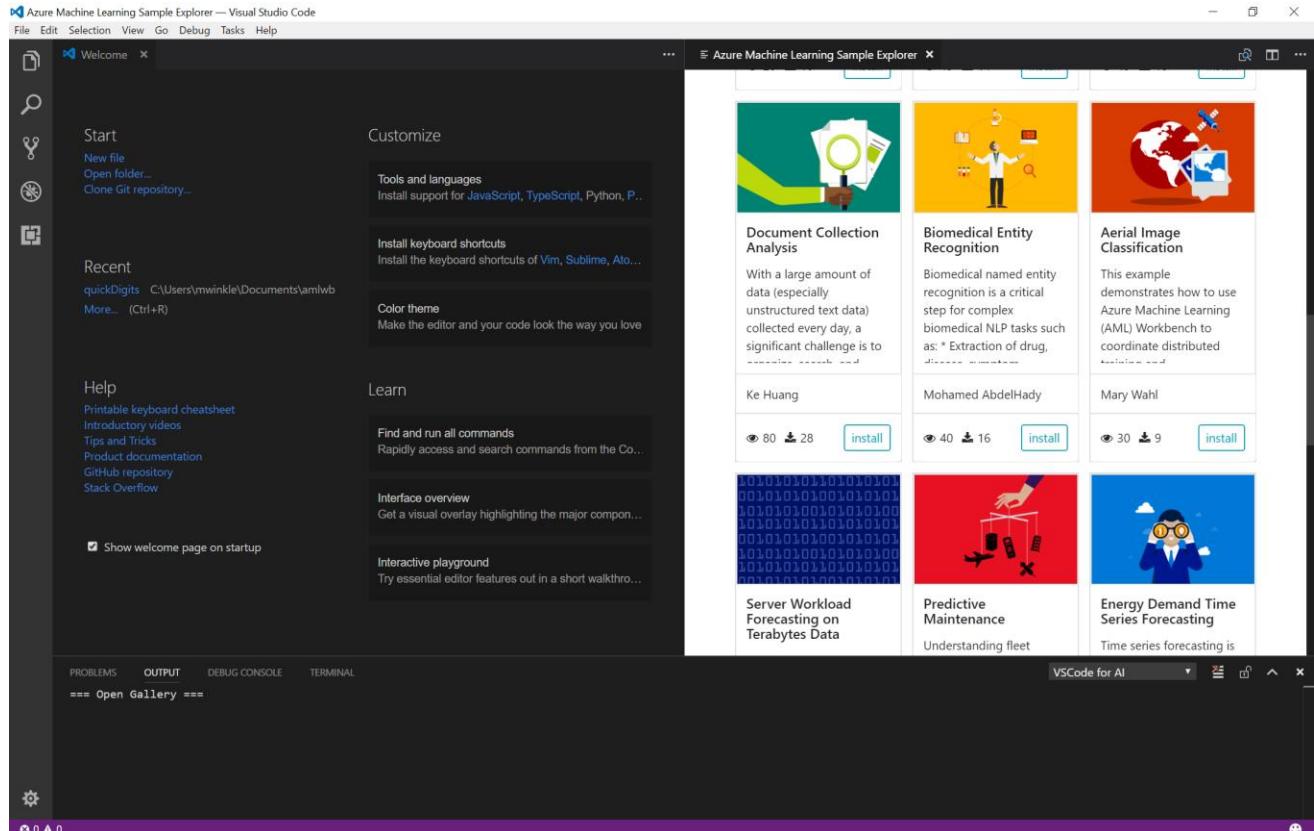
VS Code extension with deep integration to Azure ML

End to end development environment, from new project through training

Support for remote training

Job management

On top of all of the goodness of VS Code
(Python, Jupyter, Git, etc)



WORKBENCH

WHAT IS IT?

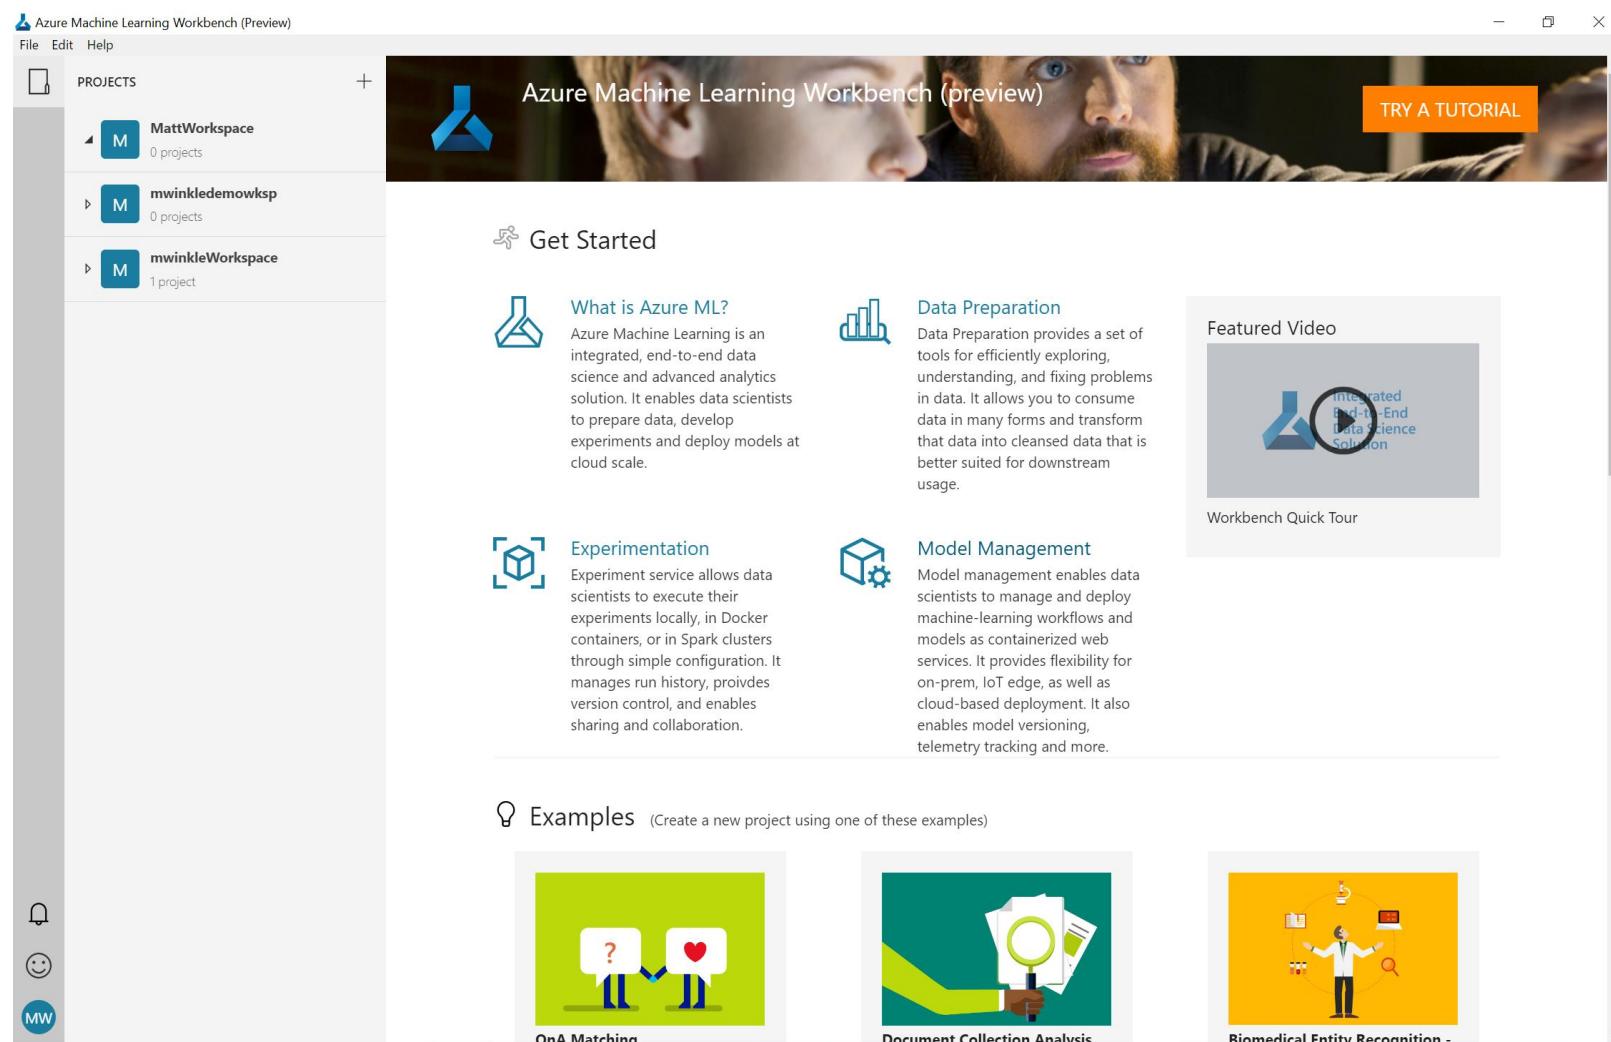
Windows and Mac based companion for AI development

Full environment set up (Python, Jupyter, etc)

Embedded notebooks

Run History and Comparison experience

New data wrangling tools



AI POWERED DATA WRANGLING

Rapidly sample, understand, and prep data

Leverage [PROSE](#) and more for intelligent, data prep by example

Extend/customize transforms and featurization through Python

Generate Python and Pyspark for execution at scale

Azure Machine Learning Workbench (Preview)

Project Dashboard prepStuff BostonWeather TrainDataV6

DATAFLOWS

TrainDataV6

	abc Hour	abc Weekday	# start statio...	# RideInitiat...	# N_DryBulb...	# N_Relative...	# N_WindSp...	
1	12AM-2AM	Thu	115	2	0.29383886255...	0.42528735632...	0.28169014084...	
2	12AM-2AM	Thu	80	1	0.29383886255...	0.42528735632...	0.28169014084...	
3	12AM-2AM	Thu	91	2	0.29383886255...	0.42528735632...	0.28169014084...	
4	12AM-2AM	Thu	105	1	0.29383886255...	0.42528735632...	0.28169014084...	
5	12AM-2AM	Thu	88	1	0.29383886255...	0.42528735632...	0.28169014084...	
6	2AM-4AM	Thu	68	1	0.30331753554...	0.40229885057...	0.33802816901...	
7	4AM-6AM	Thu	117	1	0.29857819905...	0.42528735632...	0.45070422535...	
8	8AM-10AM	Thu	67	1	0.33649289099...	0.32758620689...	0.42253521126...	
9	8AM-10AM	Thu	75	1	0.33649289099...	0.32758620689...	0.42253521126...	
10	8AM-10AM	Thu	115	1	0.33649289099...	0.32758620689...	0.42253521126...	
11	8AM-10AM	Thu	88	1	0.33649289099...	0.32758620689...	0.42253521126...	
12	8AM-10AM	Thu	90	1	0.33649289099...	0.32758620689...	0.42253521126...	
13	8AM-10AM	Thu	116	1	0.33649289099...	0.32758620689...	0.42253521126...	
14	10AM-12PM	Thu	88	1	0.37440758293...	0.24712643678...	0.49295774647...	
15	10AM-12PM	Thu	95	1	0.37440758293...	0.24712643678...	0.49295774647...	
16	10AM-12PM	Thu	116	2	0.37440758293...	0.24712643678...	0.49295774647...	
17	10AM-12PM	Thu	110	1	0.37440758293...	0.24712643678...	0.49295774647...	

INPECTORS

Top 6 values of "Hour"

Scatter Plot: start station id, RideInitiationCount

Box Plot: start station id

MACHINE LEARNING & AI PORTFOLIO

WHEN TO USE WHAT?

Build your own or consume pre-trained models?

Which experience do you want?

Deployment target

What engine(s) do you want to use?

Microsoft
ML & AI
products

Build your
own

Consume

Azure Machine Learning

Cognitive services, bots

Code first

Visual tooling

(On-prem)
ML Server

(cloud)
AML services (Preview)

(cloud)
AML Studio

On-prem
Hadoop

SQL
Server

SQL
Server

Spark

Hadoop

Azure
Batch

DSVM

Azure
Container
Service